

A/B Testing: Statistical Foundations and Practical Implementation

Diogo Ribeiro
ESMAD - Instituto Politécnico do Porto

November 19, 2025

Contents

Preface	xvii
1 Introduction to A/B Testing	1
1.1 Learning Objectives	1
1.2 What is A/B Testing?	1
1.2.1 Definition and Core Concepts	1
1.2.2 From Agricultural Fields to Digital Platforms	2
1.2.3 A/B Testing vs. Other Methodologies	3
1.3 Why A/B Testing Matters	5
1.3.1 The Business Case for Experimentation	5
1.3.2 Evidence-Based Culture	6
1.3.3 Examples from Industry Leaders	6
1.4 Basic Terminology and Concepts	7
1.4.1 Experimental Units and Assignment	7
1.4.2 Control and Treatment Groups	7
1.4.3 Metrics and Key Performance Indicators	8
1.4.4 Statistical Significance	8
1.4.5 Effect Size	9
1.5 Types of A/B Tests	9
1.5.1 Simple A/B Tests	9
1.5.2 A/B/n Tests (Multiple Variants)	9
1.5.3 Multivariate Tests (Factorial Designs)	10
1.5.4 Sequential Testing	10
1.5.5 Multi-Armed Bandits	11
1.6 Real-World Examples	11
1.6.1 Website Conversion Optimization	11
1.6.2 Email Marketing Campaigns	13
1.6.3 Product Feature Testing	15
1.6.4 Pricing Experiments	16
1.7 When A/B Testing Works (and When It Doesn't)	18
1.7.1 Ideal Scenarios for A/B Testing	18
1.7.2 When A/B Testing Struggles	18
1.8 Summary	19
1.9 Further Reading	19
1.10 Exercises	19
1.10.1 Conceptual Questions	19
1.10.2 Quantitative Problems	20
1.10.3 Programming Exercises	20

1.10.4	Critical Thinking	21
2	Statistical Foundations	23
2.1	Learning Objectives	23
2.2	Introduction	23
2.3	Probability Theory Fundamentals	24
2.3.1	Probability Spaces and Random Variables	24
2.3.2	Expected Value and Variance	24
2.3.3	The Central Limit Theorem	27
2.3.4	Law of Large Numbers	30
2.4	Key Probability Distributions for A/B Testing	31
2.4.1	Bernoulli and Binomial Distributions	31
2.4.2	Normal (Gaussian) Distribution	32
2.4.3	Student's t-Distribution	32
2.4.4	Beta Distribution	33
2.5	Statistical Inference	36
2.5.1	Population vs. Sample	36
2.5.2	Sampling Distributions	36
2.5.3	Standard Error	37
2.5.4	Confidence Intervals	37
2.6	Hypothesis Testing Framework	42
2.6.1	The Logic of Hypothesis Testing	42
2.6.2	Test Statistics and p-values	43
2.6.3	Significance Level and Decision Rule	43
2.6.4	Type I and Type II Errors	44
2.6.5	Two-Sample Z-Test for Proportions	44
2.7	Effect Size Measures	50
2.7.1	Why Effect Size Matters	50
2.7.2	Absolute vs. Relative Differences	50
2.7.3	Cohen's h for Proportions	51
2.7.4	Odds Ratio	51
2.7.5	Cohen's d for Continuous Metrics	52
2.7.6	Practical Guidelines for Effect Size Reporting	56
2.8	Frequentist vs. Bayesian Inference	57
2.8.1	Frequentist Framework	57
2.8.2	Bayesian Framework	58
2.8.3	When to Use Each Approach	59
2.9	Summary	59
2.10	Further Reading	60
2.11	Exercises	60
2.11.1	Conceptual Questions	60
2.11.2	Probability and Distributions	60
2.11.3	Hypothesis Testing	61
2.11.4	Confidence Intervals	61
2.11.5	Effect Sizes	62
2.11.6	Programming Projects	62

3	Experimental Design Principles	65
3.1	Learning Objectives	65
3.2	Introduction	65
3.3	Fundamental Principles of Experimental Design	66
3.3.1	Randomization	66
3.3.2	Control Groups	67
3.3.3	Blocking and Stratification	67
3.3.4	Blinding	68
3.4	Randomization Techniques	69
3.4.1	Simple Randomization	69
3.4.2	Block Randomization	73
3.4.3	Stratified Randomization	75
3.4.4	Cluster Randomization	80
3.5	Confounding Variables and Bias	86
3.5.1	Selection Bias	86
3.5.2	Survivorship Bias	87
3.5.3	Simpson's Paradox	87
3.5.4	Network Effects	92
3.6	Validity Considerations	92
3.6.1	Internal Validity	93
3.6.2	External Validity	93
3.6.3	Construct Validity	94
3.6.4	Statistical Conclusion Validity	95
3.7	Practical Design Decisions	95
3.7.1	Unit of Randomization	95
3.7.2	Temporal Considerations	96
3.7.3	Geographic Considerations	96
3.7.4	Platform-Specific Issues	97
3.8	Summary	97
3.9	Further Reading	97
3.10	Exercises	98
3.10.1	Conceptual Questions	98
3.10.2	Design Exercises	98
3.10.3	Implementation Exercises	98
3.10.4	Analysis Exercises	99
3.10.5	Simulation Projects	99
4	Sample Size Calculation and Statistical Power	101
4.1	Learning Objectives	101
4.2	Introduction	101
4.3	Power Analysis Fundamentals	101
4.3.1	Definition of Statistical Power	102
4.3.2	Factors Affecting Power	102
4.3.3	The Type I and Type II Error Tradeoff	103
4.3.4	Power Curves and Interpretation	103
4.4	Sample Size Formulas	107
4.4.1	Sample Size for Proportions (Conversion Rates)	107
4.4.2	Sample Size for Means (Continuous Metrics)	111

4.4.3	Sample Size for Survival Data	115
4.4.4	Sample Size for Ratio Metrics	115
4.5	Practical Considerations	116
4.5.1	Minimum Detectable Effect (MDE)	116
4.5.2	Business Significance vs. Statistical Significance	119
4.5.3	Cost-Benefit Analysis of Sample Sizes	120
4.5.4	Duration Estimation	121
4.6	Advanced Topics	122
4.6.1	Unequal Sample Sizes	122
4.6.2	Sequential Testing Implications	122
4.6.3	Multiple Metrics Considerations	123
4.7	Implementation Tools	123
4.7.1	Python Implementation	123
4.7.2	R Implementation	126
4.8	Summary	128
4.9	Further Reading	129
4.10	Exercises	129
4.10.1	Conceptual Questions	129
4.10.2	Calculation Exercises	129
4.10.3	Programming Exercises	130
4.10.4	Applied Exercises	131
4.10.5	Advanced Projects	132
5	Hypothesis Testing for A/B Tests	133
5.1	Learning Objectives	133
5.2	The Framework of Hypothesis Testing	133
5.2.1	Structure of a Hypothesis Test	133
5.2.2	One-Sided vs. Two-Sided Tests	134
5.3	Z-Test for Proportions	134
5.3.1	Two-Sample Z-Test	134
5.3.2	Assumptions and Validity	136
5.4	T-Test for Means	137
5.4.1	Two-Sample T-Test	137
5.5	Understanding P-Values	138
5.5.1	Correct Interpretation	138
5.5.2	Common Misinterpretations	138
5.5.3	Statistical Significance vs. Practical Significance	139
5.6	Chi-Squared Test for Categorical Data	139
5.6.1	Pearson's Chi-Squared Test	139
5.6.2	Fisher's Exact Test	139
5.7	Permutation Tests	140
5.7.1	Permutation Test Procedure	140
5.8	Summary	143
5.9	Further Reading	143
5.10	Exercises	143

6	Multiple Testing and False Discovery Rate	145
6.1	Learning Objectives	145
6.2	The Multiple Testing Problem	145
6.2.1	Probability of False Positives	145
6.2.2	Sources of Multiplicity in A/B Testing	146
6.3	Family-Wise Error Rate (FWER)	146
6.3.1	Bonferroni Correction	146
6.3.2	Holm-Bonferroni Method	147
6.3.3	Šidák Correction	149
6.4	False Discovery Rate (FDR)	149
6.4.1	Benjamini-Hochberg Procedure	150
6.4.2	Choosing Between FWER and FDR	151
6.5	Multiple Testing Strategies for A/B Testing	151
6.5.1	Hierarchical Testing	151
6.5.2	Pre-specification and Registration	151
6.5.3	Handling Multiple Variants	152
6.6	Practical Recommendations	152
6.7	Summary	152
6.8	Further Reading	153
6.9	Exercises	153
7	Bayesian Approaches to A/B Testing	155
7.1	Learning Objectives	155
7.2	Bayesian vs. Frequentist Paradigms	155
7.2.1	Key Philosophical Differences	155
7.2.2	Bayes' Theorem for A/B Testing	155
7.3	Conjugate Priors for Common Distributions	156
7.3.1	Beta-Binomial Model for Conversion Rates	156
7.3.2	Normal-Normal Model for Continuous Metrics	158
7.4	Bayesian A/B Testing	159
7.4.1	Comparing Two Conversion Rates	159
7.4.2	Expected Loss	160
7.5	Prior Selection	161
7.5.1	Types of Priors	161
7.5.2	Prior Sensitivity Analysis	162
7.6	Advantages of Bayesian A/B Testing	162
7.7	Limitations and Challenges	163
7.8	Bayesian vs. Frequentist: When to Use Which	163
7.9	Summary	164
7.10	Further Reading	164
7.11	Exercises	164
8	Advanced Experimental Designs	165
8.1	Learning Objectives	165
8.2	Factorial Designs	165
8.2.1	Full Factorial Designs	165
8.2.2	Statistical Model for Factorial Experiments	166
8.2.3	Fractional Factorial Designs	172

8.3	Multi-Armed Bandits	173
8.3.1	The Exploration-Exploitation Tradeoff	173
8.3.2	ϵ -Greedy Algorithm	173
8.3.3	Upper Confidence Bound (UCB) Algorithm	174
8.3.4	Thompson Sampling (Bayesian Bandit)	175
8.3.5	When to Use Bandits vs. Traditional A/B Tests	180
8.4	Sequential Testing	181
8.4.1	The Peeking Problem	181
8.4.2	Group Sequential Testing	181
8.5	Network Effects and Interference	187
8.5.1	Types of Interference	187
8.5.2	SUTVA Violation	187
8.5.3	Mitigation Strategies	187
8.6	Switchback Designs	188
8.6.1	Structure	188
8.6.2	Analysis	188
8.6.3	Practical Considerations	189
8.7	Summary	189
8.8	Further Reading	189
8.9	Exercises	190
8.9.1	Conceptual Questions	190
8.9.2	Calculation Exercises	190
8.9.3	Programming Exercises	190
8.9.4	Applied Exercises	191
8.9.5	Advanced Projects	192
9	Technical Implementation	193
9.1	Learning Objectives	193
9.2	Experimentation Infrastructure	193
9.2.1	Core Components	193
9.2.2	Architecture Patterns	193
9.3	Randomization Implementation	194
9.3.1	Hash-Based Randomization	194
9.3.2	Handling Edge Cases	195
9.4	Configuration Management	195
9.5	Logging and Data Collection	196
9.5.1	Event Schema	196
9.5.2	Data Quality Checks	198
9.6	Integration Patterns	199
9.6.1	Feature Flags Integration	199
9.6.2	Metrics Collection	200
9.7	Common Implementation Pitfalls	200
9.7.1	Assignment Bugs	200
9.7.2	Survivorship Bias	201
9.7.3	Logging Failures	201
9.8	Summary	201
9.9	Further Reading	201
9.10	Exercises	201

10 Data Analysis and Result Interpretation	203
10.1 Learning Objectives	203
10.2 Analysis Workflow	203
10.2.1 Pre-Analysis Checks	203
10.2.2 Exploratory Data Analysis	205
10.3 Handling Data Issues	206
10.3.1 Missing Data	206
10.3.2 Outliers	206
10.4 Variance Reduction Techniques	208
10.4.1 CUPED (Controlled-Experiment Using Pre-Experiment Data) . .	208
10.4.2 Stratification	209
10.5 Result Interpretation	209
10.5.1 Statistical vs. Practical Significance	209
10.5.2 Communicating Uncertainty	210
10.6 Decision Making Framework	210
10.6.1 Quantitative Criteria	210
10.6.2 Qualitative Factors	210
10.6.3 Decision Matrix	210
10.7 Summary	210
10.8 Further Reading	211
10.9 Exercises	211
11 Case Studies and Real-World Applications	213
11.1 Learning Objectives	213
11.2 Case Study 1: Google Search Results and Page Load Time	213
11.2.1 Context and Business Problem	213
11.2.2 Experimental Design	213
11.2.3 Results and Analysis	214
11.2.4 Deep Dive: Causal Mechanism	214
11.2.5 Statistical Analysis Framework	214
11.2.6 Lessons Learned	218
11.3 Case Study 2: E-commerce Checkout Optimization	218
11.3.1 Context	218
11.3.2 Experimental Design	218
11.3.3 Overall Results	219
11.3.4 Segmentation Analysis	219
11.3.5 Decision Framework	223
11.3.6 Post-Launch Monitoring	224
11.3.7 Lessons Learned	224
11.4 Case Study 3: Netflix Thumbnail Personalization	224
11.4.1 Context	224
11.4.2 Technical Challenge	224
11.4.3 Multi-Stage Experimental Approach	225
11.4.4 Experimental Design Details (Stage 2)	225
11.4.5 Results	226
11.4.6 Implementation Considerations	226
11.4.7 Lessons Learned	226
11.5 Case Study 4: Failed Experiment—Subscription Paywall	226

11.5.1	Context	227
11.5.2	Design	227
11.5.3	Results: The Good and The Bad	227
11.5.4	What Went Wrong	227
11.5.5	Correct Analysis	228
11.5.6	Lessons Learned	231
11.6	Common Anti-Patterns and How to Avoid Them	231
11.6.1	Anti-Pattern 1: P-Hacking	231
11.6.2	Anti-Pattern 2: HiPPO (Highest Paid Person's Opinion)	232
11.6.3	Anti-Pattern 3: Analysis Paralysis	232
11.6.4	Anti-Pattern 4: Novelty Effect Ignorance	233
11.7	Industry-Specific Considerations	233
11.7.1	E-commerce	233
11.7.2	Social Media	233
11.7.3	SaaS/B2B	234
11.7.4	Healthcare/EdTech	234
11.8	Building an Experimentation Culture: Meta-Lessons	235
11.8.1	Infrastructure and Tooling	235
11.8.2	Process and Governance	235
11.8.3	Culture and Education	235
11.9	Summary	235
11.10	Further Reading	236
11.11	Exercises	236
11.11.1	Case Study Analysis	236
11.11.2	Business Impact Calculation	236
11.11.3	Anti-Pattern Detection	237
11.11.4	Industry Application	237
11.11.5	Meta-Analysis Project	237
12	Ethics and Best Practices	239
12.1	Learning Objectives	239
12.2	Ethical Foundations	239
12.2.1	Core Ethical Principles	239
12.2.2	The Belmont Report	239
12.3	Informed Consent	240
12.3.1	Consent in Digital Experiments	240
12.3.2	When Explicit Consent is Required	240
12.3.3	Balancing Innovation and Consent	240
12.4	Ethical Guardrails	240
12.4.1	Minimal Risk Standard	240
12.4.2	Experimentation Review Board	241
12.4.3	Automated Guardrails	241
12.5	Privacy and Data Protection	243
12.5.1	Data Minimization	243
12.5.2	Regulatory Compliance	243
12.5.3	Internal Data Governance	244
12.6	Transparency and Communication	244
12.6.1	Internal Transparency	244

12.6.2	External Transparency	245
12.7	Building Experimentation Culture	245
12.7.1	Cultural Principles	245
12.7.2	Organizational Structure	245
12.7.3	Training and Education	246
12.8	Best Practices Checklist	246
12.8.1	Design Phase	246
12.8.2	Implementation Phase	246
12.8.3	Analysis Phase	247
12.8.4	Decision Phase	247
12.9	Common Pitfalls and How to Avoid Them	247
12.10	Summary	247
12.11	Further Reading	248
12.12	Exercises	248
Mathematical Proofs		249
.1	Mathematical Proofs	249
.1.1	Proof: Unbiasedness of Sample Mean	249
.1.2	Proof: Variance of Sample Mean	249
.1.3	Proof: Beta-Binomial Conjugacy	250
.1.4	Proof: Sample Size Formula for Proportions	250
.1.5	Proof: CUPED Variance Reduction	251
.1.6	Central Limit Theorem	252
.1.7	Bonferroni Inequality	252
.1.8	Delta Method	252
Code Examples		253
.2	Complete Code Examples	253
.2.1	Comprehensive A/B Test Analysis Class (Python)	253
.2.2	Experiment Analysis Pipeline (R)	258
.2.3	Experiment Simulation Framework	260
Statistical Tables		263
.3	Statistical Tables	263
.3.1	Standard Normal Distribution (Z) Critical Values	263
.3.2	Student's t-Distribution Critical Values ($t_{\alpha/2,df}$)	263
.3.3	Chi-Squared Distribution Critical Values ($\chi^2_{\alpha,df}$)	263
.3.4	Sample Size Requirements (Two Proportions)	264
.3.5	Minimum Detectable Effect (MDE)	264
.3.6	Bonferroni Corrected Significance Levels	265
.3.7	Power vs. Sample Size	265
.3.8	Beta Distribution Quantiles	265
.3.9	Conversion Rate Effect Size Examples	266
.3.10	Design Effect for Cluster Randomization	266
.3.11	Quick Reference Formulas	266
.3.12	Decision Matrix	267

List of Figures

List of Tables

2.1	Type I and Type II errors in hypothesis testing	44
4.1	Type I and Type II errors in hypothesis testing	103

Preface

The idea for this book crystallized during a late-night debugging session at Mysense.ai, where I serve as Lead Data Scientist. We were investigating why an A/B test that should have been straightforward—comparing two checkout button designs—had produced confusing results. The statistical significance was marginal, the effect size was tiny, yet business stakeholders were pressing for a decision. As I walked through the analysis with the team, I realized that the real challenge wasn't the mathematics or the code—it was the gap between statistical rigor and practical decision-making under uncertainty.

This experience, repeated countless times across industry and academia, revealed a fundamental problem: most practitioners learn A/B testing from one of two extremes. Either they study rigorous statistical theory without adequate grounding in real-world constraints, or they learn practical "recipes" without understanding the underlying assumptions and limitations. Both approaches leave critical gaps. The theorist may design a perfectly valid experiment that's impossible to implement in production. The practitioner may run hundreds of tests without recognizing systematic biases that invalidate their conclusions.

This book attempts to bridge that divide.

Why This Book, Why Now

The practice of experimentation has undergone a quiet revolution over the past two decades. What began as a specialized technique used primarily in pharmaceutical trials and agricultural research has become the cornerstone of decision-making at technology companies, e-commerce platforms, government agencies, and organizations of all types. Google runs tens of thousands of experiments annually. Amazon tests nearly every product change. Netflix uses experimentation to optimize everything from recommendation algorithms to thumbnail images. Even traditional industries—retail, finance, health-care—increasingly embrace data-driven experimentation.

Yet this democratization of A/B testing has created new challenges. The barrier to running an experiment has dropped dramatically—modern platforms make it trivially easy to split traffic and measure outcomes. But the barrier to running a *good* experiment, one that produces trustworthy insights and leads to sound decisions, remains high. Statistical literacy has not kept pace with experimental practice.

The consequences manifest in multiple ways. I've seen teams celebrate "winning" variants that were merely lucky noise. I've watched organizations make million-dollar decisions based on p-values they didn't understand. I've encountered experiments that violated basic statistical assumptions yet were treated as gospel. Most concerning, I've observed a growing cynicism about experimentation itself when poorly designed tests produce conflicting or irreproducible results.

These experiences, accumulated through my dual roles—as Lead Data Scientist at Mysense.ai, where I design and analyze experiments for real products and real customers, and as Researcher and Instructor at ESMAD (Instituto Politécnico do Porto), where I teach the statistical foundations to the next generation of data scientists—convinced me that a different kind of book was needed.

Who This Book Is For

I wrote this book with several overlapping audiences in mind, all united by a need to understand both the "how" and the "why" of A/B testing.

Data Scientists and Analysts working in industry will find practical guidance grounded in statistical rigor. You face daily pressures to move fast, to test everything, to "let the data decide." This book will help you move quickly *and* correctly, understanding when standard methods apply and when you need something more sophisticated. You'll learn to spot the subtle implementation issues that invalidate experiments, to communicate uncertainty clearly to stakeholders, and to push back on unreasonable demands ("Can we just peek at the results?") with solid statistical reasoning.

Graduate Students in statistics, data science, computer science, or related fields will find a modern treatment of experimental design that connects classical theory to contemporary practice. The mathematical development is rigorous—I don't shy away from proofs when they illuminate understanding—but always motivated by practical questions. You'll see how the Central Limit Theorem justifies real-world approximations, why Bayesian methods handle certain problems elegantly, and where theory meets the messiness of production systems.

Researchers in academia or industry who design experiments will discover how classical experimental design principles translate to digital contexts. If you're familiar with randomized controlled trials in clinical or social science settings, you'll learn how online experimentation differs—both the new opportunities (instant results, massive sample sizes) and new challenges (network effects, novelty effects, implementation bugs).

Product Managers and Technical Leaders who make decisions based on experimental evidence need to understand what the numbers really mean. This book will help you ask the right questions: Is this sample size adequate? Are we measuring the right thing? What assumptions underlie this analysis? You'll develop intuition for when to trust experimental results and when to dig deeper.

Practitioners Transitioning from Other Fields who have statistical training but are new to online experimentation will find familiar concepts (hypothesis testing, confidence intervals, power analysis) applied in novel contexts. You'll learn the domain-specific challenges—how to handle sample ratio mismatches, why sequential testing requires special care, when to use variance reduction techniques.

What unites all these audiences is a need for both depth and breadth—statistical foundations solid enough to support critical decisions, practical guidance specific enough to implement tomorrow, and judgment sophisticated enough to navigate the grey areas where theory doesn't provide clear answers.

What Makes This Book Different

Several principles guided the development of this book:

Theory Motivated by Practice. Every statistical concept is introduced in response to a real problem. Why do we need the Central Limit Theorem? Because we want to make probability statements about sample means. Why study multiple testing corrections? Because without them, we fool ourselves with false discoveries. The mathematics serves the practice, not the other way around.

Practice Grounded in Theory. Conversely, every practical technique is connected to its theoretical foundation. When I present a Python function for calculating sample size, I also derive the formula and explain the assumptions. When we implement Bayesian A/B testing, we understand why conjugate priors work and when they're appropriate. You won't just learn what to do—you'll understand why it works and when it doesn't.

Code as First-Class Citizen. This is not a book where code appears as an afterthought. The implementations in Python and R are production-quality, documented, and ready to adapt. I've run every code example, debugged the edge cases, and provided not just the happy path but also error handling and validation. Because I've been the person at 2 AM trying to debug a broken experiment, I know that implementation details matter.

Honest About Limitations. Statistics education often presents methods as if they solve problems cleanly. Real experimentation is messier. This book acknowledges the gaps: where assumptions are likely violated, when approximations break down, which questions statistics cannot answer. I share failures alongside successes, because learning what doesn't work is as valuable as learning what does.

Modern Methodologies. While grounded in classical statistics, this book embraces contemporary approaches. We cover Bayesian methods seriously, not as a curiosity but as a practical alternative with real advantages in certain contexts. We explore multi-armed bandits, sequential testing, and variance reduction techniques that have emerged from modern practice. We address challenges—network effects, novelty effects, metric dilution—that classical texts don't consider.

How to Use This Book

The book is designed for multiple modes of engagement:

As a Course Text: Graduate courses can proceed through chapters sequentially, using exercises to reinforce learning. I've structured the material to build progressively, with each chapter assuming mastery of previous content. Problem sets range from mathematical derivations to simulation studies to analysis of real datasets.

As a Reference: Practitioners can dive directly into specific chapters as needs arise. Comprehensive cross-referencing, a detailed index, and self-contained chapters enable targeted learning. Need to understand multiple testing corrections? Chapter 6 provides everything required. Want to implement Bayesian A/B testing? Chapter 7 is your guide.

As Professional Development: Working data scientists can deepen their understanding systematically. Perhaps you've been running A/B tests for years using standard tools but want to understand the statistical foundations more rigorously. Or you're comfortable with frequentist methods but curious about Bayesian approaches. The book supports both deepening and broadening your expertise.

Each chapter includes learning objectives, worked examples with complete solutions, production-ready code, exercises at varying difficulty levels, and suggestions for further reading. The appendices provide mathematical proofs for those seeking rigor, compre-

hensive code libraries for those seeking implementation, and statistical tables for quick reference.

A Personal Note

Writing this book has been a journey of synthesis—bringing together years of experience across industry and academia, reconciling practical constraints with theoretical ideals, and distilling complex ideas into accessible explanations. My work at Mysense.ai has provided invaluable real-world context: the pressure to move fast, the complexity of production systems, the challenge of communicating uncertainty to stakeholders who want certainty. My teaching and research at ESMAD has forced me to articulate foundations clearly, to answer "why" questions rigorously, and to stay current with methodological developments.

This dual perspective—practitioner and educator, industry and academia—has shaped every page. I've tried to write the book I wish had existed when I began my career, one that would have saved me from numerous mistakes and accelerated my understanding of this powerful methodology.

Experimentation, at its best, represents a commitment to learning from reality rather than assuming we already know. It's an acknowledgment that our intuitions, however informed, are fallible—and that careful measurement can surprise us. This intellectual humility, coupled with methodological rigor, enables genuine discovery.

My hope is that this book helps you design better experiments, analyze results more critically, make sounder decisions, and ultimately contribute to the growing culture of evidence-based decision making. Whether you're optimizing a website, evaluating a policy intervention, or conducting scientific research, the principles of rigorous experimentation apply.

The field of A/B testing continues to evolve. New methods emerge, new challenges arise, new applications appear. But the fundamentals—randomization, statistical inference, careful interpretation—endure. Master these foundations, and you'll be equipped not just for today's problems but for tomorrow's challenges.

Diogo Ribeiro
Lead Data Scientist, Mysense.ai
Researcher and Instructor, ESMAD
Instituto Politécnico do Porto
Porto, Portugal
November 19, 2025

Chapter 1

Introduction to A/B Testing

1.1 Learning Objectives

After completing this chapter, readers will be able to:

- Define A/B testing and explain its fundamental principles
- Trace the historical evolution from agricultural experiments to digital applications
- Distinguish A/B testing from other experimental methodologies
- Understand core terminology and concepts used throughout the book
- Recognize different types of A/B tests and their appropriate applications
- Implement basic A/B tests using Python and statistical libraries
- Identify real-world scenarios where A/B testing adds value

1.2 What is A/B Testing?

1.2.1 Definition and Core Concepts

Definition 1.1 (A/B Test). An A/B test is a randomized controlled experiment in which a population is randomly divided into two or more groups to compare the effects of different treatments, interventions, or variations on a measured outcome. The fundamental structure includes:

- A **control group** (A) receiving the existing or baseline treatment
- One or more **treatment groups** (B, C, etc.) receiving new variations
- **Random assignment** ensuring unbiased allocation
- **Quantifiable outcomes** measured consistently across groups

At its core, A/B testing operationalizes the scientific method for decision-making under uncertainty. Rather than relying on intuition, expert opinion, or observational

data—all valuable but potentially misleading—we directly test hypotheses through controlled experimentation. This approach answers a deceptively simple question: *Does this change actually work?*

The elegance of A/B testing lies in its simplicity. By randomly assigning subjects to groups and measuring outcomes, we isolate the causal effect of our intervention from confounding variables. If users assigned to the new checkout flow convert at a higher rate than those in the control group, and this difference exceeds what we'd expect from random variation, we have evidence that the change caused the improvement.

1.2.2 From Agricultural Fields to Digital Platforms

The intellectual foundations of A/B testing trace back over a century, evolving through several distinct phases that reveal both continuity and transformation in experimental thinking.

Early Pioneers: Fisher and Agricultural Science

In the 1920s, Sir Ronald Fisher revolutionized experimental design while working at Rothamsted Experimental Station in England. Fisher confronted a practical problem: how to determine which fertilizers, crop varieties, and farming techniques produced the best yields when countless variables—soil quality, weather, pest pressure—varied unpredictably across fields and seasons.

Fisher's solution introduced three principles that remain foundational today [?]:

1. **Randomization:** Rather than systematically assigning treatments (which might confound treatment effects with field characteristics), Fisher advocated random assignment to eliminate systematic bias.
2. **Replication:** Testing each treatment on multiple plots enabled statistical estimation of variability and increased confidence in results.
3. **Blocking:** Grouping similar experimental units (e.g., adjacent plots) and randomizing within blocks reduced noise and improved precision.

These principles solved the agricultural scientist's dilemma: how to draw valid causal inferences in the face of uncontrolled variation. They work equally well for digital experimentation, though the scale and speed have changed dramatically.

Clinical Trials and Medical Research

The mid-20th century saw experimental design principles migrate to medical research. The 1948 streptomycin tuberculosis trial is often cited as the first modern randomized controlled trial (RCT) [?]. Medical researchers recognized that comparing treatments without randomization introduced selection bias—sicker patients might systematically receive more aggressive treatments, confounding causal inference.

Clinical trials formalized concepts still central to A/B testing:

- **Blinding:** Keeping subjects (and often researchers) unaware of treatment assignment to prevent expectation effects

- **Intention-to-treat analysis:** Analyzing all randomized subjects regardless of compliance
- **Power analysis:** Calculating required sample sizes before conducting experiments
- **Interim monitoring:** Checking results during trials for ethical reasons, requiring special statistical methods

The Digital Revolution

The internet transformed experimentation. What took months or years in agricultural or medical research could now happen in hours. Early online experiments at Amazon, Google, and Microsoft in the late 1990s and early 2000s demonstrated that digital platforms offered unprecedented advantages:

- **Massive scale:** Millions of subjects available
- **Low cost:** Marginal cost of adding one more subject approaches zero
- **Rapid iteration:** Tests complete in days or weeks, not months or years
- **Precise measurement:** Every click, view, and conversion tracked automatically
- **Perfect randomization:** Computers execute random assignment flawlessly

However, digital experimentation also introduced new challenges that Fisher never contemplated:

- **Implementation bugs:** Code errors can invalidate experiments
- **Network effects:** One user's treatment may affect others
- **Novelty effects:** Users respond differently to new features initially
- **Metric dilution:** Measuring too many outcomes increases false positives
- **Continuous monitoring:** "Peeking" at results invalidates classical statistics

Modern A/B testing inherits Fisher's foundational principles while adapting them to digital contexts. This synthesis—classical rigor meets contemporary scale—defines the field today.

1.2.3 A/B Testing vs. Other Methodologies

Understanding A/B testing requires distinguishing it from alternative approaches to learning and decision-making.

Observational Studies

Observational studies analyze existing data without intervention. For example, examining whether users who engage with a particular feature have higher retention. While valuable for hypothesis generation, observational studies suffer from confounding: perhaps engaged users were already more loyal, and feature usage is a symptom rather than a cause.

Example 1.1 (Simpson’s Paradox). An e-commerce platform observes that customers who use the mobile app spend 50% more than website-only customers. However, randomizing users to receive an app download prompt shows no effect on spending. Why?

The observational data confounds app usage with customer type. High-value customers independently discovered and adopted the app. The app didn’t cause higher spending; higher-spending customers chose to use it. Randomization eliminates this confounding.

Multi-Armed Bandits

Multi-armed bandit (MAB) algorithms dynamically allocate more traffic to better-performing variants, optimizing cumulative performance during the test. Unlike A/B tests with fixed allocation, bandits adapt.

Consider testing K variants with unknown success rates $\theta_1, \dots, \theta_K$. The regret after T trials is:

$$R(T) = T \cdot \max_k \theta_k - \sum_{t=1}^T \mathbb{E}[r_t] \quad (1.1)$$

where r_t is the reward at trial t .

MAB algorithms like Thompson Sampling or Upper Confidence Bound (UCB) minimize regret by balancing exploration (testing uncertain variants) and exploitation (favoring likely winners).

Trade-offs:

- A/B tests provide clean statistical inference but sacrifice performance during testing
- MABs optimize during testing but complicate inference (non-uniform assignment probabilities)
- A/B tests suit scenarios where learning is the goal
- MABs suit scenarios where optimization during testing matters (e.g., ad selection)

We explore bandits in depth in Chapter 8.

Quasi-Experimental Methods

When randomization is infeasible, quasi-experimental designs attempt causal inference from observational data under assumptions. Methods include:

- **Difference-in-differences:** Comparing trends before/after intervention in treatment vs. control groups
- **Regression discontinuity:** Exploiting arbitrary cutoffs (e.g., test score thresholds) for assignment

- **Instrumental variables:** Using variables that affect treatment but not outcomes directly
- **Propensity score matching:** Creating matched pairs with similar characteristics

These methods complement A/B testing when experimentation is impractical but carry stronger assumptions.

1.3 Why A/B Testing Matters

1.3.1 The Business Case for Experimentation

A/B testing delivers measurable value through better decision-making, risk mitigation, and organizational learning.

Quantifiable Return on Investment

Consider a concrete example. An e-commerce company with 1 million monthly visitors and 3% conversion rate generates 30,000 orders monthly. An A/B test finds that a new checkout flow increases conversion to 3.3%—a 10% relative lift.

Annual impact:

$$\text{Additional conversions} = 1,000,000 \times 12 \times (0.033 - 0.030) = 36,000 \quad (1.2)$$

$$\text{Revenue (at \$50 AOV)} = 36,000 \times \$50 = \$1,800,000 \quad (1.3)$$

If the test cost \$10,000 to design, implement, and analyze, the ROI is 180x. Even conservative estimates assuming 5% lift still yield \$900,000 annually.

De-Risking Product Changes

Not all changes improve outcomes. Experimentation prevents costly mistakes:

Example 1.2 (Microsoft Bing). Microsoft tested displaying ads more prominently, hypothesizing increased revenue. The test showed short-term revenue gains but significant user satisfaction declines and reduced search quality metrics. Without testing, full deployment would have damaged the product and brand [?].

Failures are valuable. Learning that an idea doesn't work before full rollout saves development resources and protects user experience.

Compound Returns from Iteration

Individual experiments often yield small improvements—2-5% lifts. The power comes from compounding many improvements over time.

If a team runs one experiment weekly, each improving a key metric by 2%, annual compound growth is:

$$(1.02)^{52} - 1 \approx 1.83 = 183\% \quad (1.4)$$

This calculation oversimplifies (not all tests succeed, metrics may not be independent), but illustrates how systematic experimentation drives dramatic cumulative gains.

1.3.2 Evidence-Based Culture

Beyond individual tests, A/B testing transforms organizational decision-making.

Objective Evaluation of Ideas

Experimentation democratizes decision-making. Good ideas can come from anywhere—junior engineers, customer support, users themselves—and all receive objective evaluation. This meritocracy of ideas contrasts with authority-based or seniority-based decision-making.

Learning from Surprises

The most valuable experiments challenge assumptions. When results contradict expert opinion or conventional wisdom, we learn something fundamental about user behavior, market dynamics, or product mechanics.

Example 1.3 (Netflix Personalization). Netflix hypothesized that showing personalized movie thumbnails would increase engagement. Testing revealed the effect was far larger than anticipated—personalized images increased viewing by 20-30% for some content. This insight led to broader personalization initiatives worth billions in subscriber value [?].

1.3.3 Examples from Industry Leaders

Google: Search Result Optimization

Google runs over 10,000 experiments annually on Search alone. A famous 2009 experiment tested displaying 10, 20, or 30 search results per page. The hypothesis: more results would reduce pagination and improve user experience.

Results showed that increasing results to 30 caused a 20% increase in page load time (from 0.4s to 0.9s). This seemingly small delay significantly reduced searches per user—users abandoned searches rather than wait. The test demonstrated that even half-second delays meaningfully impact behavior at scale.

Amazon: Everything is Testable

Amazon's culture of experimentation is legendary. The company tests nearly every product change before deployment. Examples include:

- Testing "Customers who bought this also bought" recommendation algorithms
- Optimizing product page layouts
- Testing checkout flows and payment options
- Evaluating search result ranking algorithms

Amazon founder Jeff Bezos famously said, "Our success at Amazon is a function of how many experiments we do per year, per month, per week, per day."

Facebook: Continuous Optimization

Facebook treats its platform as a continuous experiment. The company tests features on small user subsets before gradual rollout. Notable examples:

- News Feed ranking algorithms optimizing for engagement
- Ad formats and placement testing
- Like button variations and reactions
- Video autoplay behavior

However, Facebook also illustrates ethical challenges—a 2014 "emotional contagion" study manipulating News Feed content sparked controversy about informed consent in experimentation.

1.4 Basic Terminology and Concepts

Precise language enables clear thinking about experiments. This section establishes terminology used throughout the book.

1.4.1 Experimental Units and Assignment

Definition 1.2 (Experimental Unit). The experimental unit is the entity randomized to treatment. Common units include:

- **Users:** Individual people (most common in digital experiments)
- **Sessions:** Individual visits or browsing sessions
- **Page views:** Individual page loads
- **Clusters:** Groups (e.g., geographic regions, schools, stores)

The choice of experimental unit affects both implementation and analysis. User-level randomization ensures consistency across sessions but requires persistent identifiers. Session-level randomization is simpler but creates inconsistent experiences for repeat visitors.

1.4.2 Control and Treatment Groups

Definition 1.3 (Control Group). The control group (often called variant A) receives the existing or baseline experience. It serves as the comparison benchmark for measuring treatment effects.

Definition 1.4 (Treatment Group). Treatment groups (B, C, etc.) receive new variations being tested. In simple A/B tests, there is one treatment group; multivariate tests have multiple.

Why maintain a control? Even if we believe a change will improve outcomes, we need the counterfactual—what would have happened without the change? External factors (seasonality, marketing campaigns, news events) affect all users. Comparing treatment to control isolates the causal effect of our intervention.

1.4.3 Metrics and Key Performance Indicators

Definition 1.5 (Metric). A metric is a quantifiable measure of user behavior or business outcome. Metrics can be:

- **Binary:** Conversion (yes/no), click (yes/no)
- **Continuous:** Revenue per user, time on site, pages viewed
- **Count:** Number of purchases, search queries, shares

Effective metrics share several properties:

1. **Measurable:** Can be tracked reliably and consistently
2. **Attributable:** Can be linked to experimental units
3. **Timely:** Observable within reasonable experimental duration
4. **Sensitive:** Responds to interventions we care about
5. **Aligned:** Correlates with long-term business goals

Metric Hierarchies

Mature experimentation programs organize metrics hierarchically:

- **Primary metric:** The main outcome for decision-making (e.g., revenue per user)
- **Secondary metrics:** Supporting measures providing context (e.g., conversion rate, average order value)
- **Guardrail metrics:** Measures we monitor but don't want to degrade (e.g., page load time, error rate)

This hierarchy prevents "metric shopping"—testing many metrics until finding significance by chance.

1.4.4 Statistical Significance

Definition 1.6 (Statistical Significance). A result is statistically significant if the observed difference between groups is unlikely to have occurred by chance alone under the null hypothesis of no effect. Conventionally, we use significance level $\alpha = 0.05$, meaning we tolerate a 5% chance of false positives.

The p-value quantifies this: $p = P(\text{observe data this extreme} \mid H_0 \text{ true})$. If $p < \alpha$, we reject the null hypothesis and conclude a statistically significant effect exists.

Critical distinction: Statistical significance differs from practical significance. A tiny effect might be statistically significant with enough data but too small to matter. Conversely, a meaningful effect might not reach significance with limited data.

1.4.5 Effect Size

Definition 1.7 (Effect Size). Effect size quantifies the magnitude of difference between groups, independent of sample size. Common measures include:

- **Absolute difference:** $\hat{p}_B - \hat{p}_A$ (e.g., 10.5% - 10.0% = 0.5 percentage points)
- **Relative lift:** $\frac{\hat{p}_B - \hat{p}_A}{\hat{p}_A}$ (e.g., 0.5% / 10.0% = 5% relative increase)
- **Cohen's d:** $d = \frac{\bar{x}_B - \bar{x}_A}{s_{\text{pooled}}}$ (standardized difference for continuous outcomes)

Effect sizes communicate practical importance and facilitate comparison across experiments. A 5% relative lift in conversion rate from 10% baseline has different business impact than the same lift from 1% baseline, even if both are statistically significant.

1.5 Types of A/B Tests

Experimentation encompasses various designs suited to different objectives and constraints.

1.5.1 Simple A/B Tests

The canonical design compares one treatment against a control:

- 50% of users see control (A)
- 50% see treatment (B)
- Measure and compare a primary metric

When to use:

- Testing a single change (button color, headline, algorithm)
- Clear hypothesis about improvement
- Want clean statistical inference
- First test of a new idea

1.5.2 A/B/n Tests (Multiple Variants)

Testing multiple variations simultaneously:

- Control (A): 50%
- Treatment 1 (B): 25%
- Treatment 2 (C): 25%

Trade-offs:

- **Efficiency:** Test multiple ideas with same traffic

- **Complexity:** More comparisons increase multiple testing burden
- **Power:** Splitting traffic reduces power for each comparison

Proper multiple testing corrections (Chapter 6) are essential to maintain valid error rates.

1.5.3 Multivariate Tests (Factorial Designs)

Factorial designs test multiple factors simultaneously. A 2^k factorial tests k binary factors:

Example 1.4 (2×2 Factorial Design). Testing email subject line (A/B) and send time (morning/evening):

	Morning	Evening
Subject A	25%	25%
Subject B	25%	25%

This design estimates:

- Main effect of subject line
- Main effect of send time
- Interaction: does subject line effect depend on time?

Advantages:

- Test multiple factors with one experiment
- Detect interactions between factors
- More efficient than separate tests

Challenges:

- Complexity grows exponentially (2^k combinations)
- Interpretation harder with many factors
- Requires larger samples for interaction detection

Chapter 8 covers factorial designs comprehensively.

1.5.4 Sequential Testing

Traditional A/B tests fix sample size in advance and analyze once at the end. Sequential testing allows monitoring during the experiment with appropriate statistical controls.

Group sequential methods: Pre-specify analysis times (e.g., 25%, 50%, 75%, 100% of planned sample) and use spending functions to control Type I error.

Advantages:

- Stop early if strong evidence emerges

- Ethical in some contexts (e.g., medical trials)
- Reduce time to decision

Challenges:

- Requires special statistical methods
- "Peeking" at results invalidates standard p-values
- More complex implementation

1.5.5 Multi-Armed Bandits

Rather than fixed allocation, bandit algorithms dynamically shift traffic toward better performers:

Algorithm 1 Simple ϵ -Greedy Bandit

```

Initialize:  $N_i = 0$ ,  $\hat{\mu}_i = 0$  for variants  $i = 1, \dots, K$ 
for each user  $t = 1, 2, \dots$  do
  With probability  $\epsilon$ : select random variant (explore)
  With probability  $1 - \epsilon$ : select  $\arg \max_i \hat{\mu}_i$  (exploit)
  Observe outcome  $r_t$ 
  Update estimates for selected variant
end for

```

When bandits outperform A/B tests:

- Many variants to test (e.g., ad creative)
- Performance during testing matters (opportunity cost high)
- Priors available from similar tests
- Willing to trade inference clarity for optimization

1.6 Real-World Examples

1.6.1 Website Conversion Optimization

E-commerce Checkout Flow

An online retailer redesigns the checkout process, reducing steps from 5 to 3 by combining shipping and billing information.

Hypothesis: Fewer steps will reduce abandonment and increase completion rate.

Metrics:

- Primary: Checkout completion rate (conversions / checkout initiations)
- Secondary: Revenue per session, average order value
- Guardrails: Form errors, customer support contacts

Results:

- Completion rate increased from 68% to 71.5% ($p < 0.01$)
- Revenue per session up 6.2%
- No increase in form errors or support contacts

Business impact: With 100,000 monthly checkout initiations, 3.5 percentage point improvement yields 3,500 additional monthly conversions. At \$80 average order value: \$280,000 monthly revenue increase, or \$3.36M annually.

Listing 1.1: Basic A/B Test Analysis for Conversion Rates

```

1 import numpy as np
2 from scipy import stats
3 import pandas as pd
4
5 def analyze_conversion_test(control_conversions, control_total,
6                             treatment_conversions, treatment_total,
7                             alpha=0.05):
8     """
9     Analyze A/B test for conversion rates.
10
11     Parameters:
12     -----
13     control_conversions: int
14         Number of conversions in control
15     control_total: int
16         Total users in control
17     treatment_conversions: int
18         Number of conversions in treatment
19     treatment_total: int
20         Total users in treatment
21     alpha: float
22         Significance level (default 0.05)
23
24     Returns:
25     -----
26     dict: Test results including statistics and interpretation
27     """
28     # Calculate proportions
29     p_control = control_conversions / control_total
30     p_treatment = treatment_conversions / treatment_total
31
32     # Pooled proportion for null hypothesis
33     p_pooled = (control_conversions + treatment_conversions) / \
34                (control_total + treatment_total)
35
36     # Standard error under null
37     se_null = np.sqrt(p_pooled * (1 - p_pooled) *
38                       (1/control_total + 1/treatment_total))
39
40     # Z-statistic
41     z_stat = (p_treatment - p_control) / se_null
42
43     # Two-sided p-value
44     p_value = 2 * (1 - stats.norm.cdf(abs(z_stat)))
45

```



```

46     # Confidence interval (using unpooled variance)
47     se_unpooled = np.sqrt(p_control*(1-p_control)/control_total +
48                             p_treatment*(1-p_treatment)/treatment_total)
49     z_crit = stats.norm.ppf(1 - alpha/2)
50     ci_lower = (p_treatment - p_control) - z_crit * se_unpooled
51     ci_upper = (p_treatment - p_control) + z_crit * se_unpooled
52
53     # Calculate effect sizes
54     absolute_lift = p_treatment - p_control
55     relative_lift = absolute_lift / p_control if p_control > 0 else np.
        nan
56
57     # Determine statistical significance
58     is_significant = p_value < alpha
59
60     results = {
61         'control_rate': p_control,
62         'treatment_rate': p_treatment,
63         'absolute_lift': absolute_lift,
64         'relative_lift': relative_lift,
65         'z_statistic': z_stat,
66         'p_value': p_value,
67         'ci_lower': ci_lower,
68         'ci_upper': ci_upper,
69         'is_significant': is_significant,
70         'alpha': alpha
71     }
72
73     return results
74
75 # Example: Checkout flow test
76 results = analyze_conversion_test(
77     control_conversions=6800,
78     control_total=10000,
79     treatment_conversions=7150,
80     treatment_total=10000
81 )
82
83 print("A/B Test Results")
84 print("=" * 50)
85 print(f"Control conversion rate: {results['control_rate']:.1%}")
86 print(f"Treatment conversion rate: {results['treatment_rate']:.1%}")
87 print(f"\nAbsolute lift: {results['absolute_lift']:.1%}")
88 print(f"Relative lift: {results['relative_lift']:.1%}")
89 print(f"\nZ-statistic: {results['z_statistic']:.3f}")
90 print(f"P-value: {results['p_value']:.4f}")
91 print(f"95% CI: [{results['ci_lower']:.1%}, {results['ci_upper']:.1%}]")
92 print(f"\nStatistically significant: {results['is_significant']}")

```

1.6.2 Email Marketing Campaigns

Subject Line Testing

An e-commerce company tests three email subject lines for a promotional campaign:

- A (Control): "Weekend Sale - 20% Off"

- B: "Flash Sale: Save 20% This Weekend Only"
- C: "Your Exclusive 20% Weekend Discount"

Experimental design:

- Randomly assign 100,000 subscribers equally to three variants
- Send simultaneously (Saturday 9 AM)
- Measure open rate, click-through rate, conversion rate

Listing 1.2: Multi-Variant Email Test Analysis

```

1 def analyze_email_test(variants_data, metric='open_rate', alpha=0.05):
2     """
3     Analyze multi-variant email test with Bonferroni correction.
4
5     Parameters:
6     -----
7     variants_data: dict
8         Dictionary mapping variant names to (successes, total) tuples
9     metric: str
10        Metric name for reporting
11    alpha: float
12        Family-wise error rate
13
14    Returns:
15    -----
16    DataFrame: Comparison of all variants with adjusted p-values
17    """
18    variant_names = list(variants_data.keys())
19    n_variants = len(variant_names)
20    n_comparisons = n_variants * (n_variants - 1) // 2
21
22    # Bonferroni correction
23    alpha_adjusted = alpha / n_comparisons
24
25    results = []
26
27    # Pairwise comparisons
28    for i in range(n_variants):
29        for j in range(i + 1, n_variants):
30            var1, var2 = variant_names[i], variant_names[j]
31
32            successes1, total1 = variants_data[var1]
33            successes2, total2 = variants_data[var2]
34
35            # Perform test
36            test_results = analyze_conversion_test(
37                successes1, total1, successes2, total2, alpha_adjusted
38            )
39
40            results.append({
41                'comparison': f'{var2}_vs_{var1}',
42                f'{var1}_{metric}': test_results['control_rate'],
43                f'{var2}_{metric}': test_results['treatment_rate'],
44                'absolute_diff': test_results['absolute_lift'],

```

```

45         'relative_lift': test_results['relative_lift'],
46         'p_value': test_results['p_value'],
47         'significant': test_results['is_significant']
48     })
49
50     df = pd.DataFrame(results)
51
52     print(f"Email_Test_Results_{metric}")
53     print("=" * 70)
54     print(f"Number_of_comparisons: {n_comparisons}")
55     print(f"Bonferroni-adjusted_alpha: {alpha_adjusted:.4f}")
56     print()
57
58     return df
59
60 # Example data
61 email_data = {
62     'A_Control': (4200, 33333),      # Open rate ~12.6%
63     'B_Flash': (4500, 33333),      # Open rate ~13.5%
64     'C_Exclusive': (4800, 33334)    # Open rate ~14.4%
65 }
66
67 results_df = analyze_email_test(email_data, metric='open_rate')
68 print(results_df.to_string(index=False))

```

1.6.3 Product Feature Testing

Recommendation Algorithm

A streaming platform tests a new content recommendation algorithm against the existing system.

Hypothesis: Machine learning-based personalized recommendations will increase viewing time compared to popularity-based recommendations.

Experimental design:

- User-level randomization (consistent experience across sessions)
- 50/50 split: 5 million users per variant
- Duration: 4 weeks (capture weekly patterns, reduce novelty effects)

Metrics:

- Primary: Hours watched per user per week
- Secondary: Content diversity (unique titles watched), completion rate
- Guardrails: Churn rate, customer satisfaction scores

Results demonstrate heterogeneous treatment effects:

User Segment	Control	Treatment	Lift
New users (< 1 month)	8.2 hrs	10.1 hrs	+23%
Established (1-6 months)	12.5 hrs	13.8 hrs	+10%
Long-term (> 6 months)	15.3 hrs	15.1 hrs	-1%
Overall	12.4 hrs	13.2 hrs	+6.5%

New algorithm benefits new and moderately engaged users but provides no value (or slight harm) to long-term users who have strong preferences. This insight suggests opportunities for segment-specific optimization.

1.6.4 Pricing Experiments

Pricing is perhaps the most sensitive area for experimentation, requiring careful ethical and business considerations.

Subscription Tier Testing

A SaaS company tests two pricing strategies:

- Control: \$29/month, \$290/year (17% annual discount)
- Treatment: \$29/month, \$249/year (28% annual discount)

Hypothesis: Larger annual discount will increase annual plan adoption without significantly reducing revenue.

Key considerations:

- **Ethics:** All users see both options; no one charged different prices for same product
- **Segment:** Test on new signups only (existing customers unaffected)
- **Duration:** 8 weeks to capture full subscription cycle

Listing 1.3: Revenue Analysis for Pricing Test

```

1 def analyze_pricing_test(control_monthly, control_annual,
2                           treatment_monthly, treatment_annual,
3                           avg_months_subscribed=12):
4     """
5     Analyze pricing experiment considering customer lifetime value.
6
7     Parameters:
8     -----
9     control_monthly, treatment_monthly: tuple
10    count, price for monthly subscribers
11    control_annual, treatment_annual: tuple
12    count, price for annual subscribers
13    avg_months_subscribed: int
14    Average subscription duration in months
15
16    Returns:
17    -----
18    dict: Revenue metrics and comparison
19    """
20    # Unpack data
21    c_monthly_count, c_monthly_price = control_monthly
22    c_annual_count, c_annual_price = control_annual
23    t_monthly_count, t_monthly_price = treatment_monthly
24    t_annual_count, t_annual_price = treatment_annual
25
26    # Calculate annual plan adoption rate
27    control_total = c_monthly_count + c_annual_count

```

```

28     treatment_total = t_monthly_count + t_annual_count
29
30     c_annual_rate = c_annual_count / control_total
31     t_annual_rate = t_annual_count / treatment_total
32
33     # Calculate average revenue per customer (annualized)
34     c_avg_revenue = (
35         (c_monthly_count * c_monthly_price * avg_months_subscribed +
36          c_annual_count * c_annual_price) / control_total
37     )
38
39     t_avg_revenue = (
40         (t_monthly_count * t_monthly_price * avg_months_subscribed +
41          t_annual_count * t_annual_price) / treatment_total
42     )
43
44     # Statistical test on annual adoption rate
45     from scipy.stats import chi2_contingency
46
47     contingency_table = [
48         [c_annual_count, c_monthly_count],
49         [t_annual_count, t_monthly_count]
50     ]
51     chi2, p_value, dof, expected = chi2_contingency(contingency_table)
52
53     results = {
54         'control_annual_rate': c_annual_rate,
55         'treatment_annual_rate': t_annual_rate,
56         'annual_adoption_lift': (t_annual_rate - c_annual_rate) /
57             c_annual_rate,
58         'control_avg_revenue': c_avg_revenue,
59         'treatment_avg_revenue': t_avg_revenue,
60         'revenue_diff': t_avg_revenue - c_avg_revenue,
61         'revenue_diff_pct': (t_avg_revenue - c_avg_revenue) /
62             c_avg_revenue,
63         'chi2_statistic': chi2,
64         'p_value': p_value,
65         'significant': p_value < 0.05
66     }
67
68     return results
69
70 # Example: Pricing test data
71 pricing_results = analyze_pricing_test(
72     control_monthly=(2800, 29),      # 2800 monthly at $29
73     control_annual=(1200, 290),      # 1200 annual at $290
74     treatment_monthly=(2400, 29),    # 2400 monthly at $29
75     treatment_annual=(1600, 249)     # 1600 annual at $249
76 )
77
78 print("Pricing Test Results")
79 print("=" * 60)
80 print(f"Control annual adoption: {pricing_results['control_annual_rate']:.1%}")
81 print(f"Treatment annual adoption: {pricing_results['treatment_annual_rate']:.1%}")
82 print(f"Lift in annual adoption: {pricing_results['annual_adoption_lift']:.1%}")

```

```

81 print()
82 print(f"Control_avg_revenue:_{pricing_results['control_avg_revenue']:.2f}")
83 print(f"Treatment_avg_revenue:_{pricing_results['treatment_avg_revenue']:.2f}")
84 print(f"Revenue_difference:_{pricing_results['revenue_diff']:+.2f}_{f"({pricing_results['revenue_diff_pct']:+.1%})}")
85
86 print()
87 print(f"Chi-squared:_{pricing_results['chi2_statistic']:.2f},_{f"p-value:_{pricing_results['p_value']:.4f}")
88
89 print(f"Statistically_significant:_{pricing_results['significant']}")

```

Results interpretation: Treatment increased annual plan adoption from 30% to 40%, but the lower price point resulted in slight decrease in average revenue per customer (from \$319.20 to \$314.10 annually). However, higher adoption of annual plans improves retention and reduces churn, making this a strategic win despite short-term revenue impact.

1.7 When A/B Testing Works (and When It Doesn't)

1.7.1 Ideal Scenarios for A/B Testing

A/B testing excels when:

1. **Sufficient traffic:** Adequate users/events to detect meaningful effects
2. **Clear metrics:** Quantifiable outcomes observable in reasonable timeframe
3. **Reversible changes:** Can easily roll back if needed
4. **Limited network effects:** Users' experiences are independent
5. **Ethical feasibility:** Randomization and experimentation are appropriate

1.7.2 When A/B Testing Struggles

Insufficient Sample Size

If your website receives 100 visitors weekly and conversion rate is 2%, detecting a 25% relative improvement (from 2% to 2.5%) requires approximately 6,300 users per variant—over a year of testing.

Long-Term Effects

Some interventions show different short-term and long-term effects. A/B tests typically run weeks to months. Evaluating impacts on 5-year customer lifetime value requires different approaches.

Rare Events

Testing impacts on events occurring once per 10,000 users requires enormous samples. Consider proxy metrics or specialized methods.

Fundamental Redesigns

Radical changes create such different experiences that incremental testing may not apply. Qualitative research, gradual rollouts, and careful monitoring complement experimentation.

Network Effects

Social networks, marketplaces, and communication platforms violate the "no interference" assumption. One user's treatment affects others. Specialized designs (Chapter 8) address this.

1.8 Summary

A/B testing provides a rigorous framework for learning from data and making evidence-based decisions. Its core principles—randomization, controlled comparison, statistical inference—enable causal conclusions that observational studies cannot support.

Key takeaways:

- A/B testing evolved from Fisher's agricultural experiments to become central to digital product development
- Randomization eliminates confounding and enables causal inference
- Proper terminology (control/treatment, metrics, significance, effect size) enables precise thinking
- Multiple test types (simple A/B, multivariate, sequential, bandits) suit different objectives
- Real-world applications span e-commerce, marketing, product features, and pricing
- Understanding both capabilities and limitations guides appropriate application

The subsequent chapters develop the statistical theory, implementation practices, and analytical techniques necessary for rigorous experimentation. We begin in Chapter 2 with probability and statistical foundations.

1.9 Further Reading

- ? : Original exposition of experimental design principles
- ? : Comprehensive modern treatment from industry leaders
- ? : Classical text bridging theory and practice
- ? : Rigorous causal inference framework

1.10 Exercises

1.10.1 Conceptual Questions

1. Understanding Randomization

- (a) Explain in your own words why randomization is essential for causal inference.
- (b) Describe a scenario where observational data would be misleading but a randomized experiment would reveal the truth.

- (c) What could go wrong if we assigned users to variants based on the first letter of their username instead of random assignment?

2. Metrics and Business Impact

- (a) A test shows 10% increase in clicks but no change in revenue. What might explain this? Would you ship the change?
- (b) Propose a complete metric hierarchy (primary, secondary, guardrails) for testing a new mobile app checkout flow.
- (c) Explain the difference between statistical significance and practical significance with a concrete example.

3. Experimental Design

- (a) When would you choose user-level over session-level randomization? Provide three scenarios for each.
- (b) A social network wants to test a new feature but worries about network effects. How would this concern affect their experimental design?
- (c) Describe the trade-offs between A/B testing and multi-armed bandits for selecting email subject lines.

1.10.2 Quantitative Problems

1. Sample Size Intuition

- (a) Your website has 10,000 daily visitors. How long must a test run to accumulate 20,000 users per variant?
- (b) If detecting a 5% effect requires 10,000 users per variant, approximately how many users are needed to detect a 2.5% effect (half as large)?
- (c) Explain why detecting small effects requires disproportionately more data.

2. Effect Size Calculations

- (a) Control: 8.2% conversion. Treatment: 9.0% conversion. Calculate absolute and relative lift.
- (b) A test increases average order value from \$45 to \$48. Express as absolute difference, relative lift, and Cohen's d (assume SD = \$15).
- (c) Why might the same 10% relative improvement have different business value at different baseline rates?

1.10.3 Programming Exercises

1. Basic Analysis

- (a) Use the provided code to analyze: Control (450 conversions / 5000 users), Treatment (520 conversions / 5000 users). Interpret all outputs.
- (b) Modify the code to calculate 90% confidence intervals instead of 95%. How does this change affect the width?

- (c) Implement a function that determines if an observed difference could plausibly be due to chance (p-value interpretation).

2. Simulation Study

- (a) Simulate 10,000 A/B tests where the null hypothesis is true (no effect). What proportion show $p < 0.05$? Why?
- (b) Simulate tests with 10% true relative lift. What proportion correctly detect this effect (statistical power)?
- (c) Create a visualization showing how power changes with sample size for a fixed effect size.

3. Real Data Analysis

- (a) Download a public A/B test dataset (e.g., from Kaggle). Perform complete analysis including data validation, statistical testing, and business recommendation.
- (b) Implement pre-analysis checks: sample ratio mismatch detection, temporal patterns, outlier identification.
- (c) Create a reusable analysis template you could apply to future experiments.

1.10.4 Critical Thinking

1. Case Analysis

- (a) Microsoft Bing's experiment (mentioned earlier) showed short-term revenue gains but long-term user satisfaction declines. Design an experimental protocol that would have detected this issue earlier.
- (b) A company runs 100 experiments simultaneously. Several show statistically significant improvements. What concerns should you have? How would you address them?
- (c) Critics argue that A/B testing leads to incremental optimization rather than breakthrough innovation. How would you respond? When is this criticism valid?

2. Ethical Considerations

- (a) A dating app wants to test whether showing fewer high-quality matches increases engagement (by creating scarcity). Is this ethical? What principles guide your answer?
- (b) When, if ever, should informed consent be required for A/B tests in commercial products?
- (c) Design ethical guidelines for experimentation in your domain of interest.

Chapter 2

Statistical Foundations

2.1 Learning Objectives

After completing this chapter, readers will be able to:

- Apply probability theory fundamentals to experimental design and analysis
- Work with probability distributions essential for A/B testing
- Understand and apply the Central Limit Theorem to justify normal approximations
- Construct and interpret confidence intervals for common metrics
- Distinguish between frequentist and Bayesian inference frameworks
- Formulate and test statistical hypotheses rigorously
- Calculate and interpret effect size measures for practical significance
- Implement statistical analyses using Python and R

2.2 Introduction

A/B testing rests on a foundation of probability theory and statistical inference. While modern software makes running experiments trivially easy, understanding the underlying mathematics is essential for:

1. **Correct Implementation:** Knowing when assumptions are violated
2. **Proper Interpretation:** Understanding what results actually mean
3. **Sound Decisions:** Distinguishing statistical significance from practical importance
4. **Effective Communication:** Explaining uncertainty to stakeholders

This chapter develops these foundations systematically, connecting abstract theory to concrete experimental practice. We emphasize not just the mechanics of calculation but the reasoning that justifies each method.

2.3 Probability Theory Fundamentals

Probability provides the language for reasoning about uncertainty—the fundamental challenge in A/B testing where we observe samples but want to make inferences about populations.

2.3.1 Probability Spaces and Random Variables

Definition 2.1 (Probability Space). A probability space is a triple $(\Omega, \mathcal{F}, P)$ where:

- Ω is the sample space (set of all possible outcomes)
- $\mathcal{F}$ is a σ -algebra of events (subsets of Ω)
- $P : \mathcal{F} \rightarrow [0, 1]$ is a probability measure satisfying:
 1. $P(\Omega) = 1$ (normalization)
 2. For disjoint events $A_1, A_2, \dots$: $P(\bigcup_{i=1}^{\infty} A_i) = \sum_{i=1}^{\infty} P(A_i)$ (countable additivity)

Example 2.1 (A/B Test Sample Space). In a simple A/B test where we observe whether a user converts:

- $\Omega = \{\text{convert}, \text{no convert}\}$
- $\mathcal{F} = \{\emptyset, \{\text{convert}\}, \{\text{no convert}\}, \Omega\}$
- $P(\{\text{convert}\}) = p$, $P(\{\text{no convert}\}) = 1 - p$

For n users, Ω becomes the set of all possible sequences of conversions and non-conversions, containing 2^n elements.

Definition 2.2 (Random Variable). A random variable X is a measurable function $X : \Omega \rightarrow \mathbb{R}$ that maps outcomes to real numbers. We distinguish:

- **Discrete random variables:** Taking countable values with probability mass function (PMF) $p(x) = P(X = x)$
- **Continuous random variables:** Taking uncountable values with probability density function (PDF) $f(x)$ where $P(a \leq X \leq b) = \int_a^b f(x)dx$

2.3.2 Expected Value and Variance

The expected value and variance are the most important summary statistics of a distribution, characterizing location and spread respectively.

Definition 2.3 (Expected Value). For discrete X :

$$\mathbb{E}[X] = \sum_x x \cdot p(x) \quad (2.1)$$

For continuous X :

$$\mathbb{E}[X] = \int_{-\infty}^{\infty} x \cdot f(x)dx \quad (2.2)$$

The expected value represents the long-run average if we could repeat the experiment infinitely many times.

Definition 2.4 (Variance and Standard Deviation).

$$\text{Var}(X) = \mathbb{E}[(X - \mathbb{E}[X])^2] = \mathbb{E}[X^2] - (\mathbb{E}[X])^2 \quad (2.3)$$

The standard deviation is $\sigma = \sqrt{\text{Var}(X)}$.

Variance measures the expected squared deviation from the mean—how spread out the distribution is. Standard deviation has the same units as X , making it more interpretable.

Essential Properties for A/B Testing

Proposition 2.1 (Linearity of Expectation). For random variables $X_1, \dots, X_n$ and constants $a_1, \dots, a_n, b$:

$$\mathbb{E}\left[b + \sum_{i=1}^n a_i X_i\right] = b + \sum_{i=1}^n a_i \mathbb{E}[X_i] \quad (2.4)$$

This holds **regardless of dependence** between variables—a remarkable fact that greatly simplifies calculations.

Proposition 2.2 (Variance of Linear Combinations). For independent random variables $X_1, \dots, X_n$ and constants $a_1, \dots, a_n, b$:

$$\text{Var}\left(b + \sum_{i=1}^n a_i X_i\right) = \sum_{i=1}^n a_i^2 \text{Var}(X_i) \quad (2.5)$$

Note: The constant b vanishes (shifting doesn't change spread), and independence is crucial.

Corollary 2.1 (Variance of Sample Mean). For $X_1, \dots, X_n$ independent with $\mathbb{E}[X_i] = \mu$ and $\text{Var}(X_i) = \sigma^2$:

$$\text{Var}(\bar{X}) = \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) = \frac{\sigma^2}{n} \quad (2.6)$$

This fundamental result shows uncertainty decreases with $\sqrt{n}$, not n . Quadrupling sample size only halves standard error.

Example 2.2 (Computing Means and Variances). Consider an A/B test where users spend time on site. Let X_i represent time spent by user i , with individual $\mathbb{E}[X_i] = 5$ minutes and $\text{Var}(X_i) = 4$ minutes².

For $n = 100$ users:

$$\mathbb{E}[\bar{X}] = \mathbb{E}\left[\frac{1}{100} \sum_{i=1}^{100} X_i\right] = \frac{1}{100} \sum_{i=1}^{100} \mathbb{E}[X_i] = 5 \text{ minutes} \quad (2.7)$$

$$\text{Var}(\bar{X}) = \frac{4}{100} = 0.04 \text{ minutes}^2 \quad (2.8)$$

$$\text{SD}(\bar{X}) = \sqrt{0.04} = 0.2 \text{ minutes} \quad (2.9)$$

The sample mean has the same expected value as individual observations, but much less variability.

Listing 2.1: Simulating Expected Value and Variance

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 from scipy import stats
5
6 sns.set_style("whitegrid")
7 np.random.seed(42)
8
9 # Simulate time spent on site (exponential distribution)
10 true_mean = 5 # minutes
11 n_users = 100
12 n_simulations = 10000
13
14 # Individual observations
15 individual_times = np.random.exponential(true_mean, size=(n_simulations
16     , n_users))
17
18 # Sample means
19 sample_means = individual_times.mean(axis=1)
20
21 # Theoretical values
22 theoretical_var_individual = true_mean**2 # For exponential
23 theoretical_var_sample_mean = theoretical_var_individual / n_users
24 theoretical_sd_sample_mean = np.sqrt(theoretical_var_sample_mean)
25
26 # Empirical values
27 empirical_mean = sample_means.mean()
28 empirical_sd = sample_means.std()
29
30 # Visualization
31 fig, axes = plt.subplots(1, 2, figsize=(14, 5))
32
33 # Individual observations
34 axes[0].hist(individual_times[0, :], bins=50, density=True, alpha=0.7,
35     label='Sample data', color='steelblue')
36 x = np.linspace(0, 25, 1000)
37 axes[0].plot(x, stats.expon.pdf(x, scale=true_mean), 'r-', linewidth=2,
38     label=f'Theoretical PDF')
39 axes[0].axvline(true_mean, color='green', linestyle='--', linewidth=2,
40     label=f'True mean = {true_mean}')
41 axes[0].set_xlabel('Time Spent (minutes)', fontsize=12)
42 axes[0].set_ylabel('Density', fontsize=12)
43 axes[0].set_title('Individual User Times', fontsize=14, fontweight='
44     bold')
45 axes[0].legend()
46 axes[0].grid(alpha=0.3)
47
48 # Sample means
49 axes[1].hist(sample_means, bins=50, density=True, alpha=0.7,
50     label=f'Simulated sample means (n={n_simulations:,})',
51     color='steelblue')
52 axes[1].axvline(empirical_mean, color='red', linestyle='-', linewidth
53     =2,
54     label=f'Empirical mean = {empirical_mean:.3f}')
55 axes[1].axvline(true_mean, color='green', linestyle='--', linewidth=2,
56     label=f'True mean = {true_mean}')
57 axes[1].set_xlabel('Sample Mean (minutes)', fontsize=12)

```

```

54 axes[1].set_ylabel('Density', fontsize=12)
55 axes[1].set_title(f'Distribution of Sample Means (n={n_users})',
56                  fontsize=14, fontweight='bold')
57 axes[1].legend()
58 axes[1].grid(alpha=0.3)
59
60 plt.tight_layout()
61 plt.savefig('expectation_variance_demo.pdf', dpi=300, bbox_inches='
    tight')
62 plt.show()
63
64 # Print comparison
65 print(f"Theoretical SD of sample mean: {theoretical_sd_sample_mean:.4f}
    ")
66 print(f"Empirical SD of sample mean: {empirical_sd:.4f}")
67 print(f"Ratio (should be ~1): {empirical_sd /
    theoretical_sd_sample_mean:.4f}")

```

2.3.3 The Central Limit Theorem

The Central Limit Theorem (CLT) is arguably the most important result in statistics. It justifies using normal approximations in A/B testing, enabling inference even when individual observations are far from normally distributed.

Theorem 2.1 (Central Limit Theorem). Let $X_1, X_2, \dots, X_n$ be independent, identically distributed random variables with $\mathbb{E}[X_i] = \mu$ and $\text{Var}(X_i) = \sigma^2 < \infty$. Define the sample mean:

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i \quad (2.10)$$

Then as $n \rightarrow \infty$:

$$\frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \xrightarrow{d} \mathcal{N}(0, 1) \quad (2.11)$$

Equivalently, for large n :

$$\bar{X}_n \approx \mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right) \quad (2.12)$$

Interpretation: The sample mean of *any* distribution (with finite variance) becomes approximately normally distributed for large sample sizes. The approximation quality improves as n increases.

Implications for A/B Testing

The CLT has profound practical consequences:

1. **Universality:** Whether testing conversions (Bernoulli), revenue (right-skewed), or engagement time (exponential), sample means are approximately normal for sufficient n .
2. **Quantifiable Uncertainty:** The standard error $\text{SE}(\bar{X}) = \sigma/\sqrt{n}$ precisely quantifies estimation uncertainty.

3. **Confidence Intervals:** Normal approximation enables constructing confidence intervals: $\bar{X} \pm z_{\alpha/2} \cdot \text{SE}(\bar{X})$.
4. **Hypothesis Testing:** Test statistics become approximately normal, enabling p-value calculation.

Example 2.3 (CLT for Conversion Rates). Users convert with probability $p = 0.05$. Individual conversions $X_i \sim \text{Bernoulli}(0.05)$ take only values 0 and 1—decidedly non-normal! The PMF is:

$$P(X_i = 0) = 0.95, \quad P(X_i = 1) = 0.05 \quad (2.13)$$

However, with $n = 1000$ users, the sample conversion rate $\hat{p} = \bar{X}_n$ is approximately:

$$\hat{p} \sim \mathcal{N}\left(0.05, \frac{0.05 \times 0.95}{1000}\right) = \mathcal{N}(0.05, 0.0000475) \quad (2.14)$$

Standard error: $\text{SE}(\hat{p}) = \sqrt{0.0000475} \approx 0.0069$ or about 0.69 percentage points.

A 95% confidence interval:

$$0.05 \pm 1.96 \times 0.0069 = [0.0365, 0.0635] \quad (2.15)$$

How Large is "Large Enough"?

A practical question: when is the normal approximation accurate? Rules of thumb:

- **For proportions:** $np \geq 5$ and $n(1 - p) \geq 5$
- **For symmetric distributions:** $n \geq 30$ often sufficient
- **For highly skewed distributions:** May need $n \geq 100$ or more

These are guidelines, not hard rules. Simulation can verify adequacy for specific cases.

Listing 2.2: Demonstrating the Central Limit Theorem

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 from scipy.stats import norm, bernoulli, expon, kstest
5
6 sns.set_style("whitegrid")
7 np.random.seed(42)
8
9 def demonstrate_clt(distribution, params, sample_sizes, n_simulations
10                     =10000):
11     """
12     Demonstrate CLT for different distributions and sample sizes.
13
14     Parameters:
15     -----
16     distribution: str
17     'bernoulli', 'exponential', or 'uniform'
18     params: dict
19     Distribution parameters
20     sample_sizes: list

```



```

20 #####List of sample sizes to demonstrate
21 #####n_simulations: int
22 #####Number of simulations per sample size
23 #####"""
24     n_plots = len(sample_sizes)
25     fig, axes = plt.subplots(2, 2, figsize=(14, 10))
26     axes = axes.flatten()
27
28     # Generate data based on distribution
29     if distribution == 'bernoulli':
30         p = params['p']
31         true_mean = p
32         true_std = np.sqrt(p * (1-p))
33         generate_sample = lambda size: bernoulli.rvs(p, size=size)
34         title_prefix = f"Bernoulli(p={p})"
35     elif distribution == 'exponential':
36         scale = params['scale']
37         true_mean = scale
38         true_std = scale
39         generate_sample = lambda size: expon.rvs(scale=scale, size=size
40         )
41         title_prefix = f"Exponential(  =1/{scale})"
42     elif distribution == 'uniform':
43         a, b = params['a'], params['b']
44         true_mean = (a + b) / 2
45         true_std = np.sqrt((b - a)**2 / 12)
46         generate_sample = lambda size: np.random.uniform(a, b, size)
47         title_prefix = f"Uniform({a},{b})"
48
49     ks_results = []
50
51     for idx, n in enumerate(sample_sizes):
52         # Simulate sample means
53         sample_means = np.array([
54             generate_sample(n).mean()
55             for _ in range(n_simulations)
56         ])
57
58         # Theoretical normal distribution for sample mean
59         theoretical_mean = true_mean
60         theoretical_std = true_std / np.sqrt(n)
61
62         # Standardize for normality test
63         standardized = (sample_means - theoretical_mean) /
64             theoretical_std
65         ks_stat, p_value = kstest(standardized, 'norm')
66         ks_results.append((n, ks_stat, p_value))
67
68         # Create theoretical normal curve
69         x_range = np.linspace(
70             theoretical_mean - 4*theoretical_std,
71             theoretical_mean + 4*theoretical_std,
72             1000
73         )
74         theoretical_pdf = norm.pdf(x_range, theoretical_mean,
75             theoretical_std)
76
77         # Plot

```

```

75     ax = axes[idx]
76     ax.hist(sample_means, bins=50, density=True, alpha=0.7,
77            color='steelblue', label='Simulated_sample_means',
78            edgecolor='black')
79     ax.plot(x_range, theoretical_pdf, 'r-', linewidth=2.5,
80            label='Theoretical_Normal', alpha=0.9)
81     ax.axvline(true_mean, color='green', linestyle='--', linewidth
82               =2,
83               label=f'True_mean={true_mean:.3f}')
84
85     ax.set_title(f'{title_prefix}, n={n}\nSE={theoretical_std
86               :.4f}, ' +
87               f'KS_p-value={p_value:.4f}',
88               fontsize=11, fontweight='bold')
89     ax.set_xlabel('Sample_Mean', fontsize=10)
90     ax.set_ylabel('Density', fontsize=10)
91     ax.legend(fontsize=9)
92     ax.grid(alpha=0.3)
93
94     plt.tight_layout()
95     plt.savefig(f'clt_demonstration_{distribution}.pdf', dpi=300,
96               bbox_inches='tight')
97     plt.show()
98
99     # Print normality test results
100    print(f"\nKolmogorov-Smirnov test for normality ({title_prefix}):")
101    print(f"{n':>6} {'KS_statistic':>15} {'p-value':>10} {'Assessment
102          ':>20}")
103    print("-" * 55)
104    for n, ks_stat, p_val in ks_results:
105        assessment = "Excellent_fit" if p_val > 0.05 else "Poor_fit"
106        print(f"{n:>6} {ks_stat:>15.4f} {p_val:>10.4f} {assessment:>20}
107              ")
108
109    # Demonstrate for different distributions
110    print("=" * 70)
111    print("CLT Demonstration: Bernoulli Distribution (conversion_rates)")
112    print("=" * 70)
113    demonstrate_clt('bernoulli', {'p': 0.05}, [10, 50, 200, 1000])
114
115    print("\n" + "=" * 70)
116    print("CLT Demonstration: Exponential Distribution (session_times)")
117    print("=" * 70)
118    demonstrate_clt('exponential', {'scale': 5}, [10, 50, 200, 1000])
119
120    print("\n" + "=" * 70)
121    print("CLT Demonstration: Uniform Distribution")
122    print("=" * 70)
123    demonstrate_clt('uniform', {'a': 0, 'b': 10}, [10, 50, 200, 1000])

```

2.3.4 Law of Large Numbers

Closely related to the CLT is the Law of Large Numbers, which formalizes the intuition that sample averages converge to population means.

Theorem 2.2 (Weak Law of Large Numbers). Let $X_1, X_2, \dots$ be independent, identically

distributed with $\mathbb{E}[X_i] = \mu$. Then for any $\epsilon > 0$:

$$P(|\bar{X}_n - \mu| > \epsilon) \rightarrow 0 \text{ as } n \rightarrow \infty \quad (2.16)$$

In words: the sample mean converges in probability to the population mean.

Distinction from CLT: The Law of Large Numbers tells us sample means get close to the truth. The CLT tells us how the sample mean is distributed around the truth, enabling quantification of uncertainty.

2.4 Key Probability Distributions for A/B Testing

Several distributions arise repeatedly in experimental analysis. Understanding their properties and relationships is essential for correct application.

2.4.1 Bernoulli and Binomial Distributions

Definition 2.5 (Bernoulli Distribution). A random variable $X \sim \text{Bernoulli}(p)$ represents a single binary trial, taking value 1 (success) with probability p and 0 (failure) with probability $1 - p$.

$$P(X = 1) = p, \quad P(X = 0) = 1 - p \quad (2.17)$$

$$\mathbb{E}[X] = p \quad (2.18)$$

$$\text{Var}(X) = p(1 - p) \quad (2.19)$$

Note that variance is maximized at $p = 0.5$ (maximum uncertainty) and approaches 0 as $p \rightarrow 0$ or $p \rightarrow 1$ (near certainty).

Definition 2.6 (Binomial Distribution). If $X_1, \dots, X_n$ are independent $\text{Bernoulli}(p)$ variables, their sum $Y = \sum_{i=1}^n X_i$ follows $\text{Binomial}(n, p)$, representing the number of successes in n trials.

$$P(Y = k) = \binom{n}{k} p^k (1 - p)^{n-k}, \quad k = 0, 1, \dots, n \quad (2.20)$$

$$\mathbb{E}[Y] = np \quad (2.21)$$

$$\text{Var}(Y) = np(1 - p) \quad (2.22)$$

Practical Application: In A/B testing:

- X_i = whether user i converts (Bernoulli)
- Y = total number of conversions (Binomial)
- $\hat{p} = Y/n$ = observed conversion rate

2.4.2 Normal (Gaussian) Distribution

The normal distribution is ubiquitous in A/B testing, justified by the CLT.

Definition 2.7 (Normal Distribution). A random variable $X \sim \mathcal{N}(\mu, \sigma^2)$ has probability density function:

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right), \quad x \in \mathbb{R} \quad (2.23)$$

Properties:

- $\mathbb{E}[X] = \mu$ (location parameter)
- $\text{Var}(X) = \sigma^2$ (scale parameter)
- Symmetric around μ
- $\approx 68\%$ of mass within $\mu \pm \sigma$, 95% within $\mu \pm 2\sigma$, 99.7% within $\mu \pm 3\sigma$

The **standard normal distribution** $Z \sim \mathcal{N}(0, 1)$ serves as reference. Any normal variable can be standardized:

$$Z = \frac{X - \mu}{\sigma} \sim \mathcal{N}(0, 1) \quad (2.24)$$

Proposition 2.3 (Properties of Normal Distribution). 1. **Linear transformation:** If $X \sim \mathcal{N}(\mu, \sigma^2)$ and $Y = aX + b$, then:

$$Y \sim \mathcal{N}(a\mu + b, a^2\sigma^2) \quad (2.25)$$

2. **Sum of independent normals:** If $X_i \sim \mathcal{N}(\mu_i, \sigma_i^2)$ independently, then:

$$\sum_{i=1}^n X_i \sim \mathcal{N}\left(\sum_{i=1}^n \mu_i, \sum_{i=1}^n \sigma_i^2\right) \quad (2.26)$$

3. **Difference of independent normals:** Crucial for A/B testing:

$$X_1 - X_2 \sim \mathcal{N}(\mu_1 - \mu_2, \sigma_1^2 + \sigma_2^2) \quad (2.27)$$

2.4.3 Student's t-Distribution

When population variance is unknown (the usual case) and must be estimated from data, the t-distribution provides more accurate inference than the normal distribution, especially for small samples.

Definition 2.8 (Student's t-Distribution). If $Z \sim \mathcal{N}(0, 1)$ and $V \sim \chi_\nu^2$ independently, then:

$$T = \frac{Z}{\sqrt{V/\nu}} \sim t_\nu \quad (2.28)$$

where ν is the degrees of freedom parameter.

Key properties:

- Symmetric around 0, like the normal distribution

- Heavier tails than normal (more probability in extremes)
- As $\nu \rightarrow \infty$, $t_\nu \rightarrow \mathcal{N}(0, 1)$
- For $\nu \geq 30$, practically indistinguishable from normal

When to use:

- Comparing means when variance is estimated from data
- Small to moderate sample sizes ($n < 30$ per group)
- Constructing confidence intervals for means
- Two-sample t-tests in A/B experiments

Example 2.4 (t-Distribution in Practice). Comparing average order value between variants with $n_1 = 25$ and $n_2 = 25$:

- Degrees of freedom: $\nu = n_1 + n_2 - 2 = 48$
- Use t_{48} critical values, not $\mathcal{N}(0, 1)$
- For 95% CI: $t_{0.025, 48} = 2.011$ vs. $z_{0.025} = 1.96$

The difference is small but matters for proper uncertainty quantification.

With $n_1 = n_2 = 1000$, we'd have $\nu = 1998$, and $t_{0.025, 1998} \approx 1.961$ — essentially identical to the normal quantile.

2.4.4 Beta Distribution

The Beta distribution is the conjugate prior for the Bernoulli/Binomial likelihood in Bayesian A/B testing (Chapter 7).

Definition 2.9 (Beta Distribution). A random variable $X \sim \text{Beta}(\alpha, \beta)$ has support on $[0, 1]$ with PDF:

$$f(x) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1-x)^{\beta-1}, \quad 0 \leq x \leq 1 \quad (2.29)$$

Properties:

$$\mathbb{E}[X] = \frac{\alpha}{\alpha + \beta} \quad (2.30)$$

$$\text{Var}(X) = \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)} \quad (2.31)$$

$$\text{Mode} = \frac{\alpha - 1}{\alpha + \beta - 2} \quad (\text{for } \alpha, \beta > 1) \quad (2.32)$$

The Beta distribution is extremely flexible:

- $\text{Beta}(1, 1) = \text{Uniform}(0, 1)$ (non-informative prior)
- $\alpha = \beta$ gives symmetric distributions
- $\alpha < \beta$ skews left, $\alpha > \beta$ skews right

- Large $\alpha + \beta$ gives narrow distributions (strong prior beliefs)

Listing 2.3: Visualizing Key Distributions

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 from scipy.stats import binom, norm, t, beta
5
6 sns.set_style("whitegrid")
7 np.random.seed(42)
8
9 fig = plt.figure(figsize=(16, 10))
10
11 # 1. Binomial distributions
12 ax1 = plt.subplot(2, 3, 1)
13 n = 100
14 for p in [0.1, 0.3, 0.5]:
15     x = np.arange(0, n+1)
16     pmf = binom.pmf(x, n, p)
17     ax1.plot(x, pmf, marker='o', markersize=2, label=f'p={p}', alpha
18             =0.7)
19 ax1.set_xlabel('Number of Successes', fontsize=11)
20 ax1.set_ylabel('Probability', fontsize=11)
21 ax1.set_title('Binomial(n=100, p) Distributions', fontsize=12,
22             fontweight='bold')
23 ax1.legend()
24 ax1.grid(alpha=0.3)
25
26 # 2. Normal distributions
27 ax2 = plt.subplot(2, 3, 2)
28 x = np.linspace(-8, 8, 1000)
29 for mu, sigma in [(0, 1), (0, 2), (2, 1)]:
30     pdf = norm.pdf(x, mu, sigma)
31     ax2.plot(x, pdf, linewidth=2.5,
32             label=f'mu={mu}, sigma={sigma}', alpha=0.7)
33 ax2.set_xlabel('x', fontsize=11)
34 ax2.set_ylabel('Density', fontsize=11)
35 ax2.set_title('Normal Distributions', fontsize=12, fontweight='bold')
36 ax2.legend()
37 ax2.grid(alpha=0.3)
38
39 # 3. t-distributions
40 ax3 = plt.subplot(2, 3, 3)
41 x = np.linspace(-4, 4, 1000)
42 ax3.plot(x, norm.pdf(x, 0, 1), linewidth=2.5,
43         label='Normal(0,1)', color='black', linestyle='--')
44 for df in [1, 3, 10, 30]:
45     pdf = t.pdf(x, df)
46     ax3.plot(x, pdf, linewidth=2, label=f'df={df}', alpha=0.7)
47 ax3.set_xlabel('x', fontsize=11)
48 ax3.set_ylabel('Density', fontsize=11)
49 ax3.set_title("Student's t-Distributions", fontsize=12, fontweight='
50             bold')
51 ax3.legend()
52 ax3.grid(alpha=0.3)
53
54 # 4. Beta distributions

```

```

52 ax4 = plt.subplot(2, 3, 4)
53 x = np.linspace(0, 1, 1000)
54 for alpha, beta_param in [(1, 1), (2, 5), (5, 2), (10, 10)]:
55     pdf = beta.pdf(x, alpha, beta_param)
56     ax4.plot(x, pdf, linewidth=2.5,
57             label=f'  $\alpha$ = $\alpha$ {alpha},  $\beta$ = $\beta$ {beta_param}', alpha=0.7)
58 ax4.set_xlabel('x', fontsize=11)
59 ax4.set_ylabel('Density', fontsize=11)
60 ax4.set_title('Beta Distributions', fontsize=12, fontweight='bold')
61 ax4.legend()
62 ax4.grid(alpha=0.3)
63
64 # 5. Normal approximation to Binomial
65 ax5 = plt.subplot(2, 3, 5)
66 n, p = 100, 0.3
67 x_binom = np.arange(0, n+1)
68 pmf = binom.pmf(x_binom, n, p)
69 ax5.bar(x_binom, pmf, alpha=0.6, label='Binomial(100, 0.3)', color='
    steelblue')
70
71 mu = n * p
72 sigma = np.sqrt(n * p * (1-p))
73 x_norm = np.linspace(0, n, 1000)
74 pdf_norm = norm.pdf(x_norm, mu, sigma)
75 ax5.plot(x_norm, pdf_norm, 'r-', linewidth=2.5,
76         label=f'Normal({mu:.0f}, {sigma:.2f})', alpha=0.9)
77 ax5.set_xlabel('Number of Successes', fontsize=11)
78 ax5.set_ylabel('Probability/Density', fontsize=11)
79 ax5.set_title('Normal Approximation to Binomial', fontsize=12,
80             fontweight='bold')
81 ax5.legend()
82 ax5.grid(alpha=0.3)
83
84 # 6. Comparison of confidence intervals: z vs t
85 ax6 = plt.subplot(2, 3, 6)
86 sample_sizes = np.arange(5, 101, 5)
87 z_critical = norm.ppf(0.975)
88 t_critical = [t.ppf(0.975, n-1) for n in sample_sizes]
89
90 ax6.plot(sample_sizes, t_critical, 'b-', linewidth=2.5,
91         label='t critical value', marker='o', markersize=4)
92 ax6.axhline(z_critical, color='red', linestyle='--', linewidth=2.5,
93         label=f'z critical value = {z_critical:.3f}')
94 ax6.set_xlabel('Sample Size', fontsize=11)
95 ax6.set_ylabel('Critical Value (two-tailed,  $\alpha$  = 0.05)', fontsize=11)
96 ax6.set_title('t vs. Normal Critical Values', fontsize=12, fontweight='
    bold')
97 ax6.legend()
98 ax6.grid(alpha=0.3)
99
100 plt.tight_layout()
101 plt.savefig('distributions_visualization.pdf', dpi=300, bbox_inches='
    tight')
102 plt.show()
103
104 # Print practical insights
105 print("\n" + "="*70)
106 print("Practical Insights on Distribution Selection")

```


3. Repeat this process infinitely many times

The distribution of all those statistics is the sampling distribution.

Key insight: Even though we observe only one sample, we can deduce properties of the sampling distribution using probability theory (especially the CLT). This enables quantifying uncertainty.

2.5.3 Standard Error

The standard error (SE) is the standard deviation of a sampling distribution—it quantifies how much a statistic varies across repeated samples.

Definition 2.11 (Standard Error). For a statistic $\hat{\theta}$ with sampling distribution variance $\text{Var}(\hat{\theta})$:

$$\text{SE}(\hat{\theta}) = \sqrt{\text{Var}(\hat{\theta})} \quad (2.33)$$

Common standard errors:

- **Sample mean:** $\text{SE}(\bar{X}) = \sigma/\sqrt{n}$
- **Sample proportion:** $\text{SE}(\hat{p}) = \sqrt{p(1-p)/n}$
- **Difference in proportions:** $\text{SE}(\hat{p}_1 - \hat{p}_2) = \sqrt{\frac{p_1(1-p_1)}{n_1} + \frac{p_2(1-p_2)}{n_2}}$

In practice, we estimate SE by plugging in sample values for unknown parameters:

$$\widehat{\text{SE}}(\hat{p}) = \sqrt{\frac{\hat{p}(1-\hat{p})}{n}} \quad (2.34)$$

2.5.4 Confidence Intervals

Confidence intervals provide a range of plausible values for a parameter, quantifying estimation uncertainty.

Definition 2.12 (Confidence Interval). A $(1-\alpha) \times 100\%$ confidence interval for parameter θ is a random interval $[L, U]$ constructed from sample data such that:

$$P(L \leq \theta \leq U) = 1 - \alpha \quad (2.35)$$

before we observe the data. Common choices: $\alpha = 0.05$ (95% CI) or $\alpha = 0.01$ (99% CI).

Correct interpretation: "If we repeated this experiment many times and constructed a 95% CI each time, approximately 95% of those intervals would contain the true parameter value."

Common misinterpretation: "There is a 95% probability the true parameter is in this specific interval." (This is a Bayesian interpretation, not frequentist!)

Confidence Interval for a Proportion

For large n , by the CLT:

$$\hat{p} \sim \mathcal{N}\left(p, \frac{p(1-p)}{n}\right) \quad (2.36)$$

A $(1 - \alpha)$ confidence interval (Wald interval):

$$\hat{p} \pm z_{\alpha/2} \sqrt{\frac{\hat{p}(1-\hat{p})}{n}} \quad (2.37)$$

where $z_{\alpha/2} = \Phi^{-1}(1 - \alpha/2)$ is the standard normal quantile. For $\alpha = 0.05$: $z_{0.025} = 1.96$.

Example 2.6 (CI for Conversion Rate). Experiment results: 85 conversions from 1,000 users.

$$\hat{p} = 85/1000 = 0.085 \quad (2.38)$$

$$\widehat{\text{SE}}(\hat{p}) = \sqrt{\frac{0.085 \times 0.915}{1000}} = 0.00882 \quad (2.39)$$

95% CI:

$$0.085 \pm 1.96 \times 0.00882 = [0.0677, 0.1023] \quad (2.40)$$

Interpretation: We are 95% confident the true conversion rate is between 6.77% and 10.23%.

Confidence Interval for Difference in Proportions

For comparing two independent proportions (the core of A/B testing):

$$(\hat{p}_2 - \hat{p}_1) \pm z_{\alpha/2} \sqrt{\frac{\hat{p}_1(1-\hat{p}_1)}{n_1} + \frac{\hat{p}_2(1-\hat{p}_2)}{n_2}} \quad (2.41)$$

If the CI contains 0, we cannot conclude a significant difference at level α . If the CI is entirely positive (or negative), we have evidence of a difference.

Listing 2.4: Computing Confidence Intervals

```

1 import numpy as np
2 from scipy import stats
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 sns.set_style("whitegrid")
7
8 def proportion_ci(conversions, n, confidence=0.95, method='wald'):
9     """
10    Calculate confidence interval for a proportion.
11
12    Parameters:
13    -----
14    conversions: int
15    Number of successes

```

```

16  #####n: int
17  #####Total number of trials
18  #####confidence: float
19  #####Confidence level (default 0.95)
20  #####method: str
21  #####'wald' (default), 'wilson', or 'agresti-coull'
22
23  #####Returns:
24  #####-----
25  #####dict with 'estimate', 'ci_lower', 'ci_upper', 'se'
26  #####"""
27      p_hat = conversions / n
28      alpha = 1 - confidence
29      z = stats.norm.ppf(1 - alpha/2)
30
31      if method == 'wald':
32          # Standard Wald interval
33          se = np.sqrt(p_hat * (1 - p_hat) / n)
34          ci_lower = p_hat - z * se
35          ci_upper = p_hat + z * se
36
37      elif method == 'wilson':
38          # Wilson score interval (better for small p or n)
39          denominator = 1 + z**2/n
40          center = (p_hat + z**2/(2*n)) / denominator
41          margin = z * np.sqrt(p_hat*(1-p_hat)/n + z**2/(4*n**2)) /
42                  denominator
43          ci_lower = center - margin
44          ci_upper = center + margin
45          se = np.sqrt(p_hat * (1 - p_hat) / n)
46          p_hat = center # Wilson estimate
47
48      elif method == 'agresti-coull':
49          # Agresti-Coull interval (add 2 successes and 2 failures)
50          n_tilde = n + z**2
51          p_tilde = (conversions + z**2/2) / n_tilde
52          se = np.sqrt(p_tilde * (1 - p_tilde) / n_tilde)
53          ci_lower = p_tilde - z * se
54          ci_upper = p_tilde + z * se
55          p_hat = p_tilde
56
57      return {
58          'estimate': p_hat,
59          'ci_lower': max(0, ci_lower), # Ensure in [0,1]
60          'ci_upper': min(1, ci_upper),
61          'se': se
62      }
63
64  def proportion_diff_ci(conv1, n1, conv2, n2, confidence=0.95):
65      """
66      #####Calculate confidence interval for difference in proportions.
67
68      #####Parameters:
69      #####-----
70      #####conv1, conv2: int
71      #####Number of conversions in each group
72      #####n1, n2: int

```

```

73 #####Sample_sizes
74 #####confidence_:float
75 #####Confidence_level
76
77 #####Returns:
78 #####-----
79 #####dict_with_difference_estimates_and_CI
80 #####"""
81     p1_hat = conv1 / n1
82     p2_hat = conv2 / n2
83     diff = p2_hat - p1_hat
84
85     se = np.sqrt(p1_hat * (1 - p1_hat) / n1 +
86                 p2_hat * (1 - p2_hat) / n2)
87
88     alpha = 1 - confidence
89     z = stats.norm.ppf(1 - alpha/2)
90
91     ci_lower = diff - z * se
92     ci_upper = diff + z * se
93
94     # Relative lift
95     relative_lift = (p2_hat - p1_hat) / p1_hat if p1_hat > 0 else np.
96         nan
97
98     return {
99         'p1': p1_hat,
100         'p2': p2_hat,
101         'difference': diff,
102         'relative_lift': relative_lift,
103         'ci_lower': ci_lower,
104         'ci_upper': ci_upper,
105         'se': se,
106         'significant': not (ci_lower <= 0 <= ci_upper)
107     }
108
109 # Example: Compare CI methods
110 print("="*70)
111 print("Comparing_Confidence_Interval_Methods")
112 print("="*70)
113
114 scenarios = [
115     (50, 500, "Moderate_p,_large_n"),
116     (5, 50, "Small_p,_small_n"),
117     (250, 500, "p_near_0.5,_large_n"),
118     (2, 100, "Very_small_p")
119 ]
120
121 for conversions, n, description in scenarios:
122     print(f"\n{description}:_{conversions}/{n}_conversions")
123     print("-" * 70)
124
125     for method in ['wald', 'wilson', 'agresti-coull']:
126         result = proportion_ci(conversions, n, method=method)
127         width = result['ci_upper'] - result['ci_lower']
128         print(f"{method:15s}:_{[{result['ci_lower']:.4f},_{result['ci_upper']:.4f}]}_")

```

```

129         f"width_{width:.4f}")
130
131 # Example: A/B test comparison
132 print("\n" + "="*70)
133 print("A/B Test: Confidence Interval for Difference")
134 print("="*70)
135
136 control_conv, control_n = 80, 1000
137 treatment_conv, treatment_n = 95, 1000
138
139 result = proportion_diff_ci(control_conv, control_n,
140                             treatment_conv, treatment_n)
141
142 print(f"\nControl: {result['p1']:.2%} conversion rate ({control_conv} / {control_n})")
143 print(f"Treatment: {result['p2']:.2%} conversion rate ({treatment_conv} / {treatment_n})")
144 print(f"\nAbsolute difference: {result['difference']:.4f} ({result['difference']*100:.2f} pp)")
145 print(f"Relative lift: {result['relative_lift']:.2%}")
146 print(f"\n95% CI for difference: [{result['ci_lower']:.4f}, {result['ci_upper']:.4f}]")
147 print(f"Significant at  $\alpha$  = 0.05: {result['significant']}")
148
149 # Visualize confidence intervals
150 fig, axes = plt.subplots(1, 2, figsize=(14, 5))
151
152 # Plot 1: CI comparison across methods
153 ax1 = axes[0]
154 methods = ['Wald', 'Wilson', 'Agresti-Coull']
155 estimates = []
156 lowers = []
157 uppers = []
158
159 for method in ['wald', 'wilson', 'agresti-coull']:
160     ci = proportion_ci(50, 500, method=method)
161     estimates.append(ci['estimate'])
162     lowers.append(ci['ci_lower'])
163     uppers.append(ci['ci_upper'])
164
165 y_pos = np.arange(len(methods))
166 errors = np.array([[est - low for est, low in zip(estimates, lowers)],
167                   [up - est for up, est in zip(uppers, estimates)]])
168
169 ax1.errorbar(estimates, y_pos, xerr=errors, fmt='o', markersize=8,
170              capsize=5, capthick=2)
171 ax1.set_yticks(y_pos)
172 ax1.set_yticklabels(methods)
173 ax1.set_xlabel('Conversion Rate', fontsize=12)
174 ax1.set_title('95% CI Methods Comparison (50/500 conversions)',
175               fontsize=12, fontweight='bold')
176 ax1.grid(axis='x', alpha=0.3)
177
178 # Plot 2: A/B test visualization
179 ax2 = axes[1]
180 groups = ['Control', 'Treatment']
181 conversion_rates = [result['p1'], result['p2']]
182 ci_control = proportion_ci(control_conv, control_n)

```

```

182 ci_treatment = proportion_ci(treatment_conv, treatment_n)
183
184 errors = np.array([
185     [ci_control['estimate'] - ci_control['ci_lower'],
186      ci_treatment['estimate'] - ci_treatment['ci_lower']],
187     [ci_control['ci_upper'] - ci_control['estimate'],
188      ci_treatment['ci_upper'] - ci_treatment['estimate']]
189 ])
190
191 x_pos = np.arange(len(groups))
192 ax2.bar(x_pos, conversion_rates, alpha=0.7, color=['steelblue', 'coral',
193 ])
194 ax2.errorbar(x_pos, conversion_rates, yerr=errors, fmt='none',
195             ecolord='black', capsize=10, capthick=2)
196 ax2.set_xticks(x_pos)
197 ax2.set_xticklabels(groups, fontsize=12)
198 ax2.set_ylabel('Conversion_Rate', fontsize=12)
199 ax2.set_title('A/B Test Comparison with 95% CIs', fontsize=12,
200             fontweight='bold')
201 ax2.grid(axis='y', alpha=0.3)
202
203 # Add significance indicator
204 if result['significant']:
205     ax2.text(0.5, max(conversion_rates)*1.1, 'Significant difference (p
206         <0.05)',
207             ha='center', fontsize=11, bbox=dict(boxstyle='round',
208             facecolor='lightgreen'))
209 else:
210     ax2.text(0.5, max(conversion_rates)*1.1, 'No significant difference
211         (p > 0.05)',
212             ha='center', fontsize=11, bbox=dict(boxstyle='round',
213             facecolor='lightyellow'))
214
215 plt.tight_layout()
216 plt.savefig('confidence_intervals_comparison.pdf', dpi=300, bbox_inches
217             = 'tight')
218 plt.show()

```

2.6 Hypothesis Testing Framework

Hypothesis testing provides a formal framework for making decisions under uncertainty. It's the workhorse methodology for A/B testing analysis.

2.6.1 The Logic of Hypothesis Testing

Hypothesis testing uses proof by contradiction:

1. Assume the null hypothesis (typically "no effect") is true
2. Calculate: "If null were true, how surprising is our observed data?"
3. If data is very surprising under the null, reject the null
4. Otherwise, fail to reject (we don't "accept" the null)

Definition 2.13 (Null and Alternative Hypotheses). • **Null hypothesis (H_0):** The default assumption, typically "no effect" or "no difference"

• **Alternative hypothesis (H_1 or H_A):** What we suspect or want to demonstrate
In A/B testing:

$$H_0 : p_{\text{treatment}} = p_{\text{control}} \quad (\text{no difference}) \quad (2.42)$$

$$H_1 : p_{\text{treatment}} \neq p_{\text{control}} \quad (\text{two-sided}) \quad (2.43)$$

$$\text{or } H_1 : p_{\text{treatment}} > p_{\text{control}} \quad (\text{one-sided}) \quad (2.44)$$

2.6.2 Test Statistics and p-values

Definition 2.14 (Test Statistic). A test statistic is a function of sample data used to test hypotheses. It quantifies how far the data deviates from what we'd expect under H_0 .

For comparing proportions:

$$z = \frac{(\hat{p}_2 - \hat{p}_1) - 0}{\text{SE}_0(\hat{p}_2 - \hat{p}_1)} \quad (2.45)$$

where SE_0 is computed under the null hypothesis of equal proportions.

Definition 2.15 (p-value). The p-value is the probability of observing a test statistic at least as extreme as what we observed, assuming H_0 is true:

$$p\text{-value} = P(\text{test statistic} \geq \text{observed value} \mid H_0 \text{ true}) \quad (2.46)$$

Interpretation:

- Small p-value (< 0.05): Data is surprising under H_0 , reject H_0
- Large p-value (> 0.05): Data is consistent with H_0 , fail to reject

Critical point: The p-value is NOT the probability that H_0 is true. It's the probability of seeing data this extreme if H_0 were true—a subtle but crucial distinction.

2.6.3 Significance Level and Decision Rule

Definition 2.16 (Significance Level). The significance level α is the threshold for rejecting H_0 . Common choices:

- $\alpha = 0.05$ (5%): Standard in most applications
- $\alpha = 0.01$ (1%): More conservative, reduces false positives
- $\alpha = 0.10$ (10%): More liberal, increases power but more false positives

Decision rule: Reject H_0 if $p\text{-value} < \alpha$

Reality	Fail to Reject H_0	Reject H_0
H_0 true	Correct decision ($1-\alpha$)	Type I Error (α)
H_1 true	Type II Error (β)	Correct decision ($1-\beta$)

Table 2.1: Type I and Type II errors in hypothesis testing

2.6.4 Type I and Type II Errors

Hypothesis testing involves two types of errors:

Definition 2.17 (Type I Error (False Positive)). Rejecting H_0 when it's actually true—declaring a difference when none exists.

$$P(\text{Type I Error}) = \alpha \quad (2.47)$$

In A/B testing: Launching a "winning" variant that's actually no better (or worse) than control.

Definition 2.18 (Type II Error (False Negative)). Failing to reject H_0 when H_1 is true—missing a real effect.

$$P(\text{Type II Error}) = \beta \quad (2.48)$$

In A/B testing: Abandoning a variant that would actually improve metrics.

Definition 2.19 (Statistical Power). Power is the probability of correctly rejecting H_0 when H_1 is true:

$$\text{Power} = 1 - \beta = P(\text{reject } H_0 \mid H_1 \text{ true}) \quad (2.49)$$

Desired power is typically 0.80 (80%) or 0.90 (90%). Chapter 4 covers power analysis in depth.

Tradeoff: Decreasing α (being more conservative about declaring significance) increases β (reduces power) for fixed sample size. The only way to decrease both is to increase sample size.

2.6.5 Two-Sample Z-Test for Proportions

The workhorse test for A/B testing with binary outcomes.

Setup:

- Control: n_1 trials, x_1 successes, $\hat{p}_1 = x_1/n_1$
- Treatment: n_2 trials, x_2 successes, $\hat{p}_2 = x_2/n_2$

Hypotheses:

$$H_0 : p_1 = p_2 \quad (2.50)$$

$$H_1 : p_1 \neq p_2 \quad (\text{two-sided}) \quad (2.51)$$

Under H_0 : Pool data to estimate common proportion:

$$\hat{p} = \frac{x_1 + x_2}{n_1 + n_2} \quad (2.52)$$

Test statistic:

$$z = \frac{\hat{p}_2 - \hat{p}_1}{\sqrt{\hat{p}(1 - \hat{p}) \left(\frac{1}{n_1} + \frac{1}{n_2} \right)}} \quad (2.53)$$

Under H_0 , $z \sim \mathcal{N}(0, 1)$ approximately (by CLT).

p-value (two-sided):

$$p = 2 \cdot P(Z \geq |z|) = 2 \cdot (1 - \Phi(|z|)) \quad (2.54)$$

where Φ is the standard normal CDF.

Example 2.7 (Z-Test for A/B Experiment). • Control: 1,200 users, 96 conversions, $\hat{p}_1 = 0.08$

• Treatment: 1,200 users, 120 conversions, $\hat{p}_2 = 0.10$

Pooled proportion:

$$\hat{p} = \frac{96 + 120}{1200 + 1200} = \frac{216}{2400} = 0.09 \quad (2.55)$$

Test statistic:

$$z = \frac{0.10 - 0.08}{\sqrt{0.09 \times 0.91 \times (1/1200 + 1/1200)}} = \frac{0.02}{0.01168} = 1.712 \quad (2.56)$$

p-value:

$$p = 2 \cdot P(Z \geq 1.712) = 2 \cdot (1 - 0.9566) = 0.0868 \quad (2.57)$$

Conclusion: With $p = 0.0868 > 0.05$, we fail to reject H_0 at the 5% level. Insufficient evidence to conclude treatment differs from control, despite the 2 percentage point observed difference.

Listing 2.5: Implementing Hypothesis Tests

```

1 import numpy as np
2 from scipy import stats
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 sns.set_style("whitegrid")
7
8 def z_test_proportions(conv1, n1, conv2, n2, alpha=0.05, alternative='
9     two-sided'):
10     """
11     Two-sample z-test for proportions.
12
13     Parameters:
14     -----
15     conv1, conv2: int
16         Number of conversions in each group
17     n1, n2: int
18         Sample sizes
19     alpha: float
20         Significance level
21     alternative: str
22         'two-sided', 'greater', or 'less'

```

```

23     """Returns:
24     """
25     dict with test results
26     """
27     p1_hat = conv1 / n1
28     p2_hat = conv2 / n2
29
30     # Pooled proportion under H0
31     p_pool = (conv1 + conv2) / (n1 + n2)
32
33     # Standard error under H0
34     se_null = np.sqrt(p_pool * (1 - p_pool) * (1/n1 + 1/n2))
35
36     # Test statistic
37     z_stat = (p2_hat - p1_hat) / se_null
38
39     # p-value
40     if alternative == 'two-sided':
41         p_value = 2 * (1 - stats.norm.cdf(abs(z_stat)))
42     elif alternative == 'greater':
43         p_value = 1 - stats.norm.cdf(z_stat)
44     elif alternative == 'less':
45         p_value = stats.norm.cdf(z_stat)
46
47     # Critical value
48     if alternative == 'two-sided':
49         z_critical = stats.norm.ppf(1 - alpha/2)
50     else:
51         z_critical = stats.norm.ppf(1 - alpha)
52
53     # Decision
54     reject_null = p_value < alpha
55
56     # Effect size (relative lift)
57     relative_lift = (p2_hat - p1_hat) / p1_hat if p1_hat > 0 else np.
        nan
58
59     return {
60         'p1': p1_hat,
61         'p2': p2_hat,
62         'p_pool': p_pool,
63         'difference': p2_hat - p1_hat,
64         'relative_lift': relative_lift,
65         'z_statistic': z_stat,
66         'p_value': p_value,
67         'alpha': alpha,
68         'z_critical': z_critical,
69         'reject_null': reject_null,
70         'significant': reject_null
71     }
72
73
74 def visualize_hypothesis_test(result):
75     """Visualize the hypothesis test result."""
76     fig, axes = plt.subplots(1, 2, figsize=(14, 5))
77
78     # Plot 1: Conversion rates with error bars
79     ax1 = axes[0]

```

```

80 groups = ['Control', 'Treatment']
81 rates = [result['p1'], result['p2']]
82 colors = ['steelblue', 'coral']
83
84 bars = ax1.bar(groups, rates, color=colors, alpha=0.7, edgecolor='
      black', linewidth=1.5)
85 ax1.set_ylabel('Conversion_Rate', fontsize=12)
86 ax1.set_title('Conversion_Rates_Comparison', fontsize=13,
      fontweight='bold')
87 ax1.grid(axis='y', alpha=0.3)
88
89 # Add value labels
90 for i, (bar, rate) in enumerate(zip(bars, rates)):
91     height = bar.get_height()
92     ax1.text(bar.get_x() + bar.get_width()/2., height,
93             f'{rate:.2%}',
94             ha='center', va='bottom', fontsize=11, fontweight='bold'
95             ')
96
97 # Add significance indicator
98 y_max = max(rates) * 1.15
99 if result['significant']:
100     ax1.plot([0, 1], [y_max, y_max], 'k-', linewidth=2)
101     ax1.text(0.5, y_max*1.02, f'p={result["p_value"]:.4f}* ',
102             ha='center', fontsize=11, fontweight='bold',
103             bbox=dict(boxstyle='round', facecolor='lightgreen',
104                     alpha=0.7))
105 else:
106     ax1.text(0.5, y_max, f'p={result["p_value"]:.4f}\n(not
107             significant)',
108             ha='center', fontsize=11,
109             bbox=dict(boxstyle='round', facecolor='lightyellow',
110                     alpha=0.7))
111
112 # Plot 2: Standard normal distribution with test statistic
113 ax2 = axes[1]
114 x = np.linspace(-4, 4, 1000)
115 y = stats.norm.pdf(x, 0, 1)
116
117 ax2.plot(x, y, 'b-', linewidth=2, label='Null_distribution_N(0,1)')
118 ax2.fill_between(x, 0, y, where=(x <= -result['z_critical']) | (x
119     >= result['z_critical']),
120     alpha=0.3, color='red', label=f'Rejection_region(
121     = {result["alpha"]})')
122
123 # Mark observed z-statistic
124 ax2.axvline(result['z_statistic'], color='darkred', linestyle='--',
125     linewidth=2.5, label=f'Observed_z={result["z_statistic"]:.3f}')
126
127 ax2.set_xlabel('z-value', fontsize=12)
128 ax2.set_ylabel('Density', fontsize=12)
129 ax2.set_title('Hypothesis_Test_Visualization', fontsize=13,
130     fontweight='bold')
131 ax2.legend(fontsize=10)
132 ax2.grid(alpha=0.3)
133
134 plt.tight_layout()

```

```

128     return fig
129
130
131 # Example: Run multiple tests
132 print("="*80)
133 print("Hypothesis Testing Examples")
134 print("="*80)
135
136 test_scenarios = [
137     {
138         'name': 'Significant_effect_(large_sample)',
139         'conv1': 80, 'n1': 1000,
140         'conv2': 120, 'n2': 1000
141     },
142     {
143         'name': 'Non-significant_effect_(large_sample)',
144         'conv1': 80, 'n1': 1000,
145         'conv2': 95, 'n2': 1000
146     },
147     {
148         'name': 'Significant_effect_(small_sample)',
149         'conv1': 8, 'n1': 50,
150         'conv2': 20, 'n2': 50
151     },
152     {
153         'name': 'Non-significant_effect_(small_sample)',
154         'conv1': 4, 'n1': 40,
155         'conv2': 8, 'n2': 40
156     }
157 ]
158
159 for i, scenario in enumerate(test_scenarios, 1):
160     print(f"\n{'='*80}")
161     print(f"Scenario_{i}: {scenario['name']}")
162     print('='*80)
163
164     result = z_test_proportions(
165         scenario['conv1'], scenario['n1'],
166         scenario['conv2'], scenario['n2']
167     )
168
169     print(f"\nControl: {result['p1']:.2%} ({scenario['conv1']}/{scenario['n1']} conversions)")
170     print(f"Treatment: {result['p2']:.2%} ({scenario['conv2']}/{scenario['n2']} conversions)")
171     print(f"\nAbsolute_difference: {result['difference']:.4f} ({result['difference']*100:.2f}pp)")
172     print(f"Relative_lift: {result['relative_lift']:.2%}")
173     print(f"\nPooled_proportion: {result['p_pool']:.4f}")
174     print(f"z-statistic: {result['z_statistic']:.4f}")
175     print(f"p-value: {result['p_value']:.6f}")
176     print(f"Critical_value_(=0.05): {result['z_critical']:.4f}")
177     print(f"\nDecision: {'REJECT_H0 - Significant difference' if result['significant'] else 'FAIL TO REJECT_H0 - No significant difference'}")
178
179 # Visualize first scenario
180 if i == 1:

```

```

181     fig = visualize_hypothesis_test(result)
182     plt.savefig(f'hypothesis_test_scenario_{i}.pdf', dpi=300,
183               bbox_inches='tight')
184     plt.show()
185
186 # Demonstrate power of test via simulation
187 print("\n" + "="*80)
188 print("Statistical_Power_Demonstration_via_Simulation")
189 print("="*80)
190
191 def simulate_power(p1, p2, n, alpha=0.05, n_sims=10000):
192     """
193     Simulate statistical power for given parameters.
194     """
195     significant_count = 0
196
197     for _ in range(n_sims):
198         # Generate data under H1 (p2 != p1)
199         conv1 = np.random.binomial(n, p1)
200         conv2 = np.random.binomial(n, p2)
201
202         # Perform test
203         result = z_test_proportions(conv1, n, conv2, n, alpha=alpha)
204         if result['significant']:
205             significant_count += 1
206
207     return significant_count / n_sims
208
209 # Test different sample sizes
210 p_control = 0.08
211 p_treatment = 0.10 # 25% relative lift
212 sample_sizes = [100, 300, 500, 1000, 2000, 5000]
213
214 print(f"\nTrue_conversion_rates: Control={p_control:.1%}, Treatment=
215       {p_treatment:.1%}")
216 print(f"True_relative_lift: {(p_treatment-p_control)/p_control:.1%}")
217 print(f"\n{'Sample_size(per_group)':>25}{'Empirical_power':>20}")
218 print("-" * 50)
219
220 powers = []
221 for n in sample_sizes:
222     power = simulate_power(p_control, p_treatment, n)
223     powers.append(power)
224     print(f"{n:>25} {power:>20.3f}")
225
226 # Plot power curve
227 fig, ax = plt.subplots(figsize=(10, 6))
228 ax.plot(sample_sizes, powers, 'bo-', linewidth=2, markersize=8)
229 ax.axhline(0.8, color='red', linestyle='--', linewidth=2, label='Target
230       power=0.80')
231 ax.set_xlabel('Sample_Size(per_group)', fontsize=12)
232 ax.set_ylabel('Statistical_Power', fontsize=12)
233 ax.set_title(f'Power_Curve(p1={p_control:.1%}, p2={p_treatment:.1%},
234       alpha=0.05)',
235             fontsize=13, fontweight='bold')
236 ax.grid(alpha=0.3)
237 ax.legend(fontsize=11)
238 plt.tight_layout()

```

```

235 plt.savefig('power_curve_simulation.pdf', dpi=300, bbox_inches='tight')
236 plt.show()

```

2.7 Effect Size Measures

Statistical significance tells us whether an effect exists; effect size tells us whether it matters. This distinction is crucial for A/B testing—a statistically significant result may be practically irrelevant, especially with large samples.

2.7.1 Why Effect Size Matters

Consider two scenarios:

1. $n = 100,000$ per variant: 5.00% vs. 5.05% conversion, $p < 0.001$ (highly significant)
2. $n = 500$ per variant: 5.0% vs. 7.5% conversion, $p = 0.08$ (not significant)

In Scenario 1, the difference is statistically significant but tiny (0.05 percentage points). Implementation costs may exceed benefits.

In Scenario 2, the difference is large (2.5 percentage points, 50% relative lift) but not statistically significant due to small sample. More data might be worthwhile.

Key principle: Always report both statistical significance (p-value) and effect size. Make decisions considering both along with practical constraints.

2.7.2 Absolute vs. Relative Differences

Definition 2.20 (Absolute Difference).

$$\text{Absolute difference} = p_2 - p_1 \quad (2.58)$$

Measured in the same units as the metric (e.g., percentage points for conversion rates).

Definition 2.21 (Relative Difference (Relative Lift)).

$$\text{Relative lift} = \frac{p_2 - p_1}{p_1} = \frac{p_2}{p_1} - 1 \quad (2.59)$$

Expressed as a percentage or decimal. Answers: "By what fraction did the metric change?"

Example 2.8 (Absolute vs. Relative). 1. Control: 2%, Treatment: 3%

- Absolute difference: $3\% - 2\% = 1$ percentage point
- Relative lift: $(3\% - 2\%)/2\% = 0.50 = 50\%$

2. Control: 20%, Treatment: 21%

- Absolute difference: $21\% - 20\% = 1$ percentage point
- Relative lift: $(21\% - 20\%)/20\% = 0.05 = 5\%$

Same absolute difference, vastly different relative lifts. Which matters more depends on context.

2.7.3 Cohen's h for Proportions

Cohen's h standardizes the difference between proportions, enabling comparison across experiments with different baseline rates.

Definition 2.22 (Cohen's h). For proportions p_1 and p_2 :

$$h = 2(\arcsin \sqrt{p_2} - \arcsin \sqrt{p_1}) \quad (2.60)$$

where $\arcsin$ is the inverse sine function (arcsine).

Interpretation guidelines (Cohen, 1988):

- $|h| < 0.2$: Small effect
- $0.2 \leq |h| < 0.5$: Medium effect
- $|h| \geq 0.5$: Large effect

Advantage: Unlike relative lift, Cohen's h behaves well for proportions near 0 or 1.

2.7.4 Odds Ratio

The odds ratio is commonly used in medical research and logistic regression contexts.

Definition 2.23 (Odds and Odds Ratio). The **odds** of an event with probability p is:

$$\text{Odds} = \frac{p}{1-p} \quad (2.61)$$

The **odds ratio** comparing two proportions:

$$\text{OR} = \frac{p_2/(1-p_2)}{p_1/(1-p_1)} = \frac{p_2(1-p_1)}{p_1(1-p_2)} \quad (2.62)$$

Interpretation:

- $\text{OR} = 1$: Equal odds (no effect)
- $\text{OR} > 1$: Higher odds in treatment
- $\text{OR} < 1$: Lower odds in treatment

Example 2.9 (Odds Ratio Calculation). Control: 10% conversion, Treatment: 15% conversion

$$\text{Odds}_{\text{control}} = 0.10/0.90 = 0.111 \quad (2.63)$$

$$\text{Odds}_{\text{treatment}} = 0.15/0.85 = 0.176 \quad (2.64)$$

$$\text{OR} = 0.176/0.111 = 1.59 \quad (2.65)$$

Interpretation: The odds of conversion are 1.59 times higher in treatment—or 59% higher odds.

Log odds ratio: Often $\log(\text{OR})$ is reported because it's symmetric around 0 and unbounded.

2.7.5 Cohen's d for Continuous Metrics

For continuous metrics (e.g., revenue, session duration), Cohen's d measures effect size in standard deviation units.

Definition 2.24 (Cohen's d).

$$d = \frac{\bar{x}_2 - \bar{x}_1}{s_{\text{pooled}}} \quad (2.66)$$

where the pooled standard deviation is:

$$s_{\text{pooled}} = \sqrt{\frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{n_1 + n_2 - 2}} \quad (2.67)$$

Interpretation guidelines:

- $|d| < 0.2$: Small effect
- $0.2 \leq |d| < 0.5$: Small to medium effect
- $0.5 \leq |d| < 0.8$: Medium to large effect
- $|d| \geq 0.8$: Large effect

Example 2.10 (Cohen's d for Revenue). Testing pricing change:

- Control: $\bar{x}_1 = \$45$, $s_1 = \$20$, $n_1 = 500$
- Treatment: $\bar{x}_2 = \$52$, $s_2 = \$22$, $n_2 = 500$

Pooled standard deviation:

$$s_{\text{pooled}} = \sqrt{\frac{499 \times 20^2 + 499 \times 22^2}{998}} = \sqrt{440.2} = 20.98 \quad (2.68)$$

Cohen's d:

$$d = \frac{52 - 45}{20.98} = \frac{7}{20.98} = 0.334 \quad (2.69)$$

This is a small to medium effect (between 0.2 and 0.5).

Listing 2.6: Computing Effect Sizes

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 from scipy import stats
5
6 sns.set_style("whitegrid")
7
8 def compute_effect_sizes(conv1, n1, conv2, n2):
9     """
10     Compute multiple effect size measures for proportion comparison.
11
12     Returns:
13     -----
14     dict with various effect size measures
15     """

```



```

16     p1 = conv1 / n1
17     p2 = conv2 / n2
18
19     # Absolute difference
20     abs_diff = p2 - p1
21
22     # Relative lift
23     rel_lift = (p2 - p1) / p1 if p1 > 0 else np.nan
24
25     # Cohen's h
26     cohens_h = 2 * (np.arcsin(np.sqrt(p2)) - np.arcsin(np.sqrt(p1)))
27
28     # Odds ratio
29     odds1 = p1 / (1 - p1) if p1 < 1 else np.inf
30     odds2 = p2 / (1 - p2) if p2 < 1 else np.inf
31     odds_ratio = odds2 / odds1 if odds1 > 0 and np.isfinite(odds1) else
        np.nan
32     log_odds_ratio = np.log(odds_ratio) if odds_ratio > 0 and np.
        isfinite(odds_ratio) else np.nan
33
34     # Number needed to treat (NNT) - for proportions
35     nnt = 1 / abs_diff if abs_diff != 0 else np.inf
36
37     return {
38         'p1': p1,
39         'p2': p2,
40         'absolute_diff': abs_diff,
41         'relative_lift': rel_lift,
42         'cohens_h': cohens_h,
43         'odds_ratio': odds_ratio,
44         'log_odds_ratio': log_odds_ratio,
45         'nnt': nnt
46     }
47
48
49 def cohens_d_continuous(mean1, std1, n1, mean2, std2, n2):
50     """
51     Compute Cohen's d for continuous metrics.
52     """
53     # Pooled standard deviation
54     pooled_std = np.sqrt(((n1-1)*std1**2 + (n2-1)*std2**2) / (n1 + n2 -
        2))
55
56     # Cohen's d
57     d = (mean2 - mean1) / pooled_std
58
59     return {
60         'mean1': mean1,
61         'mean2': mean2,
62         'pooled_std': pooled_std,
63         'cohens_d': d,
64         'interpretation': interpret_cohens_d(abs(d))
65     }
66
67
68 def interpret_cohens_h(h):
69     """Interpret Cohen's h effect size."""
70     h_abs = abs(h)

```

```

71     if h_abs < 0.2:
72         return "Small"
73     elif h_abs < 0.5:
74         return "Medium"
75     else:
76         return "Large"
77
78
79 def interpret_cohens_d(d):
80     """Interpret Cohen's d effect size."""
81     d_abs = abs(d)
82     if d_abs < 0.2:
83         return "Negligible"
84     elif d_abs < 0.5:
85         return "Small"
86     elif d_abs < 0.8:
87         return "Medium"
88     else:
89         return "Large"
90
91
92 # Example calculations
93 print("="*80)
94 print("Effect Size Examples for Proportions")
95 print("="*80)
96
97 scenarios = [
98     {
99         'name': 'Low baseline, small absolute diff',
100         'conv1': 20, 'n1': 1000,
101         'conv2': 30, 'n2': 1000
102     },
103     {
104         'name': 'Low baseline, large absolute diff',
105         'conv1': 20, 'n1': 1000,
106         'conv2': 60, 'n2': 1000
107     },
108     {
109         'name': 'High baseline, small absolute diff',
110         'conv1': 400, 'n1': 1000,
111         'conv2': 410, 'n2': 1000
112     },
113     {
114         'name': 'High baseline, large absolute diff',
115         'conv1': 400, 'n1': 1000,
116         'conv2': 500, 'n2': 1000
117     }
118 ]
119
120 for scenario in scenarios:
121     print(f"\n{'-'*80}")
122     print(f"{scenario['name']}")
123     print('-'*80)
124
125     es = compute_effect_sizes(
126         scenario['conv1'], scenario['n1'],
127         scenario['conv2'], scenario['n2']
128     )

```

```

129
130     print(f"Control: {es['p1']:.2%} ({scenario['conv1']}/{scenario['n1']})")
131     print(f"Treatment: {es['p2']:.2%} ({scenario['conv2']}/{scenario['n2']})")
132     print(f"\nEffect sizes:")
133     print(f"Absolute difference: {es['absolute_diff']:.4f} ({es['absolute_diff']*100:.2f}pp)")
134     print(f"Relative lift: {es['relative_lift']:.2%}")
135     print(f"Cohen's h: {es['cohens_h']:.4f} ({interpret_cohens_h(es['cohens_h'])})")
136     print(f"Odds ratio: {es['odds_ratio']:.4f}")
137     print(f"Log odds ratio: {es['log_odds_ratio']:.4f}")
138     print(f"NNT: {es['nnt']:.1f}")
139
140 # Example for continuous metrics
141 print("\n" + "="*80)
142 print("Effect Size Examples for Continuous Metrics")
143 print("="*80)
144
145 continuous_scenarios = [
146     {
147         'name': 'Revenue_per_user',
148         'mean1': 45, 'std1': 20, 'n1': 500,
149         'mean2': 52, 'std2': 22, 'n2': 500,
150         'unit': '$'
151     },
152     {
153         'name': 'Session_duration',
154         'mean1': 180, 'std1': 120, 'n1': 1000,
155         'mean2': 200, 'std2': 125, 'n2': 1000,
156         'unit': 'seconds'
157     }
158 ]
159
160 for scenario in continuous_scenarios:
161     print(f"\n{'-'*80}")
162     print(f"{scenario['name']}")
163     print(f{'-'*80})
164
165     result = cohens_d_continuous(
166         scenario['mean1'], scenario['std1'], scenario['n1'],
167         scenario['mean2'], scenario['std2'], scenario['n2']
168     )
169
170     print(f"Control: {scenario['unit']}{result['mean1']:.2f} (SD={scenario['std1']:.2f}, n={scenario['n1']})")
171     print(f"Treatment: {scenario['unit']}{result['mean2']:.2f} (SD={scenario['std2']:.2f}, n={scenario['n2']})")
172     print(f"\nPooled SD: {scenario['unit']}{result['pooled_std']:.2f}")
173     print(f"Cohen's d: {result['cohens_d']:.4f} ({result['interpretation']} effect)")
174
175 # Visualize effect size landscape
176 print("\n" + "="*80)
177 print("Visualizing Effect Size Landscape")
178 print("="*80)
179

```

```

180 # Create grid of p1 and p2 values
181 p1_values = np.linspace(0.01, 0.5, 50)
182 p2_values = np.linspace(0.01, 0.5, 50)
183 P1, P2 = np.meshgrid(p1_values, p2_values)
184
185 # Compute relative lift and Cohen's h for each combination
186 relative_lift = (P2 - P1) / P1
187 cohens_h_grid = 2 * (np.arcsin(np.sqrt(P2)) - np.arcsin(np.sqrt(P1)))
188
189 fig, axes = plt.subplots(1, 2, figsize=(16, 6))
190
191 # Plot 1: Relative lift heatmap
192 im1 = axes[0].contourf(P1, P2, relative_lift, levels=20, cmap='RdYlGn')
193 axes[0].plot([0, 0.5], [0, 0.5], 'k--', linewidth=2, label='No effect line')
194 axes[0].set_xlabel('Control Conversion Rate', fontsize=12)
195 axes[0].set_ylabel('Treatment Conversion Rate', fontsize=12)
196 axes[0].set_title('Relative Lift Landscape', fontsize=13, fontweight='bold')
197 plt.colorbar(im1, ax=axes[0], label='Relative Lift')
198 axes[0].legend()
199
200 # Plot 2: Cohen's h heatmap
201 im2 = axes[1].contourf(P1, P2, cohens_h_grid, levels=20, cmap='RdYlGn')
202 axes[1].plot([0, 0.5], [0, 0.5], 'k--', linewidth=2, label='No effect line')
203
204 # Add contour lines for small/medium/large thresholds
205 axes[1].contour(P1, P2, np.abs(cohens_h_grid), levels=[0.2, 0.5],
206                 colors='black', linewidths=2, linestyles=[':', '--'])
207 axes[1].set_xlabel('Control Conversion Rate', fontsize=12)
208 axes[1].set_ylabel('Treatment Conversion Rate', fontsize=12)
209 axes[1].set_title("Cohen's sh Landscape", fontsize=13, fontweight='bold')
210
211 plt.colorbar(im2, ax=axes[1], label="Cohen's sh")
212 axes[1].legend()
213
214 plt.tight_layout()
215 plt.savefig('effect_size_landscape.pdf', dpi=300, bbox_inches='tight')
216 plt.show()

```

2.7.6 Practical Guidelines for Effect Size Reporting

When reporting A/B test results, always include:

1. **Sample sizes:** $n_1 = 1,000$, $n_2 = 1,000$
2. **Observed metrics:** Control = 8.0%, Treatment = 10.0%
3. **Absolute difference:** 2.0 percentage points [95% CI: 0.5, 3.5 pp]
4. **Relative lift:** 25.0% [95% CI: 6.3%, 43.8%]
5. **Statistical significance:** $p = 0.009$ (significant at $\alpha = 0.05$)
6. **Effect size:** Cohen's $h = 0.15$ (small effect)

7. **Business impact:** Expected +200 conversions/day, \$10,000 additional revenue/-month

This comprehensive reporting enables stakeholders to judge both statistical and practical significance.

2.8 Frequentist vs. Bayesian Inference

Two philosophical frameworks underpin statistical inference. Understanding both perspectives enriches A/B testing practice.

2.8.1 Frequentist Framework

Core philosophy:

- Parameters are fixed but unknown constants
- Probability describes long-run frequencies of events
- Inference focuses on procedures with good repeated sampling properties
- Uncertainty is quantified through confidence intervals and p-values

Key concepts:

- **Confidence intervals:** "95% of such intervals contain the true parameter"
- **p-values:** "Probability of data this extreme if null is true"
- **Hypothesis testing:** Reject or fail to reject H_0 based on α

Strengths:

- Objective—no prior beliefs required
- Well-established theory and widespread acceptance
- Clear Type I and Type II error control
- Standard in regulatory contexts (FDA, etc.)

Limitations:

- Cannot make probability statements about parameters ("probability that $p > 0.05$ ")
- P-values are frequently misinterpreted
- Fixed sample size typically required for valid inference
- Doesn't naturally incorporate prior information

2.8.2 Bayesian Framework

Core philosophy:

- Parameters are random variables with probability distributions
- Probability describes degree of belief about unknowns
- Prior beliefs are updated with data via Bayes' theorem
- Inference produces posterior distributions for parameters

Theorem 2.3 (Bayes' Theorem). For parameter θ and data D :

$$\underbrace{p(\theta|D)}_{\text{Posterior}} = \frac{\overbrace{p(D|\theta)}^{\text{Likelihood}} \times \overbrace{p(\theta)}^{\text{Prior}}}{\underbrace{p(D)}_{\text{Evidence}}} \quad (2.70)$$

where $p(D) = \int p(D|\theta)p(\theta)d\theta$ normalizes the posterior.

In words:

$$\text{Posterior} \propto \text{Likelihood} \times \text{Prior} \quad (2.71)$$

Key concepts:

- **Prior distribution:** Encodes initial beliefs before seeing data
- **Likelihood:** How probable the data is under different parameter values
- **Posterior distribution:** Updated beliefs after observing data
- **Credible intervals:** "95% probability parameter is in this interval"

Strengths:

- Direct probability statements about parameters
- Natural incorporation of prior information
- Coherent framework for sequential testing
- Intuitive interpretation for decision-making

Limitations:

- Requires specifying prior distributions (though this can be an advantage)
- Computational complexity for complex models (mitigated by modern MCMC)
- Results depend on prior choice
- Less standard in regulatory contexts

2.8.3 When to Use Each Approach

Prefer frequentist when:

- Regulatory requirements demand it
- You want objectivity without prior specification
- You have fixed sample size and single analysis
- Standard null hypothesis testing suffices

Prefer Bayesian when:

- You have meaningful prior information to incorporate
- You want direct probability statements about parameters
- Sequential testing or early stopping is needed
- Decision-making requires posterior distributions

In practice: Many modern A/B testing platforms use Bayesian methods because they handle sequential testing naturally and provide intuitive outputs for business stakeholders. Chapter 7 covers Bayesian A/B testing comprehensively.

2.9 Summary

This chapter established the probabilistic and statistical foundations essential for rigorous A/B testing:

- **Probability theory** provides the mathematical framework for quantifying uncertainty. Key results include the Central Limit Theorem, which justifies normal approximations for sample means regardless of underlying distributions.
- **Key distributions**—Bernoulli, binomial, normal, t, beta, and chi-squared—model different aspects of experiments. Understanding their properties enables correct application.
- **Statistical inference** uses sample data to draw conclusions about populations. Point estimates provide best guesses; confidence intervals quantify uncertainty.
- **Hypothesis testing** offers a formal framework for decision-making. Understanding null hypotheses, p-values, Type I and Type II errors, and statistical power is essential for correct interpretation.
- **Effect sizes** measure practical significance, complementing statistical significance. Always report both—a large p-value with large effect size may warrant more data; a small p-value with tiny effect size may be practically irrelevant.
- **Frequentist and Bayesian frameworks** offer complementary perspectives. Frequentist methods dominate traditional practice; Bayesian methods offer advantages for sequential testing and direct probability statements.

These foundations enable the hypothesis testing procedures, sample size calculations, multiple testing corrections, and Bayesian methods developed in subsequent chapters. Mastery of these concepts distinguishes practitioners who merely run tests from those who design sound experiments and interpret results correctly.

2.10 Further Reading

- ? : Comprehensive treatment of statistical inference with mathematical rigor
- ? : Modern theoretical statistics with clear exposition
- ? : Mathematical statistics emphasizing practical application
- ? : Definitive text on Bayesian data analysis
- ? : Classic reference on statistical power and effect sizes
- ? : Permutation tests and resampling methods
- ? : ASA statement on p-values and their interpretation

2.11 Exercises

2.11.1 Conceptual Questions

1. Explain the distinction between a parameter and a statistic. Why does this matter for A/B testing?
2. A colleague says: "We ran the experiment and got $p = 0.03$, so there's only a 3% chance the null hypothesis is true." What is wrong with this interpretation? Provide the correct interpretation.
3. Why does the Central Limit Theorem matter for A/B testing? What would be different if it didn't hold?
4. Explain the difference between a 95% confidence interval (frequentist) and a 95% credible interval (Bayesian). When would each be appropriate?
5. An experiment shows a statistically significant difference ($p < 0.001$) with Cohen's $h = 0.05$. What does this tell you? Should you implement the change?
6. Why might Type II error be more costly than Type I error in A/B testing? Conversely, when might Type I error be more concerning?

2.11.2 Probability and Distributions

1. Prove that for $X \sim \text{Bernoulli}(p)$: $\text{Var}(X) = p(1 - p)$.
2. Show that the binomial distribution $\text{Binomial}(n, p)$ has maximum variance at $p = 0.5$ for fixed n .
3. For $X \sim \mathcal{N}(\mu, \sigma^2)$, derive the distribution of:

(a) $Y = aX + b$

(b) $Z = (X - \mu)/\sigma$

4. Simulate the CLT for an exponential distribution with $\lambda = 0.2$. Show convergence to normality for $n = 5, 20, 50, 200$. Use K-S test to assess normality.
5. A website has average session duration of 4 minutes, exponentially distributed. With 100 users, what is the approximate distribution of average session duration? Calculate a 95% confidence interval.

2.11.3 Hypothesis Testing

1. Conduct a two-sample z-test by hand (show all steps):
 - Control: 1,500 users, 120 conversions
 - Treatment: 1,500 users, 150 conversions
 - $\alpha = 0.05$

Include: hypotheses, test statistic, p-value, decision, interpretation.

2. For the scenario above, construct a 95% confidence interval for the difference in proportions. Is it consistent with your hypothesis test conclusion?
3. Verify by simulation that Type I error rate is controlled at $\alpha = 0.05$ when testing $H_0 : p_1 = p_2 = 0.1$ with $n_1 = n_2 = 500$. Run 10,000 simulations.
4. Simulate power for detecting a difference from $p_1 = 0.08$ to $p_2 = 0.10$ with $\alpha = 0.05$. Plot power vs. sample size for $n = 100, 200, \dots, 2000$ (per group).
5. Investigate the relationship between α , β , effect size, and sample size via simulation. Fix three and vary the fourth.

2.11.4 Confidence Intervals

1. Construct 90%, 95%, and 99% confidence intervals for conversion rate given 75 conversions from 1,000 users. How does interval width change with confidence level?
2. Implement and compare three CI methods for proportions: Wald, Wilson score, and Agresti-Coull. For which scenarios (n, p combinations) do they differ substantially?
3. Verify by simulation that 95% confidence intervals contain the true parameter approximately 95% of the time. Test for $p = 0.05, 0.1, 0.3, 0.5$ and $n = 50, 200, 1000$.
4. Two variants show: A ($145/2000 = 7.25\%$), B ($170/2000 = 8.50\%$). Construct:
 - (a) 95% CI for each proportion
 - (b) 95% CI for difference in proportions
 - (c) 95% CI for relative lift

2.11.5 Effect Sizes

1. Calculate all effect size measures (absolute difference, relative lift, Cohen's h , odds ratio) for:

- (a) $p_1 = 0.05, p_2 = 0.07$
- (b) $p_1 = 0.20, p_2 = 0.22$
- (c) $p_1 = 0.40, p_2 = 0.50$

Compare and interpret the measures across scenarios.

2. For continuous metrics, compute Cohen's d :

- Control: $\bar{x}_1 = 100, s_1 = 30, n_1 = 250$
- Treatment: $\bar{x}_2 = 110, s_2 = 32, n_2 = 250$

Interpret the effect size. How many additional observations would be needed to detect this effect with 80% power at $\alpha = 0.05$?

3. Create a visualization showing the "landscape" of effect sizes: For $p_1 \in [0.01, 0.5]$ and $p_2 \in [0.01, 0.5]$, plot heatmaps of (a) relative lift, (b) Cohen's h , and (c) log odds ratio. Identify regions of small/medium/large effects.
4. A company runs 100 A/B tests per year. Each test has a true effect size of Cohen's $h = 0.1$ (small). With $n = 1000$ per variant and $\alpha = 0.05$, what fraction of tests will:
 - (a) Correctly detect the effect (power)?
 - (b) Fail to detect the effect (Type II error)?
 - (c) Show a "significant" effect in the wrong direction?

Verify via simulation.

2.11.6 Programming Projects

1. **A/B Test Analysis Package:** Implement a comprehensive Python class `ABTest` that:
 - Accepts control and treatment data
 - Computes all relevant statistics (proportions, SEs, CIs)
 - Performs hypothesis tests (z-test, chi-squared test, permutation test)
 - Calculates effect sizes
 - Generates publication-quality visualizations
 - Outputs formatted report with interpretation

Include docstrings, type hints, unit tests, and example usage.

2. **Interactive Dashboard:** Create a web dashboard (using Streamlit, Dash, or Shiny) that:

- Allows users to input A/B test data
- Computes and displays all statistics
- Visualizes confidence intervals, power curves
- Provides interpretation and recommendations
- Includes sample size calculator

3. **Simulation Study:** Design and execute a comprehensive simulation investigating:

- Coverage of different CI methods across parameter space
- Power as a function of effect size, sample size, baseline rate
- Impact of unequal sample sizes on test properties
- Consequences of peeking (sequential testing without correction)

Write a report summarizing findings with visualizations.

4. **Bayesian vs. Frequentist Comparison:** Implement both frameworks for A/B testing:

- Frequentist: Confidence intervals and hypothesis tests
- Bayesian: Posterior distributions with Beta-Binomial conjugacy

Compare results across scenarios varying sample size, effect size, and prior specification. When do they agree/disagree?

Chapter 3

Experimental Design Principles

3.1 Learning Objectives

After completing this chapter, readers will be able to:

- Understand and apply the fundamental principles of experimental design
- Implement various randomization techniques for different experimental contexts
- Identify and mitigate common sources of bias and confounding
- Assess experimental validity across multiple dimensions
- Make informed design decisions for online A/B tests
- Implement randomization schemes and bias detection in Python

3.2 Introduction

Experimental design is the architecture underlying A/B testing. Poor design invalidates even the most sophisticated statistical analysis—garbage in, garbage out. Conversely, thoughtful design can dramatically increase the efficiency and reliability of experiments, enabling faster learning and better decisions.

This chapter explores the principles that separate rigorous experiments from flawed ones. We'll see how randomization protects against bias, how stratification increases power, how confounding variables mislead, and how validity considerations constrain generalization. These aren't merely theoretical concerns—they directly impact whether your experiment provides trustworthy guidance or leads you astray.

The history of experimental design offers sobering lessons. Fisher's agricultural experiments in the 1920s established randomization as essential. Medical trials in the 1940s demonstrated the necessity of control groups and blinding. The digital era has introduced new challenges: network effects, novelty effects, and implementation bugs that invalidate traditional assumptions. Modern A/B testing requires both classical wisdom and contemporary adaptations.

3.3 Fundamental Principles of Experimental Design

Four core principles underpin sound experimental design. Violate any of them, and your experiment's validity is threatened.

3.3.1 Randomization

Randomization is the cornerstone of causal inference. By randomly assigning units to treatment and control, we ensure that—in expectation—groups differ only in the treatment received, not in other characteristics that might affect outcomes.

Definition 3.1 (Randomization). Random assignment is the process of allocating experimental units to treatment conditions using a random mechanism, ensuring each unit has a known, non-zero probability of assignment to each condition.

Why randomization matters:

1. **Eliminates selection bias:** Without randomization, systematic differences between groups confound treatment effects. Users who self-select into groups differ in unmeasured ways.
2. **Balances confounders:** Randomization balances both measured and *unmeasured* confounding variables across groups. This is crucial—we can't control for variables we don't know about or can't measure.
3. **Justifies statistical inference:** Probability theory underlying hypothesis tests and confidence intervals assumes random sampling or random assignment. Without randomization, these tools lack foundation.
4. **Enables causal interpretation:** Observational studies can establish association; randomized experiments establish causation. The difference is randomization.

Example 3.1 (Randomization vs. Self-Selection). Consider testing a new premium subscription tier:

Flawed approach (self-selection): Offer the new tier to all users, compare adopters vs. non-adopters. Problem: Adopters likely differ systematically—higher engagement, more resources, different needs. Observed differences confound tier effect with user characteristics.

Correct approach (randomization): Randomly assign users to see new tier vs. old tier. Now groups are balanced in expectation. Observed differences are attributable to tier design, not pre-existing user differences.

Common randomization mistakes:

- **Alternating assignment:** "Every other user gets treatment" is NOT random. Predictability enables gaming; time-based patterns confound results.
- **Assignment by characteristics:** "Users in California get treatment" introduces geographic confounding.
- **Insufficient randomness:** Using low-quality random number generators or seeding incorrectly can introduce patterns.
- **Post-hoc selection:** Assigning users randomly but then filtering or selecting subsets for analysis reintroduces bias.

3.3.2 Control Groups

A control group provides the counterfactual: what would have happened without treatment? Without controls, we can't distinguish treatment effects from natural variation, seasonal trends, or external factors.

Definition 3.2 (Control Group). The control group receives no treatment (or the current standard treatment), providing a baseline for comparison. The treatment group receives the intervention being tested.

Example 3.2 (The Necessity of Controls). A company launches a new website design and observes a 10% increase in conversions over the next week. Success? Not necessarily:

- Was there a marketing campaign that week?
- Is it a seasonal peak (e.g., holiday shopping)?
- Did competitors change their offerings?
- Was there a temporary bug in the old design that was fixed?

Without a control group maintaining the old design, these confounds are inseparable from the design effect. With proper randomization and controls, we isolate the causal impact of the design change.

Types of controls:

- **No-treatment control:** Users receive no intervention. Appropriate when testing entirely new features.
- **Placebo control:** Users receive a neutral intervention that mimics treatment exposure without active ingredients. Less common in digital contexts but relevant for psychological effects.
- **Active control:** Users receive current standard treatment. Essential when testing improvements to existing features.
- **Multiple controls:** Sometimes multiple baseline conditions are needed to isolate specific mechanisms.

3.3.3 Blocking and Stratification

While randomization ensures balance in expectation, chance can produce imbalanced groups in any single experiment. Blocking and stratification improve balance on known covariates, increasing precision.

Definition 3.3 (Blocking). Blocking divides experimental units into homogeneous groups (blocks) based on covariates, then randomizes separately within each block. This ensures balance on the blocking variables.

Definition 3.4 (Stratification). Stratification is similar to blocking but typically refers to analysis rather than design. Units are divided into strata post-randomization, and stratum-specific effects are estimated.

When to use blocking:

- **Known important covariates:** If user type, geographic region, or device strongly affects outcomes, blocking ensures balance and increases power.
- **Small sample sizes:** With few units, chance imbalance is more likely. Blocking provides insurance.
- **Heterogeneous treatment effects:** If effects vary across subgroups, blocking enables subgroup-specific analysis.

Example 3.3 (Blocking in Practice). Testing a mobile app redesign where behavior differs dramatically between iOS and Android:

Simple randomization: 50% iOS users to treatment, 50% Android users to treatment. By chance, treatment might get 52% iOS, control 48% iOS—introducing imbalance.

Blocked randomization: Within iOS users, randomize 50% to each group. Within Android users, randomize 50% to each group. Now platform is perfectly balanced.

If iOS users have 15% conversion rate and Android users 10%, blocking eliminates this source of variance, increasing power to detect the treatment effect.

Practical considerations:

- **Number of blocks:** Too many small blocks reduce efficiency. Aim for substantial units per block.
- **Blocking variables:** Choose strong predictors of outcomes. Blocking on weak predictors adds complexity without benefit.
- **Crossover:** Can block on multiple variables (e.g., platform $\times$ country), but beware exponentially growing strata.

3.3.4 Blinding

Blinding conceals treatment assignment from participants, experimenters, or both, preventing knowledge of assignment from influencing behavior or assessment.

Definition 3.5 (Blinding). • **Single-blind:** Participants don't know their treatment assignment

- **Double-blind:** Neither participants nor experimenters know assignment
- **Triple-blind:** Participants, experimenters, and analysts all remain blind

Why blinding matters:

1. **Placebo effects:** Participants' knowledge of treatment can affect outcomes independent of treatment's direct effects.
2. **Experimenter bias:** Researchers who know assignment may unconsciously influence participants or assess outcomes differently.
3. **Demand characteristics:** Participants may alter behavior to confirm experimenters' hypotheses.

4. **Differential attrition:** If users realize they're in control, they may drop out at different rates than treatment.

Example 3.4 (Blinding in Digital Experiments). In A/B testing online, blinding is often implicit:

Naturally blinded: Users typically don't know they're in an experiment or which group they're in. A new button color appears to some users; they don't realize others see something different.

Not blinded: If treatment is a new feature prominently labeled "BETA" or "NEW!", users know they're in treatment. Their behavior may reflect curiosity, novelty-seeking, or frustration with bugs rather than the feature's inherent value.

Experimenter blinding: Data analysts should ideally remain blind to which variant is A vs. B until analysis completion, preventing p-hacking or selective reporting.

Limitations in online testing:

- Complete blinding is often impossible—users see different interfaces
- Product teams must know which variant is which to iterate
- The goal: minimize systematic bias from knowledge of assignment

3.4 Randomization Techniques

Different contexts demand different randomization approaches. We examine four primary techniques with implementation details.

3.4.1 Simple Randomization

Simple randomization assigns each unit independently to treatment or control with fixed probability, typically 50%.

Advantages:

- Conceptually simple
- Easy to implement
- Provably unbiased
- No memory or coordination needed

Disadvantages:

- Can produce imbalanced group sizes
- No guarantee of balance on covariates
- Variance in group sizes reduces power

When to use: Large experiments (thousands of units) where chance imbalance is negligible.

Listing 3.1: Simple Randomization Implementation

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from typing import List, Dict, Tuple
6
7 sns.set_style("whitegrid")
8 np.random.seed(42)
9
10 def simple_randomization(n_units: int, p_treatment: float = 0.5,
11                          seed: int = None) -> np.ndarray:
12     """
13     Simple randomization: each unit independently assigned.
14
15     Parameters:
16     -----
17     n_units: int
18         Number of units to randomize
19     p_treatment: float
20         Probability of treatment assignment (default 0.5)
21     seed: int
22         Random seed for reproducibility
23
24     Returns:
25     -----
26     np.ndarray: Binary array (0=control, 1=treatment)
27     """
28     if seed is not None:
29         np.random.seed(seed)
30
31     assignments = np.random.binomial(1, p_treatment, n_units)
32
33     return assignments
34
35
36 def evaluate_randomization(assignments: np.ndarray,
37                            covariates: pd.DataFrame = None,
38                            covariate_names: List[str] = None) -> Dict:
39     """
40     Evaluate quality of randomization.
41
42     Parameters:
43     -----
44     assignments: np.ndarray
45         Treatment assignments (0/1)
46     covariates: pd.DataFrame, optional
47         Covariates to check balance on
48     covariate_names: List[str], optional
49         Names of covariates to check
50
51     Returns:
52     -----
53     dict: Evaluation metrics
54     """
55     n_total = len(assignments)
56     n_treatment = assignments.sum()
57     n_control = n_total - n_treatment

```

```

58
59     results = {
60         'n_total': n_total,
61         'n_treatment': n_treatment,
62         'n_control': n_control,
63         'prop_treatment': n_treatment / n_total,
64         'balance_score': abs(n_treatment - n_control) / n_total
65     }
66
67     # Check covariate balance if provided
68     if covariates is not None:
69         from scipy import stats as sp_stats
70         balance_tests = {}
71
72         for col in covariate_names or covariates.columns:
73             treatment_vals = covariates.loc[assignments == 1, col]
74             control_vals = covariates.loc[assignments == 0, col]
75
76             # t-test for continuous, chi-square for categorical
77             if covariates[col].dtype in [np.float64, np.int64]:
78                 t_stat, p_val = sp_stats.ttest_ind(treatment_vals,
79                                                     control_vals)
80                 balance_tests[col] = {
81                     'mean_treatment': treatment_vals.mean(),
82                     'mean_control': control_vals.mean(),
83                     'diff': treatment_vals.mean() - control_vals.mean(),
84                     'p_value': p_val
85                 }
86
87         results['covariate_balance'] = balance_tests
88
89     return results
90
91 # Example: Simple randomization with evaluation
92 print("="*80)
93 print("Simple Randomization Example")
94 print("="*80)
95
96 n_users = 1000
97 assignments = simple_randomization(n_users, p_treatment=0.5, seed=42)
98
99 # Create synthetic covariates
100 np.random.seed(42)
101 covariates = pd.DataFrame({
102     'age': np.random.normal(35, 10, n_users),
103     'prior_purchases': np.random.poisson(3, n_users),
104     'engagement_score': np.random.normal(50, 15, n_users)
105 })
106
107 # Evaluate
108 results = evaluate_randomization(assignments, covariates)
109
110 print(f"\nTotal units: {results['n_total']}")
111 print(f"Treatment: {results['n_treatment']} ({results['prop_treatment']:.2%})")
112 print(f"Control: {results['n_control']} ({1-results['prop_treatment']:.2%})")

```

```

    ']:.2%})")
113 print(f"Balance_score:_{results['balance_score']:.4f}")
114
115 if 'covariate_balance' in results:
116     print("\nCovariate_Balance:")
117     print("-" * 80)
118     print(f"{'Covariate':<20}_{'Treatment_Mean':<15}_{'Control_Mean':<15}_{'Difference':<12}_{'p-value':<10}")
119     print("-" * 80)
120     for cov, stats in results['covariate_balance'].items():
121         print(f"{cov:<20}_{stats['mean_treatment']:<15.2f}_{stats['mean_control']:<15.2f}_{stats['diff']:<12.2f}_{stats['p_value']:<10.4f}")
122
123
124 # Simulate multiple randomizations to show variability
125 n_sims = 1000
126 treatment_props = []
127 for i in range(n_sims):
128     sim_assign = simple_randomization(1000, p_treatment=0.5, seed=i)
129     treatment_props.append(sim_assign.sum() / 1000)
130
131 # Visualize
132 fig, axes = plt.subplots(1, 2, figsize=(14, 5))
133
134 # Distribution of treatment proportions
135 axes[0].hist(treatment_props, bins=30, edgecolor='black', alpha=0.7, color='steelblue')
136 axes[0].axvline(0.5, color='red', linestyle='--', linewidth=2, label='Target_(50%)')
137 axes[0].axvline(np.mean(treatment_props), color='green', linestyle='-', linewidth=2, label=f'Mean_{(np.mean(treatment_props):.3f)}')
138
139 axes[0].set_xlabel('Treatment_Proportion', fontsize=12)
140 axes[0].set_ylabel('Frequency', fontsize=12)
141 axes[0].set_title(f'Distribution_of_Treatment_Proportions\n({n_sims} simulations, n=1000)',
142                  fontsize=12, fontweight='bold')
143 axes[0].legend()
144 axes[0].grid(alpha=0.3)
145
146 # Covariate balance check for one randomization
147 treatment_mask = assignments == 1
148 axes[1].scatter(covariates.loc[treatment_mask, 'age'],
149                covariates.loc[treatment_mask, 'engagement_score'],
150                alpha=0.5, label='Treatment', color='coral', s=20)
151 axes[1].scatter(covariates.loc[~treatment_mask, 'age'],
152                covariates.loc[~treatment_mask, 'engagement_score'],
153                alpha=0.5, label='Control', color='steelblue', s=20)
154 axes[1].set_xlabel('Age', fontsize=12)
155 axes[1].set_ylabel('Engagement_Score', fontsize=12)
156 axes[1].set_title('Covariate_Balance: Age vs Engagement', fontsize=12, fontweight='bold')
157 axes[1].legend()
158 axes[1].grid(alpha=0.3)
159
160 plt.tight_layout()
161 plt.savefig('simple_randomization.pdf', dpi=300, bbox_inches='tight')
162 plt.show()

```

3.4.2 Block Randomization

Block randomization ensures exactly balanced group sizes by randomizing in small blocks rather than one unit at a time.

Mechanism:

1. Define block size (e.g., 4 or 6)
2. Within each block, assign equal numbers to treatment and control
3. Randomize the order within block

Advantages:

- Guarantees balanced group sizes
- Protects against temporal trends
- Reduces variance in treatment effect estimates

Disadvantages:

- If block size is known and assignments observable, late assignments in a block are predictable
- Requires coordination/memory across units
- Slightly more complex implementation

Listing 3.2: Block Randomization Implementation

```

1 def block_randomization(n_units: int, block_size: int = 4,
2                           p_treatment: float = 0.5, seed: int = None) ->
3                               np.ndarray:
4     """
5     Block randomization: guarantees balance within blocks.
6
7     Parameters:
8     -----
9     n_units: int
10        Number of units to randomize
11     block_size: int
12        Size of each block (should be even for 50-50 split)
13     p_treatment: float
14        Proportion assigned to treatment within each block
15     seed: int
16        Random seed for reproducibility
17
18     Returns:
19     -----
20     np.ndarray: Binary array (0=control, 1=treatment)
21     """
22     if seed is not None:
23         np.random.seed(seed)
24
25     # Calculate number of treatment assignments per block
26     n_treatment_per_block = int(block_size * p_treatment)

```

```

27 # Create full blocks
28 n_full_blocks = n_units // block_size
29 assignments = []
30
31 for _ in range(n_full_blocks):
32     # Create block with correct proportions
33     block = np.array([1] * n_treatment_per_block +
34                     [0] * (block_size - n_treatment_per_block))
35     # Shuffle within block
36     np.random.shuffle(block)
37     assignments.extend(block)
38
39 # Handle remaining units
40 remainder = n_units % block_size
41 if remainder > 0:
42     n_treatment_remainder = int(remainder * p_treatment)
43     remainder_block = np.array([1] * n_treatment_remainder +
44                               [0] * (remainder -
45                                       n_treatment_remainder))
46     np.random.shuffle(remainder_block)
47     assignments.extend(remainder_block)
48
49 return np.array(assignments)
50
51 # Example: Compare simple vs block randomization
52 print("\n" + "="*80)
53 print("Block Randomization vs Simple Randomization")
54 print("="*80)
55
56 n_users = 1000
57 n_sims = 1000
58
59 simple_balances = []
60 block_balances = []
61
62 for i in range(n_sims):
63     # Simple randomization
64     simple_assign = simple_randomization(n_users, seed=i)
65     simple_balances.append(abs(simple_assign.sum() - n_users/2))
66
67     # Block randomization
68     block_assign = block_randomization(n_users, block_size=10, seed=i)
69     block_balances.append(abs(block_assign.sum() - n_users/2))
70
71 print(f"\nBalance (deviation from 50-50 split):")
72 print(f"{'Method':<20}{'Mean':<12}{'Std Dev':<12}{'Max':<12}")
73 print("-" * 60)
74 print(f"{'Simple':<20}{np.mean(simple_balances):<12.2f}"
75       f"{np.std(simple_balances):<12.2f}{np.max(simple_balances):<12.0f}")
76 print(f"{'Block (size=10)':<20}{np.mean(block_balances):<12.2f}"
77       f"{np.std(block_balances):<12.2f}{np.max(block_balances):<12.0f}")
78
79 # Visualize comparison
80 fig, axes = plt.subplots(1, 2, figsize=(14, 5))
81

```

```

82 axes[0].hist(simple_balances, bins=30, alpha=0.6, label='Simple',
83             color='steelblue', edgecolor='black')
84 axes[0].hist(block_balances, bins=30, alpha=0.6, label='Block_(size=10)
85             ',
86             color='coral', edgecolor='black')
87 axes[0].set_xlabel('Absolute_Deviation_from_Perfect_Balance', fontsize
88                   =12)
89 axes[0].set_ylabel('Frequency', fontsize=12)
90 axes[0].set_title('Randomization_Balance_Comparison', fontsize=12,
91                   fontweight='bold')
92 axes[0].legend()
93 axes[0].grid(alpha=0.3)
94
95 # Cumulative assignments over time
96 np.random.seed(42)
97 simple_cumsum = np.cumsum(simple_randomization(1000, seed=42))
98 block_cumsum = np.cumsum(block_randomization(1000, block_size=10, seed
99                                     =42))
100 perfect = np.arange(1, 1001) * 0.5
101
102 axes[1].plot(simple_cumsum, label='Simple_randomization', linewidth=2,
103             alpha=0.7)
104 axes[1].plot(block_cumsum, label='Block_randomization', linewidth=2,
105             alpha=0.7)
106 axes[1].plot(perfect, 'k--', label='Perfect_balance', linewidth=2)
107 axes[1].set_xlabel('User_Number', fontsize=12)
108 axes[1].set_ylabel('Cumulative_Treatment_Assignments', fontsize=12)
109 axes[1].set_title('Balance_Over_Time', fontsize=12, fontweight='bold')
110 axes[1].legend()
111 axes[1].grid(alpha=0.3)
112
113 plt.tight_layout()
114 plt.savefig('block_randomization_comparison.pdf', dpi=300, bbox_inches=
115             'tight')
116 plt.show()

```

3.4.3 Stratified Randomization

Stratified randomization divides units into strata based on important covariates, then randomizes separately within each stratum.

Difference from blocking: Blocking typically refers to small, homogeneous groups. Stratification involves larger groups defined by key characteristics (e.g., country, user segment).

Advantages:

- Guarantees balance on stratifying variables
- Increases power by reducing between-stratum variance
- Enables stratified analysis
- Protects against chance imbalances on important covariates

Disadvantages:

- Requires knowing stratifying variables before randomization


```

51         strata: np.ndarray) -> Dict:
52     """
53     Analyze stratified experiment with proper variance estimation.
54
55     Parameters:
56     -----
57     outcomes: np.ndarray
58         Outcome values
59     assignments: np.ndarray
60         Treatment assignments (0/1)
61     strata: np.ndarray
62         Stratum indicators
63
64     Returns:
65     -----
66     dict: Analysis results including stratum-specific effects
67     """
68     from scipy import stats as sp_stats
69
70     results = {
71         'strata_effects': {},
72         'overall_effect': None
73     }
74
75     unique_strata = np.unique(strata)
76     n_strata = len(unique_strata)
77
78     stratum_effects = []
79     stratum_vars = []
80     stratum_weights = []
81
82     for stratum in unique_strata:
83         stratum_mask = strata == stratum
84         stratum_outcomes = outcomes[stratum_mask]
85         stratum_assignments = assignments[stratum_mask]
86
87         # Calculate stratum-specific effect
88         treatment_outcomes = stratum_outcomes[stratum_assignments == 1]
89         control_outcomes = stratum_outcomes[stratum_assignments == 0]
90
91         if len(treatment_outcomes) > 0 and len(control_outcomes) > 0:
92             effect = treatment_outcomes.mean() - control_outcomes.mean()
93
94             # Variance of difference
95             var_treatment = treatment_outcomes.var() / len(
96                 treatment_outcomes)
97             var_control = control_outcomes.var() / len(control_outcomes)
98             var_effect = var_treatment + var_control
99
100             stratum_effects.append(effect)
101             stratum_vars.append(var_effect)
102             stratum_weights.append(stratum_mask.sum())
103
104             results['strata_effects'][stratum] = {
105                 'effect': effect,
106                 'se': np.sqrt(var_effect),

```

```

106         'n': stratum_mask.sum(),
107         'n_treatment': (stratum_assignments == 1).sum(),
108         'n_control': (stratum_assignments == 0).sum(),
109         'mean_treatment': treatment_outcomes.mean(),
110         'mean_control': control_outcomes.mean()
111     }
112
113     # Calculate overall weighted effect
114     stratum_effects = np.array(stratum_effects)
115     stratum_vars = np.array(stratum_vars)
116     stratum_weights = np.array(stratum_weights) / sum(stratum_weights)
117
118     overall_effect = np.sum(stratum_weights * stratum_effects)
119     overall_var = np.sum(stratum_weights**2 * stratum_vars)
120     overall_se = np.sqrt(overall_var)
121
122     # Test statistic
123     z_stat = overall_effect / overall_se
124     p_value = 2 * (1 - sp_stats.norm.cdf(abs(z_stat)))
125
126     results['overall_effect'] = {
127         'effect': overall_effect,
128         'se': overall_se,
129         'z_statistic': z_stat,
130         'p_value': p_value,
131         'ci_lower': overall_effect - 1.96 * overall_se,
132         'ci_upper': overall_effect + 1.96 * overall_se
133     }
134
135     return results
136
137
138     # Example: Stratified randomization with heterogeneous effects
139     print("\n" + "="*80)
140     print("Stratified Randomization Example")
141     print("="*80)
142
143     n_users = 2000
144     np.random.seed(42)
145
146     # Create strata (e.g., user segments)
147     segments = np.random.choice(['New', 'Casual', 'Power'], n_users,
148                                 p=[0.3, 0.5, 0.2])
149
150     # Stratified randomization
151     assignments = stratified_randomization(segments, p_treatment=0.5, seed
152                                           =42)
153
154     # Simulate outcomes with heterogeneous treatment effects
155     baseline_conversions = {
156         'New': 0.05,
157         'Casual': 0.10,
158         'Power': 0.20
159     }
160
161     treatment_effects = {
162         'New': 0.02,      # 40% relative lift for new users
163         'Casual': 0.01,   # 10% relative lift for casual
164         'Power': 0.01     # 5% relative lift for power users

```

```

163 }
164
165 outcomes = np.zeros(n_users)
166 for i in range(n_users):
167     segment = segments[i]
168     baseline = baseline_conversions[segment]
169     effect = treatment_effects[segment] if assignments[i] == 1 else 0
170     outcomes[i] = np.random.binomial(1, baseline + effect)
171
172 # Analyze
173 results = analyze_stratified_experiment(outcomes, assignments, segments
174 )
175
176 print("\nStratum-Specific Effects:")
177 print("-" * 90)
178 print(f"{'Stratum':<12}{'Treatment_Mean':<15}{'Control_Mean':<15}{'Effect':<12}{'SE':<10}{'n':<8}")
179 print("-" * 90)
180 for stratum, stats in results['strata_effects'].items():
181     print(f"{stratum:<12}{stats['mean_treatment']:<15.4f}{stats['mean_control']:<15.4f}"
182           f"{stats['effect']:<12.4f}{stats['se']:<10.4f}{stats['n']:<8}")
183
184 print("\nOverall Effect (Weighted):")
185 print("-" * 90)
186 overall = results['overall_effect']
187 print(f"Effect: {overall['effect']:.4f}")
188 print(f"SE: {overall['se']:.4f}")
189 print(f"95% CI: [{overall['ci_lower']:.4f}, {overall['ci_upper']:.4f}]")
190 print(f"z-statistic: {overall['z_statistic']:.4f}")
191 print(f"p-value: {overall['p_value']:.6f}")
192
193 # Visualize stratum-specific effects
194 fig, axes = plt.subplots(1, 2, figsize=(14, 5))
195
196 # Effect sizes by stratum
197 strata_names = list(results['strata_effects'].keys())
198 effects = [results['strata_effects'][s]['effect'] for s in strata_names]
199 ses = [results['strata_effects'][s]['se'] for s in strata_names]
200
201 axes[0].barh(strata_names, effects, xerr=[1.96*se for se in ses],
202             alpha=0.7, color='steelblue', capsize=5)
203 axes[0].axvline(0, color='red', linestyle='--', linewidth=2)
204 axes[0].axvline(overall['effect'], color='green', linestyle='-', linewidth=2,
205                label=f"Overall: {overall['effect']:.4f}")
206 axes[0].set_xlabel('Treatment Effect', fontsize=12)
207 axes[0].set_ylabel('Stratum', fontsize=12)
208 axes[0].set_title('Stratum-Specific Treatment Effects', fontsize=12, fontweight='bold')
209 axes[0].legend()
210 axes[0].grid(alpha=0.3)
211
212 # Conversion rates by group and stratum
213 x = np.arange(len(strata_names))

```

```

213 width = 0.35
214 control_means = [results['strata_effects'][s]['mean_control'] for s in
    strata_names]
215 treatment_means = [results['strata_effects'][s]['mean_treatment'] for s
    in strata_names]
216
217 axes[1].bar(x - width/2, control_means, width, label='Control',
218             alpha=0.7, color='steelblue')
219 axes[1].bar(x + width/2, treatment_means, width, label='Treatment',
220             alpha=0.7, color='coral')
221 axes[1].set_xlabel('Stratum', fontsize=12)
222 axes[1].set_ylabel('Conversion_Rate', fontsize=12)
223 axes[1].set_title('Conversion_Rates_by_Stratum_and_Group', fontsize=12,
224                   fontweight='bold')
225 axes[1].set_xticks(x)
226 axes[1].set_xticklabels(strata_names)
227 axes[1].legend()
228 axes[1].grid(alpha=0.3, axis='y')
229
230 plt.tight_layout()
231 plt.savefig('stratified_randomization.pdf', dpi=300, bbox_inches='tight')
232 plt.show()

```

3.4.4 Cluster Randomization

Cluster randomization assigns groups of units rather than individual units. Essential when treatment naturally applies to groups or when individual-level randomization is infeasible.

When to use:

- Treatment inherently group-level (e.g., teacher training affects entire classrooms)
- Individual contamination likely (users in same household interact)
- Technical constraints prevent individual randomization
- Network effects mean individual outcomes are not independent

Key challenge: Outcomes within clusters are correlated, inflating standard errors. Effective sample size is reduced.

Definition 3.6 (Intraclass Correlation (ICC)). The ICC measures correlation of outcomes within clusters:

$$\rho = \frac{\sigma_{\text{between}}^2}{\sigma_{\text{between}}^2 + \sigma_{\text{within}}^2} \quad (3.1)$$

where $\sigma_{\text{between}}^2$ is variance between cluster means and σ_{within}^2 is variance within clusters.

Design effect: Cluster randomization increases required sample size by the design effect:

$$\text{DE} = 1 + (m - 1)\rho \quad (3.2)$$

where m is average cluster size and ρ is ICC. Higher ICC or larger clusters require proportionally more total units.

Example 3.5 (Cluster Randomization Penalty). Testing a feature where users in same household interact:

- ICC = 0.2 (moderate correlation within households)
- Average household size = 3
- Design effect: $1 + (3 - 1) \times 0.2 = 1.4$

Must collect 40% more households (or equivalently, more total users) compared to individual randomization to achieve same power.

Listing 3.4: Cluster Randomization Implementation

```

1 def cluster_randomization(cluster_ids: np.ndarray, p_treatment: float =
2     0.5,
3     seed: int = None) -> np.ndarray:
4     """
5     Cluster randomization: randomize at cluster level.
6
7     Parameters:
8     -----
9     cluster_ids: np.ndarray
10        Array indicating cluster membership for each unit
11     p_treatment: float
12        Proportion of clusters assigned to treatment
13     seed: int
14        Random seed for reproducibility
15
16     Returns:
17     -----
18     np.ndarray: Binary array (0=control, 1=treatment)
19     """
20     if seed is not None:
21         np.random.seed(seed)
22
23     # Get unique clusters
24     unique_clusters = np.unique(cluster_ids)
25     n_clusters = len(unique_clusters)
26
27     # Randomize clusters
28     n_treatment_clusters = int(n_clusters * p_treatment)
29     treatment_clusters = np.random.choice(unique_clusters,
30        n_treatment_clusters,
31        replace=False)
32
33     # Assign all units in treatment clusters
34     assignments = np.isin(cluster_ids, treatment_clusters).astype(int)
35
36     return assignments
37
38 def calculate_icc(outcomes: np.ndarray, cluster_ids: np.ndarray) ->
39     float:
40     """
41     Calculate intraclass correlation coefficient.
42
43     Parameters:
44     """

```

```

42  """-----
43  """outcomes_: np.ndarray
44  """Outcome values
45  """cluster_ids_: np.ndarray
46  """Cluster identifiers
47
48  Returns:
49  """-----
50  """float_: ICC
51  """
52      unique_clusters = np.unique(cluster_ids)
53      n_clusters = len(unique_clusters)
54
55      # Grand mean
56      grand_mean = outcomes.mean()
57
58      # Calculate between and within variance
59      cluster_means = []
60      cluster_sizes = []
61      within_ss = 0
62
63      for cluster in unique_clusters:
64          cluster_mask = cluster_ids == cluster
65          cluster_outcomes = outcomes[cluster_mask]
66          cluster_mean = cluster_outcomes.mean()
67          cluster_size = len(cluster_outcomes)
68
69          cluster_means.append(cluster_mean)
70          cluster_sizes.append(cluster_size)
71
72          # Within-cluster sum of squares
73          within_ss += np.sum((cluster_outcomes - cluster_mean)**2)
74
75      # Between-cluster sum of squares
76      cluster_means = np.array(cluster_means)
77      cluster_sizes = np.array(cluster_sizes)
78      between_ss = np.sum(cluster_sizes * (cluster_means - grand_mean)
79                          **2)
80
81      # Mean squares
82      df_between = n_clusters - 1
83      df_within = len(outcomes) - n_clusters
84      ms_between = between_ss / df_between
85      ms_within = within_ss / df_within
86
87      # ICC calculation
88      mean_cluster_size = np.mean(cluster_sizes)
89      icc = (ms_between - ms_within) / (ms_between + (mean_cluster_size -
90          1) * ms_within)
91      icc = max(0, icc) # ICC should be non-negative
92
93      return icc
94
95  def analyze_cluster_experiment(outcomes: np.ndarray,
96                                assignments: np.ndarray,
97                                cluster_ids: np.ndarray) -> Dict:
98      """

```

```

98 Analyze_cluster-randomized_experiment_with_cluster-robust_SEs.
99
100 Parameters:
101 -----
102 outcomes: np.ndarray
103 Outcome values
104 assignments: np.ndarray
105 Treatment assignments (0/1)
106 cluster_ids: np.ndarray
107 Cluster identifiers
108
109 Returns:
110 -----
111 dict: Analysis results accounting for clustering
112 ""
113
114     from scipy import stats as sp_stats
115
116     # Calculate ICC
117     icc = calculate_icc(outcomes, cluster_ids)
118
119     # Get cluster-level summaries
120     unique_clusters = np.unique(cluster_ids)
121     cluster_means = []
122     cluster_assignments = []
123     cluster_sizes = []
124
125     for cluster in unique_clusters:
126         cluster_mask = cluster_ids == cluster
127         cluster_means.append(outcomes[cluster_mask].mean())
128         cluster_assignments.append(assignments[cluster_mask][0]) # All
129         # same
130         cluster_sizes.append(cluster_mask.sum())
131
132     cluster_means = np.array(cluster_means)
133     cluster_assignments = np.array(cluster_assignments)
134     cluster_sizes = np.array(cluster_sizes)
135
136     # Cluster-level analysis
137     treatment_cluster_means = cluster_means[cluster_assignments == 1]
138     control_cluster_means = cluster_means[cluster_assignments == 0]
139
140     effect = treatment_cluster_means.mean() - control_cluster_means.
141         mean()
142
143     # Cluster-robust standard error
144     var_treatment = treatment_cluster_means.var() / len(
145         treatment_cluster_means)
146     var_control = control_cluster_means.var() / len(
147         control_cluster_means)
148     se_cluster = np.sqrt(var_treatment + var_control)
149
150     # Naive individual-level SE (for comparison)
151     treatment_outcomes = outcomes[assignments == 1]
152     control_outcomes = outcomes[assignments == 0]
153     var_treatment_naive = treatment_outcomes.var() / len(
154         treatment_outcomes)
155     var_control_naive = control_outcomes.var() / len(control_outcomes)
156     se_naive = np.sqrt(var_treatment_naive + var_control_naive)

```

```

151
152     # Design effect
153     mean_cluster_size = np.mean(cluster_sizes)
154     design_effect = 1 + (mean_cluster_size - 1) * icc
155
156     # Test statistic
157     z_stat = effect / se_cluster
158     p_value = 2 * (1 - sp_stats.norm.cdf(abs(z_stat)))
159
160     return {
161         'effect': effect,
162         'se_cluster_robust': se_cluster,
163         'se_naive': se_naive,
164         'inflation_factor': se_cluster / se_naive,
165         'icc': icc,
166         'design_effect': design_effect,
167         'mean_cluster_size': mean_cluster_size,
168         'n_clusters': len(unique_clusters),
169         'n_treatment_clusters': (cluster_assignments == 1).sum(),
170         'n_control_clusters': (cluster_assignments == 0).sum(),
171         'z_statistic': z_stat,
172         'p_value': p_value,
173         'ci_lower': effect - 1.96 * se_cluster,
174         'ci_upper': effect + 1.96 * se_cluster
175     }
176
177
178     # Example: Cluster randomization
179     print("\n" + "="*80)
180     print("Cluster Randomization Example")
181     print("="*80)
182
183     n_clusters = 200
184     cluster_sizes = np.random.poisson(10, n_clusters) + 1 # Vary cluster
185     sizes
186     n_users = cluster_sizes.sum()
187
188     # Create cluster IDs
189     cluster_ids = np.repeat(np.arange(n_clusters), cluster_sizes)
190
191     # Randomize at cluster level
192     np.random.seed(42)
193     assignments = cluster_randomization(cluster_ids, p_treatment=0.5, seed
194     =42)
195
196     # Simulate outcomes with ICC
197     icc_true = 0.3
198     treatment_effect = 0.05
199
200     # Generate cluster-level random effects
201     cluster_effects = np.random.normal(0, np.sqrt(icc_true), n_clusters)
202
203     # Generate outcomes
204     baseline_prob = 0.10
205     outcomes = np.zeros(n_users)
206
207     for i in range(n_users):
208         cluster_id = cluster_ids[i]

```



```

207     cluster_effect = cluster_effects[cluster_id]
208     treatment_contrib = treatment_effect if assignments[i] == 1 else 0
209
210     # Probability for this user
211     prob = baseline_prob + cluster_effect + treatment_contrib
212     prob = np.clip(prob, 0, 1)
213     outcomes[i] = np.random.binomial(1, prob)
214
215 # Analyze
216 results = analyze_cluster_experiment(outcomes, assignments, cluster_ids
    )
217
218 print(f"\nCluster_Design:")
219 print(f"    Number_of_clusters: {results['n_clusters']}")
220 print(f"    Treatment_clusters: {results['n_treatment_clusters']}")
221 print(f"    Control_clusters: {results['n_control_clusters']}")
222 print(f"    Mean_cluster_size: {results['mean_cluster_size']:.1f}")
223 print(f"    Total_units: {n_users}")
224
225 print(f"\nIntraclass_Correlation:")
226 print(f"    ICC: {results['icc']:.4f} (true: {icc_true:.4f})")
227 print(f"    Design_effect: {results['design_effect']:.4f}")
228
229 print(f"\nTreatment_Effect:")
230 print(f"    Effect: {results['effect']:.4f} (true: {treatment_effect:.4f})")
231 print(f"    SE (cluster-robust): {results['se_cluster_robust']:.4f}")
232 print(f"    SE (naive): {results['se_naive']:.4f}")
233 print(f"    Inflation_factor: {results['inflation_factor']:.4f}")
234 print(f"    95% CI: [{results['ci_lower']:.4f}, {results['ci_upper']:.4f}]")
235 print(f"    z-statistic: {results['z_statistic']:.4f}")
236 print(f"    p-value: {results['p_value']:.6f}")
237
238 # Visualize clustering effect
239 fig, axes = plt.subplots(1, 2, figsize=(14, 5))
240
241 # Distribution of cluster means
242 cluster_means_by_group = {
243     'Control': [],
244     'Treatment': []
245 }
246 for cluster in np.unique(cluster_ids):
247     cluster_mask = cluster_ids == cluster
248     cluster_mean = outcomes[cluster_mask].mean()
249     cluster_assignment = assignments[cluster_mask][0]
250
251     if cluster_assignment == 1:
252         cluster_means_by_group['Treatment'].append(cluster_mean)
253     else:
254         cluster_means_by_group['Control'].append(cluster_mean)
255
256 axes[0].hist(cluster_means_by_group['Control'], bins=20, alpha=0.6,
257             label='Control_clusters', color='steelblue', edgecolor='black')
258 axes[0].hist(cluster_means_by_group['Treatment'], bins=20, alpha=0.6,
259             label='Treatment_clusters', color='coral', edgecolor='black')

```

```

260 axes[0].set_xlabel('Cluster_Mean_Outcome', fontsize=12)
261 axes[0].set_ylabel('Frequency', fontsize=12)
262 axes[0].set_title('Distribution_of_Cluster_Means', fontsize=12,
    fontweight='bold')
263 axes[0].legend()
264 axes[0].grid(alpha=0.3)
265
266 # SE inflation as function of ICC
267 icc_range = np.linspace(0, 0.5, 50)
268 cluster_size = results['mean_cluster_size']
269 inflation = np.sqrt(1 + (cluster_size - 1) * icc_range)
270
271 axes[1].plot(icc_range, inflation, linewidth=2.5, color='steelblue')
272 axes[1].axvline(results['icc'], color='red', linestyle='--', linewidth
    =2,
273               label=f"Observed_ICC={results['icc']:.3f}")
274 axes[1].axhline(results['inflation_factor'], color='green', linestyle='
    --', linewidth=2,
275               label=f"Inflation={results['inflation_factor']:.3f}")
276 axes[1].set_xlabel('Intraclass_Correlation_(ICC)', fontsize=12)
277 axes[1].set_ylabel('SE_Inflation_Factor', fontsize=12)
278 axes[1].set_title(f'SE_Inflation_(cluster_size={cluster_size:.0f})',
279               fontsize=12, fontweight='bold')
280 axes[1].legend()
281 axes[1].grid(alpha=0.3)
282
283 plt.tight_layout()
284 plt.savefig('cluster_randomization.pdf', dpi=300, bbox_inches='tight')
285 plt.show()

```

3.5 Confounding Variables and Bias

Even with randomization, various biases can threaten experimental validity. Recognizing and addressing them is crucial.

3.5.1 Selection Bias

Selection bias occurs when treatment and control groups differ systematically due to the selection mechanism.

Common sources:

1. **Self-selection:** Users choose their group rather than being randomly assigned
2. **Non-random attrition:** Differential dropout between groups
3. **Survivorship bias:** Analyzing only users who complete the experiment
4. **Sample restriction:** Post-randomization filtering creates imbalance

Example 3.6 (Selection Bias in Practice). Testing a premium feature:

Biased approach: "Offer to all users, compare those who adopt vs. don't adopt"
 Adopters likely have:

- Higher income / willingness to pay

- Different needs
- More engagement
- Different demographics

Observed differences confound adoption decision with feature effect.

Correct approach: Randomly offer feature to 50%, don't offer to other 50% (regardless of interest). Compare outcomes between groups.

3.5.2 Survivorship Bias

Survivorship bias arises when analysis excludes units that drop out, especially if dropout is differential across groups.

Example 3.7 (Survivorship Bias). Testing a complex new interface:

- Treatment (new interface): 30% of users abandon immediately due to confusion
- Control (old interface): 10% abandon
- Analysis of remaining users shows treatment has higher engagement

Conclusion: New interface is better?

Problem: You're comparing the 70% who persevered through treatment (highly engaged, patient users) with 90% of control users (more representative). The selection creates bias.

Solution: Analyze all randomized users using intention-to-treat principle, not just those who completed.

3.5.3 Simpson's Paradox

Simpson's paradox occurs when a trend appears in different groups but reverses when groups are combined—or vice versa.

Example 3.8 (Simpson's Paradox in A/B Testing). Testing checkout redesign across mobile and desktop:

Platform	Control Conv.	Treatment Conv.	Winner
Mobile	5% (100/2000)	6% (120/2000)	Treatment
Desktop	15% (1500/10000)	16% (1600/10000)	Treatment
Combined	13.3% (1600/12000)	14.3% (1720/12000)	Treatment

Treatment wins on both platforms and overall—straightforward.

Now consider a different allocation:

Platform	Control Conv.	Treatment Conv.	Winner
Mobile	5% (50/1000)	6% (180/3000)	Treatment
Desktop	15% (1650/11000)	14% (1260/9000)	Control
Combined	14.2% (1700/12000)	12.0% (1440/12000)	Control

Treatment wins on mobile but loses on desktop and overall! The reversal occurs because:

- Treatment got more mobile users (low conversion)
- Control got more desktop users (high conversion)
- Platform composition confounds treatment effect

Resolution: Proper randomization prevents such confounding by balancing group composition. Post-stratification or regression adjustment can account for it in analysis.

Listing 3.5: Detecting Simpson's Paradox

```

1 def detect_simpsons_paradox(outcomes: np.ndarray,
2                             assignments: np.ndarray,
3                             strata: np.ndarray) -> Dict:
4     """
5     Detect Simpson's paradox in stratified data.
6
7     Parameters:
8     -----
9     outcomes: np.ndarray
10            Binary outcomes
11     assignments: np.ndarray
12            Treatment assignments (0/1)
13     strata: np.ndarray
14            Stratum indicators
15
16     Returns:
17     -----
18     dict: Analysis showing overall and stratum-specific effects
19     """
20     results = {
21         'overall': {},
22         'strata': {},
23         'paradox_detected': False
24     }
25
26     # Overall effect
27     overall_treatment_rate = outcomes[assignments == 1].mean()
28     overall_control_rate = outcomes[assignments == 0].mean()
29     overall_effect = overall_treatment_rate - overall_control_rate
30
31     results['overall'] = {
32         'treatment_rate': overall_treatment_rate,
33         'control_rate': overall_control_rate,
34         'effect': overall_effect,
35         'treatment_wins': overall_effect > 0
36     }
37
38     # Stratum-specific effects
39     unique_strata = np.unique(strata)
40     stratum_treatment_wins = []
41
42     for stratum in unique_strata:
43         stratum_mask = strata == stratum
44         stratum_outcomes = outcomes[stratum_mask]
45         stratum_assignments = assignments[stratum_mask]
46

```

```

47     treatment_rate = stratum_outcomes[stratum_assignments == 1].
        mean()
48     control_rate = stratum_outcomes[stratum_assignments == 0].mean
        ()
49     effect = treatment_rate - control_rate
50
51     results['strata'][stratum] = {
52         'treatment_rate': treatment_rate,
53         'control_rate': control_rate,
54         'effect': effect,
55         'treatment_wins': effect > 0,
56         'n_treatment': (stratum_assignments == 1).sum(),
57         'n_control': (stratum_assignments == 0).sum()
58     }
59
60     stratum_treatment_wins.append(effect > 0)
61
62     # Check for paradox
63     all_strata_agree = len(set(stratum_treatment_wins)) == 1
64     overall_agrees = results['overall']['treatment_wins'] ==
        stratum_treatment_wins[0]
65
66     if all_strata_agree and not overall_agrees:
67         results['paradox_detected'] = True
68         results['paradox_type'] = 'reversal'
69     elif not all_strata_agree:
70         results['paradox_detected'] = True
71         results['paradox_type'] = 'inconsistent'
72
73     return results
74
75
76     # Example: Simpson's Paradox simulation
77     print("\n" + "="*80)
78     print("Simpson's Paradox Detection")
79     print("="*80)
80
81     # Scenario 1: No paradox
82     np.random.seed(42)
83     n_mobile = 4000
84     n_desktop = 20000
85
86     platform = np.array(['Mobile'] * n_mobile + ['Desktop'] * n_desktop)
87     assignments_balanced = simple_randomization(len(platform), seed=42)
88
89     # Generate outcomes (treatment helps on both platforms)
90     outcomes_no_paradox = np.zeros(len(platform))
91     for i in range(len(platform)):
92         if platform[i] == 'Mobile':
93             base_rate = 0.05
94             treatment_lift = 0.01
95         else:
96             base_rate = 0.15
97             treatment_lift = 0.01
98
99     prob = base_rate + (treatment_lift if assignments_balanced[i] == 1
        else 0)
100    outcomes_no_paradox[i] = np.random.binomial(1, prob)

```

```

101 results_no_paradox = detect_simpsons_paradox(outcomes_no_paradox,
102                                             assignments_balanced,
103                                             platform)
104
105 print("\nScenario 1: Balanced Randomization (No Paradox)")
106 print("-" * 80)
107 print(f"Overall: Treatment={results_no_paradox['overall']['treatment_rate']:.4f},
108       f"Control={results_no_paradox['overall']['control_rate']:.4f},
109       f"Winner={ 'Treatment' if results_no_paradox['overall']['treatment_wins'] else 'Control'}")
110
111 for stratum, stats in results_no_paradox['strata'].items():
112     print(f"{stratum}: Treatment={stats['treatment_rate']:.4f},
113           f"Control={stats['control_rate']:.4f},
114           f"Winner={ 'Treatment' if stats['treatment_wins'] else 'Control'}")
115
116 print(f"Simpson's Paradox: {results_no_paradox['paradox_detected']}")
117
118 # Scenario 2: Simpson's Paradox due to imbalanced allocation
119 # Artificially create imbalance
120 assignments_imbalanced = np.zeros(len(platform), dtype=int)
121 # More treatment in mobile (low conversion)
122 assignments_imbalanced[platform == 'Mobile'] = np.random.binomial(
123     1, 0.75, n_mobile
124 )
125 # Less treatment in desktop (high conversion)
126 assignments_imbalanced[platform == 'Desktop'] = np.random.binomial(
127     1, 0.25, n_desktop
128 )
129
130 outcomes_paradox = outcomes_no_paradox.copy() # Same true effects
131
132 results_paradox = detect_simpsons_paradox(outcomes_paradox,
133                                           assignments_imbalanced,
134                                           platform)
135
136 print("\nScenario 2: Imbalanced Allocation (Potential Paradox)")
137 print("-" * 80)
138 print(f"Overall: Treatment={results_paradox['overall']['treatment_rate']:.4f},
139       f"Control={results_paradox['overall']['control_rate']:.4f},
140       f"Winner={ 'Treatment' if results_paradox['overall']['treatment_wins'] else 'Control'}")
141
142 for stratum, stats in results_paradox['strata'].items():
143     print(f"{stratum}: Treatment={stats['treatment_rate']:.4f},
144           f"Control={stats['control_rate']:.4f},
145           f"Winner={ 'Treatment' if stats['treatment_wins'] else 'Control' },
146           f"n_treat={stats['n_treatment']}, n_ctrl={stats['n_control']}")
147
148 print(f"Simpson's Paradox: {results_paradox['paradox_detected']}")
149 if results_paradox['paradox_detected']:

```

```

149     print(f"Type: {results_paradox['paradox_type']}")
150
151 # Visualize
152 fig, axes = plt.subplots(1, 2, figsize=(14, 5))
153
154 # Scenario 1
155 strata_names = list(results_no_paradox['strata'].keys())
156 x = np.arange(len(strata_names) + 1)
157 treatment_rates = [results_no_paradox['overall']['treatment_rate']] + \
158     [results_no_paradox['strata'][s]['treatment_rate']
159     for s in strata_names]
160 control_rates = [results_no_paradox['overall']['control_rate']] + \
161     [results_no_paradox['strata'][s]['control_rate'] for s
162     in strata_names]
163 labels = ['Overall'] + strata_names
164
165 width = 0.35
166 axes[0].bar(x - width/2, control_rates, width, label='Control',
167             alpha=0.7, color='steelblue')
168 axes[0].bar(x + width/2, treatment_rates, width, label='Treatment',
169             alpha=0.7, color='coral')
170 axes[0].set_ylabel('Conversion_Rate', fontsize=12)
171 axes[0].set_title('Balanced_Allocation_(No_Paradox)', fontsize=12,
172                  fontweight='bold')
173 axes[0].set_xticks(x)
174 axes[0].set_xticklabels(labels, rotation=15)
175 axes[0].legend()
176 axes[0].grid(alpha=0.3, axis='y')
177
178 # Scenario 2
179 treatment_rates_p = [results_paradox['overall']['treatment_rate']] + \
180     [results_paradox['strata'][s]['treatment_rate'] for
181     s in strata_names]
182 control_rates_p = [results_paradox['overall']['control_rate']] + \
183     [results_paradox['strata'][s]['control_rate'] for s
184     in strata_names]
185
186 axes[1].bar(x - width/2, control_rates_p, width, label='Control',
187             alpha=0.7, color='steelblue')
188 axes[1].bar(x + width/2, treatment_rates_p, width, label='Treatment',
189             alpha=0.7, color='coral')
190 axes[1].set_ylabel('Conversion_Rate', fontsize=12)
191 axes[1].set_title("Imbalanced_Allocation_(Simpson's_Paradox)",
192                  fontsize=12, fontweight='bold')
193 axes[1].set_xticks(x)
194 axes[1].set_xticklabels(labels, rotation=15)
195 axes[1].legend()
196 axes[1].grid(alpha=0.3, axis='y')
197
198 plt.tight_layout()
199 plt.savefig('simpsons_paradox.pdf', dpi=300, bbox_inches='tight')
200 plt.show()

```

3.5.4 Network Effects

Network effects occur when one user's treatment affects another user's outcomes, violating the Stable Unit Treatment Value Assumption (SUTVA).

Definition 3.7 (SUTVA). The Stable Unit Treatment Value Assumption states that:

1. Each unit has only one potential outcome under each treatment (no hidden variations)
2. One unit's treatment does not affect another's outcome (no interference)

Common network effects in digital experiments:

- **Direct interaction:** Social features where treated users interact with control users (messaging, sharing, tagging)
- **Marketplace effects:** Two-sided platforms where treatment affecting buyers impacts sellers, and vice versa
- **Capacity constraints:** Limited resources (e.g., customer support) where treatment increasing demand affects all users
- **Competitive dynamics:** Treatment improving one seller's visibility reduces others' (zero-sum outcomes)

Consequences of network effects:

- Biased treatment effect estimates (usually toward zero)
- Invalid confidence intervals
- Spillover effects confound results
- May underestimate global impact

Solutions:

1. **Cluster randomization:** Randomize social networks, households, or geographic regions rather than individuals
2. **Geo-based experiments:** Randomize by city or region when network effects are geographically bounded
3. **Temporal rollouts:** Full rollout with before/after comparison (weaker causal inference but avoids contamination)
4. **Model-based adjustment:** Estimate and adjust for spillover effects (advanced, requires modeling assumptions)

3.6 Validity Considerations

Experimental validity assesses whether an experiment tests what it purports to test and whether conclusions generalize appropriately.

3.6.1 Internal Validity

Internal validity asks: Does the experiment actually measure the causal effect of treatment on outcomes, or are results confounded?

Threats to internal validity:

1. **Selection bias:** Non-random assignment or differential attrition
2. **History effects:** External events occurring during experiment affecting one group more than another
3. **Maturation effects:** Natural changes over time (learning, fatigue) affecting groups differently
4. **Instrumentation:** Measurement changes during experiment (e.g., bug fixes, metric definition changes)
5. **Testing effects:** Awareness of being tested affects behavior (Hawthorne effect)
6. **Regression to the mean:** Extreme values naturally regress toward average on retesting

Ensuring internal validity:

- Proper randomization
- Adequate sample size
- Blinding where feasible
- Consistent measurement
- Intention-to-treat analysis
- Pre-registration of hypotheses

3.6.2 External Validity

External validity asks: Do results generalize beyond the specific experimental context to other populations, settings, and times?

Threats to external validity:

1. **Sample non-representativeness:** Experimental sample differs from target population
2. **Setting artificiality:** Experimental conditions differ from real-world deployment
3. **Temporal specificity:** Results hold only at specific times (seasonality, trends)
4. **Treatment variation:** Experimental treatment differs from how it would be deployed at scale

Example 3.9 (External Validity Challenges). Testing a new mobile app design:

- **Sample:** Only iOS users, English language, US residents
- **Setting:** Short 1-week test during holiday season
- **Treatment:** Prototype with potential bugs, incomplete features

Questions:

- Will effects hold for Android users? Non-English speakers? Other countries?
- Are holiday patterns representative of typical usage?
- Will polished version perform differently than prototype?

These questions concern external validity—whether specific experimental findings generalize.

Improving external validity:

- Representative sampling
- Realistic treatment implementation
- Sufficient duration to capture normal patterns
- Multiple experiments across contexts
- Replication studies

3.6.3 Construct Validity

Construct validity asks: Do we measure the theoretical construct we intend to measure?

Common construct validity issues:

1. **Metric doesn't capture intent:** Using clicks as proxy for user satisfaction
2. **Metric captures unintended constructs:** Revenue increase might reflect price increase, not value increase
3. **Missing important dimensions:** Focusing on engagement while ignoring user well-being
4. **Confounded metrics:** Single metric conflates multiple constructs

Example 3.10 (Construct Validity). Goal: Improve user satisfaction with search results

Weak construct validity:

- Metric: Number of search result clicks
- Problem: More clicks might indicate *worse* results (users can't find answer, keep clicking)

Better construct validity:

- Metrics: Time to first click (faster is better), probability of return to search (lower is better), explicit satisfaction ratings
- These more directly measure the intended construct (satisfaction)

3.6.4 Statistical Conclusion Validity

Statistical conclusion validity asks: Are statistical inferences correct? Do we properly quantify uncertainty?

Threats to statistical conclusion validity:

1. **Low statistical power:** Insufficient sample size to detect real effects
2. **Violated assumptions:** Using tests that assume normality, independence, etc. when assumptions don't hold
3. **Fishing and p-hacking:** Multiple testing, selective reporting, post-hoc hypotheses
4. **Improper error estimation:** Ignoring clustering, using wrong standard errors
5. **Unreliable measurement:** High measurement error inflates variance, reduces power

3.7 Practical Design Decisions

Translating experimental design principles into actual A/B tests requires numerous practical decisions.

3.7.1 Unit of Randomization

The unit of randomization is the entity assigned to treatment or control. Common choices:

- **User:** Most common; tracks individual across sessions
- **Session:** Each visit randomized independently (rarely appropriate)
- **Device:** When user identity is unavailable or unreliable
- **Household/Account:** When multiple users share accounts
- **Geographic unit:** City, region, or country for market-level interventions
- **Temporal unit:** Time periods (week, day) for platform-wide changes

Considerations:

1. **Treatment application level:** Unit must be at least as coarse as treatment application
2. **Outcome measurement level:** Ideally aligned with measurement
3. **Network effects:** Cluster at level that contains interference
4. **Sample size:** More granular units provide larger samples but risk contamination

3.7.2 Temporal Considerations

Time affects experiments in multiple ways:

1. **Duration:**

- Too short: Miss delayed effects, affected by day-of-week patterns
- Too long: External events confound results, delayed decisions
- Typical: 1-4 weeks, capturing full weekly cycles

2. **Seasonality:**

- Holiday periods atypical
- Back-to-school, tax season, etc. for relevant domains
- Consider running in representative periods

3. **Novelty effects:**

- Initial curiosity inflates treatment metrics
- May fade after days/weeks
- Distinguish short-term from long-term effects

4. **Learning effects:**

- Users adapt to new interfaces
- Performance may improve over time
- Consider learning curves in duration decisions

3.7.3 Geographic Considerations

Geographic variation affects experimental design:

1. **Regulatory constraints:** Some features unavailable in certain regions
2. **Market maturity:** Effects may differ between mature and emerging markets
3. **Cultural differences:** UI conventions, color meanings, interaction patterns vary
4. **Infrastructure:** Network speed, device types affect experience
5. **Time zones:** Usage patterns differ; ensure balanced sampling across times

Approaches:

- **Stratify by region:** Ensure balance across key markets
- **Run region-specific tests:** Estimate heterogeneous effects
- **Stage rollouts geographically:** Test in one market before expanding

Platform	Issues	Considerations
Web	Session continuity, cookie deletion	Use server-side randomization
Mobile apps	App store delays, version fragmentation	Staged rollouts, feature flags
Desktop software	Infrequent updates, offline usage	Long experiment duration
APIs	Rate limiting, caching	Cache invalidation, proper logging
Email	Spam filters, client rendering	Check deliverability, test across clients

3.7.4 Platform-Specific Issues

Different platforms present unique challenges:

3.8 Summary

Experimental design principles form the foundation for trustworthy A/B testing:

- **Randomization** eliminates selection bias and enables causal inference. It's non-negotiable for valid experiments.
- **Control groups** provide the counterfactual, distinguishing treatment effects from background variation.
- **Blocking and stratification** improve precision by ensuring balance on important covariates.
- **Different randomization techniques**—simple, block, stratified, cluster—suit different contexts. Choose based on sample size, covariate importance, and interference patterns.
- **Biases**—selection, survivorship, confounding—threaten validity even with randomization. Recognize and mitigate them through design and analysis choices.
- **Validity** encompasses internal (causal correctness), external (generalizability), construct (measurement appropriateness), and statistical conclusion (inference correctness) dimensions. Strong experiments excel on all four.
- **Practical decisions**—unit of randomization, duration, geographic scope, platform specifics—require balancing statistical ideals with operational constraints.

Sound experimental design precedes and enables sound statistical analysis. Invest time in design; it pays dividends in reliable insights and confident decisions.

3.9 Further Reading

- ? : Fisher's foundational work on experimental design and randomization
- ? : Comprehensive modern treatment of experimental design
- ? : Causal inference framework for experimental and observational studies
- ? : Practical guide to online controlled experiments
- ? : Detailed treatment of cluster randomized trials
- ? : Network interference in experiments

3.10 Exercises

3.10.1 Conceptual Questions

1. Explain why randomization is essential for causal inference. What assumptions does randomization satisfy that observational studies typically violate?
2. A colleague proposes assigning even-numbered user IDs to treatment and odd-numbered to control. What's wrong with this approach?
3. When would you prefer stratified randomization over simple randomization? What are the costs and benefits?
4. Describe three scenarios where cluster randomization is necessary despite its statistical efficiency loss.
5. Explain Simpson's paradox in your own words. Provide an original example from a domain you're familiar with.
6. Distinguish between internal and external validity. Can an experiment have high internal validity but low external validity? Provide an example.

3.10.2 Design Exercises

1. Design an A/B test for a two-sided marketplace (e.g., Airbnb, Uber) where treatment affects one side but you care about both sides' outcomes. What design challenges arise? How would you address them?
2. You're testing a social feature where users can tag friends. Describe the network effects you'd expect. How would you design the randomization to account for these?
3. Design a long-duration experiment (6+ months) for a subscription service. What temporal factors would you consider? How would you handle seasonality?
4. Compare designs for testing the same feature on web, iOS app, and Android app. What platform-specific challenges arise? Would you run separate experiments or one combined experiment?
5. You want to test a pricing change but worry about fairness (charging different prices to different users). Design an experiment that balances statistical validity with ethical concerns.

3.10.3 Implementation Exercises

1. Implement a hash-based randomization function that:
 - Takes user ID and experiment name as input
 - Returns deterministic assignment (same user always gets same variant)
 - Supports arbitrary split ratios
 - Allows easy "salting" for independent experiments

2. Create a stratified randomization system that:
 - Accepts multiple stratifying variables
 - Handles continuous and categorical variables
 - Ensures balance within strata
 - Provides diagnostic plots and balance tests
3. Write a bias detection tool that:
 - Checks for selection bias via covariate balance tests
 - Detects potential Simpson's paradox across strata
 - Identifies temporal trends suggesting history effects
 - Generates report with visualizations and recommendations
4. Implement cluster randomization with ICC estimation:
 - Randomize at cluster level
 - Estimate ICC from pilot data
 - Calculate design effect and required sample size
 - Perform cluster-robust inference

3.10.4 Analysis Exercises

1. Given experimental data with known selection bias (provided dataset), use multiple methods to detect and quantify the bias. How would you correct for it in analysis?
2. Simulate a scenario with network effects. Compare results from individual randomization (ignoring interference) vs. cluster randomization. Quantify the bias introduced by ignoring network effects.
3. Create synthetic data exhibiting Simpson's paradox. Show reversal of treatment effect sign when aggregating vs. stratifying. Explain which analysis is correct and why.
4. Analyze a multi-platform experiment (web + mobile) where treatment was randomly assigned but platform composition differs between groups. Demonstrate stratified analysis accounting for this imbalance.
5. Given survival data (time-to-conversion) with differential attrition, compare intention-to-treat analysis vs. per-protocol analysis vs. survivor analysis. Discuss when each is appropriate.

3.10.5 Simulation Projects

1. **Randomization comparison:** Simulate 10,000 experiments comparing simple, block, and stratified randomization. Compare:
 - Group size balance

- Covariate balance
- Power to detect effects
- Type I error rates

Write report with visualizations and recommendations.

2. **Network effects impact:** Simulate social network with varying levels of connectivity. Compare:

- Individual randomization (ignoring network)
- Cluster randomization by network communities
- Bias as function of network density and treatment effect strength

3. **Validity threats:** Implement experiments with various validity threats:

- Selection bias (non-random attrition)
- History effects (external shock during experiment)
- Instrumentation (metric definition change mid-experiment)
- Regression to mean (selecting extreme performers)

Quantify how each threat affects conclusions. Demonstrate mitigation strategies.

4. **Optimal stratification:** Simulate experiments with multiple potential stratifying variables. Investigate:

- How many strata provide optimal power?
- Which variables are worth stratifying on?
- Trade-off between stratification benefit and complexity

Chapter 4

Sample Size Calculation and Statistical Power

4.1 Learning Objectives

After completing this chapter, readers will be able to:

- Calculate required sample sizes for different types of A/B tests
- Understand the relationship between sample size, effect size, and statistical power
- Apply power analysis to determine experiment feasibility
- Create and interpret power curves
- Account for practical constraints in sample size determination
- Implement sample size calculators in Python and R
- Make informed decisions about minimum detectable effects

4.2 Introduction

"How many users do we need?" is perhaps the most common question in A/B testing. Too few, and we waste time running inconclusive experiments. Too many, and we delay decisions unnecessarily or expose users to potentially inferior variants longer than needed.

Sample size calculation transforms this practical question into a rigorous statistical framework. By specifying what we want to detect and how confident we need to be, we determine exactly how much data is required. This chapter develops the theory and practice of sample size determination, from fundamental formulas to real-world constraints that complicate the textbook approach.

4.3 Power Analysis Fundamentals

Statistical power is the probability of detecting an effect when it truly exists. Understanding power is essential for designing effective experiments.

4.3.1 Definition of Statistical Power

Definition 4.1 (Statistical Power). Statistical power is the probability of correctly rejecting the null hypothesis when the alternative hypothesis is true:

$$\text{Power} = 1 - \beta = P(\text{Reject } H_0 \mid H_1 \text{ true}) \quad (4.1)$$

where β is the probability of a Type II error (false negative).

Power answers the question: "If there really is an effect of size δ , what's the probability our experiment will detect it?"

Example 4.1 (Interpreting Power). Power = 0.80 means:

- If the true treatment effect is δ , we have an 80% chance of obtaining $p < \alpha$
- Equivalently, 20% chance of missing the effect (Type II error)
- If we ran 100 such experiments where an effect exists, 80 would correctly declare significance

4.3.2 Factors Affecting Power

Four interrelated factors determine statistical power:

1. **Effect Size (δ)**: Larger effects are easier to detect
 - Power increases with δ
 - Detecting 10% lift requires less data than 1% lift
2. **Sample Size (n)**: More data increases power
 - Power increases with $\sqrt{n}$
 - Doubling power requires roughly 4× sample size
3. **Significance Level (α)**: Higher α increases power
 - Relaxing from $\alpha = 0.01$ to 0.05 increases power
 - Tradeoff: more false positives
4. **Variance (σ^2)**: Lower variance increases power
 - Variance reduction techniques boost power without more data
 - Covered in Chapter 8

Theorem 4.1 (Power Relationships). For comparing two proportions with equal sample sizes:

$$\text{Power} = \Phi \left(\frac{\delta \sqrt{n}}{2\sqrt{p(1-p)}} - z_{1-\alpha/2} \right) \quad (4.2)$$

where $p = (p_1 + p_2)/2$, $\delta = p_2 - p_1$, and Φ is the standard normal CDF.
Key insights:

- Power increases with effect size δ
- Power increases with $\sqrt{n}$ (not linearly!)
- Power decreases as α becomes more stringent
- Power depends on baseline rate p (highest at $p = 0.5$)

4.3.3 The Type I and Type II Error Tradeoff

Sample size determination involves balancing two types of errors:

Reality	Fail to Reject H_0	Reject H_0
H_0 true	Correct decision ($1-\alpha$)	Type I Error (α)
H_1 true	Type II Error (β)	Correct decision ($1-\beta$)

Table 4.1: Type I and Type II errors in hypothesis testing

Conventional choices:

- $\alpha = 0.05$: Willing to accept 5% false positive rate
- Power = 0.80: Willing to accept 20% false negative rate

Why power = 0.80 is standard: This represents a 4:1 tradeoff between Type II and Type I errors ($\beta/\alpha = 0.20/0.05 = 4$). In most contexts, missing a real effect is considered 4× worse than claiming a spurious one.

When to use higher power:

- High implementation costs (difficult to change later)
- Regulatory requirements
- Mission-critical features
- One-shot opportunities

4.3.4 Power Curves and Interpretation

Power curves visualize how power depends on various parameters, providing intuition about experimental design tradeoffs.

Listing 4.1: Power Curves for Proportions

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 from scipy import stats
5 from typing import Tuple, List
6
7 sns.set_style("whitegrid")
8
9 def calculate_power_proportions(p1: float, p2: float, n: int,
10                                alpha: float = 0.05,
11                                alternative: str = 'two-sided') ->
                                float:

```

```

12     """
13     Calculate statistical power for two-sample proportion test.
14
15     Parameters:
16     -----
17     p1: float
18         Control proportion
19     p2: float
20         Treatment proportion
21     n: int
22         Sample size per group
23     alpha: float
24         Significance level
25     alternative: str
26         'two-sided' or 'one-sided'
27
28     Returns:
29     -----
30     float: Statistical power
31     """
32     # Effect size
33     delta = p2 - p1
34
35     # Pooled proportion under H0
36     p_bar = (p1 + p2) / 2
37
38     # Standard error under H0
39     se_null = np.sqrt(2 * p_bar * (1 - p_bar) / n)
40
41     # Standard error under H1
42     se_alt = np.sqrt((p1 * (1 - p1) + p2 * (1 - p2)) / n)
43
44     # Critical value
45     if alternative == 'two-sided':
46         z_crit = stats.norm.ppf(1 - alpha/2)
47     else:
48         z_crit = stats.norm.ppf(1 - alpha)
49
50     # Non-centrality parameter
51     ncp = delta / se_alt
52
53     # Power calculation
54     if alternative == 'two-sided':
55         power = 1 - stats.norm.cdf(z_crit - ncp) + stats.norm.cdf(-
56             z_crit - ncp)
57     else:
58         power = 1 - stats.norm.cdf(z_crit - ncp)
59
60     return power
61
62 def plot_power_curves(p1: float = 0.10, alpha: float = 0.05,
63                       save_path: str = None):
64     """
65     Generate comprehensive power curve visualizations.
66
67     Parameters:
68     -----

```

```

69 p1: float
70 Baseline_conversion_rate
71 alpha: float
72 Significance_level
73 save_path: str
74 Optional_path_to_save_figure
75 """
76 fig, axes = plt.subplots(2, 2, figsize=(16, 12))
77
78 # 1. Power vs Sample Size for different effect sizes
79 ax1 = axes[0, 0]
80 sample_sizes = np.arange(100, 10001, 100)
81 relative_lifts = [0.05, 0.10, 0.20, 0.30]
82
83 for lift in relative_lifts:
84     p2 = p1 * (1 + lift)
85     powers = [calculate_power_proportions(p1, p2, n, alpha) for n
86               in sample_sizes]
87     ax1.plot(sample_sizes, powers, linewidth=2.5,
88             label=f'{lift:.0%}_relative_lift', alpha=0.8)
89
90 ax1.axhline(0.8, color='red', linestyle='--', linewidth=2,
91            label='Target_power=0.80')
92 ax1.set_xlabel('Sample_Size(per_group)', fontsize=12)
93 ax1.set_ylabel('Statistical_Power', fontsize=12)
94 ax1.set_title(f'Power vs Sample Size (baseline={p1:.1%}, alpha={
95               alpha})',
96               fontsize=13, fontweight='bold')
97 ax1.legend(fontsize=10)
98 ax1.grid(alpha=0.3)
99 ax1.set_ylim([0, 1])
100
101 # 2. Power vs Effect Size for different sample sizes
102 ax2 = axes[0, 1]
103 relative_lifts = np.linspace(0.01, 0.50, 100)
104 sample_sizes_to_plot = [500, 1000, 2000, 5000]
105
106 for n in sample_sizes_to_plot:
107     powers = []
108     for lift in relative_lifts:
109         p2 = p1 * (1 + lift)
110         power = calculate_power_proportions(p1, p2, n, alpha)
111         powers.append(power)
112     ax2.plot(relative_lifts * 100, powers, linewidth=2.5,
113             label=f'n={n:,}', alpha=0.8)
114
115 ax2.axhline(0.8, color='red', linestyle='--', linewidth=2,
116            label='Target_power=0.80')
117 ax2.set_xlabel('Relative_Lift(%)', fontsize=12)
118 ax2.set_ylabel('Statistical_Power', fontsize=12)
119 ax2.set_title(f'Power vs Effect Size (baseline={p1:.1%}, alpha={
120               alpha})',
121               fontsize=13, fontweight='bold')
122 ax2.legend(fontsize=10)
123 ax2.grid(alpha=0.3)
124 ax2.set_ylim([0, 1])
125
126 # 3. Power vs Baseline Rate for fixed relative lift

```



```

179 plt.tight_layout()
180
181 if save_path:
182     plt.savefig(save_path, dpi=300, bbox_inches='tight')
183 plt.show()
184
185 # Print key insights
186 print("="*80)
187 print("Power Analysis Insights")
188 print("="*80)
189
190 # Calculate some specific examples
191 print(f"\nBaseline conversion rate: {p1:.1%}")
192 print(f"Significance level: {alpha}")
193 print("\nSample size required for 80% power:\n")
194
195 for lift in [0.05, 0.10, 0.20]:
196     p2 = p1 * (1 + lift)
197     # Binary search
198     n_low, n_high = 10, 100000
199     while n_high - n_low > 1:
200         n_mid = (n_low + n_high) // 2
201         power = calculate_power_proportions(p1, p2, n_mid, alpha)
202         if power < 0.80:
203             n_low = n_mid
204         else:
205             n_high = n_mid
206
207     print(f"{lift:>5.0%} lift ({p1:.1%} to {p2:.1%}): {n_high} users/group")
208
209 # Generate power curves
210 print("Generating power curve visualizations...\n")
211 plot_power_curves(p1=0.10, alpha=0.05, save_path='power_curves.pdf')

```

4.4 Sample Size Formulas

Sample size formulas translate desired power, significance level, and effect size into required data volume. We derive formulas for the most common metrics.

4.4.1 Sample Size for Proportions (Conversion Rates)

The most common A/B test compares conversion rates between two groups.

Two-Sample Test for Proportions

Test: $H_0 : p_1 = p_2$ vs. $H_1 : p_1 \neq p_2$

Assumptions:

- Two independent samples
- Equal sample sizes $n_1 = n_2 = n$
- Two-sided test

Theorem 4.2 (Sample Size for Proportions). The required sample size per group is:

$$n = \frac{(z_{1-\alpha/2} + z_{1-\beta})^2 \cdot [p_1(1-p_1) + p_2(1-p_2)]}{(p_2 - p_1)^2} \quad (4.3)$$

where:

- $z_{1-\alpha/2}$ is the $(1 - \alpha/2)$ quantile of standard normal (e.g., 1.96 for $\alpha = 0.05$)
- $z_{1-\beta}$ is the $(1 - \beta)$ quantile of standard normal (e.g., 0.84 for power = 0.80)
- p_1, p_2 are the control and treatment proportions

Derivation (sketch). Under H_0 , the test statistic follows:

$$Z_0 = \frac{\hat{p}_2 - \hat{p}_1}{\sqrt{\bar{p}(1-\bar{p})(1/n + 1/n)}} \sim \mathcal{N}(0, 1) \quad (4.4)$$

Under H_1 with true difference $\delta = p_2 - p_1$:

$$Z_1 = \frac{\hat{p}_2 - \hat{p}_1 - \delta}{\sqrt{p_1(1-p_1)/n + p_2(1-p_2)/n}} \sim \mathcal{N}(0, 1) \quad (4.5)$$

For power $1 - \beta$, we require:

$$P(|Z_0| > z_{1-\alpha/2} \mid H_1) = 1 - \beta \quad (4.6)$$

This occurs when:

$$\frac{\delta}{\sqrt{(p_1(1-p_1) + p_2(1-p_2))/n}} = z_{1-\alpha/2} + z_{1-\beta} \quad (4.7)$$

Solving for n yields the formula above. $\square$

Example 4.2 (Sample Size Calculation). An e-commerce site wants to test a new check-out flow:

- Current conversion rate: $p_1 = 0.10$ (10%)
- Minimum detectable effect: 10% relative lift $\rightarrow p_2 = 0.11$ (11%)
- Significance level: $\alpha = 0.05$
- Desired power: $1 - \beta = 0.80$

$$n = \frac{(1.96 + 0.84)^2 \cdot [0.10(0.90) + 0.11(0.89)]}{(0.11 - 0.10)^2} \quad (4.8)$$

$$= \frac{7.84 \cdot [0.09 + 0.0979]}{0.0001} \quad (4.9)$$

$$= \frac{7.84 \cdot 0.1879}{0.0001} \quad (4.10)$$

$$= 14,732 \quad (4.11)$$

Required: 14,732 users per group (29,464 total).

At 1,000 users/day, this experiment takes 30 days.

Listing 4.2: Sample Size Calculator for Proportions

```

1 def sample_size_proportions(p1: float, p2: float,
2                             alpha: float = 0.05,
3                             power: float = 0.80,
4                             ratio: float = 1.0) -> dict:
5     """
6     Calculate required sample size for two-proportion test.
7
8     Parameters:
9     -----
10    p1: float
11        Control proportion
12    p2: float
13        Treatment proportion
14    alpha: float
15        Significance level (default 0.05)
16    power: float
17        Statistical power (default 0.80)
18    ratio: float
19        Ratio n2/n1 (default 1.0 for equal allocation)
20
21    Returns:
22    -----
23    dict: Sample sizes and related statistics
24    """
25    from scipy import stats
26
27    # Critical values
28    z_alpha = stats.norm.ppf(1 - alpha/2) # Two-sided
29    z_beta = stats.norm.ppf(power)
30
31    # Effect size
32    delta = p2 - p1
33    if delta == 0:
34        raise ValueError("p1 and p2 must be different")
35
36    # Sample size calculation
37    if ratio == 1.0:
38        # Equal allocation (more efficient)
39        numerator = (z_alpha + z_beta)**2 * (p1*(1-p1) + p2*(1-p2))
40        denominator = delta**2
41        n1 = numerator / denominator
42        n2 = n1
43    else:
44        # Unequal allocation
45        p_avg = (p1 + ratio * p2) / (1 + ratio)
46        numerator = (z_alpha * np.sqrt((1+ratio)*p_avg*(1-p_avg)) +
47                    z_beta * np.sqrt(p1*(1-p1) + ratio*p2*(1-p2)))**2
48        denominator = ratio * delta**2
49        n1 = numerator / denominator
50        n2 = ratio * n1
51
52    # Round up
53    n1 = int(np.ceil(n1))
54    n2 = int(np.ceil(n2))
55
56    # Calculate actual power with these sample sizes
57    se = np.sqrt(p1*(1-p1)/n1 + p2*(1-p2)/n2)

```

```

58     ncp = delta / se
59     actual_power = 1 - stats.norm.cdf(z_alpha - ncp) + stats.norm.cdf(-
        z_alpha - ncp)
60
61     return {
62         'n1': n1,
63         'n2': n2,
64         'n_total': n1 + n2,
65         'p1': p1,
66         'p2': p2,
67         'absolute_difference': delta,
68         'relative_lift': delta / p1,
69         'alpha': alpha,
70         'requested_power': power,
71         'actual_power': actual_power,
72         'z_alpha': z_alpha,
73         'z_beta': z_beta
74     }
75
76
77 def sample_size_table(p1: float, relative_lifts: List[float],
78                       alpha: float = 0.05, power: float = 0.80):
79     """
80     Generate table of sample sizes for different effect sizes.
81
82     Parameters:
83     -----
84     p1: float
85         Baseline conversion rate
86     relative_lifts: List[float]
87         List of relative lifts to evaluate
88     alpha: float
89         Significance level
90     power: float
91         Desired power
92     """
93     print("="*90)
94     print(f"Sample Size Requirements (baseline={p1:.1%},    = {alpha
95         }, power={power:.0%})")
96     print("="*90)
97     print(f"{'Relative Lift':<15} {'Absolute':<12} {'p2':<10}"
98         f"{'n_per_group':<15} {'Total_n':<12} {'Duration (days)':>15}"
99         ")")
100     print("-"*90)
101
102     for lift in relative_lifts:
103         p2 = p1 * (1 + lift)
104         result = sample_size_proportions(p1, p2, alpha, power)
105
106         # Assume 1000 users per day per group
107         duration = np.ceil(result['n1'] / 1000)
108
109         print(f"{'lift':>14.1%} {result['absolute_difference']:>11.4f}"
110             f"{'p2':>9.3f} {result['n1']:>14,} {result['n_total']:>11,}"
111             f"{'duration':>14.0f}")

```

```

112     print("Assumes 1,000 users per day per group")
113     print()
114
115
116 # Example usage
117 print("\n" + "="*80)
118 print("Sample Size Calculations for A/B Testing")
119 print("="*80 + "\n")
120
121 # Example 1: E-commerce checkout
122 print("Example 1: E-commerce Checkout Optimization")
123 print("-"*80)
124 result = sample_size_proportions(p1=0.10, p2=0.11, alpha=0.05, power
    =0.80)
125 print(f"Baseline conversion: {result['p1']:.1%}")
126 print(f"Target conversion: {result['p2']:.1%}")
127 print(f"Relative lift: {result['relative_lift']:.1%}")
128 print(f"Required n per group: {result['n1']:,}")
129 print(f"Total required n: {result['n_total']:,}")
130 print(f"Actual power: {result['actual_power']:.3f}")
131 print()
132
133 # Example 2: Generate table
134 print("Example 2: Sample Size Table for Different Effect Sizes")
135 print("-"*80)
136 sample_size_table(p1=0.10,
137                   relative_lifts=[0.05, 0.10, 0.15, 0.20, 0.25, 0.30],
138                   alpha=0.05,
139                   power=0.80)
140
141 # Example 3: Compare different power levels
142 print("Example 3: Impact of Power Requirements")
143 print("-"*80)
144 p1, p2 = 0.10, 0.12
145 print(f"Baseline: {p1:.1%}, Target: {p2:.1%} ({(p2-p1)/p1:.0%} lift)\n"
    )
146 print(f"{'Power':<10} {'n per group':<15} {'Total n':<12}")
147 print("-"*40)
148 for power in [0.70, 0.75, 0.80, 0.85, 0.90]:
149     result = sample_size_proportions(p1, p2, power=power)
150     print(f"{'power':<10.0%} {result['n1']:<14,} {result['n_total']:<11,}"
    )

```

4.4.2 Sample Size for Means (Continuous Metrics)

When testing continuous outcomes like revenue, session duration, or engagement scores, we use formulas for comparing means.

Theorem 4.3 (Sample Size for Means). For testing $H_0 : \mu_1 = \mu_2$ vs. $H_1 : \mu_1 \neq \mu_2$ with equal sample sizes:

$$n = \frac{2(z_{1-\alpha/2} + z_{1-\beta})^2 \sigma^2}{(\mu_2 - \mu_1)^2} \quad (4.12)$$

where:

- σ^2 is the common variance (assumed equal in both groups)

- μ_1, μ_2 are the control and treatment means

For unequal variances (Welch's t-test):

$$n_1 = \frac{(z_{1-\alpha/2} + z_{1-\beta})^2(\sigma_1^2 + \sigma_2^2/r)}{(\mu_2 - \mu_1)^2} \quad (4.13)$$

where $r = n_2/n_1$ is the allocation ratio.

Cohen's d effect size:

Often we specify effect size in standardized units:

$$d = \frac{\mu_2 - \mu_1}{\sigma} \quad (4.14)$$

Then:

$$n = \frac{2(z_{1-\alpha/2} + z_{1-\beta})^2}{d^2} \quad (4.15)$$

Cohen's conventions:

- Small effect: $d = 0.2$
- Medium effect: $d = 0.5$
- Large effect: $d = 0.8$

Example 4.3 (Sample Size for Revenue). Testing pricing change:

- Current average revenue per user: \$50
- Standard deviation: \$30
- Minimum detectable increase: \$5 (10%)
- $\alpha = 0.05$, power = 0.80

$$n = \frac{2(1.96 + 0.84)^2 \cdot 30^2}{5^2} \quad (4.16)$$

$$= \frac{2 \cdot 7.84 \cdot 900}{25} \quad (4.17)$$

$$= \frac{14,112}{25} \quad (4.18)$$

$$= 565 \quad (4.19)$$

Required: 565 users per group (1,130 total).

Cohen's d: $d = 5/30 = 0.167$ (small effect).

Listing 4.3: Sample Size Calculator for Means

```

1 def sample_size_means(mu1: float, mu2: float, sigma: float,
2     alpha: float = 0.05,
3     power: float = 0.80,
4     ratio: float = 1.0) -> dict:
5     """

```

```

6  """Calculate required sample size for two-sample t-test.
7
8  """Parameters:
9  """-----
10 """mu1: float
11 """Control mean
12 """mu2: float
13 """Treatment mean
14 """sigma: float
15 """Common standard deviation
16 """alpha: float
17 """Significance level
18 """power: float
19 """Statistical power
20 """ratio: float
21 """Ratio n2/n1
22
23 """Returns:
24 """-----
25 """dict: Sample sizes and effect size measures
26 """
27 from scipy import stats
28
29 # Critical values
30 z_alpha = stats.norm.ppf(1 - alpha/2)
31 z_beta = stats.norm.ppf(power)
32
33 # Effect size
34 delta = mu2 - mu1
35 if delta == 0:
36     raise ValueError("mu1 and mu2 must be different")
37
38 cohens_d = delta / sigma
39
40 # Sample size
41 if ratio == 1.0:
42     n1 = 2 * (z_alpha + z_beta)**2 * sigma**2 / delta**2
43     n2 = n1
44 else:
45     n1 = (1 + 1/ratio) * (z_alpha + z_beta)**2 * sigma**2 / delta
46         **2
47     n2 = ratio * n1
48
49 n1 = int(np.ceil(n1))
50 n2 = int(np.ceil(n2))
51
52 # Actual power
53 se = sigma * np.sqrt(1/n1 + 1/n2)
54 ncp = delta / se
55 actual_power = 1 - stats.norm.cdf(z_alpha - ncp) + stats.norm.cdf(-
56     z_alpha - ncp)
57
58 return {
59     'n1': n1,
60     'n2': n2,
61     'n_total': n1 + n2,
62     'mu1': mu1,
63     'mu2': mu2,

```

```

62     'sigma': sigma,
63     'absolute_difference': delta,
64     'relative_difference': delta / mu1,
65     'cohens_d': cohens_d,
66     'effect_size_interpretation': interpret_cohens_d(abs(cohens_d))
67     ,
68     'alpha': alpha,
69     'requested_power': power,
70     'actual_power': actual_power
71 }
72
73 def interpret_cohens_d(d: float) -> str:
74     """Interpret Cohen's d effect size."""
75     if d < 0.2:
76         return "Negligible"
77     elif d < 0.5:
78         return "Small"
79     elif d < 0.8:
80         return "Medium"
81     else:
82         return "Large"
83
84
85 # Example usage
86 print("\n" + "="*80)
87 print("Sample Size for Continuous Metrics")
88 print("="*80 + "\n")
89
90 # Example 1: Revenue per user
91 print("Example 1: Revenue Per User")
92 print("-"*80)
93 result = sample_size_means(mu1=50, mu2=55, sigma=30, alpha=0.05, power
94                             =0.80)
95 print(f"Control mean: {result['mu1']:.2f}")
96 print(f"Treatment mean: {result['mu2']:.2f}")
97 print(f"Standard dev: {result['sigma']:.2f}")
98 print(f"Absolute diff: {result['absolute_difference']:.2f}")
99 print(f"Relative diff: {result['relative_difference']:.1%}")
100 print(f"Cohen's d: {result['cohens_d']:.3f} ({result['effect_size_interpretation']})")
101 print(f"Required n/group: {result['n1']:,}")
102 print(f"Total required n: {result['n_total']:,}")
103 print(f"Actual power: {result['actual_power']:.3f}")
104 print()
105
106 # Example 2: Session duration
107 print("Example 2: Session Duration (seconds)")
108 print("-"*80)
109 result = sample_size_means(mu1=180, mu2=200, sigma=120, alpha=0.05,
110                             power=0.80)
111 print(f"Control mean: {result['mu1']:.0f}s")
112 print(f"Treatment mean: {result['mu2']:.0f}s")
113 print(f"Increase: {result['absolute_difference']:.0f}s ({result['relative_difference']:.1%})")
114 print(f"Cohen's d: {result['cohens_d']:.3f} ({result['effect_size_interpretation']})")
115 print(f"Required n/group: {result['n1']:,}")

```

```

114 print()
115
116 # Example 3: Effect size table
117 print("Example 3: Sample Size by Effect Size (sigma=30)")
118 print("-"*80)
119 print(f"{'Cohen\'s d':<12}{'Interpretation':<15}{'Delta mu':<10}"
120       f"{'n per group':<15}{'Total n':<12}")
121 print("-"*70)
122
123 for d in [0.1, 0.2, 0.3, 0.5, 0.8]:
124     delta = d * 30 # sigma = 30
125     result = sample_size_means(mu1=100, mu2=100+delta, sigma=30)
126     print(f"{'d':<12.1f}{'result['effect_size_interpretation']':<15}"
127           f"{'delta':<10.1f}{'result['n1']':<14,}{'result['n_total']':<11,}"
128           ")

```

4.4.3 Sample Size for Survival Data

For time-to-event outcomes (e.g., time to purchase, churn time), sample size depends on the expected number of events rather than just sample size.

Theorem 4.4 (Sample Size for Survival Analysis). For comparing two survival curves using the log-rank test:

$$n = \frac{(z_{1-\alpha/2} + z_{1-\beta})^2}{p_1 p_2 (\log(HR))^2} \quad (4.20)$$

where:

- HR is the hazard ratio (treatment/control)
- p_1, p_2 are proportions in each group
- The formula gives total number of *events* needed

Total sample size depends on event rate:

$$N = \frac{n}{\text{event rate}} \quad (4.21)$$

4.4.4 Sample Size for Ratio Metrics

Ratio metrics (e.g., revenue per user, clicks per session) require special handling due to their non-normal distributions.

Delta method approximation:

For ratio $R = X/Y$ where X and Y are correlated:

$$\text{Var}(R) \approx \frac{1}{\mu_Y^2} [\text{Var}(X) + R^2 \text{Var}(Y) - 2R \text{Cov}(X, Y)] \quad (4.22)$$

Sample size calculation uses this variance in the standard means formula, but often requires pilot data to estimate covariances.

Practical approach: Bootstrap or simulate from pilot data to estimate variance, then use means formula.

4.5 Practical Considerations

Textbook formulas provide starting points, but real-world experiments require additional considerations.

4.5.1 Minimum Detectable Effect (MDE)

The minimum detectable effect is the smallest effect size your experiment can reliably detect given sample size, power, and significance level constraints.

Definition 4.2 (Minimum Detectable Effect). Given fixed n , α , and power, the MDE is the effect size δ satisfying:

$$\text{Power}(\delta, n, \alpha) = 1 - \beta \quad (4.23)$$

Why MDE matters:

- Experiments can only detect effects $\geq$ MDE reliably
- Smaller effects may exist but will appear non-significant
- MDE should align with minimum practically important difference

Example 4.4 (MDE Calculation). Given:

- Available sample: 5,000 users per group
- Baseline conversion: 10%
- $\alpha = 0.05$, power = 0.80

Solve for MDE by inverting sample size formula:

$$\delta = \sqrt{\frac{(z_{1-\alpha/2} + z_{1-\beta})^2 [p_1(1-p_1) + p_2(1-p_2)]}{n}} \quad (4.24)$$

With $p_2 = p_1 + \delta$ and $n = 5000$, we solve iteratively to find MDE ≈ 0.0165 (1.65 percentage points), corresponding to 16.5% relative lift.

Interpretation: This experiment can reliably detect effects of 16.5% or larger, but will likely miss smaller effects.

Listing 4.4: MDE Calculator

```

1 def calculate_mde(p1: float, n: int,
2                   alpha: float = 0.05,
3                   power: float = 0.80) -> dict:
4     """
5     Calculate minimum detectable effect given sample size.
6
7     Parameters:
8     -----
9     p1: float
10    Baseline conversion rate
11    n: int
12    Sample size per group
13    alpha: float

```



```

14 #####Significance_level
15 #####power_:float
16 #####Desired_power
17
18 #####Returns:
19 #####-----
20 #####dict_:MDE_statistics
21 #####"""
22     from scipy import stats
23
24     z_alpha = stats.norm.ppf(1 - alpha/2)
25     z_beta = stats.norm.ppf(power)
26
27     # Binary search for MDE
28     delta_low, delta_high = 0.0001, 1.0
29     epsilon = 0.0001
30
31     while delta_high - delta_low > epsilon:
32         delta_mid = (delta_low + delta_high) / 2
33         p2 = p1 + delta_mid
34
35         # Calculate power for this effect size
36         actual_power = calculate_power_proportions(p1, p2, n, alpha)
37
38         if actual_power < power:
39             delta_low = delta_mid
40         else:
41             delta_high = delta_mid
42
43     mde_absolute = delta_high
44     mde_relative = mde_absolute / p1
45
46     return {
47         'p1': p1,
48         'n_per_group': n,
49         'mde_absolute': mde_absolute,
50         'mde_relative': mde_relative,
51         'p2': p1 + mde_absolute,
52         'alpha': alpha,
53         'power': power
54     }
55
56
57 def mde_sensitivity_analysis(p1: float, sample_sizes: List[int],
58                             alpha: float = 0.05, power: float = 0.80):
59     """
60     #####Analyze how MDE changes with sample size.
61     #####"""
62     print("="*80)
63     print(f"MDE Sensitivity Analysis (baseline={p1:.1%}, n={alpha
64         }, power={power:.0%})")
65     print("="*80)
66     print(f"{'n_per_group':<15}{'MDE(abs)':<12}{'MDE(rel)':<12}"
67         f"{'p2':<10}{'Feasibility':<20}")
68     print("-"*80)
69
70     for n in sample_sizes:
71         result = calculate_mde(p1, n, alpha, power)

```

```

71     feasibility = assess_mde_feasibility(result['mde_relative'])
72
73     print(f"{n:<14,}\t{result['mde_absolute']:<11.4f}\t"
74           f"{result['mde_relative']:<11.1%}\t{result['p2']:<9.3f}\t{feasibility:<20}")
75
76     print("="*80 + "\n")
77
78
79 def assess_mde_feasibility(mde_relative: float) -> str:
80     """Assess whether MDE is practically feasible."""
81     if mde_relative < 0.05:
82         return "Very ambitious"
83     elif mde_relative < 0.10:
84         return "Ambitious"
85     elif mde_relative < 0.20:
86         return "Moderate"
87     else:
88         return "Conservative"
89
90
91 # Example usage
92 print("\n" + "="*80)
93 print("Minimum Detectable Effect (MDE) Analysis")
94 print("="*80 + "\n")
95
96 # Example 1: What can we detect with available traffic?
97 print("Example 1: MDE for Given Sample Size")
98 print("-"*80)
99 result = calculate_mde(p1=0.10, n=5000, alpha=0.05, power=0.80)
100 print(f"Baseline conversion: {result['p1']:.1%}")
101 print(f"Sample size per group: {result['n_per_group']:,}")
102 print(f"MDE (absolute): {result['mde_absolute']:.4f} ({result['mde_absolute']*100:.2f}pp)")
103 print(f"MDE (relative): {result['mde_relative']:.1%}")
104 print(f"Detectable p2: {result['p2']:.3f}")
105 print()
106
107 # Example 2: Sensitivity analysis
108 print("Example 2: How MDE Changes with Sample Size")
109 print("-"*80)
110 mde_sensitivity_analysis(
111     p1=0.10,
112     sample_sizes=[1000, 2000, 5000, 10000, 20000, 50000],
113     alpha=0.05,
114     power=0.80
115 )
116
117 # Visualization
118 fig, axes = plt.subplots(1, 2, figsize=(14, 5))
119
120 # Plot 1: MDE vs sample size
121 sample_sizes = np.logspace(2.5, 5, 50)
122 mdes_abs = []
123 mdes_rel = []
124
125 for n in sample_sizes:
126     n_int = int(n)

```

```

127     result = calculate_mde(p1=0.10, n=n_int, alpha=0.05, power=0.80)
128     mdes_abs.append(result['mde_absolute'])
129     mdes_rel.append(result['mde_relative'])
130
131 axes[0].plot(sample_sizes, np.array(mdes_abs)*100, linewidth=2.5, color
132             ='steelblue')
133 axes[0].set_xlabel('Sample_Size(per_group)', fontsize=12)
134 axes[0].set_ylabel('MDE_(percentage_points)', fontsize=12)
135 axes[0].set_title('Minimum_Detectable_Effect_vs_Sample_Size\n(baseline_
136                   =10%, alpha=0.05, power=0.80)',
137                   fontsize=12, fontweight='bold')
138 axes[0].set_xscale('log')
139 axes[0].grid(alpha=0.3)
140
141 axes[1].plot(sample_sizes, np.array(mdes_rel)*100, linewidth=2.5, color
142             ='coral')
143 axes[1].set_xlabel('Sample_Size(per_group)', fontsize=12)
144 axes[1].set_ylabel('MDE_(relative_lift_%)', fontsize=12)
145 axes[1].set_title('Relative_MDE_vs_Sample_Size\n(baseline_alpha=10%,
146                   =0.05, power=0.80)',
147                   fontsize=12, fontweight='bold')
148 axes[1].set_xscale('log')
149 axes[1].grid(alpha=0.3)
150
151 plt.tight_layout()
152 plt.savefig('mde_analysis.pdf', dpi=300, bbox_inches='tight')
153 plt.show()

```

4.5.2 Business Significance vs. Statistical Significance

Statistical significance ($p < 0.05$) doesn't imply business significance. Sample size planning must consider minimum *practically important* differences.

Framework for determining MDE:

1. **Business value:** What's the smallest improvement worth implementing?
 - Consider implementation costs
 - Account for risks
 - Factor in opportunity costs
2. **Statistical feasibility:** Can we detect that difference with available traffic?
 - Calculate required sample size
 - Estimate experiment duration
 - Assess if timeline is acceptable
3. **Reconciliation:** If desired MDE requires impractical sample size:
 - Increase traffic allocation
 - Run experiment longer
 - Accept lower power
 - Reconsider if test is worthwhile

Example 4.5 (Business vs. Statistical Significance). E-commerce company testing checkout redesign:

Business analysis:

- Current conversion: 10%
- Average order value: \$100
- Monthly users: 100,000
- Implementation cost: \$50,000

Break-even calculation:

- Need $\$50,000 / (\$100 \times 100,000/\text{month}) = 5\%$ conversion lift to break even in 1 month
- Absolute lift: 0.5 percentage points
- This becomes minimum practically important difference

Sample size for 0.5pp lift:

- $n = 14,732$ per group (from earlier calculation)
- At 50,000 users/month per group: 9 days
- Feasible!

If implementation cost were \$500,000, break-even requires 50% lift—likely infeasible.

4.5.3 Cost-Benefit Analysis of Sample Sizes

Larger samples increase power but have costs: longer experiments, more users exposed to potentially inferior variants, delayed decisions.

Cost factors:

- **Opportunity cost:** Delayed rollout of winning variant
- **Exposure cost:** Users experiencing losing variant
- **Resource cost:** Engineering time, infrastructure

Benefit factors:

- **Reduced Type II error:** Lower risk of missing beneficial changes
- **More precise estimates:** Better understanding of effect magnitude
- **Subgroup analysis:** Sufficient power for heterogeneous effects

Optimal sample size balances these factors. Standard power = 0.80 represents conventional wisdom, but context-specific optimization may justify different choices.

4.5.4 Duration Estimation

Sample size determines experiment duration given traffic volume.

$$\text{Duration (days)} = \frac{n_{\text{total}}}{\text{users per day} \times \text{allocation fraction}} \quad (4.25)$$

Considerations:

- **Day-of-week effects:** Run full weekly cycles
- **Seasonality:** Avoid atypical periods
- **Minimum duration:** At least 1 week to capture patterns
- **Maximum duration:** Product changes, external events limit feasibility

Example 4.6 (Duration Calculation). Required sample: 20,000 per group (40,000 total)
Scenario 1: High traffic site

- Daily users: 50,000
- 50/50 split: 25,000 per group per day
- Duration: $20,000 / 25,000 = 0.8$ days
- **Decision:** Run for 1 week to capture weekly patterns

Scenario 2: Medium traffic site

- Daily users: 2,000
- 50/50 split: 1,000 per group per day
- Duration: $20,000 / 1,000 = 20$ days
- **Decision:** Acceptable, run for 3 weeks (21 days)

Scenario 3: Low traffic site

- Daily users: 200
- 50/50 split: 100 per group per day
- Duration: $20,000 / 100 = 200$ days
- **Decision:** Infeasible! Consider:
 - Accept larger MDE (smaller sample)
 - Lower power requirement
 - Don't run experiment (too uncertain)

4.6 Advanced Topics

4.6.1 Unequal Sample Sizes

Sometimes unequal allocation is optimal or necessary.

Reasons for unequal allocation:

- Limited traffic to treatment variant
- Risk mitigation (keep most users on proven control)
- Cost asymmetry (treatment is expensive)

Theorem 4.5 (Optimal Allocation Ratio). For fixed total sample size $N = n_1 + n_2$, power is maximized when:

$$\frac{n_2}{n_1} = \frac{\sigma_2}{\sigma_1} \quad (4.26)$$

For proportions with similar rates, equal allocation (1:1) is nearly optimal.

For highly unequal allocation, efficiency loss is:

$$\text{Efficiency} = \frac{4r}{(1+r)^2} \quad (4.27)$$

where $r = n_2/n_1$.

Example 4.7 (Allocation Efficiency). • 1:1 (50/50): Efficiency = 1.00 (optimal)

- 2:1 (67/33): Efficiency = 0.89 (11% loss)
- 3:1 (75/25): Efficiency = 0.75 (25% loss)
- 9:1 (90/10): Efficiency = 0.36 (64% loss!)

Practical recommendation: Avoid allocation more extreme than 2:1 unless necessary. Efficiency loss outweighs most benefits.

4.6.2 Sequential Testing Implications

Classical sample size formulas assume:

- Fixed sample size determined in advance
- Single analysis at experiment conclusion

Sequential testing (peeking at results, stopping early) inflates Type I error rate. Naïve peeking can increase α from 0.05 to 0.30!

Solutions:

- **Alpha spending:** Adjust significance thresholds for multiple looks
- **Sequential probability ratio test (SPRT):** Optimal sequential procedure
- **Group sequential design:** Pre-specified analysis points with adjusted α

Sample sizes for sequential testing are **not** determined by classical formulas. Chapter 8 covers sequential methods.

4.6.3 Multiple Metrics Considerations

Testing multiple metrics simultaneously requires adjustment for multiple comparisons.

Approaches:

1. **Primary metric only:** Design for power on single most important metric
 - Simple, clear decision rule
 - Risk: Miss effects on other metrics
2. **Family-wise error rate (FWER):** Control probability of any false positive
 - Bonferroni: Increase sample size by factor of $\sqrt{m}$ for m metrics
 - Very conservative
3. **False discovery rate (FDR):** Control expected proportion of false discoveries
 - More powerful than FWER
 - Appropriate when some false positives acceptable
4. **Guardrail metrics:** Primary metric at $\alpha = 0.05$, secondary metrics at $\alpha = 0.10$ or 0.20
 - Practical compromise
 - Flags potential issues without requiring huge samples

Practical recommendation: Design for adequate power on primary metric. Use secondary metrics as diagnostics with less stringent thresholds. See Chapter 6 for comprehensive treatment.

4.7 Implementation Tools

4.7.1 Python Implementation

We've developed comprehensive functions. Here's a unified calculator:

Listing 4.5: Comprehensive Sample Size Calculator

```

1 class ABTestCalculator:
2     """
3     Comprehensive A/B test sample size and power calculator.
4     """
5
6     def __init__(self, alpha: float = 0.05, power: float = 0.80):
7         """
8         Initialize calculator with default parameters.
9
10        Parameters:
11        -----
12        alpha: float
13        Significance level (default 0.05)
14        power: float
15        Statistical power (default 0.80)
16        """

```

```

17         self.alpha = alpha
18         self.power = power
19
20     def sample_size_proportions(self, p1: float, p2: float,
21                                ratio: float = 1.0) -> dict:
22         """Calculate sample size for proportion test."""
23         return sample_size_proportions(p1, p2, self.alpha, self.power,
24                                         ratio)
25
26     def sample_size_means(self, mu1: float, mu2: float, sigma: float,
27                           ratio: float = 1.0) -> dict:
28         """Calculate sample size for means test."""
29         return sample_size_means(mu1, mu2, sigma, self.alpha, self.
30                                 power, ratio)
31
32     def calculate_power(self, p1: float, p2: float, n: int) -> float:
33         """Calculate power for given parameters."""
34         return calculate_power_proportions(p1, p2, n, self.alpha)
35
36     def calculate_mde(self, p1: float, n: int) -> dict:
37         """Calculate minimum detectable effect."""
38         return calculate_mde(p1, n, self.alpha, self.power)
39
40     def recommend_sample_size(self, p1: float, desired_lift: float,
41                               traffic_per_day: int,
42                               max_duration_days: int = 30) -> dict:
43         """
44         Provide sample size recommendation with practical constraints.
45
46         Parameters:
47         -----
48         p1: float
49             Baseline conversion rate
50         desired_lift: float
51             Desired relative lift to detect
52         traffic_per_day: int
53             Available users per day
54         max_duration_days: int
55             Maximum acceptable experiment duration
56
57         Returns:
58         -----
59         dict: Recommendation with feasibility assessment
60         """
61         p2 = p1 * (1 + desired_lift)
62
63         # Calculate required sample size
64         result = self.sample_size_proportions(p1, p2)
65
66         # Calculate duration
67         users_per_group_per_day = traffic_per_day / 2
68         duration_days = np.ceil(result['n1'] / users_per_group_per_day)
69
70         # Feasibility assessment
71         feasible = duration_days <= max_duration_days
72
73         # If not feasible, calculate achievable MDE
74         if not feasible:

```



```

73         max_n = int(users_per_group_per_day * max_duration_days)
74         achievable_mde = self.calculate_mde(p1, max_n)
75     else:
76         achievable_mde = None
77
78     return {
79         'required_n_per_group': result['n1'],
80         'required_total_n': result['n_total'],
81         'duration_days': int(duration_days),
82         'duration_weeks': int(np.ceil(duration_days / 7)),
83         'feasible': feasible,
84         'achievable_mde': achievable_mde,
85         'recommendation': self._generate_recommendation(
86             feasible, duration_days, max_duration_days,
87             achievable_mde
88         ),
89         'details': result
90     }
91
92     def _generate_recommendation(self, feasible: bool, duration: float,
93                                 max_duration: float, achievable_mde:
94                                     dict) -> str:
95         """Generate human-readable recommendation."""
96         if feasible:
97             return (f"Experiment is FEASIBLE. Run for {int(np.ceil(
98                 duration/7))} weeks"
99                 f"({int(duration)} days) to achieve desired power.")
100         else:
101             if achievable_mde:
102                 return (f"Experiment NOT FEASIBLE within {max_duration}
103                     days."
104                     f"Can only detect {achievable_mde['mde_relative']:.1%} lift"
105                     f"with available traffic. Consider: (1) Accepting larger MDE,"
106                     f"(2) Running longer, (3) Increasing traffic allocation,"
107                     f"(4) Not running experiment.")
108             else:
109                 return "Experiment NOT FEASIBLE. Insufficient traffic."
110
111 # Example: Comprehensive analysis
112 print("\n" + "="*80)
113 print("Comprehensive A/B Test Sample Size Analysis")
114 print("="*80 + "\n")
115
116 # Initialize calculator
117 calc = ABTestCalculator(alpha=0.05, power=0.80)
118
119 # Scenario: E-commerce checkout optimization
120 print("Scenario: E-commerce Checkout Redesign")
121 print("-"*80)
122 print("Current conversion rate: 10%")
123 print("Desired lift to detect: 10% (relative)")
124 print("Daily traffic: 10,000 users")
125 print("Maximum acceptable duration: 30 days")
126 print()

```

```

124
125 recommendation = calc.recommend_sample_size(
126     p1=0.10,
127     desired_lift=0.10,
128     traffic_per_day=10000,
129     max_duration_days=30
130 )
131
132 print("ANALYSIS RESULTS:")
133 print("-"*80)
134 print(f"Required sample size: {recommendation['required_n_per_group']:,} per group")
135 print(f"Total required: {recommendation['required_total_n']:,} users")
136 print(f"Estimated duration: {recommendation['duration_days']} days ({recommendation['duration_weeks']} weeks)")
137 print(f"Feasibility: {' FEASIBLE' if recommendation['feasible'] else ' NOT FEASIBLE'}")
138 print()
139 print("RECOMMENDATION:")
140 print("-"*80)
141 print(recommendation['recommendation'])
142 print()
143
144 # Additional scenarios
145 print("\nComparing Different Traffic Levels:")
146 print("-"*80)
147
148 for daily_traffic in [1000, 5000, 10000, 50000]:
149     rec = calc.recommend_sample_size(
150         p1=0.10,
151         desired_lift=0.10,
152         traffic_per_day=daily_traffic,
153         max_duration_days=30
154     )
155
156     print(f"\nDaily traffic: {daily_traffic:,}")
157     print(f"Duration: {rec['duration_days']} days")
158     print(f"Feasible: {'Yes' if rec['feasible'] else 'No'}")
159     if not rec['feasible'] and rec['achievable_mde']:
160         print(f"Achievable MDE: {rec['achievable_mde']['mde_relative']:.1%}")

```

4.7.2 R Implementation

R provides excellent power analysis tools through the `pwr` package.

Listing 4.6: Sample Size Calculation in R

```

1 # Install and load required packages
2 install.packages("pwr")
3 library(pwr)
4
5 # 1. Sample size for two proportions
6 # Testing 10% vs 11% conversion
7 p1 <- 0.10
8 p2 <- 0.11
9

```

```

10 # Calculate effect size (Cohen's h)
11 h <- ES.h(p1, p2)
12 cat(sprintf("Cohen's h: %.4f\n", h))
13
14 # Calculate required sample size
15 power_result <- pwr.2p.test(
16   h = h,
17   sig.level = 0.05,
18   power = 0.80,
19   alternative = "two.sided"
20 )
21
22 print(power_result)
23 cat(sprintf("\nRequired n per group: %.0f\n", ceiling(power_result$n)))
24 cat(sprintf("Total n: %.0f\n", ceiling(power_result$n * 2)))
25
26 # 2. Sample size for means
27 # Testing $50 vs $55 with SD = $30
28 mu1 <- 50
29 mu2 <- 55
30 sigma <- 30
31
32 # Cohen's d
33 d <- (mu2 - mu1) / sigma
34 cat(sprintf("\nCohen's d: %.4f\n", d))
35
36 # Calculate required sample size
37 power_result_means <- pwr.t.test(
38   d = d,
39   sig.level = 0.05,
40   power = 0.80,
41   type = "two.sample",
42   alternative = "two.sided"
43 )
44
45 print(power_result_means)
46
47 # 3. Power curve
48 p1_vals <- seq(0.05, 0.30, 0.05)
49 lift <- 0.10
50 sample_sizes <- c(500, 1000, 2000, 5000)
51
52 power_matrix <- matrix(NA, nrow = length(p1_vals), ncol = length(sample
53   _sizes))
54 colnames(power_matrix) <- paste0("n=", sample_sizes)
55 rownames(power_matrix) <- paste0("p1=", p1_vals)
56
57 for (i in seq_along(p1_vals)) {
58   p1 <- p1_vals[i]
59   p2 <- p1 * (1 + lift)
60   h <- ES.h(p1, p2)
61
62   for (j in seq_along(sample_sizes)) {
63     n <- sample_sizes[j]
64     result <- pwr.2p.test(
65       h = h,
66       n = n,
67       sig.level = 0.05,

```

```

67     alternative = "two.sided"
68   )
69   power_matrix[i, j] <- result$power
70 }
71 }
72
73 print(round(power_matrix, 3))
74
75 # 4. Plot power curve
76 plot_power_curve <- function(p1, lift_range, n, alpha = 0.05) {
77   powers <- numeric(length(lift_range))
78
79   for (i in seq_along(lift_range)) {
80     p2 <- p1 * (1 + lift_range[i])
81     h <- ES.h(p1, p2)
82     result <- pwr.2p.test(h = h, n = n, sig.level = alpha)
83     powers[i] <- result$power
84   }
85
86   plot(lift_range * 100, powers,
87        type = "l", lwd = 2,
88        xlab = "Relative Lift (%)",
89        ylab = "Statistical Power",
90        main = sprintf("Power Curve (baseline = %.1f%%, n = %d)", p1 *
91                        100, n))
92   abline(h = 0.80, col = "red", lty = 2)
93   grid()
94 }
95
96 # Generate power curve
97 lift_range <- seq(0.05, 0.50, 0.01)
98 plot_power_curve(p1 = 0.10, lift_range = lift_range, n = 5000)

```

4.8 Summary

Sample size determination transforms experimental planning from guesswork to rigorous calculation:

- **Power analysis** connects sample size, effect size, significance level, and power. Specifying any three determines the fourth.
- **Standard conventions** ($\alpha = 0.05$, power = 0.80) represent reasonable defaults but should be adjusted for context.
- **Sample size formulas** differ by metric type. Proportions and means have closed-form solutions; other metrics require simulation or approximation.
- **Minimum detectable effect** (MDE) defines the smallest effect reliably detectable. MDE must be both statistically feasible and business-relevant.
- **Practical constraints**—traffic volume, acceptable duration, implementation costs—often dominate pure statistical considerations.
- **Advanced topics**—unequal allocation, sequential testing, multiple metrics—require specialized methods beyond classical formulas.

- **Implementation tools** in Python and R enable rapid calculation and sensitivity analysis.

Proper sample size calculation ensures experiments are neither underpowered (wasting time and resources) nor overpowered (unnecessarily delaying decisions). Master these techniques to design efficient, informative experiments.

4.9 Further Reading

- [?:](#) Definitive reference on statistical power analysis
- [?:](#) Sample size methods for clinical research
- [?:](#) Practical guidance for online experiments
- [?:](#) Documentation for R's pwr package

4.10 Exercises

4.10.1 Conceptual Questions

1. Explain the relationship between Type I error, Type II error, power, and significance level. How are they connected?
2. Why is $\text{power} = 0.80$ a common standard? Under what circumstances would you use $\text{power} = 0.90$? $\text{Power} = 0.70$?
3. A colleague says: "Let's just run the test for a month and see what happens." What's wrong with this approach?
4. Explain why doubling sample size doesn't double statistical power. What's the actual relationship?
5. Describe three scenarios where unequal allocation (e.g., 75/25 split) might be justified despite efficiency loss.
6. What's the difference between minimum detectable effect and minimum practically important difference? Why do both matter?

4.10.2 Calculation Exercises

1. Calculate required sample size for testing:
 - Baseline conversion: 5%
 - Target conversion: 6% (20% relative lift)
 - $\alpha = 0.05$, $\text{power} = 0.80$

Show all work. How long would this experiment take with 10,000 users/day?

2. Repeat previous calculation for $\text{power} = 0.90$. By what factor does required sample size increase?

3. Given 3,000 users per group available, baseline conversion 15%, calculate the minimum detectable effect (MDE) for $\alpha = 0.05$, power = 0.80.
4. Calculate required sample size for continuous metric:
 - Control mean: \$75, SD: \$40
 - Target mean: \$85 (13.3% increase)
 - $\alpha = 0.05$, power = 0.80
5. Compare sample size requirements for detecting 10% relative lift at baseline rates of 2%, 10%, 30%, and 50%. What pattern emerges?

4.10.3 Programming Exercises

1. Implement a comprehensive power analysis function that:
 - Accepts multiple input formats (absolute difference, relative lift, Cohen's h)
 - Supports both one-sided and two-sided tests
 - Handles unequal allocation ratios
 - Provides clear warnings about assumptions
 - Includes input validation
2. Create an interactive sample size calculator using Streamlit or Shiny that:
 - Takes baseline metrics and desired lift as input
 - Shows power curves dynamically
 - Estimates experiment duration given traffic
 - Provides feasibility recommendations
 - Exports results as PDF report
3. Write a function to perform sensitivity analysis:
 - How does power change if true effect is 50% smaller/larger than assumed?
 - How does power change if variance is 20% higher than expected?
 - Generate visualization showing sensitivity regions
4. Implement simulation-based power calculation for:
 - Ratio metrics (revenue per user)
 - Count metrics (purchases per customer)
 - Mixture distributions (users who convert vs. don't convert)

Compare simulation results to closed-form approximations.

4.10.4 Applied Exercises

1. Design experiments for these scenarios with full justification:

Scenario A: SaaS company testing onboarding flow

- Current trial-to-paid conversion: 8%
- 500 new trials per week
- Implementation cost: \$100,000
- Average customer lifetime value: \$5,000

Scenario B: News website testing article layout

- Current average time on page: 2.5 minutes (SD: 3.0 minutes)
- 1 million pageviews per day
- Easy to implement and reverse

Scenario C: Mobile app testing premium feature

- Current premium conversion: 2%
- 10,000 daily active users
- Major engineering effort (\$500K)

For each: Specify desired lift, calculate sample size, estimate duration, assess feasibility, provide recommendation.

2. A stakeholder insists on detecting 2% relative lift with 95% power. Your analysis shows this requires 100,000 users per group but only 50,000 total are available. Prepare a memo explaining:
 - Why their requirement is infeasible
 - What can be achieved with available sample
 - Alternative approaches (sequential testing, focus on larger effects, etc.)
 - Recommended path forward
3. Review a past A/B test that failed to reach significance. Using actual baseline metrics and observed difference:
 - Calculate what power the experiment actually had
 - Determine what sample size would have been needed for 80% power
 - Assess whether observed effect was below MDE
 - Recommend whether to rerun with larger sample or abandon

4.10.5 Advanced Projects

1. **Sequential testing sample size:** Implement group sequential design with:
 - O'Brien-Fleming spending function
 - 3-5 planned interim analyses
 - Calculate required sample size accounting for multiple looks
 - Compare to fixed sample size
2. **Multiple metrics:** Design experiment testing 1 primary and 4 secondary metrics:
 - Determine sample size for each metric individually
 - Apply Bonferroni, Holm, and FDR corrections
 - Compare required sample sizes
 - Recommend practical approach
3. **Variance reduction impact:** Simulate experiments with and without CUPED variance reduction:
 - Generate correlated pre/post data
 - Compare power with/without covariate adjustment
 - Quantify sample size savings
 - Visualize power gains
4. **Bayesian sample size determination:** Implement Bayesian approach:
 - Specify prior distribution for effect size
 - Calculate probability of success for various sample sizes
 - Compare to frequentist power analysis
 - Discuss advantages and limitations

Chapter 5

Hypothesis Testing for A/B Tests

5.1 Learning Objectives

After completing this chapter, readers will be able to:

- Conduct z-tests and t-tests for A/B testing scenarios
- Interpret p-values correctly and avoid common misinterpretations
- Apply chi-squared tests for categorical outcomes
- Use permutation tests as non-parametric alternatives
- Understand assumptions underlying different test procedures
- Implement hypothesis tests in Python and R

5.2 The Framework of Hypothesis Testing

Hypothesis testing provides a formal procedure for making decisions under uncertainty. While we introduced the concept earlier, this chapter provides rigorous treatment of test procedures for A/B testing.

5.2.1 Structure of a Hypothesis Test

1. State hypotheses:

- Null hypothesis H_0 : No effect or no difference
- Alternative hypothesis H_1 : An effect exists

2. Choose significance level: α (typically 0.05 or 0.01)

3. Calculate test statistic: Function of data measuring evidence against H_0

4. Determine p-value or critical region:

- P-value: Probability of observing data this extreme under H_0
- Critical region: Values of test statistic leading to rejection

5. **Make decision:**

- Reject H_0 if p-value $< \alpha$ (or test statistic in critical region)
- Fail to reject H_0 otherwise

6. **Interpret in context:** Translate statistical conclusion to practical meaning

5.2.2 One-Sided vs. Two-Sided Tests

Two-sided test:

$$H_0 : \mu_1 = \mu_2 \quad \text{vs.} \quad H_1 : \mu_1 \neq \mu_2 \quad (5.1)$$

Tests for any difference (positive or negative). Most common in A/B testing.

One-sided test:

$$H_0 : \mu_1 \geq \mu_2 \quad \text{vs.} \quad H_1 : \mu_1 < \mu_2 \quad (5.2)$$

Tests specifically for improvement (or degradation). More powerful but only detects effects in specified direction.

When to use one-sided:

- Only care about improvements (would not implement if treatment is worse)
- Regulatory or business constraints dictate directionality
- Preregistered directional hypothesis

Caution: One-sided tests are often misused. Default to two-sided unless strong justification exists.

5.3 Z-Test for Proportions

The z-test compares proportions between two groups using normal approximation.

5.3.1 Two-Sample Z-Test

For control proportion $\hat{p}_1 = x_1/n_1$ and treatment proportion $\hat{p}_2 = x_2/n_2$:

$$H_0 : p_1 = p_2 \quad \text{vs.} \quad H_1 : p_1 \neq p_2 \quad (5.3)$$

Under H_0 , we use the pooled proportion:

$$\hat{p}_{\text{pooled}} = \frac{x_1 + x_2}{n_1 + n_2} \quad (5.4)$$

Test statistic:

$$Z = \frac{\hat{p}_2 - \hat{p}_1}{\sqrt{\hat{p}_{\text{pooled}}(1 - \hat{p}_{\text{pooled}})(1/n_1 + 1/n_2)}} \quad (5.5)$$

Under H_0 : $Z \sim \mathcal{N}(0, 1)$ approximately.

P-value (two-sided): $p = 2(1 - \Phi(|Z|))$ where Φ is the standard normal CDF.

Decision rule: Reject H_0 if $|Z| > z_{\alpha/2}$ (equivalently, if $p < \alpha$).

Example 5.1 (Conversion Rate Test).Control: $n_1 = 2000, x_1 = 160, \hat{p}_1 = 0.080$ Treatment: $n_2 = 2000, x_2 = 190, \hat{p}_2 = 0.095$

Pooled proportion:

$$\hat{p}_{\text{pooled}} = \frac{160 + 190}{2000 + 2000} = \frac{350}{4000} = 0.0875 \quad (5.6)$$

Standard error:

$$SE = \sqrt{0.0875 \times 0.9125 \times (1/2000 + 1/2000)} = 0.00894 \quad (5.7)$$

Test statistic:

$$Z = \frac{0.095 - 0.080}{0.00894} = 1.68 \quad (5.8)$$

P-value: $p = 2(1 - \Phi(1.68)) = 2(1 - 0.9535) = 0.093$ At $\alpha = 0.05$, we fail to reject H_0 . The observed difference is not statistically significant.

Listing 5.1: Z-Test Implementation

```

1 import numpy as np
2 from scipy.stats import norm
3
4 def z_test_proportions(x1, n1, x2, n2, alpha=0.05, two_sided=True):
5     """
6     Perform z-test for comparing two proportions.
7
8     Parameters:
9     -----
10    x1, x2: int
11    Number of successes in each group
12    n1, n2: int
13    Sample sizes
14    alpha: float
15    Significance level
16    two_sided: bool
17    Perform two-sided test (default True)
18
19    Returns:
20    -----
21    dict: Test results including statistic, p-value, decision
22    """
23    # Sample proportions
24    p1_hat = x1 / n1
25    p2_hat = x2 / n2
26
27    # Pooled proportion under null hypothesis
28    p_pooled = (x1 + x2) / (n1 + n2)
29
30    # Standard error under null
31    se = np.sqrt(p_pooled * (1 - p_pooled) * (1/n1 + 1/n2))
32
33    # Test statistic

```

```

34     z_stat = (p2_hat - p1_hat) / se
35
36     # P-value
37     if two_sided:
38         p_value = 2 * (1 - norm.cdf(abs(z_stat)))
39     else:
40         p_value = 1 - norm.cdf(z_stat)
41
42     # Decision
43     reject_null = p_value < alpha
44
45     # Confidence interval (using unpooled variance)
46     se_unpooled = np.sqrt(p1_hat*(1-p1_hat)/n1 + p2_hat*(1-p2_hat)/n2)
47     z_crit = norm.ppf(1 - alpha/2) if two_sided else norm.ppf(1 - alpha
48 )
49     ci_lower = (p2_hat - p1_hat) - z_crit * se_unpooled
50     ci_upper = (p2_hat - p1_hat) + z_crit * se_unpooled
51
52     return {
53         'p1_hat': p1_hat,
54         'p2_hat': p2_hat,
55         'difference': p2_hat - p1_hat,
56         'z_statistic': z_stat,
57         'p_value': p_value,
58         'alpha': alpha,
59         'reject_null': reject_null,
60         'ci_lower': ci_lower,
61         'ci_upper': ci_upper
62     }
63
64 # Example
65 result = z_test_proportions(x1=160, n1=2000, x2=190, n2=2000)
66
67 print(f"Control conversion: {result['p1_hat']:.3f}")
68 print(f"Treatment conversion: {result['p2_hat']:.3f}")
69 print(f"Difference: {result['difference']:.3f}")
70 print(f"Z-statistic: {result['z_statistic']:.3f}")
71 print(f"P-value: {result['p_value']:.4f}")
72 print(f"Reject null: {result['reject_null']}")
73 print(f"95% CI: [{result['ci_lower']:.4f}, {result['ci_upper']:.4f}]")

```

5.3.2 Assumptions and Validity

The z-test relies on:

1. **Independence:** Observations are independent within and between groups
2. **Random sampling:** Groups formed by randomization
3. **Large samples:** Normal approximation valid

Rule of thumb: $n_i p_i > 5$ and $n_i(1 - p_i) > 5$ for $i = 1, 2$.

When assumptions are violated, consider:

- Exact tests (Fisher's exact test) for small samples

- Permutation tests (robust to distributional assumptions)
- Adjustments for clustering or correlation

5.4 T-Test for Means

When comparing continuous outcomes with unknown variance, the t-test is appropriate.

5.4.1 Two-Sample T-Test

For means $\bar{X}_1$ and $\bar{X}_2$ from groups with variances s_1^2 and s_2^2 :

$$H_0 : \mu_1 = \mu_2 \quad \text{vs.} \quad H_1 : \mu_1 \neq \mu_2 \quad (5.9)$$

Equal variances (pooled t-test):

Pooled variance:

$$s_p^2 = \frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{n_1 + n_2 - 2} \quad (5.10)$$

Test statistic:

$$t = \frac{\bar{X}_2 - \bar{X}_1}{s_p \sqrt{1/n_1 + 1/n_2}} \quad (5.11)$$

Distribution under H_0 : $t \sim t_{n_1+n_2-2}$

Unequal variances (Welch's t-test):

Test statistic:

$$t = \frac{\bar{X}_2 - \bar{X}_1}{\sqrt{s_1^2/n_1 + s_2^2/n_2}} \quad (5.12)$$

Degrees of freedom (Welch-Satterthwaite):

$$\nu = \frac{(s_1^2/n_1 + s_2^2/n_2)^2}{\frac{(s_1^2/n_1)^2}{n_1-1} + \frac{(s_2^2/n_2)^2}{n_2-1}} \quad (5.13)$$

When to use which:

- Welch's t-test is more robust and generally recommended
- Use pooled when variances are truly equal (rare in practice)

Example 5.2 (Revenue Per User).

Control: $n_1 = 500, \bar{X}_1 = 45.2, s_1 = 28.3$

Treatment: $n_2 = 500, \bar{X}_2 = 48.7, s_2 = 31.1$

Using Welch's t-test:

Standard error:

$$SE = \sqrt{\frac{28.3^2}{500} + \frac{31.1^2}{500}} = \sqrt{1.603 + 1.934} = 1.88 \quad (5.14)$$

Test statistic:

$$t = \frac{48.7 - 45.2}{1.88} = 1.86 \quad (5.15)$$

Degrees of freedom ≈ 982 .

For t_{982} , $t \approx \mathcal{N}(0, 1)$, so $p \approx 0.063$.

Not significant at $\alpha = 0.05$.

Listing 5.2: T-Test in R

```

1 # Simulate A/B test data
2 set.seed(123)
3 control <- rnorm(500, mean = 45.2, sd = 28.3)
4 treatment <- rnorm(500, mean = 48.7, sd = 31.1)
5
6 # Welch's t-test (default in R)
7 test_result <- t.test(treatment, control, var.equal = FALSE)
8
9 print(test_result)
10
11 # Extract key components
12 cat("\nTest statistic:", test_result$statistic, "\n")
13 cat("P-value:", test_result$p.value, "\n")
14 cat("95% CI:", test_result$conf.int, "\n")
15
16 # Pooled t-test (for comparison)
17 test_pooled <- t.test(treatment, control, var.equal = TRUE)
18 cat("\nPooled t-test p-value:", test_pooled$p.value, "\n")

```

5.5 Understanding P-Values

Definition 5.1 (P-Value). The p-value is the probability, under the null hypothesis, of observing a test statistic at least as extreme as the one actually observed.

5.5.1 Correct Interpretation

A p-value of 0.03 means:

- ✓ If H_0 were true, we would observe data this extreme in 3% of experiments
- ✓ The data are somewhat inconsistent with H_0

5.5.2 Common Misinterpretations

1. **WRONG:** "The probability that H_0 is true is 0.03"

P-values are NOT the probability of hypotheses. They assume H_0 is true.

2. **WRONG:** "The probability that the result occurred by chance is 0.03"

Everything has some probability; this confuses conditional probabilities.

3. **WRONG:** "A p-value of 0.049 is meaningfully different from 0.051"

The $\alpha = 0.05$ threshold is arbitrary. Treat p-values as continuous measures of evidence, not binary outcomes.

4. **WRONG:** "A non-significant result means no effect exists"

Absence of evidence is not evidence of absence. May lack power to detect small effects.

5.5.3 Statistical Significance vs. Practical Significance

Statistical significance ($p < \alpha$) does not imply practical importance.

Example 5.3 (Large Sample Paradox). With $n = 100,000$ per variant, a conversion rate increase from 10.00% to 10.05% (0.05 percentage points) may be highly statistically significant ($p < 0.001$) but economically negligible.

Conversely, with $n = 100$, a 5 percentage point increase may not reach significance despite clear practical importance.

Recommendation: Report effect sizes and confidence intervals alongside p-values. Make decisions based on practical significance, informed by statistical uncertainty.

5.6 Chi-Squared Test for Categorical Data

The chi-squared test assesses association between categorical variables.

5.6.1 Pearson's Chi-Squared Test

For 2×2 contingency table:

	Converted	Not Converted	Total
Control	O_{11}	O_{12}	n_1
Treatment	O_{21}	O_{22}	n_2
Total	c_1	c_2	n

Expected counts under independence:

$$E_{ij} = \frac{(\text{row } i \text{ total}) \times (\text{column } j \text{ total})}{n} \quad (5.16)$$

Test statistic:

$$\chi^2 = \sum_{i,j} \frac{(O_{ij} - E_{ij})^2}{E_{ij}} \quad (5.17)$$

Under H_0 of independence: $\chi^2 \sim \chi^2_{(r-1)(c-1)}$ where r and c are number of rows and columns.

For 2×2 tables: $df = 1$.

Proposition 5.1 (Equivalence of Z-Test and Chi-Squared Test). For 2×2 tables, $\chi^2 = Z^2$ where Z is the z-test statistic for proportions. The tests are equivalent for two-sided alternatives.

5.6.2 Fisher's Exact Test

When sample sizes are small or expected cell counts < 5 , use Fisher's exact test instead of chi-squared.

Fisher's test calculates the exact probability of the observed table and all more extreme tables under H_0 , without relying on asymptotics.

Listing 5.3: Chi-Squared and Fisher's Exact Tests

```

1 from scipy.stats import chi2_contingency, fisher_exact
2 import numpy as np
3
4 # Contingency table
5 # Rows: Control, Treatment
6 # Columns: Converted, Not Converted
7 table = np.array([[160, 1840],
8                   [190, 1810]])
9
10 # Chi-squared test
11 chi2_stat, p_chi2, dof, expected = chi2_contingency(table)
12
13 print("Chi-Squared Test:")
14 print(f"Chi-squared statistic: {chi2_stat:.3f}")
15 print(f"P-value: {p_chi2:.4f}")
16 print(f"Degrees of freedom: {dof}")
17 print("\nExpected frequencies:")
18 print(expected)
19
20 # Fisher's exact test
21 oddsratio, p_fisher = fisher_exact(table)
22
23 print("\nFisher's Exact Test:")
24 print(f"Odds ratio: {oddsratio:.3f}")
25 print(f"P-value: {p_fisher:.4f}")
26
27 # Verify equivalence with z-test
28 from scipy.stats import norm
29
30 x1, n1 = table[0, 0], table[0].sum()
31 x2, n2 = table[1, 0], table[1].sum()
32
33 p1 = x1 / n1
34 p2 = x2 / n2
35 p_pooled = (x1 + x2) / (n1 + n2)
36
37 z = (p2 - p1) / np.sqrt(p_pooled * (1 - p_pooled) * (1/n1 + 1/n2))
38 p_z = 2 * (1 - norm.cdf(abs(z)))
39
40 print(f"\nZ-test p-value: {p_z:.4f}")
41 print(f"z^2 = {z**2:.3f}, chi2 = {chi2_stat:.3f}")
42 print(f"Difference: {abs(z**2 - chi2_stat):.6f}")

```

5.7 Permutation Tests

Permutation tests offer a non-parametric alternative requiring minimal distributional assumptions.

5.7.1 Permutation Test Procedure

1. Calculate observed test statistic T_{obs} (e.g., difference in means)
2. Under H_0 of no effect, treatment labels are exchangeable

3. Generate permutation distribution:

- Randomly shuffle treatment assignments
- Calculate test statistic for permuted data
- Repeat many times (e.g., 10,000)

4. P-value: proportion of permutations where $|T| \geq |T_{\text{obs}}|$

Advantages:

- Exact finite-sample validity under randomization
- No distributional assumptions
- Works for complex test statistics

Disadvantages:

- Computationally intensive
- Less familiar to stakeholders

Listing 5.4: Permutation Test Implementation

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 def permutation_test(group1, group2, n_permutations=10000,
5                     statistic='mean_diff', seed=None):
6     """
7     Perform permutation test for two groups.
8
9     Parameters:
10    -----
11    group1, group2: array-like
12    Data for each group
13    n_permutations: int
14    Number of permutations
15    statistic: str
16    Test statistic ('mean_diff', 'median_diff', 't_stat')
17    seed: int
18    Random seed
19
20    Returns:
21    -----
22    dict: Test results
23    """
24    if seed is not None:
25        np.random.seed(seed)
26
27    group1 = np.array(group1)
28    group2 = np.array(group2)
29    combined = np.concatenate([group1, group2])
30    n1 = len(group1)
31
32    # Define test statistic function
33    def calc_statistic(data1, data2):

```

```

34     if statistic == 'mean_diff':
35         return np.mean(data2) - np.mean(data1)
36     elif statistic == 'median_diff':
37         return np.median(data2) - np.median(data1)
38     elif statistic == 't_stat':
39         from scipy.stats import ttest_ind
40         return ttest_ind(data2, data1, equal_var=False).statistic
41
42     # Observed test statistic
43     obs_stat = calc_statistic(group1, group2)
44
45     # Permutation distribution
46     perm_stats = []
47     for _ in range(n_permutations):
48         # Shuffle combined data
49         permuted = np.random.permutation(combined)
50         perm_group1 = permuted[:n1]
51         perm_group2 = permuted[n1:]
52         perm_stat = calc_statistic(perm_group1, perm_group2)
53         perm_stats.append(perm_stat)
54
55     perm_stats = np.array(perm_stats)
56
57     # P-value (two-sided)
58     p_value = np.mean(np.abs(perm_stats) >= np.abs(obs_stat))
59
60     return {
61         'observed_statistic': obs_stat,
62         'permutation_statistics': perm_stats,
63         'p_value': p_value,
64         'n_permutations': n_permutations
65     }
66
67 # Example with simulated data
68 np.random.seed(42)
69 control = np.random.exponential(scale=50, size=500)
70 treatment = np.random.exponential(scale=55, size=500)
71
72 result = permutation_test(control, treatment, n_permutations=10000)
73
74 print(f"Observed difference in means: {result['observed_statistic']:.3f}")
75 print(f"P-value: {result['p_value']:.4f}")
76
77 # Visualize permutation distribution
78 plt.figure(figsize=(10, 6))
79 plt.hist(result['permutation_statistics'], bins=50, density=True,
80         alpha=0.7, label='Permutation distribution')
81 plt.axvline(result['observed_statistic'], color='red',
82         linestyle='--', linewidth=2, label='Observed statistic')
83 plt.axvline(-result['observed_statistic'], color='red',
84         linestyle='--', linewidth=2)
85 plt.xlabel('Test Statistic (Difference in Means)')
86 plt.ylabel('Density')
87 plt.title('Permutation Test: Null Distribution')
88 plt.legend()
89 plt.grid(alpha=0.3)
90 plt.tight_layout()

```

```
91 plt.savefig('permutation_distribution.pdf')
92 plt.show()
```

5.8 Summary

Hypothesis testing provides a rigorous framework for making decisions from experimental data:

- Z-tests compare proportions using normal approximation
- T-tests compare means with unknown variances
- P-values measure consistency of data with null hypothesis
- Statistical significance differs from practical significance
- Chi-squared tests assess categorical associations
- Permutation tests offer non-parametric alternatives

Proper application requires understanding assumptions, limitations, and correct interpretation. The next chapter addresses challenges arising from multiple comparisons.

5.9 Further Reading

- ? : Theoretical foundations
- ? : Comprehensive treatment of permutation tests
- ? : ASA statement on p-values

5.10 Exercises

1. Test Implementation

- (a) Implement a function performing z-test, t-test, chi-squared test, and Fisher's exact test. Compare results for various scenarios.
- (b) Create visualizations showing p-value distributions under null and alternative hypotheses.

2. P-Value Interpretation

- (a) Explain why "probability null hypothesis is true" is incorrect interpretation.
- (b) Simulate experiments under null hypothesis. What proportion have $p < 0.05$?

3. Power and Sample Size

- (a) For given effect size, simulate experiments with various sample sizes. Plot empirical power vs. theoretical power.

- (b) Demonstrate how increasing sample size can make trivial effects statistically significant.

4. **Permutation Tests**

- (a) Compare permutation test p-values to parametric test p-values for: normal data, exponential data, heavy-tailed data.
- (b) Implement a permutation test for ratio of means instead of difference. When might this be useful?

5. **Real Data Analysis**

- (a) Analyze a real A/B test dataset using multiple test procedures. Compare conclusions.
- (b) Investigate how outliers affect different test procedures.

Chapter 6

Multiple Testing and False Discovery Rate

6.1 Learning Objectives

After completing this chapter, readers will be able to:

- Understand why multiple testing inflates false positive rates
- Apply family-wise error rate (FWER) correction methods
- Use false discovery rate (FDR) control procedures
- Choose appropriate correction methods for different scenarios
- Handle multiple metrics and variants in A/B testing

6.2 The Multiple Testing Problem

When conducting many hypothesis tests simultaneously, the probability of at least one false positive increases dramatically—even when all null hypotheses are true.

6.2.1 Probability of False Positives

For a single test at level $\alpha = 0.05$:

$$P(\text{Type I error}) = 0.05 \quad (6.1)$$

For m independent tests, all with true null hypotheses:

$$P(\text{at least one Type I error}) = 1 - (1 - \alpha)^m \quad (6.2)$$

Example 6.1 (Inflation of False Positive Rate).

Tests	FWER ($\alpha = 0.05$)	FWER ($\alpha = 0.01$)
1	0.050	0.010
5	0.226	0.049
10	0.401	0.096
20	0.642	0.182
50	0.923	0.395
100	0.994	0.634

With 20 tests, probability of at least one false positive exceeds 64%!

6.2.2 Sources of Multiplicity in A/B Testing

1. **Multiple Metrics:** Testing conversion, revenue, engagement simultaneously
2. **Multiple Variants:** A/B/C/D tests with $\binom{k}{2}$ pairwise comparisons
3. **Multiple Subgroups:** Analyzing mobile vs. desktop, new vs. returning users
4. **Multiple Time Points:** Continuous monitoring, interim analyses
5. **Multiple Experiments:** Running many tests in parallel

6.3 Family-Wise Error Rate (FWER)

Definition 6.1 (Family-Wise Error Rate). The FWER is the probability of making at least one Type I error among all tests:

$$\text{FWER} = P(\text{Reject at least one true } H_0) \quad (6.3)$$

FWER control ensures strong protection against false positives but may reduce power.

6.3.1 Bonferroni Correction

The Bonferroni method is the simplest FWER control procedure.

Theorem 6.1 (Bonferroni Correction). For m tests, reject null hypothesis i if $p_i < \alpha/m$. This guarantees $\text{FWER} \leq \alpha$ regardless of test dependence.

Proof: By Boole's inequality:

$$\text{FWER} = P\left(\bigcup_{i=1}^m \{p_i < \alpha/m\}\right) \leq \sum_{i=1}^m P(p_i < \alpha/m) = m \cdot \frac{\alpha}{m} = \alpha \quad (6.4)$$

Example 6.2 (Bonferroni in Practice). Testing 5 metrics at overall level $\alpha = 0.05$:

Bonferroni threshold: $\alpha^* = 0.05/5 = 0.01$

Reject individual tests only if $p < 0.01$.

Observed p-values: 0.008, 0.025, 0.045, 0.003, 0.120

Reject: Tests 1 and 4 (p-values 0.008 and 0.003)

Without correction, would reject tests 1, 2, 3, and 4.

Advantages:

- Simple to implement
- Valid under arbitrary dependence
- Conservative (strong Type I error control)

Disadvantages:

- Often too conservative
- Reduces power substantially
- Treats all tests equally (no prioritization)

6.3.2 Holm-Bonferroni Method

The Holm-Bonferroni procedure is uniformly more powerful than Bonferroni while maintaining FWER control.

Algorithm 2 Holm-Bonferroni Procedure

```

Order p-values:  $p_{(1)} \leq p_{(2)} \leq \dots \leq p_{(m)}$ 
for  $i = 1$  to  $m$  do
  if  $p_{(i)} \leq \alpha / (m - i + 1)$  then
    Reject  $H_{(i)}$ 
  else
    Stop; retain  $H_{(i)}, H_{(i+1)}, \dots, H_{(m)}$ 
  end if
end for

```

Example 6.3 (Holm-Bonferroni Application). Five p-values: 0.008, 0.025, 0.045, 0.003, 0.120

Ordered: $p_{(1)} = 0.003, p_{(2)} = 0.008, p_{(3)} = 0.025, p_{(4)} = 0.045, p_{(5)} = 0.120$

Thresholds ($\alpha = 0.05$):

$$\text{Test 1: } p_{(1)} = 0.003 \leq 0.05/5 = 0.010 \quad \checkmark \text{ Reject} \quad (6.5)$$

$$\text{Test 2: } p_{(2)} = 0.008 \leq 0.05/4 = 0.0125 \quad \checkmark \text{ Reject} \quad (6.6)$$

$$\text{Test 3: } p_{(3)} = 0.025 > 0.05/3 = 0.0167 \quad \times \text{ Stop} \quad (6.7)$$

Reject tests 1 and 2; retain 3, 4, 5.

Bonferroni would reject only tests 1 and 4, missing test 2.

Listing 6.1: Bonferroni and Holm-Bonferroni Implementation

```

1 import numpy as np
2
3 def bonferroni_correction(p_values, alpha=0.05):
4     """
5     Apply Bonferroni correction.
6
7     Parameters:
8     -----
9     p_values: array-like
10             Observed p-values
11     alpha: float
12             Family-wise error rate
13
14     Returns:
15     -----
16     dict: Rejection decisions and adjusted p-values
17     """
18     p_values = np.array(p_values)
19     m = len(p_values)
20
21     # Adjusted significance level
22     alpha_bonf = alpha / m
23
24     # Adjusted p-values (capped at 1)

```

```

25     adjusted_p = np.minimum(p_values * m, 1.0)
26
27     # Rejection decisions
28     reject = p_values < alpha_bonf
29
30     return {
31         'reject': reject,
32         'adjusted_p_values': adjusted_p,
33         'threshold': alpha_bonf,
34         'num_rejected': np.sum(reject)
35     }
36
37 def holm_bonferroni(p_values, alpha=0.05):
38     """
39     Apply Holm-Bonferroni procedure.
40
41     Parameters:
42     -----
43     p_values: array-like
44             Observed p-values
45     alpha: float
46             Family-wise error rate
47
48     Returns:
49     -----
50     dict: Rejection decisions and adjusted p-values
51     """
52     p_values = np.array(p_values)
53     m = len(p_values)
54
55     # Sort p-values and track original indices
56     sorted_indices = np.argsort(p_values)
57     sorted_p = p_values[sorted_indices]
58
59     # Apply Holm procedure
60     reject_sorted = np.zeros(m, dtype=bool)
61     for i in range(m):
62         if sorted_p[i] <= alpha / (m - i):
63             reject_sorted[i] = True
64         else:
65             break # Stop at first non-rejection
66
67     # Map back to original order
68     reject = np.zeros(m, dtype=bool)
69     reject[sorted_indices] = reject_sorted
70
71     # Adjusted p-values
72     adjusted_p = np.zeros(m)
73     for i in range(m):
74         adjusted_p[sorted_indices[i]] = min(
75             max(sorted_p[:i+1] * (m - np.arange(i+1))),
76             1.0
77         )
78
79     return {
80         'reject': reject,
81         'adjusted_p_values': adjusted_p,
82         'num_rejected': np.sum(reject)

```



```

83     }
84
85 # Example
86 p_values = [0.008, 0.025, 0.045, 0.003, 0.120]
87
88 bonf_result = bonferroni_correction(p_values)
89 holm_result = holm_bonferroni(p_values)
90
91 print("Original p-values:", p_values)
92 print("\nBonferroni:")
93 print("  Threshold:", bonf_result['threshold'])
94 print("  Reject:", bonf_result['reject'])
95 print("  Adjusted p-values:", bonf_result['adjusted_p_values'])
96
97 print("\nHolm-Bonferroni:")
98 print("  Reject:", holm_result['reject'])
99 print("  Adjusted p-values:", holm_result['adjusted_p_values'])

```

6.3.3 Šidák Correction

For independent tests, the Šidák correction is less conservative than Bonferroni.

$$\alpha_{\text{Šidák}} = 1 - (1 - \alpha)^{1/m} \quad (6.8)$$

This achieves exact FWER = α under independence.

For $\alpha = 0.05$ and $m = 10$:

$$\text{Bonferroni: } \alpha^* = 0.05/10 = 0.005 \quad (6.9)$$

$$\text{Šidák: } \alpha^* = 1 - (1 - 0.05)^{1/10} \approx 0.00512 \quad (6.10)$$

Slightly more powerful but requires independence assumption.

6.4 False Discovery Rate (FDR)

FWER control is conservative when testing many hypotheses. The false discovery rate offers a less stringent alternative.

Definition 6.2 (False Discovery Rate). The FDR is the expected proportion of false positives among rejected hypotheses:

$$\text{FDR} = \mathbb{E} \left[\frac{\text{False Positives}}{\text{Total Rejections}} \right] \quad (6.11)$$

where the ratio is defined as 0 if no hypotheses are rejected.

FDR control is particularly useful when:

- Testing many hypotheses (e.g., $m > 20$)
- Some false positives are acceptable
- Power is important (exploratory analysis)

6.4.1 Benjamini-Hochberg Procedure

The Benjamini-Hochberg (BH) procedure controls FDR at level α .

Algorithm 3 Benjamini-Hochberg Procedure

Order p-values: $p_{(1)} \leq p_{(2)} \leq \dots \leq p_{(m)}$

Find largest k such that $p_{(k)} \leq \frac{k}{m}\alpha$

Reject hypotheses $H_{(1)}, \dots, H_{(k)}$

Theorem 6.2 (BH FDR Control). Under independence or positive dependence, the BH procedure controls FDR at level α .

Example 6.4 (Benjamini-Hochberg Application). Ten p-values at $\alpha = 0.10$:

0.001, 0.008, 0.039, 0.041, 0.042, 0.060, 0.074, 0.205, 0.212, 0.216

Rank	P-value	Threshold	Decision
1	0.001	$1/10 \times 0.10 = 0.010$	Reject
2	0.008	$2/10 \times 0.10 = 0.020$	Reject
3	0.039	$3/10 \times 0.10 = 0.030$	Fail
4	0.041	$4/10 \times 0.10 = 0.040$	Fail

Largest k with $p_{(k)} \leq k/m \times \alpha$: $k = 2$

Reject tests 1 and 2.

With Bonferroni ($\alpha^* = 0.01$): reject only test 1.

Listing 6.2: FDR Control in R

```

1 # Benjamini-Hochberg procedure
2 p_values <- c(0.001, 0.008, 0.039, 0.041, 0.042,
3               0.060, 0.074, 0.205, 0.212, 0.216)
4
5 # Using built-in p.adjust
6 bh_adjusted <- p.adjust(p_values, method = "BH")
7 bonf_adjusted <- p.adjust(p_values, method = "bonferroni")
8 holm_adjusted <- p.adjust(p_values, method = "holm")
9
10 # Compare methods
11 comparison <- data.frame(
12   Original = p_values,
13   BH = bh_adjusted,
14   Bonferroni = bonf_adjusted,
15   Holm = holm_adjusted
16 )
17
18 print(comparison)
19
20 # Rejection decisions at alpha = 0.10
21 alpha <- 0.10
22 cat("\nRejections at alpha =", alpha, "\n")
23 cat("\t\tBH:", sum(bh_adjusted < alpha), "\n")
24 cat("\t\tBonferroni:", sum(bonf_adjusted < alpha), "\n")
25 cat("\t\tHolm:", sum(holm_adjusted < alpha), "\n")

```

6.4.2 Choosing Between FWER and FDR

Consideration	FWER	FDR
Number of tests	Few ($m < 20$)	Many ($m > 20$)
False positive cost	High	Moderate
Power priority	Lower	Higher
Setting	Confirmatory	Exploratory
Common use	Primary analysis	Screening

6.5 Multiple Testing Strategies for A/B Testing

6.5.1 Hierarchical Testing

Designate metrics hierarchically:

1. **Primary Metric:** Decision-making metric (no correction needed)
2. **Secondary Metrics:** Supporting evidence (apply correction)
3. **Guardrail Metrics:** Monitor for harm (separate testing)

Example:

- Primary: Conversion rate
- Secondary: Revenue per user, engagement time
- Guardrails: Page load time, error rate

Test primary metric at $\alpha = 0.05$. If significant, test secondary metrics with Bonferroni or FDR correction.

6.5.2 Pre-specification and Registration

Pre-specify:

- Which metrics will be tested
- Primary vs. secondary distinction
- Multiple testing correction method
- Subgroups of interest

This prevents:

- Post-hoc hypothesis generation
- P-hacking and data dredging
- Cherry-picking significant results

6.5.3 Handling Multiple Variants

For k variants, there are $\binom{k}{2}$ pairwise comparisons.

Strategies:

1. **Control Only:** Compare each treatment to control (most common)
Number of tests: $k - 1$
Less correction needed than all pairwise comparisons
2. **Dunnett's Test:** Specialized procedure for comparing multiple treatments to control
More powerful than Bonferroni for this specific structure
3. **All Pairwise:** Compare all pairs
Number of tests: $\binom{k}{2}$
Requires substantial correction
4. **Best Treatment:** Simply select best-performing variant
No hypothesis testing—exploratory approach

6.6 Practical Recommendations

1. **Limit Metrics:** Test only truly important metrics
2. **Hierarchical Structure:** Use primary/secondary framework
3. **Pre-register:** Specify hypotheses before seeing data
4. **Choose Correction Wisely:**
 - Few tests, high stakes: FWER (Bonferroni/Holm)
 - Many tests, exploratory: FDR (Benjamini-Hochberg)
5. **Report All Tests:** Document all tests conducted, not just significant ones
6. **Distinguish Confirmatory and Exploratory:** Use different standards
7. **Consider Effect Sizes:** Don't rely solely on p-values

6.7 Summary

Multiple testing is pervasive in A/B testing and requires careful attention:

- Without correction, false positive rates increase dramatically
- FWER methods (Bonferroni, Holm) provide strong error control
- FDR methods (Benjamini-Hochberg) offer more power for many tests
- Hierarchical testing and pre-specification prevent p-hacking

- Method choice depends on number of tests and acceptable error rates

Proper multiple testing control maintains experimental integrity while preserving statistical power.

6.8 Further Reading

- ? : Original FDR paper
- ? : Comprehensive multiple testing treatment
- ? : Practical applications

6.9 Exercises

1. Error Rate Calculations

- Derive the formula for FWER with m independent tests.
- Simulate 10,000 experiments each testing 20 true nulls. Verify empirical FWER matches theory.

2. Method Comparisons

- Compare power of Bonferroni, Holm, and BH for $m = 5, 10, 20, 50$ tests. Plot results.
- Under what conditions does FDR control provide substantially more power than FWER?

3. A/B Testing Applications

- Design a multiple testing strategy for an A/B/C/D test with 5 metrics.
- Analyze a dataset with 15 metrics. Compare conclusions under different correction methods.

4. Simulation Studies

- Simulate mixture of null and alternative hypotheses. Evaluate FWER and FDR empirically.
- Investigate robustness of BH procedure to positive and negative dependence among tests.

Chapter 7

Bayesian Approaches to A/B Testing

7.1 Learning Objectives

After completing this chapter, readers will be able to:

- Apply Bayesian inference to A/B testing
- Construct and interpret posterior distributions
- Use conjugate priors for analytical solutions
- Calculate probability of superiority and expected loss
- Implement Bayesian A/B tests in Python and R
- Understand advantages and limitations of Bayesian approaches

7.2 Bayesian vs. Frequentist Paradigms

The Bayesian framework treats parameters as random variables and updates beliefs using observed data.

7.2.1 Key Philosophical Differences

Aspect	Frequentist	Bayesian
Parameters	Fixed constants	Random variables
Probability	Long-run frequency	Degree of belief
Inference	Confidence intervals	Credible intervals
Prior info	Not incorporated	Explicitly modeled
Interpretation	Indirect (via procedures)	Direct (probability)
Sequential testing	Problematic	Natural

7.2.2 Bayes' Theorem for A/B Testing

$$p(\theta|D) = \frac{p(D|\theta) \cdot p(\theta)}{p(D)} \propto p(D|\theta) \cdot p(\theta) \quad (7.1)$$

where:

- θ : Parameter of interest (e.g., conversion rate)
- D : Observed data
- $p(\theta)$: Prior distribution (before seeing data)
- $p(D|\theta)$: Likelihood (probability of data given parameter)
- $p(\theta|D)$: Posterior distribution (after seeing data)

The posterior combines prior beliefs with observed evidence.

7.3 Conjugate Priors for Common Distributions

Conjugate priors yield posterior distributions in the same family, enabling analytical solutions.

7.3.1 Beta-Binomial Model for Conversion Rates

For binary outcomes (conversions), the natural model is:

Likelihood:

$$X|p \sim \text{Binomial}(n, p) \quad (7.2)$$

Prior:

$$p \sim \text{Beta}(\alpha, \beta) \quad (7.3)$$

The Beta distribution has PDF:

$$f(p; \alpha, \beta) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} p^{\alpha-1} (1-p)^{\beta-1}, \quad p \in (0, 1) \quad (7.4)$$

Properties:

- $\mathbb{E}[p] = \frac{\alpha}{\alpha + \beta}$
- $\text{Mode}[p] = \frac{\alpha-1}{\alpha + \beta - 2}$ for $\alpha, \beta > 1$
- $\text{Var}(p) = \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)}$

Posterior:

Given x conversions from n trials:

Theorem 7.1 (Beta-Binomial Conjugacy). If $p \sim \text{Beta}(\alpha, \beta)$ and $X|p \sim \text{Binomial}(n, p)$, then:

$$p|X = x \sim \text{Beta}(\alpha + x, \beta + n - x) \quad (7.5)$$

Interpretation:

- Prior $\text{Beta}(\alpha, \beta)$ represents $\alpha - 1$ prior successes and $\beta - 1$ prior failures
- Posterior adds observed data: x successes and $n - x$ failures
- Total: $\alpha + x - 1$ successes, $\beta + n - x - 1$ failures

Example 7.1 (Uninformative Prior). Use $\text{Beta}(1, 1)$ (uniform on $[0, 1]$) when no prior information exists.

After observing 80 conversions from 1000 users:

$$p|D \sim \text{Beta}(1 + 80, 1 + 920) = \text{Beta}(81, 921) \quad (7.6)$$

Posterior mean: $\mathbb{E}[p|D] = \frac{81}{81+921} = 0.0808$

95% credible interval: $[0.0649, 0.0992]$ (using beta quantiles)

Listing 7.1: Bayesian Inference for Conversion Rate

```

1 import numpy as np
2 from scipy.stats import beta
3 import matplotlib.pyplot as plt
4
5 def bayesian_conversion_inference(conversions, trials,
6                                   prior_alpha=1, prior_beta=1):
7     """
8     Bayesian inference for conversion rate using Beta-Binomial model.
9
10    Parameters:
11    -----
12    conversions: int
13    Number of conversions
14    trials: int
15    Number of trials
16    prior_alpha, prior_beta: float
17    Beta prior parameters
18
19    Returns:
20    -----
21    dict: Posterior statistics
22    """
23    # Posterior parameters
24    post_alpha = prior_alpha + conversions
25    post_beta = prior_beta + trials - conversions
26
27    # Posterior statistics
28    post_mean = post_alpha / (post_alpha + post_beta)
29    post_mode = (post_alpha - 1) / (post_alpha + post_beta - 2) \
30                if post_alpha > 1 and post_beta > 1 else None
31    post_var = (post_alpha * post_beta) / \
32                ((post_alpha + post_beta)**2 * (post_alpha + post_beta +
33                1))
34
35    # Credible intervals
36    ci_95 = beta.ppf([0.025, 0.975], post_alpha, post_beta)
37    ci_90 = beta.ppf([0.05, 0.95], post_alpha, post_beta)
38
39    return {
40        'posterior_alpha': post_alpha,
41        'posterior_beta': post_beta,
42        'posterior_mean': post_mean,
43        'posterior_mode': post_mode,
44        'posterior_std': np.sqrt(post_var),
45        'ci_95': ci_95,
46        'ci_90': ci_90,

```

```

46         'observed_rate': conversions / trials
47     }
48
49 # Example
50 result = bayesian_conversion_inference(conversions=80, trials=1000)
51
52 print("Posterior_Beta({}, {})".format(result['posterior_alpha'],
53                                     result['posterior_beta']))
54 print(f"Posterior_mean: {result['posterior_mean']:.4f}")
55 print(f"Posterior_std: {result['posterior_std']:.4f}")
56 print(f"95%_Credible_Interval: [{result['ci_95'][0]:.4f}, {result['ci_95'][1]:.4f}]")
57
58
59 # Visualize posterior
60 p_values = np.linspace(0.05, 0.12, 1000)
61 posterior_pdf = beta.pdf(p_values, result['posterior_alpha'],
62                          result['posterior_beta'])
63
64 plt.figure(figsize=(10, 6))
65 plt.plot(p_values, posterior_pdf, linewidth=2, label='Posterior')
66 plt.axvline(result['posterior_mean'], color='red', linestyle='--',
67            label=f"Mean={result['posterior_mean']:.4f}")
68 plt.fill_between(p_values, 0, posterior_pdf,
69                 where=(p_values >= result['ci_95'][0]) &
70                 (p_values <= result['ci_95'][1]),
71                 alpha=0.3, label='95%_Credible_Interval')
72 plt.xlabel('Conversion_Rate')
73 plt.ylabel('Density')
74 plt.title('Posterior_Distribution_for_Conversion_Rate')
75 plt.legend()
76 plt.grid(alpha=0.3)
77 plt.tight_layout()
78 plt.savefig('bayesian_posterior.pdf')
79 plt.show()

```

7.3.2 Normal-Normal Model for Continuous Metrics

For continuous outcomes with known variance:

Likelihood:

$$Y_i | \mu \sim \mathcal{N}(\mu, \sigma^2) \quad (7.7)$$

Prior:

$$\mu \sim \mathcal{N}(\mu_0, \tau^2) \quad (7.8)$$

Posterior:

Given data $\mathbf{y} = (y_1, \dots, y_n)$ with sample mean $\bar{y}$:

Theorem 7.2 (Normal-Normal Conjugacy).

$$\mu | \mathbf{y} \sim \mathcal{N} \left(\frac{\frac{\mu_0}{\tau^2} + \frac{n\bar{y}}{\sigma^2}}{\frac{1}{\tau^2} + \frac{n}{\sigma^2}}, \left(\frac{1}{\tau^2} + \frac{n}{\sigma^2} \right)^{-1} \right) \quad (7.9)$$

The posterior mean is a precision-weighted average of prior mean and sample mean.

7.4 Bayesian A/B Testing

For comparing two variants, we need the posterior distribution of the difference or ratio.

7.4.1 Comparing Two Conversion Rates

Given:

$$p_A \sim \text{Beta}(\alpha_A, \beta_A) \quad (7.10)$$

$$p_B \sim \text{Beta}(\alpha_B, \beta_B) \quad (7.11)$$

We want to know:

$$P(p_B > p_A | D) = \int_0^1 \int_0^{p_B} f(p_A | D) f(p_B | D) dp_A dp_B \quad (7.12)$$

No closed form, but easily computed via:

1. Numerical integration

2. Monte Carlo simulation:

- Sample $p_A^{(1)}, \dots, p_A^{(S)}$ from $\text{Beta}(\alpha_A + x_A, \beta_A + n_A - x_A)$
- Sample $p_B^{(1)}, \dots, p_B^{(S)}$ from $\text{Beta}(\alpha_B + x_B, \beta_B + n_B - x_B)$
- Estimate $P(p_B > p_A | D) \approx \frac{1}{S} \sum_{s=1}^S \mathbb{I}(p_B^{(s)} > p_A^{(s)})$

Example 7.2 (Bayesian A/B Test).

Variant A: 80 conversions from 1000 users (7.13)

Variant B: 95 conversions from 1000 users (7.14)

Using uniform priors $\text{Beta}(1, 1)$:

$$p_A | D \sim \text{Beta}(81, 921) \quad (7.15)$$

$$p_B | D \sim \text{Beta}(96, 906) \quad (7.16)$$

Via Monte Carlo: $P(p_B > p_A | D) \approx 0.92$

Interpretation: There is a 92% probability that variant B has higher conversion rate than A.

Listing 7.2: Bayesian A/B Test in R

```

1 # Bayesian A/B test comparison
2 bayesian_ab_test <- function(x_a, n_a, x_b, n_b,
3                             prior_alpha = 1, prior_beta = 1,
4                             n_simulations = 100000) {
5   # Posterior parameters
6   post_alpha_a <- prior_alpha + x_a
7   post_beta_a <- prior_beta + n_a - x_a
8
9   post_alpha_b <- prior_alpha + x_b
10  post_beta_b <- prior_beta + n_b - x_b
11

```

```

12  # Monte Carlo simulation
13  p_a_samples <- rbeta(n_simulations, post_alpha_a, post_beta_a)
14  p_b_samples <- rbeta(n_simulations, post_alpha_b, post_beta_b)
15
16  # Probability B > A
17  prob_b_better <- mean(p_b_samples > p_a_samples)
18
19  # Expected lift
20  lift <- p_b_samples / p_a_samples - 1
21  expected_lift <- mean(lift)
22  lift_ci <- quantile(lift, c(0.025, 0.975))
23
24  # Difference distribution
25  diff <- p_b_samples - p_a_samples
26  diff_mean <- mean(diff)
27  diff_ci <- quantile(diff, c(0.025, 0.975))
28
29  list(
30    prob_b_better = prob_b_better,
31    expected_lift = expected_lift,
32    lift_ci = lift_ci,
33    difference_mean = diff_mean,
34    difference_ci = diff_ci,
35    posterior_a = list(alpha = post_alpha_a, beta = post_beta_a),
36    posterior_b = list(alpha = post_alpha_b, beta = post_beta_b)
37  )
38 }
39
40 # Example
41 result <- bayesian_ab_test(x_a = 80, n_a = 1000, x_b = 95, n_b = 1000)
42
43 cat("Probability B > A:", result$prob_b_better, "\n")
44 cat("Expected relative lift:", result$expected_lift, "\n")
45 cat("95% CI for lift:", result$lift_ci, "\n")
46 cat("Mean difference:", result$difference_mean, "\n")
47 cat("95% CI for difference:", result$difference_ci, "\n")

```

7.4.2 Expected Loss

Instead of binary decisions, we can quantify the expected loss from choosing the wrong variant.

Definition 7.1 (Expected Loss). The expected loss from choosing variant A is:

$$\mathcal{L}(A) = \mathbb{E}[\max(0, p_B - p_A) | D] \quad (7.17)$$

Similarly for choosing B:

$$\mathcal{L}(B) = \mathbb{E}[\max(0, p_A - p_B) | D] \quad (7.18)$$

Decision rule: Choose the variant with smaller expected loss.

Interpretation: $\mathcal{L}(A)$ is the expected opportunity cost (in terms of conversion rate) if we choose A but B is actually better.

```

1 def expected_loss(x_a, n_a, x_b, n_b, prior_alpha=1, prior_beta=1,
2                   n_simulations=100000):
3     """
4     Calculate expected loss for both variants.
5
6     Returns:
7     -----
8     dict: Expected losses and recommendation
9     """
10    # Posterior parameters
11    post_alpha_a = prior_alpha + x_a
12    post_beta_a = prior_beta + n_a - x_a
13    post_alpha_b = prior_alpha + x_b
14    post_beta_b = prior_beta + n_b - x_b
15
16    # Simulate from posteriors
17    p_a = np.random.beta(post_alpha_a, post_beta_a, n_simulations)
18    p_b = np.random.beta(post_alpha_b, post_beta_b, n_simulations)
19
20    # Expected loss calculations
21    loss_a = np.mean(np.maximum(0, p_b - p_a))
22    loss_b = np.mean(np.maximum(0, p_a - p_b))
23
24    # Relative expected loss
25    rel_loss_a = loss_a / np.mean(p_a)
26    rel_loss_b = loss_b / np.mean(p_b)
27
28    recommendation = 'A' if loss_a < loss_b else 'B'
29
30    return {
31        'loss_a': loss_a,
32        'loss_b': loss_b,
33        'relative_loss_a': rel_loss_a,
34        'relative_loss_b': rel_loss_b,
35        'recommendation': recommendation,
36        'loss_ratio': min(loss_a, loss_b) / max(loss_a, loss_b)
37    }
38
39 # Example
40 result = expected_loss(x_a=80, n_a=1000, x_b=95, n_b=1000)
41
42 print(f"Expected loss choosing A: {result['loss_a']:.6f}")
43 print(f"Expected loss choosing B: {result['loss_b']:.6f}")
44 print(f"Relative loss A: {result['relative_loss_a']:.2%}")
45 print(f"Relative loss B: {result['relative_loss_b']:.2%}")
46 print(f"Recommendation: Variant {result['recommendation']}")

```

7.5 Prior Selection

Prior choice affects posterior, especially with limited data.

7.5.1 Types of Priors

1. Uninformative (Flat) Priors

Beta(1, 1) (uniform) or Jeffrey's prior Beta(0.5, 0.5)

Lets data dominate; minimal assumptions

2. Weakly Informative Priors

Beta(2, 20) for expected conversion around 10%

Provides gentle regularization without strong assumptions

3. Informative Priors

Based on historical data or expert knowledge

Example: Beta(50, 450) from previous experiments showing 10% conversion

4. Empirical Bayes

Use data itself to estimate prior (controversial but practical)

7.5.2 Prior Sensitivity Analysis

Always check how conclusions depend on prior choice.

Example 7.3 (Prior Sensitivity). Test with 10 conversions from 100 users:

Prior	Post. Mean	Post. Std	95% CI
Beta(1, 1)	0.108	0.031	[0.058, 0.175]
Beta(5, 45)	0.106	0.026	[0.063, 0.162]
Beta(10, 90)	0.105	0.023	[0.066, 0.154]
Beta(20, 180)	0.103	0.020	[0.069, 0.145]

With small data, stronger priors pull estimates toward prior mean.

7.6 Advantages of Bayesian A/B Testing

1. Direct Probability Statements

"95% probability that B is better" vs. "reject null with $p = 0.03$ "

More intuitive for stakeholders

2. Natural Sequential Testing

Can stop anytime without inflating error rates

No need for α -spending functions

3. Incorporation of Prior Information

Leverage historical data and expert knowledge

Particularly valuable for rare events

4. Coherent Decision Framework

Expected loss provides principled stopping rules

Balances statistical and business considerations

5. Full Uncertainty Quantification

Entire posterior distribution available

Can compute any quantity of interest

7.7 Limitations and Challenges

1. Prior Dependence

Results depend on prior, especially with little data
Requires prior specification and sensitivity analysis

2. Computational Complexity

Complex models may require MCMC or variational inference
Not always analytical like conjugate examples

3. Interpretation Challenges

"Probability $B > A$ " requires careful framing
Can be misinterpreted as certainty

4. Less Familiar

Stakeholders more familiar with p-values and confidence intervals
Requires education and buy-in

5. No Error Rate Guarantees

Unlike frequentist methods, no Type I/II error control
Different philosophical framework for errors

7.8 Bayesian vs. Frequentist: When to Use Which

Scenario	Frequentist	Bayesian
Fixed sample size	✓	✓
Sequential/adaptive		✓
Prior information available		✓
Regulatory requirement	✓	
Simple communication	✓	
Probability statements needed		✓
Multiple metrics/variants	✓	✓

Recommendation: Both frameworks are valuable. Choose based on:

- Organizational culture and stakeholder preferences
- Regulatory requirements
- Prior information availability
- Sequential testing needs
- Computational resources

Many organizations use hybrid approaches: frequentist for primary analysis, Bayesian for exploration.

7.9 Summary

Bayesian A/B testing offers a powerful alternative to frequentist methods:

- Treats parameters as random variables with probability distributions
- Updates beliefs via Bayes' theorem
- Conjugate priors enable analytical solutions
- Provides direct probability statements
- Handles sequential testing naturally
- Requires prior specification and sensitivity analysis

The choice between Bayesian and frequentist approaches depends on context, but understanding both enriches the practitioner's toolkit.

7.10 Further Reading

- ? : Comprehensive Bayesian data analysis
- ? : Bayesian analysis for psychology (accessible)
- ? : Bayesian A/B testing for practitioners

7.11 Exercises

1. Beta-Binomial Calculations

- Derive the posterior for Beta-Binomial model from first principles.
- Show that $\text{Beta}(1, 1)$ is the uniform distribution on $[0, 1]$.

2. Bayesian A/B Test

- Implement complete Bayesian A/B test including posterior visualization, probability calculations, and expected loss.
- Compare Bayesian and frequentist conclusions for various sample sizes.

3. Prior Sensitivity

- Conduct sensitivity analysis varying prior from uninformative to strongly informative.
- At what sample size do priors matter little?

4. Sequential Testing

- Simulate sequential Bayesian testing with stopping rule based on probability threshold.
- Compare to frequentist sequential testing in terms of sample size and error rates.

Chapter 8

Advanced Experimental Designs

8.1 Learning Objectives

After completing this chapter, readers will be able to:

- Design and analyze factorial experiments testing multiple factors simultaneously
- Implement multi-armed bandit algorithms for dynamic traffic allocation
- Conduct sequential testing with appropriate statistical controls
- Handle network effects and interference in experimental settings
- Apply switchback designs for marketplace and platform experiments
- Choose appropriate advanced designs for specific business contexts

8.2 Factorial Designs

Factorial designs test multiple factors simultaneously, providing efficiency gains and enabling the detection of interaction effects that would be missed in separate experiments.

8.2.1 Full Factorial Designs

A 2^k factorial design tests k binary factors with all 2^k combinations.

Example 8.1 (2×2 Factorial Design). Testing email subject line (A vs. B) and send time (morning vs. evening):

	Morning	Evening
Subject A	Cell 1 (n=250)	Cell 2 (n=250)
Subject B	Cell 3 (n=250)	Cell 4 (n=250)

With 1,000 total users, we can estimate:

- Main effect of subject line: Effect of A vs. B averaged across times
- Main effect of send time: Effect of morning vs. evening averaged across subjects
- Interaction: Does subject line effect depend on time of day?

8.2.2 Statistical Model for Factorial Experiments

For a two-factor design:

$$Y_{ijk} = \mu + \alpha_i + \beta_j + (\alpha\beta)_{ij} + \epsilon_{ijk} \quad (8.1)$$

where:

- μ : Grand mean
- α_i : Main effect of factor A (level i)
- β_j : Main effect of factor B (level j)
- $(\alpha\beta)_{ij}$: Interaction effect
- ϵ_{ijk} : Random error, $\epsilon_{ijk} \sim \mathcal{N}(0, \sigma^2)$

Main Effects:

$$\text{Effect of A} = \bar{Y}_{A \text{ high}} - \bar{Y}_{A \text{ low}} \quad (8.2)$$

$$\text{Effect of B} = \bar{Y}_{B \text{ high}} - \bar{Y}_{B \text{ low}} \quad (8.3)$$

Interaction Effect:

$$\text{Interaction} = (\bar{Y}_{11} - \bar{Y}_{10}) - (\bar{Y}_{01} - \bar{Y}_{00}) \quad (8.4)$$

An interaction exists when the effect of one factor depends on the level of another factor.

Theorem 8.1 (Efficiency of Factorial Designs). A 2^k factorial design with n observations per cell provides the same precision for estimating main effects as k separate experiments each with $2^{k-1} \times n$ observations, while also estimating all interaction effects.

Listing 8.1: Comprehensive Factorial Design Analysis

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from scipy import stats
6 import statsmodels.api as sm
7 from statsmodels.formula.api import ols
8 from typing import Dict, Tuple
9
10 def simulate_factorial_experiment(
11     n_per_cell: int,
12     main_effect_A: float,
13     main_effect_B: float,
14     interaction: float,
15     baseline: float = 0.10,
16     seed: int = None
17 ) -> pd.DataFrame:
18     """
19     Simulate a 2x2 factorial experiment with binary outcome.
20 
```

```

21  """Parameters:
22  """-----
23  """n_per_cell: int
24  """Sample size per cell
25  """main_effect_A: float
26  """Main effect of factor A (additive on probability scale)
27  """main_effect_B: float
28  """Main effect of factor B
29  """interaction: float
30  """Interaction effect (additional effect when both A and B are
    present)
31  """baseline: float
32  """Baseline conversion rate (A=0, B=0)
33  """seed: int
34  """Random seed
35
36  """Returns:
37  """-----
38  """pd.DataFrame: Simulated experiment data
39  """
40      if seed is not None:
41          np.random.seed(seed)
42
43      # Define cell probabilities
44      # Cell (0,0): baseline
45      # Cell (1,0): baseline + main_effect_A
46      # Cell (0,1): baseline + main_effect_B
47      # Cell (1,1): baseline + main_effect_A + main_effect_B +
        interaction
48
49      probabilities = {
50          (0, 0): baseline,
51          (1, 0): baseline + main_effect_A,
52          (0, 1): baseline + main_effect_B,
53          (1, 1): baseline + main_effect_A + main_effect_B + interaction
54      }
55
56      data_list = []
57
58      for factor_A in [0, 1]:
59          for factor_B in [0, 1]:
60              p = probabilities[(factor_A, factor_B)]
61              conversions = np.random.binomial(1, p, n_per_cell)
62
63              cell_data = pd.DataFrame({
64                  'factor_A': factor_A,
65                  'factor_B': factor_B,
66                  'conversion': conversions,
67                  'user_id': range(
68                      (factor_A * 2 + factor_B) * n_per_cell,
69                      (factor_A * 2 + factor_B + 1) * n_per_cell
70                  )
71              })
72              data_list.append(cell_data)
73
74      data = pd.concat(data_list, ignore_index=True)
75
76      # Add categorical labels

```

```

77     data['factor_A_label'] = data['factor_A'].map({0: 'A0', 1: 'A1'})
78     data['factor_B_label'] = data['factor_B'].map({0: 'B0', 1: 'B1'})
79
80     return data
81
82 def analyze_factorial_experiment(data: pd.DataFrame) -> Dict:
83     """
84     Comprehensive analysis of 2x2 factorial experiment.
85
86     Parameters:
87     -----
88     data: pd.DataFrame
89           Must contain columns: factor_A, factor_B, conversion
90
91     Returns:
92     -----
93     dict: Analysis results including effects, tests, and diagnostics
94     """
95     # Cell means
96     cell_means = data.groupby(['factor_A', 'factor_B'])['conversion'].agg([
97         'mean', 'count', 'std'
98     ]).reset_index()
99
100    # Extract cell means
101    y00 = cell_means[(cell_means['factor_A']==0) &
102                     (cell_means['factor_B']==0)]['mean'].values[0]
103    y10 = cell_means[(cell_means['factor_A']==1) &
104                     (cell_means['factor_B']==0)]['mean'].values[0]
105    y01 = cell_means[(cell_means['factor_A']==0) &
106                     (cell_means['factor_B']==1)]['mean'].values[0]
107    y11 = cell_means[(cell_means['factor_A']==1) &
108                     (cell_means['factor_B']==1)]['mean'].values[0]
109
110    # Calculate effects
111    main_effect_A = ((y10 + y11) / 2) - ((y00 + y01) / 2)
112    main_effect_B = ((y01 + y11) / 2) - ((y00 + y10) / 2)
113    interaction_effect = (y11 - y10) - (y01 - y00)
114
115    # ANOVA using statsmodels
116    formula = 'conversion ~ C(factor_A) * C(factor_B)'
117    model = ols(formula, data=data).fit()
118    anova_table = sm.stats.anova_lm(model, typ=2)
119
120    # Effect sizes (Cohen's d for binary factors)
121    overall_std = data['conversion'].std()
122    cohens_d_A = main_effect_A / overall_std
123    cohens_d_B = main_effect_B / overall_std
124    cohens_d_interaction = interaction_effect / overall_std
125
126    # Confidence intervals for effects (using bootstrap)
127    n_bootstrap = 1000
128    bootstrap_effects = {
129        'main_A': [],
130        'main_B': [],
131        'interaction': []
132    }
133

```

```

134 np.random.seed(42)
135 for _ in range(n_bootstrap):
136     boot_sample = data.sample(n=len(data), replace=True)
137     boot_means = boot_sample.groupby(
138         ['factor_A', 'factor_B']
139     )['conversion'].mean()
140
141     b_y00 = boot_means[(0, 0)]
142     b_y10 = boot_means[(1, 0)]
143     b_y01 = boot_means[(0, 1)]
144     b_y11 = boot_means[(1, 1)]
145
146     bootstrap_effects['main_A'].append(
147         ((b_y10 + b_y11) / 2) - ((b_y00 + b_y01) / 2)
148     )
149     bootstrap_effects['main_B'].append(
150         ((b_y01 + b_y11) / 2) - ((b_y00 + b_y10) / 2)
151     )
152     bootstrap_effects['interaction'].append(
153         (b_y11 - b_y10) - (b_y01 - b_y00)
154     )
155
156     ci_main_A = np.percentile(bootstrap_effects['main_A'], [2.5, 97.5])
157     ci_main_B = np.percentile(bootstrap_effects['main_B'], [2.5, 97.5])
158     ci_interaction = np.percentile(bootstrap_effects['interaction'],
159                                     [2.5, 97.5])
160
161     return {
162         'cell_means': {
163             'A0_B0': y00,
164             'A1_B0': y10,
165             'A0_B1': y01,
166             'A1_B1': y11
167         },
168         'effects': {
169             'main_effect_A': main_effect_A,
170             'main_effect_B': main_effect_B,
171             'interaction': interaction_effect
172         },
173         'effect_sizes': {
174             'cohens_d_A': cohens_d_A,
175             'cohens_d_B': cohens_d_B,
176             'cohens_d_interaction': cohens_d_interaction
177         },
178         'confidence_intervals': {
179             'main_A_ci': ci_main_A,
180             'main_B_ci': ci_main_B,
181             'interaction_ci': ci_interaction
182         },
183         'anova_table': anova_table,
184         'model': model
185     }
186
187 def visualize_factorial_experiment(data: pd.DataFrame, results: Dict):
188     """Create comprehensive visualizations for factorial experiment."""
189     fig, axes = plt.subplots(2, 2, figsize=(14, 12))
190
191     # 1. Cell means with error bars

```

```

191 ax = axes[0, 0]
192 cell_stats = data.groupby(['factor_A_label', 'factor_B_label'])['
193     'conversion'
194 ].agg(['mean', 'sem'])
195
196 x_pos = np.arange(4)
197 cell_labels = ['A0_B0', 'A1_B0', 'A0_B1', 'A1_B1']
198 means = [results['cell_means'][label] for label in cell_labels]
199
200 # Calculate SEM for error bars
201 sems = []
202 for a in [0, 1]:
203     for b in [0, 1]:
204         subset = data[(data['factor_A']==a) & (data['factor_B']==b)
205             ]
206         sems.append(subset['conversion'].sem())
207
208 ax.bar(x_pos, means, yerr=[1.96*s for s in sems],
209     capsize=5, alpha=0.7, color=['C0', 'C1', 'C0', 'C1'])
210 ax.set_xticks(x_pos)
211 ax.set_xticklabels(cell_labels)
212 ax.set_ylabel('Conversion_Rate')
213 ax.set_title('Cell_Means_with_95%_CI')
214 ax.grid(axis='y', alpha=0.3)
215
216 # 2. Interaction plot
217 ax = axes[0, 1]
218 for b_val, b_label in [(0, 'B0'), (1, 'B1')]:
219     subset = data[data['factor_B'] == b_val]
220     means_by_A = subset.groupby('factor_A')['conversion'].mean()
221     ax.plot([0, 1], means_by_A.values, marker='o',
222         markersize=10, linewidth=2, label=b_label)
223
224 ax.set_xticks([0, 1])
225 ax.set_xticklabels(['A0', 'A1'])
226 ax.set_xlabel('Factor_A')
227 ax.set_ylabel('Conversion_Rate')
228 ax.set_title('Interaction_Plot')
229 ax.legend(title='Factor_B')
230 ax.grid(alpha=0.3)
231
232 # 3. Main effects visualization
233 ax = axes[1, 0]
234 effects = results['effects']
235 effect_names = ['Main_Effect_A', 'Main_Effect_B', 'Interaction']
236 effect_values = [
237     effects['main_effect_A'],
238     effects['main_effect_B'],
239     effects['interaction']
240 ]
241
242 cis = [
243     results['confidence_intervals']['main_A_ci'],
244     results['confidence_intervals']['main_B_ci'],
245     results['confidence_intervals']['interaction_ci']
246 ]
247
248 y_pos = np.arange(len(effect_names))

```

```

248     colors = ['green' if ci[0] > 0 or ci[1] < 0 else 'gray'
249               for ci in cis]
250
251     ax.barh(y_pos, effect_values, color=colors, alpha=0.7)
252
253     for i, (val, ci) in enumerate(zip(effect_values, cis)):
254         ax.plot([ci[0], ci[1]], [i, i], 'k-', linewidth=2)
255         ax.plot([ci[0], ci[0]], [i-0.1, i+0.1], 'k-', linewidth=2)
256         ax.plot([ci[1], ci[1]], [i-0.1, i+0.1], 'k-', linewidth=2)
257
258     ax.axvline(x=0, color='red', linestyle='--', linewidth=1)
259     ax.set_yticks(y_pos)
260     ax.set_yticklabels(effect_names)
261     ax.set_xlabel('Effect_Size_(Absolute)')
262     ax.set_title('Effects_with_95%_Bootstrap_CI')
263     ax.grid(axis='x', alpha=0.3)
264
265     # 4. Heatmap of cell means
266     ax = axes[1, 1]
267     heatmap_data = np.array([
268         [results['cell_means']['A0_B0'], results['cell_means']['A0_B1'],
269         [results['cell_means']['A1_B0'], results['cell_means']['A1_B1']
270     ])
271
272     im = ax.imshow(heatmap_data, cmap='RdYlGn', aspect='auto')
273     ax.set_xticks([0, 1])
274     ax.set_yticks([0, 1])
275     ax.set_xticklabels(['B0', 'B1'])
276     ax.set_yticklabels(['A0', 'A1'])
277     ax.set_xlabel('Factor_B')
278     ax.set_ylabel('Factor_A')
279     ax.set_title('Heatmap_of_Conversion_Rates')
280
281     # Add text annotations
282     for i in range(2):
283         for j in range(2):
284             text = ax.text(j, i, f'{heatmap_data[i,j]:.3f}',
285                             ha="center", va="center", color="black",
286                             fontsize=14, fontweight='bold')
287
288     plt.colorbar(im, ax=ax)
289
290     plt.tight_layout()
291     return fig
292
293     # Example: Simulate and analyze factorial experiment
294     np.random.seed(42)
295     data = simulate_factorial_experiment(
296         n_per_cell=500,
297         main_effect_A=0.02,      # Subject B increases conversion by 2pp
298         main_effect_B=0.015,    # Evening increases conversion by 1.5pp
299         interaction=0.01,       # Additional 1pp when both are combined
300         baseline=0.10
301     )
302
303     results = analyze_factorial_experiment(data)

```

```

304
305 print("=== Factorial Experiment Analysis ===\n")
306 print("Cell Means:")
307 for cell, mean in results['cell_means'].items():
308     print(f"    {cell}: {mean:.4f}")
309
310 print("\nEstimated Effects:")
311 print(f"    Main Effect A: {results['effects']['main_effect_A']:.4f}")
312 print(f"    Main Effect B: {results['effects']['main_effect_B']:.4f}")
313 print(f"    Interaction: {results['effects']['interaction']:.4f}")
314
315 print("\nEffect Sizes (Cohen's d):")
316 print(f"    Main Effect A: {results['effect_sizes']['cohens_d_A']:.3f}")
317 print(f"    Main Effect B: {results['effect_sizes']['cohens_d_B']:.3f}")
318 print(f"    Interaction: {results['effect_sizes']['cohens_d_interaction']:.3f}")
319
320 print("\n95% Confidence Intervals:")
321 print(f"    Main A: [{results['confidence_intervals']['main_A_ci'][0]:.4f}, "
322         f"{results['confidence_intervals']['main_A_ci'][1]:.4f}]")
323 print(f"    Main B: [{results['confidence_intervals']['main_B_ci'][0]:.4f}, "
324         f"{results['confidence_intervals']['main_B_ci'][1]:.4f}]")
325 print(f"    Interaction: [{results['confidence_intervals']['interaction_ci'][0]:.4f}, "
326         f"{results['confidence_intervals']['interaction_ci'][1]:.4f}]")
327
328 print("\nANOVA Table:")
329 print(results['anova_table'])
330
331 # Create visualizations
332 fig = visualize_factorial_experiment(data, results)
333 plt.savefig('factorial_analysis.pdf', bbox_inches='tight')
334 plt.show()

```

8.2.3 Fractional Factorial Designs

When testing many factors, full factorial designs become impractical. A 2^k design requires 2^k experimental conditions.

Problem: Testing 6 factors requires $2^6 = 64$ cells. With 100 users per cell, need 6,400 users.

Solution: Fractional factorial designs test a carefully chosen subset of conditions.

Example 8.2 (Half-Fraction Design). A 2^{5-1} (read: "two to the five minus one") design tests 5 factors using only $2^4 = 16$ conditions instead of 32.

We sacrifice the ability to estimate certain high-order interactions (which are often negligible) to gain efficiency.

Confounding: In fractional designs, some effects cannot be separately estimated (they are "aliased").

For 2^{5-1} with defining relation $I = ABCDE$:

- Main effect A is aliased with 4-way interaction BCDE

- Two-way interaction AB is aliased with CDE

Resolution: Describes which effects are confounded:

- Resolution III: Main effects clear of each other but aliased with 2-way interactions
- Resolution IV: Main effects clear; 2-way interactions aliased with other 2-way interactions
- Resolution V: Main effects and 2-way interactions clear

Practical Recommendation: Use Resolution IV or V designs when possible. Assume high-order interactions are negligible.

8.3 Multi-Armed Bandits

Multi-armed bandits (MAB) dynamically allocate traffic to variants, balancing exploration (learning about variants) and exploitation (assigning users to better-performing variants).

8.3.1 The Exploration-Exploitation Tradeoff

Traditional A/B test: Fixed allocation (e.g., 50/50) throughout experiment.

Multi-armed bandit: Adaptive allocation that shifts traffic toward better variants as data accumulates.

Regret: Cumulative opportunity cost from suboptimal assignments:

$$R(T) = \sum_{t=1}^T (\mu^* - \mu_{A_t}) \quad (8.5)$$

where:

- T : Total number of users
- μ^* : True mean of best variant
- A_t : Variant assigned to user t
- μ_{A_t} : True mean of assigned variant

Goal: Minimize cumulative regret while maintaining valid inference.

8.3.2 ϵ -Greedy Algorithm

Simplest bandit algorithm: with probability ϵ , explore randomly; otherwise, exploit current best.

Algorithm 4 ϵ -Greedy Bandit

```

Initialize:  $N_i = 0, \hat{\mu}_i = 0$  for all variants  $i = 1, \dots, K$ 
for each user  $t = 1, 2, \dots, T$  do
  Generate  $u \sim \text{Uniform}(0, 1)$ 
  if  $u < \epsilon$  then
    Select random variant  $A_t$  (explore)
  else
    Select  $A_t = \arg \max_i \hat{\mu}_i$  (exploit)
  end if
  Observe reward  $r_t$ 
  Update:  $N_{A_t} \leftarrow N_{A_t} + 1$ 
  Update:  $\hat{\mu}_{A_t} \leftarrow \frac{(N_{A_t}-1)\hat{\mu}_{A_t} + r_t}{N_{A_t}}$ 
end for

```

Parameter Selection:

- Fixed $\epsilon \in [0.05, 0.20]$: Simple but continues exploring indefinitely
- Decaying $\epsilon_t = \min(1, \epsilon_0/t)$: Reduces exploration over time

8.3.3 Upper Confidence Bound (UCB) Algorithm

UCB addresses ϵ -greedy's limitation by explicitly balancing exploitation and exploration using confidence intervals.

Algorithm 5 UCB1 Algorithm

```

Initialize:  $N_i = 0, \hat{\mu}_i = 0$  for all variants  $i$ 
Play each variant once to initialize
for each user  $t = K + 1, K + 2, \dots, T$  do
  Select variant:
    
$$A_t = \arg \max_i \left[ \hat{\mu}_i + \sqrt{\frac{2 \ln t}{N_i}} \right]$$

  Observe reward  $r_t$ 
  Update  $N_{A_t}$  and  $\hat{\mu}_{A_t}$ 
end for

```

The UCB term $\sqrt{\frac{2 \ln t}{N_i}}$ represents uncertainty: variants with fewer observations get bonus exploration.

Theorem (Auer et al., 2002): UCB1 achieves logarithmic regret:

$$R(T) \leq 8 \sum_{i: \mu_i < \mu^*} \frac{\ln T}{\Delta_i} + \left(1 + \frac{\pi^2}{3}\right) \sum_{i=1}^K \Delta_i \quad (8.6)$$

where $\Delta_i = \mu^* - \mu_i$ is the gap between optimal and variant i .

8.3.4 Thompson Sampling (Bayesian Bandit)

Thompson sampling maintains posterior distributions over variant success rates and samples from them to make allocation decisions.

For Bernoulli rewards (conversion rates):

Algorithm 6 Thompson Sampling for Bernoulli Bandits

Initialize: $\alpha_i = 1, \beta_i = 1$ for all variants i (Beta prior)

for each user $t = 1, 2, \dots, T$ **do**

for each variant $i = 1, \dots, K$ **do**

 Sample $\theta_i \sim \text{Beta}(\alpha_i, \beta_i)$

end for

 Select variant $A_t = \arg \max_i \theta_i$

 Observe reward $r_t \in \{0, 1\}$

 Update: $\alpha_{A_t} \leftarrow \alpha_{A_t} + r_t$

 Update: $\beta_{A_t} \leftarrow \beta_{A_t} + (1 - r_t)$

end for

Intuition: Variants with higher uncertainty (wider posteriors) have higher probability of being sampled, naturally balancing exploration and exploitation.

Listing 8.2: Comprehensive Bandit Implementations and Comparison

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from typing import List, Tuple
4 from abc import ABC, abstractmethod
5
6 class Bandit(ABC):
7     """Abstract base class for bandit algorithms."""
8
9     def __init__(self, n_arms: int):
10         self.n_arms = n_arms
11         self.counts = np.zeros(n_arms)
12         self.values = np.zeros(n_arms)
13         self.total_reward = 0
14
15     @abstractmethod
16     def select_arm(self, t: int) -> int:
17         """Select an arm to pull."""
18         pass
19
20     def update(self, arm: int, reward: float):
21         """Update statistics after observing reward."""
22         self.counts[arm] += 1
23         self.values[arm] += (reward - self.values[arm]) / self.counts[
24             arm]
25         self.total_reward += reward
26
27 class EpsilonGreedy(Bandit):
28     """Epsilon-greedy bandit."""
29
30     def __init__(self, n_arms: int, epsilon: float = 0.1,
31                 decay: bool = False):
32         super().__init__(n_arms)

```

```

32     self.epsilon = epsilon
33     self.decay = decay
34
35     def select_arm(self, t: int) -> int:
36         epsilon = self.epsilon / t if self.decay else self.epsilon
37
38         if np.random.random() < epsilon:
39             return np.random.randint(0, self.n_arms)
40         else:
41             return np.argmax(self.values)
42
43 class UCB1(Bandit):
44     """UCB1 algorithm."""
45
46     def __init__(self, n_arms: int):
47         super().__init__(n_arms)
48
49     def select_arm(self, t: int) -> int:
50         # Initially pull each arm once
51         for arm in range(self.n_arms):
52             if self.counts[arm] == 0:
53                 return arm
54
55         # Calculate UCB for each arm
56         ucb_values = self.values + np.sqrt(
57             2 * np.log(t) / self.counts
58         )
59         return np.argmax(ucb_values)
60
61 class ThompsonSampling(Bandit):
62     """Thompson Sampling with Beta-Bernoulli model."""
63
64     def __init__(self, n_arms: int):
65         super().__init__(n_arms)
66         self.alpha = np.ones(n_arms) # Successes + 1 (prior)
67         self.beta = np.ones(n_arms)  # Failures + 1 (prior)
68
69     def select_arm(self, t: int) -> int:
70         # Sample from posterior for each arm
71         samples = np.random.beta(self.alpha, self.beta)
72         return np.argmax(samples)
73
74     def update(self, arm: int, reward: float):
75         super().update(arm, reward)
76         self.alpha[arm] += reward
77         self.beta[arm] += (1 - reward)
78
79 def simulate_bandit_experiment(
80     bandit: Bandit,
81     true_rates: List[float],
82     n_trials: int
83 ) -> dict:
84     """
85     Simulate bandit experiment.
86
87     Parameters:
88     -----
89     bandit: Bandit

```

```

90  #####Bandit_algorithm_instance
91  #####true_rates: list
92  #####True_conversion_rate_for_each_arm
93  #####n_trials: int
94  #####Number_of_users
95
96  #####Returns:
97  #####-----
98  #####dict: Simulation_results
99  #####"""
100     n_arms = len(true_rates)
101     true_rates = np.array(true_rates)
102     optimal_rate = np.max(true_rates)
103     optimal_arm = np.argmax(true_rates)
104
105     # Track metrics over time
106     arm_selections = np.zeros((n_trials, n_arms))
107     cumulative_reward = np.zeros(n_trials)
108     cumulative_regret = np.zeros(n_trials)
109     optimal_selections = np.zeros(n_trials)
110
111     for t in range(n_trials):
112         # Select arm
113         arm = bandit.select_arm(t + 1)
114         arm_selections[t, arm] = 1
115
116         # Observe reward
117         reward = np.random.binomial(1, true_rates[arm])
118
119         # Update bandit
120         bandit.update(arm, reward)
121
122         # Track metrics
123         regret = optimal_rate - true_rates[arm]
124
125         cumulative_reward[t] = bandit.total_reward
126         cumulative_regret[t] = (
127             cumulative_regret[t-1] + regret if t > 0 else regret
128         )
129         optimal_selections[t] = (arm == optimal_arm)
130
131     # Cumulative arm selections
132     cumulative_selections = np.cumsum(arm_selections, axis=0)
133
134     return {
135         'arm_selections': arm_selections,
136         'cumulative_selections': cumulative_selections,
137         'final_proportions': cumulative_selections[-1] / n_trials,
138         'cumulative_reward': cumulative_reward,
139         'cumulative_regret': cumulative_regret,
140         'optimal_selection_rate': np.cumsum(optimal_selections) /
141             np.arange(1, n_trials + 1),
142         'final_values': bandit.values
143     }
144
145 def compare_bandit_algorithms(
146     true_rates: List[float],
147     n_trials: int = 10000,

```

```

148     n_simulations: int = 100
149 ):
150     """
151     Compare different bandit algorithms through simulation.
152
153     Parameters:
154     -----
155     true_rates: list
156     True conversion rate for each variant
157     n_trials: int
158     Number of trials per simulation
159     n_simulations: int
160     Number of simulation replications
161     """
162     n_arms = len(true_rates)
163
164     algorithms = {
165         'epsilon-greedy(0.1)': lambda: EpsilonGreedy(n_arms, 0.1),
166         'epsilon-greedy(decay)': lambda: EpsilonGreedy(n_arms, 1.0,
167                                                         decay=True),
168         'UCB1': lambda: UCB1(n_arms),
169         'ThompsonSampling': lambda: ThompsonSampling(n_arms)
170     }
171
172     results = {name: [] for name in algorithms.keys()}
173
174     for sim in range(n_simulations):
175         np.random.seed(sim)
176
177         for name, create_bandit in algorithms.items():
178             bandit = create_bandit()
179             result = simulate_bandit_experiment(bandit, true_rates, n_
180                                                 trials)
181             results[name].append(result)
182
183     # Aggregate results
184     aggregated = {}
185     for name in algorithms.keys():
186         regrets = [r['cumulative_regret'] for r in results[name]]
187         optimal_rates = [r['optimal_selection_rate'] for r in results[
188                         name]]
189
190         aggregated[name] = {
191             'mean_regret': np.mean(regrets, axis=0),
192             'std_regret': np.std(regrets, axis=0),
193             'mean_optimal_rate': np.mean(optimal_rates, axis=0)
194         }
195
196     # Visualization
197     fig, axes = plt.subplots(2, 2, figsize=(14, 10))
198
199     # 1. Cumulative regret
200     ax = axes[0, 0]
201     for name, data in aggregated.items():
202         ax.plot(data['mean_regret'], label=name, linewidth=2)
203         ax.fill_between(
204             range(n_trials),
205             data['mean_regret'] - data['std_regret'],

```

```

204         data['mean_regret'] + data['std_regret'],
205         alpha=0.2
206     )
207     ax.set_xlabel('Trial')
208     ax.set_ylabel('Cumulative Regret')
209     ax.set_title('Cumulative Regret Over Time')
210     ax.legend()
211     ax.grid(alpha=0.3)
212
213     # 2. Optimal arm selection rate
214     ax = axes[0, 1]
215     for name, data in aggregated.items():
216         ax.plot(data['mean_optimal_rate'], label=name, linewidth=2)
217     ax.set_xlabel('Trial')
218     ax.set_ylabel('Proportion Selecting Optimal Arm')
219     ax.set_title('Convergence to Optimal Arm')
220     ax.legend()
221     ax.grid(alpha=0.3)
222     ax.set_ylim([0, 1])
223
224     # 3. Final regret comparison
225     ax = axes[1, 0]
226     final_regrets = [data['mean_regret'][-1] for data in aggregated.
227                     values()]
228     final_stds = [data['std_regret'][-1] for data in aggregated.values
229                  ()]
230
231     x_pos = np.arange(len(algorithms))
232     ax.bar(x_pos, final_regrets, yerr=final_stds, capsize=5, alpha=0.7)
233     ax.set_xticks(x_pos)
234     ax.set_xticklabels(algorithms.keys(), rotation=15, ha='right')
235     ax.set_ylabel('Final Cumulative Regret')
236     ax.set_title(f'Final Regret after {n_trials} Trials')
237     ax.grid(axis='y', alpha=0.3)
238
239     # 4. Arm selection distribution (Thompson Sampling)
240     ax = axes[1, 1]
241     ts_results = results['Thompson Sampling']
242     final_props = np.array([r['final_proportions'] for r in ts_results
243                           ])
244     mean_props = np.mean(final_props, axis=0)
245
246     x_pos = np.arange(n_arms)
247     ax.bar(x_pos, mean_props, alpha=0.7)
248     ax.axhline(y=1/n_arms, color='red', linestyle='--',
249               label='Equal allocation')
250     ax.set_xticks(x_pos)
251     ax.set_xticklabels([f'Arm_{i}\n({rate:.1%})'
252                        for i, rate in enumerate(true_rates)])
253     ax.set_ylabel('Proportion of Selections')
254     ax.set_title('Thompson Sampling: Final Allocation')
255     ax.legend()
256     ax.grid(axis='y', alpha=0.3)
257
258     plt.tight_layout()
259     return fig, aggregated
260
261 # Example comparison

```

```

259 true_rates = [0.08, 0.10, 0.12]
260 fig, results = compare_bandit_algorithms(
261     true_rates,
262     n_trials=10000,
263     n_simulations=100
264 )
265
266 print("=== Bandit Algorithm Comparison ===")
267 print(f"True rates: {true_rates}")
268 print(f"Optimal arm: Arm {np.argmax(true_rates)} ({max(true_rates)
269     :.1%})\n")
270
271 for name, data in results.items():
272     print(f"{name}:")
273     print(f"Final regret: {data['mean_regret'][-1]:.1f}")
274     print(f"Final std regret: {data['std_regret'][-1]:.1f}")
275     print(f"Final optimal rate: {data['mean_optimal_rate'][-1]:.3f}\n")
276
277 plt.savefig('bandit_comparison.pdf', bbox_inches='tight')
278 plt.show()

```

8.3.5 When to Use Bandits vs. Traditional A/B Tests

Use bandits when:

- **Optimizing during the test is valuable:** E.g., content recommendation, ad selection
- **Cost of poor experiences is high:** Want to minimize exposure to inferior variants
- **Rapid learning is possible:** Quick feedback enables adaptation
- **Continuous optimization desired:** Not a one-time decision
- **Many variants:** Efficient at eliminating poor options

Use traditional A/B when:

- **Clean inference is critical:** Need p-values, confidence intervals for decision making
- **Regulatory requirements:** Medical devices, financial products
- **Long-term effects matter:** User habituation, retention effects
- **Stakeholder communication:** Fixed allocation easier to explain
- **Implementation is permanent:** Making a one-time product decision

Hybrid Approach: Use bandits for traffic allocation during experiment, then analyze using causal inference methods (e.g., inverse propensity weighting) for final decision.

8.4 Sequential Testing

Sequential testing allows stopping experiments early when sufficient evidence accumulates, potentially saving time and resources.

8.4.1 The Peeking Problem

Naive peeking: Continuously checking p-values and stopping when $p < 0.05$ inflates Type I error far above 5%.

Example 8.3 (Inflation from Peeking). Simulation with no true effect (H_0 true):

- Fixed sample (n=1000): Type I error = 5.0%
- Check every 100 users, stop if $p < 0.05$: Type I error = 29%
- Check every 50 users: Type I error = 38%

Problem: Multiple testing without correction.

Solution: Group sequential methods with adjusted boundaries.

8.4.2 Group Sequential Testing

Pre-specify:

- Information fractions: e.g., analyze at 25%, 50%, 75%, 100% of planned sample
- Spending function: how to allocate Type I error across looks
- Stopping boundaries: thresholds for stopping

O'Brien-Fleming Boundaries

Conservative early stopping; save most α for final analysis.

For K equally spaced looks at information fractions $1/K, 2/K, \dots, K/K$:

$$c_k = c \sqrt{\frac{K}{k}} \quad (8.7)$$

where $c \approx 2$ for $\alpha = 0.05$ (determined numerically to control FWER).

Early boundaries are stringent:

- At 25% information: $c_1 = c\sqrt{4} = 2c \approx 4$ (need $|Z| > 4$)
- At 50% information: $c_2 = c\sqrt{2} \approx 2.83$
- At 100% information: $c_K = c \approx 2$

Advantage: Preserves power well; if don't stop early, final test nearly as powerful as fixed design.

Pocock Boundaries

Equal thresholds at all looks: $c_1 = c_2 = \dots = c_K$.

For $\alpha = 0.05$ and $K = 5$ looks: $c \approx 2.41$.

More aggressive early stopping than O'Brien-Fleming.

Disadvantage: If no early stop, final test less powerful due to α spent early.

Listing 8.3: Group Sequential Testing Implementation

```

1 import numpy as np
2 from scipy import stats
3 from typing import List, Tuple
4 import matplotlib.pyplot as plt
5
6 class GroupSequentialTest:
7     """
8     Implement group sequential testing with O'Brien-Fleming or Pocock
9     boundaries.
10    """
11
12    def __init__(self,
13                 max_n: int,
14                 n_looks: int,
15                 alpha: float = 0.05,
16                 beta: float = 0.20,
17                 boundary_type: str = 'obf'):
18        """
19        Parameters:
20        -----
21        max_n: int
22            Maximum sample size per group
23        n_looks: int
24            Number of interim analyses (including final)
25        alpha: float
26            Type I error rate
27        beta: float
28            Type II error rate (1 - power)
29        boundary_type: str
30            'obf' (O'Brien-Fleming) or 'pocock'
31        """
32        self.max_n = max_n
33        self.n_looks = n_looks
34        self.alpha = alpha
35        self.beta = beta
36        self.boundary_type = boundary_type
37
38        # Information fractions
39        self.info_fracs = np.array([
40            (k+1) / n_looks for k in range(n_looks)
41        ])
42
43        # Calculate boundaries
44        self.boundaries = self._calculate_boundaries()
45
46    def _calculate_boundaries(self) -> np.ndarray:
47        """Calculate stopping boundaries."""
48        if self.boundary_type == 'obf':
49            # O'Brien-Fleming: conservative early, lenient late

```

```

49         # Use Lan-DeMets spending function approximation
50         z_alpha = stats.norm.ppf(1 - self.alpha/2)
51         boundaries = z_alpha * np.sqrt(1 / self.info_frac)
52
53     elif self.boundary_type == 'pocock':
54         # Pocock: equal boundaries
55         # Approximate for 2-sided test
56         if self.n_looks == 1:
57             c = stats.norm.ppf(1 - self.alpha/2)
58         elif self.n_looks == 2:
59             c = 2.178
60         elif self.n_looks == 3:
61             c = 2.289
62         elif self.n_looks == 4:
63             c = 2.361
64         elif self.n_looks == 5:
65             c = 2.413
66         else:
67             # Approximation for larger K
68             c = stats.norm.ppf(1 - self.alpha/(2*self.n_looks))
69
70         boundaries = np.ones(self.n_looks) * c
71     else:
72         raise ValueError("boundary_type must be 'obf' or 'pocock'")
73
74     return boundaries
75
76     def analyze_look(self,
77                     n_current: int,
78                     control_conversions: int,
79                     treatment_conversions: int,
80                     control_n: int,
81                     treatment_n: int) -> dict:
82         """
83         Analyze data at a specific look.
84
85         Returns:
86         -----
87         dict with keys:
88         'z_stat': Test statistic
89         'boundary': Critical value for this look
90         'stop_efficacy': Stop for efficacy
91         'stop_futility': Stop for futility
92         'continue': Continue to next look
93         """
94         # Calculate test statistic
95         p1 = control_conversions / control_n
96         p2 = treatment_conversions / treatment_n
97         p_pooled = (control_conversions + treatment_conversions) / (
98             control_n + treatment_n
99         )
100
101         se = np.sqrt(p_pooled * (1 - p_pooled) * (1/control_n + 1/
102             treatment_n))
103         z_stat = (p2 - p1) / se
104
105         # Determine which look we're at
106         look_idx = int(np.round(n_current / self.max_n * self.n_looks))

```

```

- 1
look_idx = np.clip(look_idx, 0, self.n_looks - 1)

boundary = self.boundaries[look_idx]

# Stopping decisions
stop_efficacy = abs(z_stat) >= boundary

# Simple futility boundary (can be more sophisticated)
futility_boundary = 0.5
stop_futility = abs(z_stat) < futility_boundary and look_idx >=
    1

return {
    'look': look_idx + 1,
    'n_current': n_current,
    'z_stat': z_stat,
    'boundary': boundary,
    'stop_efficacy': stop_efficacy,
    'stop_futility': stop_futility,
    'continue': not (stop_efficacy or stop_futility),
    'p1': p1,
    'p2': p2
}

def visualize_boundaries(self):
    """Visualize stopping boundaries."""
    fig, ax = plt.subplots(figsize=(10, 6))

    # Plot boundaries
    looks = np.arange(1, self.n_looks + 1)
    ax.plot(looks, self.boundaries, 'ro-', linewidth=2,
            markersize=10, label='Efficacy boundary')
    ax.plot(looks, -self.boundaries, 'ro-', linewidth=2, markersize
            =10)

    # Add futility boundary
    ax.axhline(y=0.5, color='orange', linestyle='--',
               label='Futility boundary')
    ax.axhline(y=-0.5, color='orange', linestyle='--')
    ax.axhline(y=0, color='gray', linestyle='-', alpha=0.3)

    # Styling
    ax.set_xlabel('Analysis Number', fontsize=12)
    ax.set_ylabel('Z-statistic Threshold', fontsize=12)
    ax.set_title(f'{self.boundary_type.upper()} Boundaries'
                 f'(K={self.n_looks}, alpha={self.alpha})', fontsize
                 =14)
    ax.set_xticks(looks)
    ax.grid(alpha=0.3)
    ax.legend(fontsize=11)

    # Add information fractions
    for i, (look, frac) in enumerate(zip(looks, self.info_fracs)):
        ax.text(look, self.boundaries[i] + 0.3,
                f'{frac:.1%}', ha='center', fontsize=10)

    plt.tight_layout()

```

```

160         return fig
161
162     def simulate_sequential_experiment(
163         true_effect: float,
164         baseline_rate: float,
165         max_n_per_group: int,
166         n_looks: int,
167         boundary_type: str,
168         n_simulations: int = 1000
169     ) -> dict:
170         """
171         Simulate group sequential experiments.
172
173         Returns type I error, power, and average sample size.
174         """
175         gst = GroupSequentialTest(
176             max_n=max_n_per_group,
177             n_looks=n_looks,
178             boundary_type=boundary_type
179         )
180
181         stopped_for_efficacy = 0
182         stopped_for_futility = 0
183         total_n = []
184
185         p1 = baseline_rate
186         p2 = baseline_rate + true_effect
187
188         for sim in range(n_simulations):
189             np.random.seed(sim)
190
191             stopped = False
192
193             for look in range(1, n_looks + 1):
194                 n_current = int(max_n_per_group * look / n_looks)
195
196                 # Generate data
197                 control = np.random.binomial(1, p1, n_current)
198                 treatment = np.random.binomial(1, p2, n_current)
199
200                 result = gst.analyze_look(
201                     n_current=n_current,
202                     control_conversions=control.sum(),
203                     treatment_conversions=treatment.sum(),
204                     control_n=n_current,
205                     treatment_n=n_current
206                 )
207
208                 if result['stop_efficacy']:
209                     stopped_for_efficacy += 1
210                     total_n.append(n_current * 2)
211                     stopped = True
212                     break
213                 elif result['stop_futility']:
214                     stopped_for_futility += 1
215                     total_n.append(n_current * 2)
216                     stopped = True
217                     break

```

```

218         if not stopped:
219             total_n.append(max_n_per_group * 2)
220
221     return {
222         'type_i_error' if true_effect == 0 else 'power':
223             stopped_for_efficacy / n_simulations,
224         'stopped_for_futility_rate': stopped_for_futility / n_
225             simulations,
226         'average_sample_size': np.mean(total_n),
227         'median_sample_size': np.median(total_n),
228         'sample_size_reduction': 1 - np.mean(total_n) / (max_n_per_
229             group * 2)
230     }
231
232 # Example usage
233 print("=== Group Sequential Testing ===\n")
234
235 # Visualize boundaries
236 gst_obf = GroupSequentialTest(max_n=5000, n_looks=5, boundary_type='obf
237 ')
238 gst_pocock = GroupSequentialTest(max_n=5000, n_looks=5, boundary_type='
239 pocock')
240
241 print("O'Brien-Fleming Boundaries:")
242 print(gst_obf.boundaries)
243 print("\nPocock Boundaries:")
244 print(gst_pocock.boundaries)
245
246 fig_obf = gst_obf.visualize_boundaries()
247 plt.savefig('obf_boundaries.pdf')
248
249 fig_pocock = gst_pocock.visualize_boundaries()
250 plt.savefig('pocock_boundaries.pdf')
251
252 # Simulate to check Type I error control
253 print("\n=== Simulation Results (1000 replications) ===")
254 print("\nUnder null (true_effect=0):")
255
256 for boundary_type in ['obf', 'pocock']:
257     result = simulate_sequential_experiment(
258         true_effect=0.0,
259         baseline_rate=0.10,
260         max_n_per_group=5000,
261         n_looks=5,
262         boundary_type=boundary_type,
263         n_simulations=1000
264     )
265
266     print(f"\n{boundary_type.upper()}:")
267     print(f"Type I error: {result['type_i_error']:.3f} "
268           f"(target: 0.05)")
269     print(f"Avg sample size: {result['average_sample_size']:.0f} "
270           f"({result['sample_size_reduction']:.1%} reduction)")
271
272 print("\nUnder alternative (true_effect=0.02):")
273
274 for boundary_type in ['obf', 'pocock']:

```

```

272     result = simulate_sequential_experiment(
273         true_effect=0.02,
274         baseline_rate=0.10,
275         max_n_per_group=5000,
276         n_looks=5,
277         boundary_type=boundary_type,
278         n_simulations=1000
279     )
280
281     print(f"\n{boundary_type.upper()}:")
282     print(f"    Power: {result['power']:.3f}")
283     print(f"    Avg sample size: {result['average_sample_size']:.0f}")
284     print(f"    ({result['sample_size_reduction']:.1%} reduction)")

```

8.5 Network Effects and Interference

When users interact, standard A/B testing assumptions break down. One user's treatment can affect another user's outcome.

8.5.1 Types of Interference

1. **Direct Network Effects:** User A's treatment directly affects User B

Example: Social network feature where treated users send more friend requests, affecting control users

2. **Marketplace Effects:** Two-sided platforms with supply-demand dynamics

Example: Ride-sharing app where increasing driver supply in treatment affects wait times for all users

3. **Competition Effects:** Users compete for limited resources

Example: Bidding platform where treated users' behavior affects prices for everyone

Consequence: Treatment effect estimates are biased. May underestimate (positive spillover) or overestimate (negative spillover) true effects.

8.5.2 SUTVA Violation

The Stable Unit Treatment Value Assumption (SUTVA) requires:

1. No interference: User i 's outcome depends only on their own treatment
2. Consistency: Treatment is well-defined and consistently applied

Network effects violate SUTVA.

8.5.3 Mitigation Strategies

1. Cluster Randomization

Randomize groups of connected users together.

Example: For social network experiment, randomize entire friend groups or geographic regions.

Trade-off: Reduces interference but decreases effective sample size.

2. Ego Network Randomization

Randomize both focal user and their immediate neighbors to same treatment.

Advantage: Reduces spillover for randomized pairs.

Limitation: Difficult with overlapping networks.

3. Temporal Switchbacks

Alternate entire system between control and treatment over time periods.

Example 8.4 (Switchback Design). Week 1: Monday (Control), Tuesday (Treatment), Wednesday (Control), ...

Week 2: Monday (Treatment), Tuesday (Control), ...

Compare outcomes between control and treatment periods.

Analysis: Must account for temporal correlation and carryover effects.

8.6 Switchback Designs

Switchback (crossover) designs are particularly useful for marketplace and platform experiments.

8.6.1 Structure

Alternate between control and treatment conditions over time, typically:

- Geographic units (cities, regions)
- Time periods (hours, days, weeks)

City	Week 1	Week 2	Week 3	Week 4
New York	Control	Treatment	Control	Treatment
San Francisco	Treatment	Control	Treatment	Control
Chicago	Control	Treatment	Control	Treatment
Boston	Treatment	Control	Treatment	Control

8.6.2 Analysis

Use mixed effects model or difference-in-differences:

$$Y_{it} = \alpha_i + \beta T_{it} + \gamma_t + \epsilon_{it} \quad (8.8)$$

where:

- α_i : Unit fixed effect
- T_{it} : Treatment indicator
- γ_t : Time fixed effect
- β : Treatment effect (estimand)

Key Assumption: No carryover effects, or sufficient washout period between switches.

8.6.3 Practical Considerations

- **Switch frequency:** Balance temporal correlation vs. statistical power
- **Washout periods:** Allow for treatment effects to dissipate
- **Temporal trends:** Control for day-of-week effects, seasonality
- **Synchronization:** Coordinate switches across units

8.7 Summary

Advanced experimental designs extend A/B testing capabilities:

- **Factorial designs** efficiently test multiple factors and detect interactions
- **Multi-armed bandits** optimize allocation during experiments, minimizing regret
- **Sequential testing** allows early stopping while controlling error rates
- **Network-aware designs** handle interference when users interact
- **Switchback designs** address marketplace and platform challenges

Each design involves trade-offs between efficiency, complexity, and validity. Choose designs based on:

- Business objectives (optimization vs. inference)
- User interaction structure (independent vs. networked)
- Resource constraints (sample size, time)
- Stakeholder requirements (interpretability, regulatory)

8.8 Further Reading

- ?: Comprehensive treatment of factorial designs
- ?: Theoretical foundations of multi-armed bandits
- ?: Group sequential methods for clinical trials
- ?: Design and analysis with network effects
- ?: Cluster randomized trials

8.9 Exercises

8.9.1 Conceptual Questions

1. Explain when a 2^3 full factorial design would be preferable to three separate A/B tests, and when separate tests might be better.
2. Describe a scenario where Thompson Sampling would be preferable to a traditional A/B test, and explain your reasoning.
3. Why does "peeking" at A/B test results inflate Type I error? How do group sequential methods address this?
4. Give an example of a product feature that would violate SUTVA due to network effects. Propose an experimental design to test this feature.

8.9.2 Calculation Exercises

1. For a 2^4 factorial design with 200 observations per cell:
 - (a) How many total observations are needed?
 - (b) How many parameters can be estimated (main effects + interactions)?
 - (c) What would be the effective sample size for estimating each main effect?
2. Calculate the UCB1 selection criterion for three arms with:
 - Arm 1: $\hat{\mu}_1 = 0.10$, $N_1 = 100$
 - Arm 2: $\hat{\mu}_2 = 0.12$, $N_2 = 50$
 - Arm 3: $\hat{\mu}_3 = 0.11$, $N_3 = 150$

At time $t = 301$, which arm should be selected?

3. For O'Brien-Fleming boundaries with $K = 4$ looks and $\alpha = 0.05$, if the first look yields $Z = 2.5$, should you stop? What if the second look yields $Z = 2.8$?

8.9.3 Programming Exercises

1. Implement a complete factorial experiment simulator:
 - Generate data for 2^k designs with specified effects
 - Analyze using ANOVA
 - Create interaction plots
 - Calculate power for detecting interactions
2. Compare ϵ -greedy, UCB1, and Thompson Sampling:
 - Simulate 10 arms with various true rates
 - Run 10,000 trials for each algorithm
 - Plot cumulative regret

- Analyze convergence properties

3. Implement group sequential testing:

- Create O'Brien-Fleming and Pocock boundaries
- Simulate experiments under null and alternative
- Verify Type I error control
- Measure average sample size savings

8.9.4 Applied Exercises

1. **E-commerce Factorial Design:** Design a 2^3 experiment testing:

- Product image size (small vs. large)
- Price display (with vs. without discount)
- Reviews prominence (top vs. bottom)

Specify sample size, analysis plan, and expected interactions.

2. **Content Recommendation Bandit:** A news app wants to optimize article headlines dynamically:

- 5 headline variants per article
- Goal: maximize click-through rate
- New headlines added continuously

Recommend a bandit algorithm and justify your choice.

3. **Marketplace Experiment:** A ride-sharing platform wants to test dynamic pricing:

- Treatment affects all riders and drivers in treated regions
- Spillover between nearby regions
- Time-varying demand

Design an experiment accounting for interference. Specify randomization unit, analysis method, and duration.

4. **Sequential Testing Strategy:** Your company typically runs A/B tests for 2 weeks with 10,000 users per variant. Design a sequential testing procedure that could reduce average duration by 30% while maintaining $\alpha = 0.05$ and power = 80%.

8.9.5 Advanced Projects

1. **Bandit with Delayed Feedback:** Extend Thompson Sampling to handle rewards observed with delay (e.g., 7-day retention). Simulate and compare to standard Thompson Sampling.
2. **Network Interference Simulation:** Create a social network graph (1000 nodes) and simulate an experiment with spillover effects. Compare:
 - Individual randomization (biased)
 - Cluster randomization (less power)
 - Ego-network randomization

Measure bias and power for each approach.

3. **Adaptive Factorial Design:** Implement a factorial design that drops unpromising factor levels based on interim data, while controlling error rates.

Chapter 9

Technical Implementation

9.1 Learning Objectives

After completing this chapter, readers will be able to:

- Design robust experimentation infrastructure
- Implement randomization and assignment systems
- Build reliable logging and data collection pipelines
- Integrate A/B testing into applications
- Debug common implementation issues

9.2 Experimentation Infrastructure

A robust A/B testing system requires careful technical implementation.

9.2.1 Core Components

1. **Randomization Service:** Assigns users to variants
2. **Configuration Management:** Stores experiment definitions
3. **Logging System:** Records assignments and events
4. **Analysis Pipeline:** Computes metrics and statistical tests
5. **Monitoring Dashboard:** Tracks experiment health

9.2.2 Architecture Patterns

Client-side randomization:

- Randomization happens in browser/app
- Lightweight, no server dependency
- Risk of manipulation

Server-side randomization:

- Centralized assignment service
- More secure and consistent
- Requires infrastructure

Hybrid approach:

- Server assigns, client implements
- Balance of benefits

9.3 Randomization Implementation

9.3.1 Hash-Based Randomization

Use deterministic hashing for consistent assignments across sessions.

Listing 9.1: Hash-Based User Assignment

```

1 import hashlib
2
3 def assign_variant(user_id, experiment_id, num_variants=2,
4                   traffic_allocation=1.0, salt=''):
5     """
6     Assign user to variant using deterministic hashing.
7
8     Parameters:
9     -----
10    user_id: str
11            Unique user identifier
12    experiment_id: str
13            Unique experiment identifier
14    num_variants: int
15            Number of variants
16    traffic_allocation: float
17            Fraction of users to include (0.0-1.0)
18    salt: str
19            Random salt for security
20
21    Returns:
22    -----
23    variant: int or None
24            Assigned variant (0 to num_variants-1) or None if not allocated
25    """
26    # Create hash input
27    hash_input = f"{user_id}:{experiment_id}:{salt}".encode('utf-8')
28
29    # Generate hash
30    hash_output = hashlib.md5(hash_input).hexdigest()
31
32    # Convert to integer and normalize to [0, 1]
33    hash_int = int(hash_output, 16)
34    hash_fraction = (hash_int % 10000) / 10000.0
35

```

```

36     # Check traffic allocation
37     if hash_fraction >= traffic_allocation:
38         return None # User not in experiment
39
40     # Assign to variant
41     variant_fraction = hash_fraction / traffic_allocation
42     variant = int(variant_fraction * num_variants)
43
44     return variant
45
46 # Example usage
47 user_id = "user_12345"
48 experiment_id = "button_color_test"
49
50 # Consistent assignment
51 for _ in range(3):
52     variant = assign_variant(user_id, experiment_id, num_variants=2)
53     print(f"User_{user_id}_assigned_to_variant:{variant}")
54
55 # Different experiments give different assignments
56 for exp_id in ["exp_1", "exp_2", "exp_3"]:
57     variant = assign_variant(user_id, exp_id)
58     print(f"Experiment_{exp_id}:_variant_{variant}")

```

Key properties:

- Deterministic: same user always gets same variant
- Independent: different experiments independent
- Uniform: variants equally likely
- Secure: hard to game with salt

9.3.2 Handling Edge Cases

1. **New vs. returning users:** Ensure consistent assignment
2. **Logged out users:** Use cookies with careful handling
3. **Multi-device users:** Decision: device-level or user-level?
4. **Bots and crawlers:** Filter before assignment
5. **Internal traffic:** Exclude or separate analysis

9.4 Configuration Management

Store experiment definitions in a centralized system.

Listing 9.2: Experiment Configuration Schema

```

1 import json
2 from datetime import datetime
3
4 class ExperimentConfig:

```

```

5      """Experiment_configuration_manager."""
6
7      def __init__(self, config_path):
8          with open(config_path, 'r') as f:
9              self.config = json.load(f)
10
11      def get_experiment(self, experiment_id):
12          """Retrieve_experiment_configuration."""
13          return self.config.get('experiments', {}).get(experiment_id)
14
15      def is_active(self, experiment_id):
16          """Check_if_experiment_is_currently_active."""
17          exp = self.get_experiment(experiment_id)
18          if not exp:
19              return False
20
21          now = datetime.now()
22          start = datetime.fromisoformat(exp.get('start_date'))
23          end = datetime.fromisoformat(exp.get('end_date'))
24
25          return start <= now <= end and exp.get('status') == 'active'
26
27      # Example configuration file (JSON)
28      example_config = {
29          "experiments": {
30              "checkout_button_test": {
31                  "name": "Checkout_Button_Color_Test",
32                  "description": "Testing_red_vs_blue_checkout_button",
33                  "owner": "product_team@example.com",
34                  "start_date": "2025-01-01T00:00:00",
35                  "end_date": "2025-01-31T23:59:59",
36                  "status": "active",
37                  "variants": [
38                      {"id": 0, "name": "control", "allocation": 0.5},
39                      {"id": 1, "name": "treatment", "allocation": 0.5}
40                  ],
41                  "metrics": {
42                      "primary": "conversion_rate",
43                      "secondary": ["revenue_per_user", "bounce_rate"],
44                      "guardrails": ["page_load_time", "error_rate"]
45                  },
46                  "targeting": {
47                      "countries": ["US", "CA", "UK"],
48                      "platforms": ["web", "mobile_web"]
49                  }
50              }
51          }
52      }

```

9.5 Logging and Data Collection

9.5.1 Event Schema

Design a comprehensive event schema capturing all necessary data.

Listing 9.3: Event Logging System

```

1 import uuid
2 from datetime import datetime
3 import json
4
5 class ExperimentLogger:
6     """Log experiment assignments and events."""
7
8     def __init__(self, storage_backend):
9         self.storage = storage_backend
10
11     def log_assignment(self, user_id, experiment_id, variant,
12                       timestamp=None, metadata=None):
13         """Log user assignment to experiment variant."""
14         event = {
15             'event_id': str(uuid.uuid4()),
16             'event_type': 'assignment',
17             'timestamp': timestamp or datetime.utcnow().isoformat(),
18             'user_id': user_id,
19             'experiment_id': experiment_id,
20             'variant': variant,
21             'metadata': metadata or {}
22         }
23
24         self.storage.write(event)
25         return event
26
27     def log_conversion(self, user_id, experiment_id, conversion_type,
28                       value=None, timestamp=None, metadata=None):
29         """Log conversion event."""
30         event = {
31             'event_id': str(uuid.uuid4()),
32             'event_type': 'conversion',
33             'timestamp': timestamp or datetime.utcnow().isoformat(),
34             'user_id': user_id,
35             'experiment_id': experiment_id,
36             'conversion_type': conversion_type,
37             'value': value,
38             'metadata': metadata or {}
39         }
40
41         self.storage.write(event)
42         return event
43
44 # Example: In-memory storage for demonstration
45 class InMemoryStorage:
46     def __init__(self):
47         self.events = []
48
49     def write(self, event):
50         self.events.append(event)
51
52     def query(self, event_type=None, experiment_id=None):
53         results = self.events
54         if event_type:
55             results = [e for e in results if e['event_type'] == event_type]
56         if experiment_id:

```

```

57         results = [e for e in results
58                     if e.get('experiment_id') == experiment_id]
59     return results
60
61 # Usage
62 storage = InMemoryStorage()
63 logger = ExperimentLogger(storage)
64
65 # Log assignment
66 logger.log_assignment('user_001', 'exp_123', variant=1,
67                      metadata={'device': 'mobile'})
68
69 # Log conversion
70 logger.log_conversion('user_001', 'exp_123', 'purchase',
71                      value=49.99, metadata={'product_id': 'SKU_789'})
72
73 # Query events
74 assignments = storage.query(event_type='assignment',
75                             experiment_id='exp_123')
76 print(f"Found {len(assignments)} assignment events")

```

9.5.2 Data Quality Checks

Implement automated checks:

1. **Sample Ratio Mismatch (SRM):** Verify allocation matches expectations
2. **Assignment Consistency:** Users don't switch variants
3. **Timestamp Validity:** Events in reasonable order
4. **Data Completeness:** No missing critical fields
5. **Metric Sanity:** Values within expected ranges

Listing 9.4: Data Quality Validation

```

1 def validate_experiment_data(assignments_df, events_df):
2     """
3     Run data quality checks on experiment data.
4
5     Returns:
6     -----
7     dict: Validation results with issues flagged
8     """
9     issues = []
10
11     # Check 1: Sample Ratio Mismatch
12     variant_counts = assignments_df['variant'].value_counts()
13     expected_ratio = 1.0 / len(variant_counts)
14     actual_ratios = variant_counts / len(assignments_df)
15
16     from scipy.stats import chisquare
17     expected = [len(assignments_df) * expected_ratio] * len(variant_counts)
18     chi2, p_value = chisquare(variant_counts.values, expected)

```

```

19
20     if p_value < 0.01:
21         issues.append({
22             'type': 'SRM',
23             'severity': 'high',
24             'message': f'Sample_ratio_mismatch_detected_(p={p_value:.4f}',
25             'details': dict(variant_counts)
26         })
27
28     # Check 2: Multiple assignments
29     multi_assigned = assignments_df.groupby('user_id')['variant'].
30         unique()
31     multi_assigned_users = (multi_assigned > 1).sum()
32
33     if multi_assigned_users > 0:
34         issues.append({
35             'type': 'consistency',
36             'severity': 'high',
37             'message': f'{multi_assigned_users}_users_assigned_to_
38                 multiple_variants',
39         })
40
41     # Check 3: Events before assignment
42     merged = events_df.merge(assignments_df, on='user_id', how='left')
43     merged['assignment_timestamp'] = pd.to_datetime(merged['assignment_
44         timestamp'])
45     merged['event_timestamp'] = pd.to_datetime(merged['event_timestamp'
46         ])
47
48     early_events = (merged['event_timestamp'] <
49         merged['assignment_timestamp']).sum()
50
51     if early_events > 0:
52         issues.append({
53             'type': 'timing',
54             'severity': 'medium',
55             'message': f'{early_events}_events_occurred_before_
56                 assignment'
57         })
58
59     return {
60         'valid': len(issues) == 0,
61         'issues': issues,
62         'checks_run': ['SRM', 'consistency', 'timing']
63     }

```

9.6 Integration Patterns

9.6.1 Feature Flags Integration

Combine experimentation with feature flag systems.

Listing 9.5: Feature Flag with A/B Testing

```

1 class FeatureFlagService:

```

```

2      """Combined feature flags and A/B testing."""
3
4      def __init__(self, experiment_service):
5          self.experiment_service = experiment_service
6          self.permanent_flags = {}
7
8      def is_enabled(self, flag_name, user_id, default=False):
9          """
10         """
11         """Check if feature is enabled for user.
12         Supports:
13         """
14         """Boolean flags (on/off for all users)
15         """
16         """Percentage rollouts
17         """
18         """A/B test variants
19         """
20         """
21         # Check permanent flags first
22         if flag_name in self.permanent_flags:
23             return self.permanent_flags[flag_name]
24
25         # Check if it's an A/B test
26         if self.experiment_service.is_experiment(flag_name):
27             variant = self.experiment_service.get_variant(flag_name,
28                                                         user_id)
29             return variant == 1 # Treat variant 1 as "enabled"
30
31         return default
32
33     def get_variant(self, experiment_name, user_id):
34         """Get experiment variant for user."""
35         return self.experiment_service.get_variant(experiment_name,
36                                                     user_id)

```

9.6.2 Metrics Collection

Integrate with existing analytics systems.

1. **Passive tracking:** Automatic event capture
2. **Active instrumentation:** Explicit metric logging
3. **Computed metrics:** Derived from raw events
4. **Real-time aggregation:** Pre-compute for dashboards

9.7 Common Implementation Pitfalls

9.7.1 Assignment Bugs

- **Issue:** Inconsistent assignments across sessions
- **Cause:** Not using stable user IDs
- **Fix:** Hash-based deterministic assignment

9.7.2 Survivorship Bias

- **Issue:** Only analyzing users who complete flows
- **Cause:** Filtering out non-converters
- **Fix:** Intent-to-treat analysis

9.7.3 Logging Failures

- **Issue:** Missing assignment or event data
- **Cause:** Client-side logging dropped
- **Fix:** Server-side logging, retry logic, monitoring

9.8 Summary

Robust implementation is critical for trustworthy A/B testing:

- Use deterministic hash-based randomization
- Maintain centralized configuration management
- Implement comprehensive logging with quality checks
- Integrate with feature flags and analytics
- Monitor for common implementation issues

Proper infrastructure enables scaling experimentation across the organization.

9.9 Further Reading

- ?: Implementation best practices
- ?: Experiment system architecture
- Open source platforms: Optimizely, GrowthBook, Unleash

9.10 Exercises

1. Implement complete hash-based randomization with traffic allocation and salting
2. Build logging system with SRM detection and automated alerts
3. Design experiment configuration schema supporting complex targeting
4. Create feature flag service integrating A/B tests and gradual rollouts

Chapter 10

Data Analysis and Result Interpretation

10.1 Learning Objectives

After completing this chapter, readers will be able to:

- Conduct end-to-end analysis of A/B test results
- Handle missing data and outliers appropriately
- Perform variance reduction techniques
- Interpret results for stakeholders
- Make sound business decisions from statistical evidence

10.2 Analysis Workflow

A systematic approach ensures thorough and accurate analysis.

10.2.1 Pre-Analysis Checks

Before analyzing results, verify data quality:

1. **Sample Ratio Mismatch:** Check allocation matches design
2. **Assignment Consistency:** Verify users don't switch variants
3. **Novelty Effects:** Examine temporal patterns
4. **External Validity Threats:** Check for anomalous events
5. **Data Completeness:** Assess missing data patterns

Listing 10.1: Pre-Analysis Data Quality Checks

```

1 import pandas as pd
2 import numpy as np
3 from scipy.stats import chisquare
4 import matplotlib.pyplot as plt
5
6 def preanalysis_checks(data):
7     """
8     Run comprehensive pre-analysis checks.
9
10    Parameters:
11    -----
12    data: DataFrame
13    Must contain: user_id, variant, timestamp, conversion
14
15    Returns:
16    -----
17    dict: Check results and warnings
18    """
19    warnings = []
20
21    # 1. Sample Ratio Mismatch
22    variant_counts = data['variant'].value_counts()
23    expected = [len(data) / len(variant_counts)] * len(variant_counts)
24    chi2, srm_p_value = chisquare(variant_counts.values, expected)
25
26    if srm_p_value < 0.01:
27        warnings.append(f"SRM detected: p-value={srm_p_value:.4f}")
28
29    # 2. Temporal patterns
30    data['date'] = pd.to_datetime(data['timestamp']).dt.date
31    daily_data = data.groupby(['date', 'variant'])['conversion'].mean()
32        .unstack()
33
34    # Check for systematic time trends
35    daily_variance = daily_data.var()
36    if (daily_variance > 0.01).any():
37        warnings.append("High temporal variance detected")
38
39    # 3. Multiple assignments
40    multi_assigned = data.groupby('user_id')['variant'].nunique()
41    n_multi = (multi_assigned > 1).sum()
42
43    if n_multi > 0:
44        warnings.append(f"{n_multi} users with multiple variant assignments")
45
46    # 4. Missing data
47    missing_pct = data.isnull().sum() / len(data) * 100
48    if (missing_pct > 5).any():
49        warnings.append(f"High missing data: {dict(missing_pct[missing_pct > 5])}")
50
51    return {
52        'srm_p_value': srm_p_value,
53        'variant_counts': dict(variant_counts),
54        'warnings': warnings,
55        'passed': len(warnings) == 0
56    }

```



```

55     }
56
57 # Example usage
58 # data = pd.read_csv('experiment_data.csv')
59 # checks = preanalysis_checks(data)
60 # print(f"Pre-analysis checks: {'PASSED' if checks['passed'] else '
    WARNINGS'}")
61 # for warning in checks['warnings']:
62 #     print(f"    - {warning}")

```

10.2.2 Exploratory Data Analysis

Visualize and summarize data before formal testing.

Listing 10.2: Exploratory Analysis Dashboard

```

1 def create_eda_dashboard(data):
2     """Create exploratory analysis visualizations."""
3     fig, axes = plt.subplots(2, 2, figsize=(14, 10))
4
5     # 1. Overall conversion rates
6     conversion_by_variant = data.groupby('variant')['conversion'].agg([
7         ('mean', 'mean'),
8         ('std', 'std'),
9         ('count', 'count')
10    ])
11
12    axes[0, 0].bar(conversion_by_variant.index,
13                  conversion_by_variant['mean'],
14                  yerr=conversion_by_variant['std'] /
15                      np.sqrt(conversion_by_variant['count']))
16    axes[0, 0].set_xlabel('Variant')
17    axes[0, 0].set_ylabel('Conversion Rate')
18    axes[0, 0].set_title('Conversion Rate by Variant')
19    axes[0, 0].grid(alpha=0.3)
20
21    # 2. Time series
22    data['date'] = pd.to_datetime(data['timestamp']).dt.date
23    daily_cr = data.groupby(['date', 'variant'])['conversion'].mean().
        unstack()
24
25    for variant in daily_cr.columns:
26        axes[0, 1].plot(daily_cr.index, daily_cr[variant],
27                        marker='o', label=f'Variant_{variant}')
28
29    axes[0, 1].set_xlabel('Date')
30    axes[0, 1].set_ylabel('Daily Conversion Rate')
31    axes[0, 1].set_title('Conversion Rate Over Time')
32    axes[0, 1].legend()
33    axes[0, 1].grid(alpha=0.3)
34    plt.setp(axes[0, 1].xaxis.get_majorticklabels(), rotation=45)
35
36    # 3. Distribution of continuous metric (if available)
37    if 'value' in data.columns:
38        for variant in data['variant'].unique():
39            variant_data = data[data['variant'] == variant]['value'].
                dropna()

```

```

40         axes[1, 0].hist(variant_data, bins=50, alpha=0.5,
41                          label=f'Variant_{variant}')
42
43     axes[1, 0].set_xlabel('Value')
44     axes[1, 0].set_ylabel('Frequency')
45     axes[1, 0].set_title('Distribution of Values by Variant')
46     axes[1, 0].legend()
47     axes[1, 0].grid(alpha=0.3)
48
49     # 4. Sample size accumulation
50     data_sorted = data.sort_values('timestamp')
51     data_sorted['cumulative'] = range(1, len(data_sorted) + 1)
52
53     for variant in data['variant'].unique():
54         variant_data = data_sorted[data_sorted['variant'] == variant]
55         axes[1, 1].plot(variant_data['timestamp'],
56                        variant_data['cumulative'],
57                        label=f'Variant_{variant}')
58
59     axes[1, 1].set_xlabel('Time')
60     axes[1, 1].set_ylabel('Cumulative Sample Size')
61     axes[1, 1].set_title('Sample Size Accumulation')
62     axes[1, 1].legend()
63     axes[1, 1].grid(alpha=0.3)
64
65     plt.tight_layout()
66     return fig

```

10.3 Handling Data Issues

10.3.1 Missing Data

Missing completely at random (MCAR): Missing mechanism independent of data

Missing at random (MAR): Missing depends on observed data

Missing not at random (MNAR): Missing depends on unobserved data

Strategies:

1. **Complete case analysis:** Drop missing (if MCAR and limited missingness)
2. **Imputation:** Fill with mean, median, or model predictions
3. **Multiple imputation:** Account for imputation uncertainty
4. **Sensitivity analysis:** Test robustness to missing data assumptions

10.3.2 Outliers

Extreme values can bias estimates, especially for means.

Detection methods:

- Visual inspection (box plots, scatter plots)
- Statistical rules (e.g., beyond 3 standard deviations)

- Domain knowledge (impossible or implausible values)

Handling approaches:

1. **Winsorization:** Cap at percentiles (e.g., 1st and 99th)
2. **Transformation:** Log transform to reduce skewness
3. **Robust statistics:** Use median instead of mean
4. **Separate analysis:** Analyze with and without outliers

Listing 10.3: Outlier Treatment

```

1 def winsorize_data(data, column, lower_percentile=1, upper_percentile
  =99):
2     """
3     Winsorize data by capping at specified percentiles.
4
5     Parameters:
6     -----
7     data: DataFrame
8     column: str
9     lower_percentile, upper_percentile: float
10    Percentile thresholds
11
12    Returns:
13    -----
14    DataFrame: Data with winsorized column
15    """
16    lower_bound = np.percentile(data[column], lower_percentile)
17    upper_bound = np.percentile(data[column], upper_percentile)
18
19    data[f'{column}_winsorized'] = data[column].clip(
20        lower=lower_bound,
21        upper=upper_bound
22    )
23
24    n_capped = ((data[column] < lower_bound) |
25                (data[column] > upper_bound)).sum()
26
27    print(f"Winsorized {n_capped} values ({n_capped/len(data)*100:.2f
28          }%)")
29    print(f"Bounds: [{lower_bound:.2f}, {upper_bound:.2f}]")
30
31    return data
32
33
34 # Compare analyses
35 def compare_outlier_treatment(data, metric_col):
36     """Compare results with different outlier treatments."""
37     results = {}
38
39     # Original data
40     for variant in data['variant'].unique():
41         variant_data = data[data['variant'] == variant][metric_col]
42         results[f'variant_{variant}_original'] = {

```

```

43         'mean': variant_data.mean(),
44         'median': variant_data.median(),
45         'std': variant_data.std()
46     }
47
48     # Winsorized data
49     data_winsorized = winsorize_data(data.copy(), metric_col)
50     for variant in data['variant'].unique():
51         variant_data = data_winsorized[
52             data_winsorized['variant'] == variant
53         ][f'{metric_col}_winsorized']
54         results[f'variant_{variant}_winsorized'] = {
55             'mean': variant_data.mean(),
56             'median': variant_data.median(),
57             'std': variant_data.std()
58         }
59
60     return pd.DataFrame(results).T

```

10.4 Variance Reduction Techniques

Reducing variance increases power without additional sample size.

10.4.1 CUPED (Controlled-Experiment Using Pre-Experiment Data)

Use pre-experiment covariates to reduce variance.

$$Y_{\text{adjusted}} = Y - \theta(X - \mathbb{E}[X]) \quad (10.1)$$

where X is pre-experiment covariate and $\theta = \frac{\text{Cov}(Y, X)}{\text{Var}(X)}$.

Variance reduction:

$$\text{Var}(Y_{\text{adjusted}}) = \text{Var}(Y)(1 - \rho^2) \quad (10.2)$$

where $\rho = \text{Cor}(Y, X)$.

Listing 10.4: CUPED Implementation

```

1  cuped_adjustment <- function(y, x) {
2      # y: outcome metric
3      # x: pre-experiment covariate
4
5      # Estimate theta
6      theta <- cov(y, x) / var(x)
7
8      # Adjust outcome
9      y_adjusted <- y - theta * (x - mean(x))
10
11     # Calculate variance reduction
12     variance_reduction <- 1 - (var(y_adjusted) / var(y))
13
14     list(
15         y_adjusted = y_adjusted,
16         theta = theta,

```

```

17     variance_reduction = variance_reduction,
18     original_var = var(y),
19     adjusted_var = var(y_adjusted)
20   )
21 }
22
23 # Example
24 set.seed(123)
25 n <- 1000
26
27 # Pre-experiment metric (e.g., baseline conversion)
28 x_pre <- rbinom(n, 1, 0.10)
29
30 # Experiment outcome (correlated with baseline)
31 y <- rbinom(n, 1, 0.10 + 0.2 * x_pre)
32
33 # Apply CUPED
34 result <- cuped_adjustment(y, x_pre)
35
36 cat("Variance_reduction:", result$variance_reduction, "\n")
37 cat("Original_var:", result$original_var, "\n")
38 cat("Adjusted_var:", result$adjusted_var, "\n")

```

10.4.2 Stratification

Analyze within strata and combine results.

Benefits:

- Increases precision
- Ensures balance on stratification variables
- Enables subgroup analysis

10.5 Result Interpretation

10.5.1 Statistical vs. Practical Significance

Report both:

1. **Statistical significance:** P-value and confidence interval
2. **Effect size:** Absolute and relative differences
3. **Business impact:** Expected revenue/user impact
4. **Confidence interval:** Range of plausible effects

Example 10.1 (Interpretation Template). "The treatment increased conversion rate from 8.0% to 8.5%, a relative increase of 6.25% (95% CI: [2%, 11%], $p = 0.01$). This translates to an estimated 500 additional conversions per month and \$25,000 in additional revenue (95% CI: [\$10,000, \$40,000])."

10.5.2 Communicating Uncertainty

Always quantify uncertainty:

- Report confidence intervals, not just point estimates
- Visualize uncertainty with error bars
- Discuss what we cannot conclude
- Consider robustness under different assumptions

10.6 Decision Making Framework

10.6.1 Quantitative Criteria

1. **Statistical significance:** $p < 0.05$ (or chosen α)
2. **Practical significance:** Effect exceeds minimum meaningful threshold
3. **Guardrail metrics:** No degradation on critical metrics
4. **Cost-benefit:** Expected value exceeds implementation cost

10.6.2 Qualitative Factors

- **Strategic alignment:** Fits product vision
- **User experience:** Beyond measured metrics
- **Technical debt:** Long-term maintainability
- **Competitive dynamics:** Market positioning
- **Brand considerations:** Perception and values

10.6.3 Decision Matrix

Statistical	Practical	Decision	Action
Significant	Significant	Ship	Implement
Significant	Not significant	Hold	More investigation
Not significant	Significant	Iterate	Refine and retest
Not significant	Not significant	Kill	Abandon

10.7 Summary

Thorough analysis and interpretation are critical:

- Conduct pre-analysis quality checks
- Handle missing data and outliers appropriately

- Apply variance reduction techniques
- Report both statistical and practical significance
- Communicate uncertainty clearly
- Make decisions combining quantitative and qualitative factors

Good analysis transforms data into actionable insights.

10.8 Further Reading

- ? : CUPED methodology
- ? : Interpretation best practices
- ? : Missing data analysis

10.9 Exercises

1. Implement complete analysis pipeline including quality checks and variance reduction
2. Compare power with and without CUPED for varying correlations
3. Create stakeholder report template with visualizations
4. Build decision framework incorporating business constraints

Chapter 11

Case Studies and Real-World Applications

11.1 Learning Objectives

After completing this chapter, readers will be able to:

- Learn from real-world A/B testing successes and failures with detailed statistical analysis
- Understand common pitfalls through case examples with actual data
- Apply quantitative frameworks to similar business problems
- Recognize industry-specific considerations and adapt methods accordingly
- Conduct post-mortem analysis of experiments using provided templates

11.2 Case Study 1: Google Search Results and Page Load Time

11.2.1 Context and Business Problem

In 2009, Google tested whether displaying more search results per page (20 or 30 instead of 10) would improve user engagement. The hypothesis seemed reasonable: fewer clicks to pagination should improve user experience.

11.2.2 Experimental Design

Design: A/B/C test with three variants

- Control: 10 results per page (baseline)
- Variant A: 20 results per page
- Variant B: 30 results per page

Randomization: User-level randomization, 33.3% allocation each

Duration: 3 weeks

Sample size: Millions of users (Google traffic scale)

Metrics:

- Primary: Searches per user session
- Secondary: Page load time, click-through rate, user satisfaction

11.2.3 Results and Analysis

Metric	10 Results	20 Results	30 Results
Page load time (ms)	400	650	900
Load time increase	–	+63%	+125%
Searches per user	5.2	4.9	4.6
Relative change	–	-5.8%	-11.5%
Satisfaction score	7.8/10	7.5/10	7.2/10

Statistical significance: All differences highly significant ($p < 0.001$) due to massive sample size.

Correlation analysis: Strong negative correlation between page load time and searches per user:

$$\text{Elasticity} = \frac{\Delta \text{Searches}}{\Delta \text{Load time}} \approx -1.0 \quad (11.1)$$

A 100ms increase in load time corresponded to approximately 1% decrease in searches.

11.2.4 Deep Dive: Causal Mechanism

Why did more results decrease engagement?

1. **Performance degradation:** Rendering more results increased page weight and computation
2. **Perceived value:** Users valued speed over choice
3. **Cognitive overload:** More options paradoxically made decision harder
4. **User behavior:** Most users (>90%) found results in first 10 regardless

11.2.5 Statistical Analysis Framework

Listing 11.1: Analyzing Load Time vs. Engagement Trade-offs

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 from scipy import stats
5 import seaborn as sns
6
7 # Simulated data based on Google's published results
```

```

8 np.random.seed(42)
9
10 def simulate_google_experiment(n_users=100000):
11     """Simulate Google search results experiment."""
12
13     # Allocate users to variants
14     variants = np.random.choice(['10_results', '20_results', '30_
15                                 results'],
16                                 size=n_users)
17
18     # Simulate page load times (milliseconds)
19     load_times = {
20         '10_results': np.random.normal(400, 50),
21         '20_results': np.random.normal(650, 70),
22         '30_results': np.random.normal(900, 90)
23     }
24
25     # Simulate searches per user (Poisson with rate dependent on load
26     # time)
27     # Model: slower load time -> lower engagement
28     base_searches = {
29         '10_results': 5.2,
30         '20_results': 4.9,
31         '30_results': 4.6
32     }
33
34     data = []
35     for i, variant in enumerate(variants):
36         load_time = load_times[variant]
37
38         # Add user-level noise
39         lambda_param = base_searches[variant] * np.random.lognormal(0,
40                               0.3)
41         searches = np.random.poisson(lambda_param)
42
43         # Satisfaction (0-10 scale, correlated with searches and load
44         # time)
45         satisfaction = np.clip(
46             7.8 - (load_time - 400) / 200 + np.random.normal(0, 1),
47             0, 10
48         )
49
50         data.append({
51             'user_id': i,
52             'variant': variant,
53             'load_time_ms': load_time,
54             'searches': searches,
55             'satisfaction': satisfaction
56         })
57
58     return pd.DataFrame(data)
59
60 # Generate data
61 data = simulate_google_experiment(n_users=100000)
62
63 # Aggregate statistics by variant
64 summary = data.groupby('variant').agg({
65     'load_time_ms': ['mean', 'std'],

```

```

62     'searches': ['mean', 'std'],
63     'satisfaction': ['mean', 'std']
64 }).round(2)
65
66 print("=== Experiment Summary ===")
67 print(summary)
68
69 # Statistical tests
70 variants = ['10_results', '20_results', '30_results']
71
72 print("\n=== Pairwise Comparisons (Two-sample t-tests) ===")
73 for i in range(len(variants)):
74     for j in range(i+1, len(variants)):
75         v1, v2 = variants[i], variants[j]
76
77         searches_1 = data[data['variant']==v1]['searches']
78         searches_2 = data[data['variant']==v2]['searches']
79
80         t_stat, p_value = stats.ttest_ind(searches_1, searches_2)
81
82         diff = searches_2.mean() - searches_1.mean()
83         rel_diff = diff / searches_1.mean() * 100
84
85         print(f"\n{v1} vs {v2}:")
86         print(f"Mean difference: {diff:.3f} searches ({rel_diff:+.1f}%)"
87               f"\nt-statistic: {t_stat:.2f}"
88               f"p-value: {p_value:.2e}")
89
90 # Regression: searches ~ load_time
91 from sklearn.linear_model import LinearRegression
92
93 X = data[['load_time_ms']].values
94 y = data['searches'].values
95
96 model = LinearRegression()
97 model.fit(X, y)
98
99 print("\n=== Regression: Searches ~ Load Time ===")
100 print(f"Slope: {model.coef_[0]:.6f} searches per ms")
101 print(f"Interpretation: Each 100ms increase -> "
102       f"{model.coef_[0]*100:.2f} fewer searches")
103
104 # Visualization
105 fig, axes = plt.subplots(2, 2, figsize=(14, 10))
106
107 # 1. Load time by variant
108 ax = axes[0, 0]
109 data.boxplot(column='load_time_ms', by='variant', ax=ax)
110 ax.set_xlabel('Variant')
111 ax.set_ylabel('Page Load Time (ms)')
112 ax.set_title('Load Time Distribution by Variant')
113 plt.sca(ax)
114 plt.xticks(range(1, 4), ['10_results', '20_results', '30_results'])
115
116 # 2. Searches by variant
117 ax = axes[0, 1]

```

```

118 variant_means = data.groupby('variant')['searches'].agg(['mean', 'sem',
119 ])
119 variant_order = ['10_results', '20_results', '30_results']
120 means = [variant_means.loc[v, 'mean'] for v in variant_order]
121 sems = [variant_means.loc[v, 'sem'] for v in variant_order]
122
123 x = np.arange(3)
124 ax.bar(x, means, yerr=[1.96*s for s in sems], capsize=5, alpha=0.7)
125 ax.set_xticks(x)
126 ax.set_xticklabels(['10_results', '20_results', '30_results'])
127 ax.set_ylabel('Searches_per_User')
128 ax.set_title('Mean_Searches_by_Variant_(95%_CI)')
129 ax.grid(axis='y', alpha=0.3)
130
131 # 3. Scatter: Load time vs searches
132 ax = axes[1, 0]
133 for variant, color in zip(variant_order, ['C0', 'C1', 'C2']):
134     subset = data[data['variant']==variant].sample(1000)
135     ax.scatter(subset['load_time_ms'], subset['searches'],
136               alpha=0.3, s=20, label=variant.replace('_', ' '), color=
               color)
137
138 # Add regression line
139 X_line = np.linspace(300, 1000, 100).reshape(-1, 1)
140 y_line = model.predict(X_line)
141 ax.plot(X_line, y_line, 'r--', linewidth=2, label='Regression_fit')
142
143 ax.set_xlabel('Page_Load_Time_(ms)')
144 ax.set_ylabel('Searches_per_User')
145 ax.set_title('Load_Time_vs._Engagement')
146 ax.legend()
147 ax.grid(alpha=0.3)
148
149 # 4. Satisfaction scores
150 ax = axes[1, 1]
151 data.boxplot(column='satisfaction', by='variant', ax=ax)
152 ax.set_xlabel('Variant')
153 ax.set_ylabel('Satisfaction_Score_(0-10)')
154 ax.set_title('User_Satisfaction_by_Variant')
155 plt.sca(ax)
156 plt.xticks(range(1, 4), ['10_results', '20_results', '30_results'])
157
158 plt.tight_layout()
159 plt.savefig('google_case_study.pdf', bbox_inches='tight')
160 plt.show()
161
162 # Calculate business impact
163 print("\n===_Business_Impact_Estimation_===")
164 baseline_searches = data[data['variant']=='10_results']['searches'].
    mean()
165
166 for variant in ['20_results', '30_results']:
167     variant_searches = data[data['variant']==variant]['searches'].mean
        ()
168     loss = (variant_searches - baseline_searches) / baseline_searches *
        100
169
170 # At Google scale (billions of searches/day)

```

```

171     daily_searches_billions = 8.5 # Approximate 2009 volume
172     annual_loss_millions = daily_searches_billions * 1000 * loss / 100
        * 365
173
174     print(f"\n{variant}:")
175     print(f"    Relative loss: {loss:.1f}%")
176     print(f"    Estimated annual search loss: {annual_loss_millions:.0f}M
        searches")
177     print(f"    Decision: DO NOT SHIP")

```

11.2.6 Lessons Learned

1. **Performance is a feature:** Page load time directly impacts user behavior at all scales, but especially at massive scale
2. **Guardrail metrics:** Always monitor technical performance metrics alongside engagement metrics
3. **User psychology:** More choice isn't always better; cognitive load and perceived value matter
4. **Scale amplification:** At Google's traffic volume, small performance differences have billion-search impacts
5. **Regression analysis:** Quantifying relationship between technical metrics and business outcomes enables forecasting

Actionable framework:

- Measure page load time as standard guardrail
- Set performance budgets (e.g., < 500ms) before testing
- Estimate elasticity of engagement with respect to performance
- Calculate ROI of performance optimization

11.3 Case Study 2: E-commerce Checkout Optimization

11.3.1 Context

An e-commerce company with \$500M annual revenue hypothesized that simplifying checkout from 5 steps to 3 steps would increase conversion.

11.3.2 Experimental Design

Hypothesis: Reducing friction in checkout process will increase completion rate.

Design: A/B test

- Control: 5-step checkout (address, shipping, payment, review, confirm)

- Treatment: 3-step checkout (combined address/shipping, payment/review, confirm)

Randomization: User-level, 50/50 split

Sample size: 50,000 users per variant (calculated for 80% power to detect 5% relative lift)

Duration: 14 days (to capture two full weeks including weekends)

Metrics:

- Primary: Checkout completion rate (among initiated)
- Secondary: Overall conversion rate, revenue per user, abandonment points
- Guardrails: Customer support contacts, return rate

11.3.3 Overall Results

Metric	Control	Treatment	Lift
Checkout initiation	12.0%	12.1%	+0.8%
Checkout completion	68.0%	71.5%	+5.1%**
Overall conversion	8.16%	8.65%	+6.0%**
Revenue per user	\$4.25	\$4.50	+5.9%**
Avg order value	\$52.08	\$52.02	-0.1%
Support contacts	2.1%	2.3%	+9.5%

** Statistically significant at $\alpha = 0.05$

11.3.4 Segmentation Analysis

Critical insight came from examining heterogeneous treatment effects:

By customer type:

Segment	Control	Treatment	Lift	Size
New users	6.5%	7.8%	+20.0%**	60%
Returning users	9.2%	9.0%	-2.2%	40%

By device:

Device	Control	Treatment	Lift	Size
Mobile	5.8%	6.9%	+19.0%**	55%
Desktop	11.2%	11.0%	-1.8%	45%

Listing 11.2: Heterogeneous Treatment Effect Analysis

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 from scipy import stats
5 import seaborn as sns
6
7 def simulate_checkout_experiment(n_users=100000):

```

```

8      """Simulate_e-commerce_checkout_experiment_with_segments."""
9
10     np.random.seed(42)
11
12     # User characteristics
13     is_new = np.random.binomial(1, 0.60, n_users) # 60% new users
14     is_mobile = np.random.binomial(1, 0.55, n_users) # 55% mobile
15
16     # Treatment assignment
17     treatment = np.random.binomial(1, 0.5, n_users)
18
19     # Baseline conversion probabilities
20     baseline_conv = np.where(
21         is_new,
22         np.where(is_mobile, 0.058, 0.070), # New users: mobile vs
23         np.where(is_mobile, 0.085, 0.112) # Returning: mobile vs
24     )
25
26     # Treatment effects (heterogeneous)
27     treatment_effect = np.where(
28         is_new,
29         np.where(is_mobile, 0.011, -0.001), # New users benefit, esp
30         np.where(is_mobile, 0.005, -0.002) # Returning slightly worse
31     )
32
33     # Realized conversion
34     conv_prob = baseline_conv + treatment * treatment_effect
35     converted = np.random.binomial(1, conv_prob)
36
37     # Revenue (conditional on conversion)
38     revenue = converted * np.random.lognormal(np.log(50), 0.6)
39
40     data = pd.DataFrame({
41         'user_id': range(n_users),
42         'treatment': treatment,
43         'is_new': is_new,
44         'is_mobile': is_mobile,
45         'converted': converted,
46         'revenue': revenue
47     })
48
49     # Add labels
50     data['variant'] = data['treatment'].map({0: 'Control', 1: '
51         Treatment'})
52     data['user_type'] = data['is_new'].map({1: 'New', 0: 'Returning'})
53     data['device'] = data['is_mobile'].map({1: 'Mobile', 0: 'Desktop'})
54
55     return data
56
57 # Generate data
58 data = simulate_checkout_experiment(n_users=100000)
59
60 # Overall analysis
61 print("===OverallResults===")
62 overall = data.groupby('variant').agg({

```



```

62     'converted': ['sum', 'mean', 'sem'],
63     'revenue': ['sum', 'mean']
64 ))
65
66 control_rate = overall.loc['Control', ('converted', 'mean')]
67 treatment_rate = overall.loc['Treatment', ('converted', 'mean')]
68 lift = (treatment_rate - control_rate) / control_rate * 100
69
70 print(f"Control_conversion_rate:_{control_rate:.4f}")
71 print(f"Treatment_conversion_rate:_{treatment_rate:.4f}")
72 print(f"Relative_lift:_{lift:+.1f}%")
73
74 # Statistical test
75 control_conv = data[data['treatment']==0]['converted']
76 treatment_conv = data[data['treatment']==1]['converted']
77
78 z_stat = (treatment_conv.mean() - control_conv.mean()) / np.sqrt(
79     control_conv.var() / len(control_conv) +
80     treatment_conv.var() / len(treatment_conv)
81 )
82 p_value = 2 * (1 - stats.norm.cdf(abs(z_stat)))
83
84 print(f"\nZ-statistic:_{z_stat:.3f}")
85 print(f"P-value:_{p_value:.4f}")
86
87 # Segmented analysis
88 print("\n===_Heterogeneous_Treatment_Effects_===")
89
90 for segment_var, segment_name in [('user_type', 'User_Type'),
91                                   ('device', 'Device')]:
92     print(f"\nBy_{segment_name}:")
93
94     segment_results = data.groupby([segment_var, 'variant']).agg({
95         'converted': ['mean', 'count']
96     }).reset_index()
97
98     # Pivot for easier comparison
99     pivot = segment_results.pivot(
100         index=segment_var,
101         columns='variant',
102         values=('converted', 'mean')
103     )
104
105     for segment in pivot.index:
106         control_rate = pivot.loc[segment, ('Control')]
107         treatment_rate = pivot.loc[segment, ('Treatment')]
108         lift = (treatment_rate - control_rate) / control_rate * 100
109
110         # Test within segment
111         seg_data = data[data[segment_var]==segment]
112         seg_control = seg_data[seg_data['treatment']==0]['converted']
113         seg_treatment = seg_data[seg_data['treatment']==1]['converted']
114
115         _, p = stats.ttest_ind(seg_treatment, seg_control)
116         sig = '**' if p < 0.05 else '_'
117
118         print(f"_{segment}:_{control_rate:.4f}_->_{treatment_rate:.4f}_
            _")

```

```

119         f"({lift:+.1f}%) {sig}")
120
121 # Visualization
122 fig, axes = plt.subplots(2, 2, figsize=(14, 10))
123
124 # 1. Overall conversion rates
125 ax = axes[0, 0]
126 overall_data = data.groupby('variant')['converted'].agg(['mean', 'sem',
127 ])
128 x = np.arange(2)
129 ax.bar(x, overall_data['mean'], yerr=1.96*overall_data['sem'],
130       capsize=5, alpha=0.7, color=['C0', 'C1'])
131 ax.set_xticks(x)
132 ax.set_xticklabels(['Control', 'Treatment'])
133 ax.set_ylabel('Conversion_Rate')
134 ax.set_title('Overall_Conversion_Rate_(95%_CI)')
135 ax.set_ylim([0, 0.10])
136 for i, (idx, row) in enumerate(overall_data.iterrows()):
137     ax.text(i, row['mean'] + 0.003, f"{row['mean']:.3f}",
138           ha='center', fontweight='bold')
139 ax.grid(axis='y', alpha=0.3)
140
141 # 2. By user type
142 ax = axes[0, 1]
143 user_type_data = data.groupby(['user_type', 'variant'])['converted'].
144     mean().unstack()
145 user_type_data.plot(kind='bar', ax=ax, color=['C0', 'C1'], alpha=0.7)
146 ax.set_ylabel('Conversion_Rate')
147 ax.set_title('Conversion_Rate_by_User_Type')
148 ax.set_xticklabels(ax.get_xticklabels(), rotation=0)
149 ax.legend(title='Variant')
150 ax.grid(axis='y', alpha=0.3)
151
152 # 3. By device
153 ax = axes[1, 0]
154 device_data = data.groupby(['device', 'variant'])['converted'].mean().
155     unstack()
156 device_data.plot(kind='bar', ax=ax, color=['C0', 'C1'], alpha=0.7)
157 ax.set_ylabel('Conversion_Rate')
158 ax.set_title('Conversion_Rate_by_Device')
159 ax.set_xticklabels(ax.get_xticklabels(), rotation=0)
160 ax.legend(title='Variant')
161 ax.grid(axis='y', alpha=0.3)
162
163 # 4. Treatment effect by segment
164 ax = axes[1, 1]
165
166 segments = []
167 effects = []
168
169 for user_type in ['New', 'Returning']:
170     for device in ['Mobile', 'Desktop']:
171         seg_data = data[(data['user_type']==user_type) &
172             (data['device']==device)]
173
174         control_mean = seg_data[seg_data['treatment']==0]['converted'].
175             mean()
176         treatment_mean = seg_data[seg_data['treatment']==1]['converted']

```

```

    ].mean()
173
    effect = (treatment_mean - control_mean) / control_mean * 100
174
175
    label = f"{user_type}\n{device}"
176
    segments.append(label)
177
    effects.append(effect)
178
179
180 colors = ['green' if e > 0 else 'red' for e in effects]
181 x_pos = np.arange(len(segments))
182
183 ax.barh(x_pos, effects, color=colors, alpha=0.7)
184 ax.axvline(x=0, color='black', linestyle='-', linewidth=0.8)
185 ax.set_yticks(x_pos)
186 ax.set_yticklabels(segments)
187 ax.set_xlabel('Relative Lift (%)')
188 ax.set_title('Treatment Effect by Segment')
189 ax.grid(axis='x', alpha=0.3)
190
191 plt.tight_layout()
192 plt.savefig('checkout_segmentation.pdf', bbox_inches='tight')
193 plt.show()
194
195 # Revenue impact calculation
196 print("\n=== Business Impact ===")
197 annual_visitors = 10_000_000 # 10M annual visitors
198 current_revenue = 500_000_000 # $500M annually
199
200 baseline_conv = control_rate
201 new_conv = treatment_rate
202 revenue_per_conversion = data[data['converted']==1]['revenue'].mean()
203
204 incremental_conversions = annual_visitors * (new_conv - baseline_conv)
205 incremental_revenue = incremental_conversions * revenue_per_conversion
206
207 print(f"Baseline conversions: {annual_visitors*baseline_conv:,.0f}")
208 print(f"New conversions: {annual_visitors*new_conv:,.0f}")
209 print(f"Incremental conversions: {incremental_conversions:,.0f}")
210 print(f"Avg revenue per conversion: ${revenue_per_conversion:.2f}")
211 print(f"Incremental annual revenue: ${incremental_revenue:,.0f}")
212 print(f"Relative revenue increase: {incremental_revenue/current_revenue
    *100:.2f}%")
213
214 # Recommendation
215 print("\n=== Recommendation ===")
216 print("SHIP TO ALL USERS")
217 print("Rationale:")
218 print("- Statistically significant overall lift")
219 print("- Strong benefit for strategically important segments (new,
    mobile)")
220 print("- Small negative effect on returning/desktop acceptable")
221 print("- Projected $30M+ incremental annual revenue")
222 print("- Support contact increase minimal and manageable")

```

11.3.5 Decision Framework

The team faced a trade-off: overall benefit with segment-specific harm.

Decision criteria:

1. **Overall impact:** +6.0% conversion, highly significant
 2. **Strategic alignment:** New users and mobile are growth priorities
 3. **Magnitude:** +20% for new/mobile vs. -2% for returning/desktop
 4. **Business value:** \$30M+ projected incremental annual revenue
 5. **Reversibility:** Can iterate if returning user experience degrades further
- Decision:** Ship to all users, monitor returning user metrics closely.

11.3.6 Post-Launch Monitoring

After launch, team implemented:

- Weekly cohort analysis of returning users
- Customer feedback analysis (surveys, support tickets)
- Funnel tracking for abandonment points
- Quarterly re-testing to validate continued benefit

11.3.7 Lessons Learned

1. **Segment analysis is critical:** Averages hide important heterogeneity
2. **Strategic weighting:** Not all segments are equal; weight by business priorities
3. **Trade-offs are normal:** Few changes benefit everyone universally
4. **Quantify business impact:** Translate statistical findings to revenue/profit
5. **Monitor post-launch:** Continue learning after shipping

11.4 Case Study 3: Netflix Thumbnail Personalization

11.4.1 Context

Netflix hypothesized that personalizing video thumbnails based on user viewing history and preferences would increase engagement (click-through rate and watch time).

11.4.2 Technical Challenge

Consistency requirement: Users see the same videos multiple times (homepage, search, recommendations). Showing different thumbnails each time would be confusing.

Solution: Deterministic assignment— $\text{hash}(\text{user_id}, \text{video_id}) \rightarrow \text{thumbnail_variant}$
This ensures:

- Same user always sees same thumbnail for given video
- Different users see different thumbnails (personalization)
- Randomization for experimentation

11.4.3 Multi-Stage Experimental Approach

Netflix used an iterative experimental strategy:

Stage 1: Validate thumbnail selection matters

- Test: Random thumbnails vs. editorial picks
- Result: Editorial picks had 30% higher engagement
- Conclusion: Thumbnail selection is important lever

Stage 2: Test personalization hypothesis

- Test: Editorial picks (best average) vs. personalized selection
- Personalization: Select thumbnail based on user's genre preferences and viewing history
- Result: Personalized thumbnails 20% higher engagement than editorial
- Conclusion: Personalization works

Stage 3: Optimize personalization model

- Test: Multiple ML models for personalization
 - Collaborative filtering
 - Content-based filtering
 - Neural network (deep learning)
 - Ensemble
- Result: Neural network approach best (additional 5% lift)
- Conclusion: Ship neural network model

11.4.4 Experimental Design Details (Stage 2)

Randomization: Video-level stratified randomization

- For each video, select N candidate thumbnails (N=3-10)
- Randomly assign users to personalization (treatment) vs. best-on-average (control)
- Within treatment: use model to select best thumbnail for each user

Metrics:

- Primary: Click-through rate (CTR) - play rate when video is visible
- Secondary: Watch time, completion rate, satisfaction rating
- Guardrails: Customer support, cancellations

Duration: 4 weeks (to account for learning effects)

Sample: 10M users, 5M control / 5M treatment

11.4.5 Results

Metric	Control	Treatment	Lift
CTR (overall)	3.2%	3.84%	+20.0%**
Watch time (hrs/user/week)	12.5	13.1	+4.8%**
Completion rate	68%	70%	+2.9%**
Satisfaction (1-5)	4.1	4.2	+2.4%**

** $p < 0.001$

Heterogeneity by video type:

Personalization worked especially well for:

- Ensemble casts (different thumbnails feature different actors)
- Genre-spanning content (show different aspects to different users)
- Foreign language content (different cultural framing)

11.4.6 Implementation Considerations

1. **Infrastructure:** Real-time personalization requires low-latency model serving
2. **Quality control:** Filter inappropriate thumbnails (spoilers, misleading images)
3. **Fallback logic:** When model fails, default to editorial pick
4. **Cold start:** New users see editorial picks until preferences learned
5. **Monitoring:** Track model performance, A/A tests for bias

11.4.7 Lessons Learned

1. **Iterative experimentation:** Build incrementally, validating each assumption
2. **Technical-statistical integration:** Implementation details (hashing, persistence) affect experimental validity
3. **Personalization value:** Context-specific optimization often outperforms one-size-fits-all
4. **Heterogeneous effects:** Personalization benefits vary by content type
5. **Production readiness:** Consider operational constraints before launch

11.5 Case Study 4: Failed Experiment—Subscription Paywall

Not all experiments succeed. Learning from failures is as valuable as from successes.

11.5.1 Context

A digital media platform with freemium model (free articles with monthly limit, premium unlimited access) wanted to increase paid subscriptions.

Hypothesis: More aggressive paywall messaging and earlier paywall triggers would increase conversion to paid subscriptions.

11.5.2 Design

- Control: Paywall after 5 free articles/month, soft messaging
- Treatment: Paywall after 3 free articles/month, aggressive CTA

Primary metric: Subscription conversion rate (

Duration: 3 weeks

Sample: 200,000 users per variant

11.5.3 Results: The Good and The Bad

Metric	Control	Treatment	Abs. Diff	Rel. Lift
<i>Positive:</i>				
Subscription conversion	1.20%	1.38%	+0.18pp	+15.0%**
<i>Negative:</i>				
Bounce rate	35%	49%	+14pp	+40.0%**
Pages per user	8.2	5.1	-3.1	-37.8%**
Return rate (7-day)	42%	31%	-11pp	-26.2%**
<i>Net effect:</i>				
Net new subscribers	1200	1140	-60	-5.0%

** $p < 0.001$

11.5.4 What Went Wrong

1. **Metric misalignment:** Optimized subscription conversion rate (local metric) at expense of overall user base (global metric)
2. **Missing guardrails:** Didn't adequately monitor engagement and retention metrics
3. **Time horizon mismatch:** Immediate conversions vs. long-term user value
4. **User experience:** Aggressive tactics damaged brand perception

Root cause analysis:

The aggressive paywall converted users faster but drove away many more users who would have converted later:

- Fewer users reached the paywall (40% bounce increase)
- Lower return rate meant fewer conversion opportunities
- Net effect: slightly fewer total conversions despite higher rate

11.5.5 Correct Analysis

Listing 11.3: Analyzing Funnel Effects

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4
5 def analyze_paywall_funnel():
6     """Analyze the full funnel, not just conversion rate."""
7
8     # Monthly visitors
9     monthly_visitors = 100000
10
11     # Control funnel
12     control_funnel = {
13         'visitors': monthly_visitors,
14         'bounce_rate': 0.35,
15         'engaged': None, # calculated
16         'reach_paywall_rate': 0.15, # % of engaged users
17         'paywall_conversion': 0.012,
18         'new_subscribers': None # calculated
19     }
20
21     control_funnel['engaged'] = control_funnel['visitors'] * (
22         1 - control_funnel['bounce_rate']
23     )
24     control_funnel['reach_paywall'] = control_funnel['engaged'] * (
25         control_funnel['reach_paywall_rate']
26     )
27     control_funnel['new_subscribers'] = control_funnel['reach_paywall']
28         * (
29         control_funnel['paywall_conversion']
30     )
31
32     # Treatment funnel
33     treatment_funnel = {
34         'visitors': monthly_visitors,
35         'bounce_rate': 0.49,
36         'engaged': None,
37         'reach_paywall_rate': 0.20, # Higher due to earlier trigger
38         'paywall_conversion': 0.0138, # Higher rate
39         'new_subscribers': None
40     }
41
42     treatment_funnel['engaged'] = treatment_funnel['visitors'] * (
43         1 - treatment_funnel['bounce_rate']
44     )
45     treatment_funnel['reach_paywall'] = treatment_funnel['engaged'] * (
46         treatment_funnel['reach_paywall_rate']
47     )
48     treatment_funnel['new_subscribers'] = treatment_funnel['reach_
49         paywall'] * (
50         treatment_funnel['paywall_conversion']
51     )
52
53     # Print analysis
54     print("=== Funnel Analysis ===\n")

```



```

53 print("Control:")
54 print(f"1. Visitors: {control_funnel['visitors']:, .0f}")
55 print(f"2. Engaged (not bounced): {control_funnel['engaged']:, .0f}")
56     f"({(1-control_funnel['bounce_rate'])*100:.1f}%)"
57 print(f"3. Reached paywall: {control_funnel['reach_paywall']:, .0f}")
58     f"({control_funnel['reach_paywall_rate']*100:.1f}%)"
59 print(f"4. NEW SUBSCRIBERS: {control_funnel['new_subscribers']:, .0f}")
60     f"({control_funnel['paywall_conversion']*100:.2f}% conversion)"
61
62
63 print("\nTreatment:")
64 print(f"1. Visitors: {treatment_funnel['visitors']:, .0f}")
65 print(f"2. Engaged (not bounced): {treatment_funnel['engaged']:, .0f}")
66     f"({(1-treatment_funnel['bounce_rate'])*100:.1f}%)"
67 print(f"3. Reached paywall: {treatment_funnel['reach_paywall']:, .0f}")
68     f"({treatment_funnel['reach_paywall_rate']*100:.1f}%)"
69 print(f"4. NEW SUBSCRIBERS: {treatment_funnel['new_subscribers']:, .0f}")
70     f"({treatment_funnel['paywall_conversion']*100:.2f}% conversion)"
71
72 print("\n=== Key Insight ===")
73 subscriber_diff = (treatment_funnel['new_subscribers'] -
74                    control_funnel['new_subscribers'])
75 rel_diff = subscriber_diff / control_funnel['new_subscribers'] *
76            100
77 print(f"Net subscriber change: {subscriber_diff:, .0f} ({rel_diff:+.1f}%)")
78 print("\nProblem: Higher conversion rate, but fewer total subscribers!")
79 print("Reason: Aggressive paywall drove away more users than it converted.")
80
81 # Visualization
82 fig, axes = plt.subplots(1, 2, figsize=(14, 6))
83
84 # Funnel visualization
85 ax = axes[0]
86
87 stages = ['Visitors', 'Engaged', 'Reached\nPaywall', 'Subscribed']
88 control_values = [
89     control_funnel['visitors'],
90     control_funnel['engaged'],
91     control_funnel['reach_paywall'],
92     control_funnel['new_subscribers']
93 ]
94 treatment_values = [
95     treatment_funnel['visitors'],
96     treatment_funnel['engaged'],
97     treatment_funnel['reach_paywall'],
98     treatment_funnel['new_subscribers']

```

```

99     ]
100
101     x = np.arange(len(stages))
102     width = 0.35
103
104     ax.bar(x - width/2, control_values, width, label='Control',
105            alpha=0.8, color='C0')
106     ax.bar(x + width/2, treatment_values, width, label='Treatment',
107            alpha=0.8, color='C1')
108
109     ax.set_ylabel('Users')
110     ax.set_title('Subscription Funnel Comparison')
111     ax.set_xticks(x)
112     ax.set_xticklabels(stages)
113     ax.legend()
114     ax.grid(axis='y', alpha=0.3)
115     ax.set_yscale('log')
116
117     # Highlight the problem
118     for i, (c_val, t_val) in enumerate(zip(control_values, treatment_
119                                         values)):
120         if i == len(stages) - 1: # Final stage
121             if t_val < c_val:
122                 ax.text(i, max(c_val, t_val) * 1.5, 'Problem!',
123                        ha='center', color='red', fontweight='bold',
124                        fontsize=12)
125
126     # Conversion rates vs. absolute numbers
127     ax = axes[1]
128
129     metrics = ['Conversion\nRate', 'Total\nSubscribers']
130     control_rel = [100, 100] # Baseline
131     treatment_rel = [
132         (treatment_funnel['paywall_conversion'] /
133          control_funnel['paywall_conversion']) * 100,
134         (treatment_funnel['new_subscribers'] /
135          control_funnel['new_subscribers']) * 100
136     ]
137
138     x = np.arange(len(metrics))
139     width = 0.35
140
141     ax.bar(x - width/2, control_rel, width, label='Control (baseline)',
142            alpha=0.8, color='C0')
143     bars = ax.bar(x + width/2, treatment_rel, width, label='Treatment',
144                   alpha=0.8, color='C1')
145
146     # Color bars by good/bad
147     bars[0].set_color('green') # Higher rate is green
148     bars[1].set_color('red')   # Fewer subscribers is red
149
150     ax.axhline(y=100, color='black', linestyle='--', alpha=0.5)
151     ax.set_ylabel('Relative to Control (%)')
152     ax.set_title('The Optimization Paradox')
153     ax.set_xticks(x)
154     ax.set_xticklabels(metrics)
155     ax.legend()
156     ax.grid(axis='y', alpha=0.3)

```

```
156
157     # Add value labels
158     for i, (val, bar) in enumerate(zip(treatment_rel, bars)):
159         ax.text(bar.get_x() + bar.get_width()/2, val + 2,
160                 f'{val:.0f}%', ha='center', fontweight='bold')
161
162     plt.tight_layout()
163     plt.savefig('paywall_failure_analysis.pdf', bbox_inches='tight')
164     plt.show()
165
166 analyze_paywall_funnel()
```

11.5.6 Lessons Learned

1. **Holistic metrics:** Measure the entire funnel, not isolated steps
2. **Guardrail importance:** Protect critical health metrics (engagement, retention)
3. **Local vs. global optimization:** Improving one metric can harm overall goal
4. **Customer lifetime value:** Consider long-term impact beyond immediate conversion
5. **Document failures:** Failed experiments provide valuable learning

Corrective actions:

- Revised success criteria to include engagement and retention
- Established mandatory guardrail metrics for all experiments
- Implemented funnel analysis as standard practice
- Created decision framework weighting multiple objectives

11.6 Common Anti-Patterns and How to Avoid Them

11.6.1 Anti-Pattern 1: P-Hacking

Definition: Manipulating analysis until finding statistical significance.

Manifestations:

- Testing many metrics without correction
- Stopping experiment early when results look good
- Trying different segmentations until finding significance
- Removing "outliers" selectively

Example: A team ran an email subject line test. Primary metric (open rate) showed no effect. They then tested click rate, conversion rate, revenue per user, time to conversion, and found one significant result (time to conversion, $p = 0.04$) and declared success.

Why it's wrong: Testing 5 metrics with $\alpha = 0.05$ gives $1 - (0.95)^5 = 22.6\%$ false positive rate.

Prevention:

- Pre-register hypotheses and analysis plan
- Apply Bonferroni or FDR correction for multiple metrics
- Use sequential testing methods if monitoring continuously
- Document all analyses, including non-significant results

11.6.2 Anti-Pattern 2: HiPPO (Highest Paid Person's Opinion)

Definition: Overriding experimental evidence based on executive intuition or seniority.

Example: An e-commerce A/B test showed that removing a product feature decreased conversion by 8% ($p < 0.001$). VP insisted on shipping anyway because "customers told me they wanted this."

Why it's problematic: Undermines experimentation culture, leads to poor decisions, demoralizes data teams.

Prevention:

- Establish decision-making framework upfront
- Educate leadership on statistical methodology
- Create "escape valves" for overriding data (with documentation)
- Balance quantitative data with qualitative insights systematically
- Foster culture where challenging authority with data is encouraged

11.6.3 Anti-Pattern 3: Analysis Paralysis

Definition: Running experiments indefinitely without making decisions.

Example: Team ran checkout test for 8 weeks (planned duration: 2 weeks). Results were conclusive after 2 weeks, but team kept running "to be sure," delaying feature launch.

Prevention:

- Set clear success criteria before starting
- Establish decision deadlines
- Accept that uncertainty is inherent
- Define minimum confidence levels
- Implement decision frameworks (when to ship, when to iterate, when to kill)

11.6.4 Anti-Pattern 4: Novelty Effect Ignorance

Definition: Ignoring that users may respond differently initially than long-term.

Example: Redesigned navigation initially showed +15% engagement. After 4 weeks, effect degraded to +2

Prevention:

- Run experiments long enough to capture habituation (typically 2-4 weeks minimum)
- Analyze early vs. late periods separately
- Use cohort analysis to distinguish novelty from sustained effects
- For major changes, maintain long-term holdout groups

11.7 Industry-Specific Considerations

11.7.1 E-commerce

Unique challenges:

- High-value, low-frequency transactions (variance)
- Seasonality (holidays, sales events)
- Multi-device journeys (attribution)
- Returning customer effects

Best practices:

- Use pre-experiment covariates (CUPED) to reduce variance
- Stratify by customer type (new/returning) and device
- Run experiments across seasonal periods or control for calendar effects
- Monitor cart abandonment as key guardrail
- Track long-term metrics (LTV, repeat purchase rate)

11.7.2 Social Media

Unique challenges:

- Network effects and interference
- Viral content creates non-independence
- User-generated content variability
- Engagement vs. time spent vs. well-being trade-offs

Best practices:

- Use cluster randomization (geo, time-based)
- Account for network structure in analysis
- Monitor both engagement and well-being metrics
- Use switchback designs for platform-level changes
- Be cautious of metric gaming (e.g., optimizing for clicks vs. value)

11.7.3 SaaS/B2B

Unique challenges:

- Long sales cycles
- Account-level vs. user-level randomization
- Low sample sizes (enterprise customers)
- Contractual obligations limiting experimentation

Best practices:

- Use hierarchical models for account/user nesting
- Focus on leading indicators (trial engagement, feature adoption)
- Stratify by account size/value
- Supplement with qualitative research (customer interviews)
- Run longer experiments to capture full conversion cycle

11.7.4 Healthcare/EdTech

Unique challenges:

- Ethical constraints on randomization
- Long-term outcomes more important than short-term
- Regulatory compliance (HIPAA, FERPA, IRB)
- Equity and access considerations

Best practices:

- Prioritize equipoise (clinical uncertainty) for ethical randomization
- Use stepped-wedge designs (everyone gets treatment eventually)
- Implement robust informed consent processes
- Monitor equity metrics across demographic groups
- Focus on validated outcome measures
- Consider long-term follow-up studies

11.8 Building an Experimentation Culture: Meta-Lessons

Beyond individual experiments, successful organizations build systematic experimentation capabilities:

11.8.1 Infrastructure and Tooling

- **Experimentation platform:** Centralized system for randomization, tracking, analysis
- **Metric pipeline:** Automated computation of key metrics with proper variance estimation
- **Analysis tools:** Self-service dashboards with correct statistical methods
- **Experiment registry:** Database of all experiments for meta-analysis

11.8.2 Process and Governance

- **Hypothesis template:** Structured approach to defining experiments
- **Review process:** Statistical review before launch
- **Decision framework:** Clear criteria for ship/no-ship/iterate
- **Documentation:** Consistent recording of design, results, decisions

11.8.3 Culture and Education

- **Leadership buy-in:** Executives champion data-driven decisions
- **Training programs:** Educate PMs, engineers, analysts on methodology
- **Celebrate learning:** Reward well-designed experiments, not just "wins"
- **Transparent communication:** Share results widely, including failures

11.9 Summary

Real-world case studies provide invaluable lessons:

- **Google:** Technical performance directly impacts user behavior; always monitor guardrails
- **E-commerce:** Segment analysis reveals heterogeneous effects; weight by strategic priorities
- **Netflix:** Iterative experimentation validates assumptions incrementally; personalization adds value
- **Paywall failure:** Local optimization can harm global objectives; measure full funnel

- **Anti-patterns:** Avoid p-hacking, HiPPO, analysis paralysis; establish processes to prevent
- **Industry context:** Adapt methods to specific challenges (network effects, seasonality, regulations)

Learning from both successes and failures builds experimentation expertise. The most mature organizations view experiments as learning opportunities, documenting and sharing insights to improve continuously.

11.10 Further Reading

- **?:** Extensive case studies from Microsoft, Amazon, Google, with detailed analysis
- **?:** Exploration-exploitation trade-offs in practice
- Company engineering blogs:
 - Netflix TechBlog: experimentation at scale
 - Airbnb Engineering: experiment analysis frameworks
 - Booking.com: 1,000+ experiments per year
 - Spotify Research: sequential testing and bandits

11.11 Exercises

11.11.1 Case Study Analysis

1. **Post-Mortem Template:** Create a structured template for analyzing experiments. Include sections for: hypothesis, design, implementation, results, decision, lessons learned. Apply to a recent experiment from your organization.
2. **Replication:** Reproduce the Google case study analysis using the provided code. Modify parameters to explore sensitivity: What if load time increased by 50ms instead of 500ms? What if engagement elasticity were -0.5 instead of -1.0?
3. **Segment Discovery:** Using the e-commerce checkout code, add a third segment dimension (e.g., product category, geographic region). Analyze three-way interactions. When do you stop slicing data?

11.11.2 Business Impact Calculation

1. For your organization:
 - Define key business metrics (revenue, LTV, retention)
 - Map experiment metrics to business outcomes
 - Create calculator estimating business impact from experimental results
 - Include confidence intervals and break-even analysis
2. **Paywall Re-design:** The failed paywall experiment. Design a better version that balances conversion rate and engagement. Specify success criteria and guardrails.

11.11.3 Anti-Pattern Detection

1. **Audit:** Review 5 recent experiments from your organization. Identify any anti-patterns (p-hacking, missing guardrails, insufficient duration, etc.). Propose corrective actions.
2. **Process Design:** Design a review checklist to prevent common anti-patterns. Include statistical, technical, and business considerations. Specify who reviews, when, and with what authority.

11.11.4 Industry Application

1. Choose your industry and design an experiment addressing industry-specific challenges:
 - E-commerce: Multi-device attribution problem
 - Social media: Network effects mitigation
 - SaaS: Long-sales cycle with leading indicators
 - Healthcare: Ethical randomization with stepped-wedge

Specify design, analysis plan, and decision criteria.

11.11.5 Meta-Analysis Project

1. **Experiment Database:** Create a structured database of past experiments with fields: hypothesis, design, sample size, duration, primary metric, result, decision. Analyze:
 - What proportion of experiments are "successful"?
 - Do success rates differ by experiment type, team, or domain?
 - What factors predict successful experiments?
 - Are sample size and duration adequate on average?

Chapter 12

Ethics and Best Practices

12.1 Learning Objectives

After completing this chapter, readers will be able to:

- Navigate ethical considerations in human subjects experimentation
- Apply principles of informed consent and transparency
- Build organizational experimentation culture
- Establish governance frameworks
- Follow industry best practices

12.2 Ethical Foundations

A/B testing involves human subjects and carries ethical responsibilities.

12.2.1 Core Ethical Principles

1. **Beneficence:** Maximize benefits, minimize harms
2. **Non-maleficence:** Do no harm
3. **Autonomy:** Respect individual agency
4. **Justice:** Ensure fair distribution of benefits and burdens
5. **Transparency:** Be open about practices

12.2.2 The Belmont Report

While focused on medical research, the Belmont Report's principles apply to digital experimentation:

- **Respect for persons:** Treat individuals as autonomous agents
- **Beneficence:** Maximize possible benefits and minimize possible harms
- **Justice:** Fair distribution of research benefits and burdens

12.3 Informed Consent

12.3.1 Consent in Digital Experiments

Unlike medical trials, digital A/B tests typically operate under:

General consent: Users agree to terms of service that may include experimentation

Challenges:

- Users rarely read terms of service
- Consent may not be truly informed
- Ongoing vs. one-time consent
- Right to opt out

12.3.2 When Explicit Consent is Required

- Testing potentially harmful interventions
- Collecting sensitive personal data beyond normal use
- Significantly changing core functionality
- Research published academically
- Regulated industries (healthcare, finance, education)

12.3.3 Balancing Innovation and Consent

Example 12.1 (Acceptable Implicit Consent). Testing button colors, layout variations, recommendation algorithms—changes within normal product optimization scope.

Example 12.2 (Requiring Explicit Consent). Testing emotional manipulation, significantly degraded experiences, collection of biometric data.

12.4 Ethical Guardrails

12.4.1 Minimal Risk Standard

Experiments should pose no more than minimal risk—risks not greater than encountered in daily life or routine tests.

Risk assessment questions:

1. Could this experiment cause physical, psychological, or economic harm?
2. Are vulnerable populations disproportionately affected?
3. Is there potential for manipulation or deception?
4. Could this damage user trust or brand reputation?
5. Are there privacy or data security concerns?

12.4.2 Experimentation Review Board

Large organizations should establish review processes:

Composition:

- Product leaders
- Legal/compliance representatives
- Ethics officer or designated ethicist
- Data scientists
- Privacy/security experts

Responsibilities:

- Review high-risk experiments
- Establish ethical guidelines
- Provide guidance on edge cases
- Monitor incidents and update policies

12.4.3 Automated Guardrails

Implement technical controls:

Listing 12.1: Automated Ethical Guardrails

```

1 class ExperimentEthicsChecker:
2     """Automated ethics and safety checks."""
3
4     def __init__(self):
5         self.high_risk_keywords = [
6             'children', 'minors', 'emotion', 'manipulation',
7             'addiction', 'health', 'safety', 'security'
8         ]
9         self.sensitive_metrics = [
10            'mental_health', 'addiction_score', 'emotional_state'
11        ]
12
13     def review_experiment(self, experiment_config):
14         """
15         Automated ethics review of experiment configuration.
16
17         Returns:
18         -----
19         dict: Review results and recommendations
20         """
21         flags = []
22
23         # Check for high-risk keywords in description
24         description = experiment_config.get('description', '').lower()
25         for keyword in self.high_risk_keywords:
26             if keyword in description:
27                 flags.append({

```

```

28         'severity': 'high',
29         'type': 'high_risk_keyword',
30         'message': f"High-risk_keyword_detected: '{keyword}'\n",
31         'action': 'Requires_human_review'
32     })
33
34     # Check for sensitive metrics
35     metrics = experiment_config.get('metrics', {})
36     all_metrics = (metrics.get('primary', []) +
37                   metrics.get('secondary', []))
38
39     for metric in all_metrics:
40         if metric in self.sensitive_metrics:
41             flags.append({
42                 'severity': 'high',
43                 'type': 'sensitive_metric',
44                 'message': f"Sensitive_metric: '{metric}'\n",
45                 'action': 'Requires_ethics_board_approval'
46             })
47
48     # Check for degraded experiences
49     variants = experiment_config.get('variants', [])
50     for variant in variants:
51         if variant.get('intentionally_degraded', False):
52             flags.append({
53                 'severity': 'medium',
54                 'type': 'degraded_experience',
55                 'message': f"Variant_{variant['name']}\nintentionally_degraded",
56                 'action': 'Document_justification'
57             })
58
59     # Check target population
60     targeting = experiment_config.get('targeting', {})
61     if 'age_max' in targeting and targeting['age_max'] < 18:
62         flags.append({
63             'severity': 'critical',
64             'type': 'minors_targeted',
65             'message': 'Experiment_targets_minors',
66             'action': 'Ethics_board_approval_and_parental_consent_required'
67         })
68
69     return {
70         'approved': len([f for f in flags if f['severity'] in
71                        ['high', 'critical']]) == 0,
72         'flags': flags,
73         'recommendation': self._get_recommendation(flags)
74     }
75
76     def _get_recommendation(self, flags):
77         """Generate_recommendation_based_on_flags."""
78         if not flags:
79             return "No_ethical_concerns_detected. Approved_for_launch."
80
81         critical = [f for f in flags if f['severity'] == 'critical']
82         high = [f for f in flags if f['severity'] == 'high']

```

```

83         if critical:
84             return "BLOCKED: Critical ethical issues. Ethics board
85                 review required."
86         elif high:
87             return "CAUTION: High-risk experiment. Manager approval
88                 required."
89         else:
90             return "PROCEED WITH CAUTION: Document ethical
91                 considerations."
92
93 # Example usage
94 experiment = {
95     'description': 'Testing notification frequency to increase
96     engagement',
97     'metrics': {
98         'primary': 'daily_active_users',
99         'secondary': ['session_duration', 'notification_clicks']
100     },
101     'variants': [
102         {'name': 'control', 'notification_freq': 1},
103         {'name': 'treatment', 'notification_freq': 10}
104     ]
105 }
106
107 checker = ExperimentEthicsChecker()
108 result = checker.review_experiment(experiment)
109
110 print(f"Approved: {result['approved']}")
111 print(f"Recommendation: {result['recommendation']}")
112 for flag in result['flags']:
113     print(f"{flag['message']}")

```

12.5 Privacy and Data Protection

12.5.1 Data Minimization

Collect only data necessary for the experiment:

- Define metrics clearly and collect only those
- Aggregate early when possible
- Delete raw data after analysis
- Anonymize or pseudonymize user identifiers

12.5.2 Regulatory Compliance

GDPR (General Data Protection Regulation):

- Right to access: Users can request their data
- Right to erasure: "Right to be forgotten"

- Data protection by design
- Lawful basis for processing

CCPA (California Consumer Privacy Act):

- Right to know what data is collected
- Right to delete personal information
- Right to opt out of sale
- Non-discrimination for exercising rights

12.5.3 Internal Data Governance

1. **Access controls:** Limit who can view individual-level data
2. **Audit trails:** Log data access and analysis
3. **Retention policies:** Delete data after specified period
4. **Anonymization:** Remove personally identifiable information
5. **Differential privacy:** Add noise to protect individual privacy

12.6 Transparency and Communication

12.6.1 Internal Transparency

Experiment registry:

- Centralized catalog of all experiments
- Pre-registration of hypotheses
- Documentation of methods and results
- Sharing of negative results

Benefits:

- Prevents duplicate testing
- Enables meta-analysis
- Builds institutional knowledge
- Reduces publication bias

12.6.2 External Transparency

User-facing:

- Clear terms of service
- Privacy policies mentioning experimentation
- Opt-out mechanisms where appropriate
- Transparency reports

Academic:

- Publishing findings in peer-reviewed venues
- Open-sourcing methodologies
- Sharing datasets (when appropriate)
- Contributing to field knowledge

12.7 Building Experimentation Culture

12.7.1 Cultural Principles

1. **Data-informed, not data-driven:** Balance data with judgment
2. **Embrace failure:** Learn from experiments that don't succeed
3. **Democratize experimentation:** Enable all teams to test
4. **Question assumptions:** Challenge conventional wisdom
5. **Iterate rapidly:** Small, frequent experiments over big bets
6. **Share learnings:** Disseminate insights across organization

12.7.2 Organizational Structure

Centralized model:

- Dedicated experimentation team
- Ensures methodological rigor
- Can become bottleneck

Decentralized model:

- Each team runs own experiments
- Faster iteration
- Risk of inconsistent standards

Hybrid model (recommended):

- Central platform and guidelines
- Distributed execution
- Centers of excellence provide support

12.7.3 Training and Education

Invest in capability building:

- Statistical literacy training
- Experimentation best practices workshops
- Case study discussions
- Mentorship programs
- Documentation and runbooks

12.8 Best Practices Checklist**12.8.1 Design Phase**

- ☐ Clear hypothesis stated
- ☐ Primary metric identified
- ☐ Sample size calculated
- ☐ Guardrail metrics defined
- ☐ Ethical review completed
- ☐ Success criteria pre-specified

12.8.2 Implementation Phase

- ☐ Randomization mechanism tested
- ☐ Instrumentation validated
- ☐ SRM monitoring configured
- ☐ Logging working correctly
- ☐ Configuration documented

12.8.3 Analysis Phase

- ☐ Pre-analysis quality checks passed
- ☐ Multiple testing corrections applied
- ☐ Subgroup analysis planned
- ☐ Confidence intervals reported
- ☐ Effect sizes calculated
- ☐ Business impact estimated

12.8.4 Decision Phase

- ☐ Results communicated clearly
- ☐ Statistical and practical significance assessed
- ☐ Stakeholders aligned
- ☐ Implementation plan created
- ☐ Learnings documented

12.9 Common Pitfalls and How to Avoid Them

Pitfall	Consequence	Prevention
No pre-specified hypothesis	P-hacking, false positives	Pre-register experiments
Peeking at results	Inflated Type I error	Sequential testing methods
Ignoring guardrails	Harm to user experience	Monitor health metrics
Small sample sizes	Underpowered tests	Power analysis upfront
Short duration	Missing long-term effects	Run for full business cycles
No documentation	Lost institutional knowledge	Experiment registry

12.10 Summary

Ethical and effective experimentation requires:

- Commitment to ethical principles and user welfare
- Robust governance and review processes
- Privacy protection and regulatory compliance
- Transparency internally and externally
- Strong organizational culture supporting experimentation

- Continuous learning and improvement

A/B testing is powerful but carries responsibility. Practitioners must balance innovation with ethics, rigor with pragmatism, and data with judgment.

12.11 Further Reading

- ? : Controversial study highlighting ethical issues
- ? : Ethics of tech experimentation
- ? : Building experimentation culture
- IRB guidelines: Institutional Review Board frameworks

12.12 Exercises

1. Create an ethics review framework for your organization
2. Analyze a controversial experiment (e.g., Facebook emotional contagion study)
3. Design an experimentation training program
4. Build an experiment registry system
5. Draft transparency report for external publication

"The goal is to turn data into information, and information into insight."
— Carly Fiorina

Mathematical Proofs

.1 Mathematical Proofs

This appendix provides rigorous mathematical proofs for key theorems referenced throughout the book.

.1.1 Proof: Unbiasedness of Sample Mean

Theorem .1. Let $X_1, X_2, \dots, X_n$ be independent and identically distributed random variables with $\mathbb{E}[X_i] = \mu$. Then the sample mean $\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$ is an unbiased estimator of μ .

Proof. By linearity of expectation:

$$\mathbb{E}[\bar{X}] = \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n X_i\right] \tag{1}$$

$$= \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i] \tag{2}$$

$$= \frac{1}{n} \sum_{i=1}^n \mu \tag{3}$$

$$= \frac{1}{n} \cdot n\mu \tag{4}$$

$$= \mu \tag{5}$$

Therefore, $\bar{X}$ is an unbiased estimator of μ . □

.1.2 Proof: Variance of Sample Mean

Theorem .2. For independent random variables $X_1, \dots, X_n$ with $\text{Var}(X_i) = \sigma^2$:

$$\text{Var}(\bar{X}) = \frac{\sigma^2}{n} \tag{6}$$

Proof. Using independence:

$$\text{Var}(\bar{X}) = \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \quad (7)$$

$$= \frac{1}{n^2} \text{Var}\left(\sum_{i=1}^n X_i\right) \quad (8)$$

$$= \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) \quad (\text{by independence}) \quad (9)$$

$$= \frac{1}{n^2} \sum_{i=1}^n \sigma^2 \quad (10)$$

$$= \frac{1}{n^2} \cdot n\sigma^2 \quad (11)$$

$$= \frac{\sigma^2}{n} \quad (12)$$

□

.1.3 Proof: Beta-Binomial Conjugacy

Theorem .3. If $p \sim \text{Beta}(\alpha, \beta)$ and $X|p \sim \text{Binomial}(n, p)$, then:

$$p|X \sim \text{Beta}(\alpha + X, \beta + n - X) \quad (13)$$

Proof. By Bayes' theorem:

$$f(p|X) \propto f(X|p) \cdot f(p) \quad (14)$$

The likelihood is:

$$f(X|p) = \binom{n}{X} p^X (1-p)^{n-X} \quad (15)$$

The prior is:

$$f(p) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} p^{\alpha-1} (1-p)^{\beta-1} \quad (16)$$

The posterior is proportional to:

$$f(p|X) \propto p^X (1-p)^{n-X} \cdot p^{\alpha-1} (1-p)^{\beta-1} \quad (17)$$

$$= p^{X+\alpha-1} (1-p)^{n-X+\beta-1} \quad (18)$$

This is the kernel of a $\text{Beta}(\alpha + X, \beta + n - X)$ distribution. □

.1.4 Proof: Sample Size Formula for Proportions

Theorem .4. For comparing two proportions p_1 and p_2 with significance level α and power $1 - \beta$:

$$n = \frac{(z_{\alpha/2} + z_{\beta})^2 [p_1(1-p_1) + p_2(1-p_2)]}{(p_2 - p_1)^2} \quad (19)$$

Proof. Under $H_0 : p_1 = p_2 = p$, using pooled proportion:

$$Z_0 = \frac{\hat{p}_2 - \hat{p}_1}{\sqrt{2p(1-p)/n}} \sim \mathcal{N}(0, 1) \quad (20)$$

Under $H_1 : p_1 \neq p_2$:

$$Z_1 = \frac{(\hat{p}_2 - \hat{p}_1) - (p_2 - p_1)}{\sqrt{[p_1(1-p_1) + p_2(1-p_2)]/n}} \sim \mathcal{N}(0, 1) \quad (21)$$

Rejection region: $|\hat{p}_2 - \hat{p}_1| > z_{\alpha/2} \sqrt{2p(1-p)/n}$

For power $1 - \beta$, we require:

$$P\left(|\hat{p}_2 - \hat{p}_1| > z_{\alpha/2} \sqrt{2p(1-p)/n} \mid H_1\right) = 1 - \beta \quad (22)$$

Assuming $p_2 > p_1$ and approximating $p \approx (p_1 + p_2)/2$:

$$\frac{(p_2 - p_1) - z_{\alpha/2} \sqrt{2p(1-p)/n}}{\sqrt{[p_1(1-p_1) + p_2(1-p_2)]/n}} = z_\beta \quad (23)$$

Solving for n yields the formula above (with slight approximation). $\square$

.1.5 Proof: CUPED Variance Reduction

Theorem .5. For adjusted outcome $Y^* = Y - \theta(X - \mathbb{E}[X])$ with $\theta = \frac{\text{Cov}(Y, X)}{\text{Var}(X)}$:

$$\text{Var}(Y^*) = \text{Var}(Y)(1 - \rho^2) \quad (24)$$

where $\rho = \text{Cor}(Y, X)$.

Proof.

$$\text{Var}(Y^*) = \text{Var}(Y - \theta(X - \mathbb{E}[X])) \quad (25)$$

$$= \text{Var}(Y) + \theta^2 \text{Var}(X) - 2\theta \text{Cov}(Y, X) \quad (26)$$

$$= \text{Var}(Y) + \frac{\text{Cov}(Y, X)^2}{\text{Var}(X)^2} \text{Var}(X) - 2 \frac{\text{Cov}(Y, X)}{\text{Var}(X)} \text{Cov}(Y, X) \quad (27)$$

$$= \text{Var}(Y) + \frac{\text{Cov}(Y, X)^2}{\text{Var}(X)} - 2 \frac{\text{Cov}(Y, X)^2}{\text{Var}(X)} \quad (28)$$

$$= \text{Var}(Y) - \frac{\text{Cov}(Y, X)^2}{\text{Var}(X)} \quad (29)$$

$$= \text{Var}(Y) \left(1 - \frac{\text{Cov}(Y, X)^2}{\text{Var}(Y)\text{Var}(X)}\right) \quad (30)$$

$$= \text{Var}(Y)(1 - \rho^2) \quad (31)$$

$\square$

.1.6 Central Limit Theorem

Theorem .6 (Lindeberg-Lévy CLT). Let $X_1, X_2, \dots$ be i.i.d. random variables with $\mathbb{E}[X_i] = \mu$ and $\text{Var}(X_i) = \sigma^2 < \infty$. Then:

$$\frac{\sqrt{n}(\bar{X}_n - \mu)}{\sigma} \xrightarrow{d} \mathcal{N}(0, 1) \quad (32)$$

The full proof requires characteristic functions and is beyond scope, but the key idea is that standardized sums of i.i.d. random variables converge in distribution to the standard normal distribution.

.1.7 Bonferroni Inequality

Theorem .7 (Bonferroni Correction). For m hypothesis tests, each at level α/m , the family-wise error rate satisfies:

$$\text{FWER} = P(\text{Reject at least one true } H_0) \leq \alpha \quad (33)$$

Proof. Let R_i be the event that test i rejects a true null hypothesis. By Boole's inequality:

$$\text{FWER} = P\left(\bigcup_{i=1}^m R_i\right) \quad (34)$$

$$\leq \sum_{i=1}^m P(R_i) \quad (35)$$

$$= \sum_{i=1}^m \frac{\alpha}{m} \quad (36)$$

$$= \alpha \quad (37)$$

□

.1.8 Delta Method

Theorem .8 (Delta Method). Let $\sqrt{n}(\hat{\theta} - \theta) \xrightarrow{d} \mathcal{N}(0, \sigma^2)$ and g be differentiable at θ with $g'(\theta) \neq 0$. Then:

$$\sqrt{n}(g(\hat{\theta}) - g(\theta)) \xrightarrow{d} \mathcal{N}(0, [g'(\theta)]^2 \sigma^2) \quad (38)$$

Application: Useful for transformations like log odds ratios or relative lifts.

Proof Sketch. By Taylor expansion around θ :

$$g(\hat{\theta}) \approx g(\theta) + g'(\theta)(\hat{\theta} - \theta) \quad (39)$$

Therefore:

$$\sqrt{n}(g(\hat{\theta}) - g(\theta)) \approx g'(\theta)\sqrt{n}(\hat{\theta} - \theta) \quad (40)$$

Since $\sqrt{n}(\hat{\theta} - \theta) \xrightarrow{d} \mathcal{N}(0, \sigma^2)$ and multiplication by constant preserves normality, the result follows. □

Code Examples

.2 Complete Code Examples

This appendix provides complete, production-ready code implementations for common A/B testing tasks.

.2.1 Comprehensive A/B Test Analysis Class (Python)

Listing 2: Automated Ethical Guardrails

```
1 import numpy as np
2 import pandas as pd
3 from scipy import stats
4 from dataclasses import dataclass
5 from typing import Optional, Dict, List
6
7 @dataclass
8 class ABTestResult:
9     """Container for A/B test results."""
10     control_mean: float
11     treatment_mean: float
12     difference: float
13     relative_lift: float
14     standard_error: float
15     confidence_interval: tuple
16     p_value: float
17     statistically_significant: bool
18     test_statistic: float
19     sample_size_control: int
20     sample_size_treatment: int
21
22 class ABTestAnalyzer:
23     """
24     Comprehensive A/B test analysis.
25
26     Supports:
27     - Proportions (conversion rates)
28     - Continuous metrics (revenue, time, etc.)
29     - Multiple testing corrections
30     - Subgroup analysis
31     """
32
33     def __init__(self, alpha=0.05):
34         self.alpha = alpha
35
36     def test_proportions(self, control_conversions, control_total,
```

```

37         treatment_conversions, treatment_total,
38         method='z_test'):
39
40     """
41     Test difference in proportions.
42
43     Parameters:
44     -----
45     control_conversions: int
46     Number of successes in control
47     control_total: int
48     Total observations in control
49     treatment_conversions: int
50     Number of successes in treatment
51     treatment_total: int
52     Total observations in treatment
53     method: str
54     'z_test', 'chi_squared', or 'fisher_exact'
55
56     Returns:
57     -----
58     ABTestResult: Test results
59     """
60     p1 = control_conversions / control_total
61     p2 = treatment_conversions / treatment_total
62
63     if method == 'z_test':
64         # Pooled proportion
65         p_pool = ((control_conversions + treatment_conversions) /
66                  (control_total + treatment_total))
67
68         # Standard error under null
69         se_null = np.sqrt(p_pool * (1 - p_pool) *
70                          (1/control_total + 1/treatment_total))
71
72         # Test statistic
73         z_stat = (p2 - p1) / se_null
74         p_value = 2 * (1 - stats.norm.cdf(abs(z_stat)))
75
76         # Confidence interval (unpooled)
77         se_unpooled = np.sqrt(p1*(1-p1)/control_total +
78                               p2*(1-p2)/treatment_total)
79         z_crit = stats.norm.ppf(1 - self.alpha/2)
80         ci = (p2 - p1 - z_crit * se_unpooled,
81              p2 - p1 + z_crit * se_unpooled)
82
83         test_stat = z_stat
84         se = se_unpooled
85
86     elif method == 'chi_squared':
87         table = [[control_conversions, control_total - control_
88                  conversions],
89                  [treatment_conversions, treatment_total - treatment
90                   _conversions]]
91         chi2, p_value, dof, expected = stats.chi2_contingency(table
92 )
93         test_stat = chi2
94
95     # CI from z-test

```

```

92         se = np.sqrt(p1*(1-p1)/control_total +
93                     p2*(1-p2)/treatment_total)
94         z_crit = stats.norm.ppf(1 - self.alpha/2)
95         ci = (p2 - p1 - z_crit * se, p2 - p1 + z_crit * se)
96
97     elif method == 'fisher_exact':
98         table = [[control_conversions, control_total - control_
99                   conversions],
100                 [treatment_conversions, treatment_total - treatment
101                   _conversions]]
102         _, p_value = stats.fisher_exact(table)
103         test_stat = None
104
105     # CI from z-test
106     se = np.sqrt(p1*(1-p1)/control_total +
107                 p2*(1-p2)/treatment_total)
108     z_crit = stats.norm.ppf(1 - self.alpha/2)
109     ci = (p2 - p1 - z_crit * se, p2 - p1 + z_crit * se)
110
111     return ABTestResult(
112         control_mean=p1,
113         treatment_mean=p2,
114         difference=p2 - p1,
115         relative_lift=(p2 - p1) / p1 if p1 > 0 else np.nan,
116         standard_error=se,
117         confidence_interval=ci,
118         p_value=p_value,
119         statistically_significant=p_value < self.alpha,
120         test_statistic=test_stat,
121         sample_size_control=control_total,
122         sample_size_treatment=treatment_total
123     )
124
125     def test_means(self, control_data, treatment_data,
126                   equal_variance=False):
127         """
128         Test difference in means.
129
130         Parameters:
131         control_data: array-like
132         Control group values
133         treatment_data: array-like
134         Treatment group values
135         equal_variance: bool
136         Assume equal variances (default False, uses Welch's t-test)
137
138         Returns:
139         ABTestResult: Test results
140         """
141         control_data = np.array(control_data)
142         treatment_data = np.array(treatment_data)
143
144         n1 = len(control_data)
145         n2 = len(treatment_data)
146         mean1 = np.mean(control_data)
147         mean2 = np.mean(treatment_data)

```

```

148     var1 = np.var(control_data, ddof=1)
149     var2 = np.var(treatment_data, ddof=1)
150
151     # T-test
152     t_stat, p_value = stats.ttest_ind(treatment_data, control_data,
153                                     equal_var=equal_variance)
154
155     # Standard error and CI
156     se = np.sqrt(var1/n1 + var2/n2)
157
158     if equal_variance:
159         df = n1 + n2 - 2
160     else:
161         # Welch-Satterthwaite degrees of freedom
162         df = ((var1/n1 + var2/n2)**2 /
163              ((var1/n1)**2/(n1-1) + (var2/n2)**2/(n2-1)))
164
165     t_crit = stats.t.ppf(1 - self.alpha/2, df)
166     ci = (mean2 - mean1 - t_crit * se, mean2 - mean1 + t_crit * se)
167
168     return ABTestResult(
169         control_mean=mean1,
170         treatment_mean=mean2,
171         difference=mean2 - mean1,
172         relative_lift=(mean2 - mean1) / mean1 if mean1 > 0 else np.
173             nan,
174         standard_error=se,
175         confidence_interval=ci,
176         p_value=p_value,
177         statistically_significant=p_value < self.alpha,
178         test_statistic=t_stat,
179         sample_size_control=n1,
180         sample_size_treatment=n2
181     )
182
183     def calculate_sample_size_proportions(self, p1, p2, power=0.80):
184         """
185         Calculate required sample size for proportion test.
186
187         Parameters:
188         p1: float
189             Control proportion
190         p2: float
191             Treatment proportion
192         power: float
193             Desired statistical power
194
195         Returns:
196         dict: Sample size requirements
197         """
198         z_alpha = stats.norm.ppf(1 - self.alpha/2)
199         z_beta = stats.norm.ppf(power)
200
201         variance = p1*(1-p1) + p2*(1-p2)
202         effect = (p2 - p1)**2
203
204

```

```

205     n_per_group = int(np.ceil(
206         ((z_alpha + z_beta)**2 * variance) / effect
207     ))
208
209     return {
210         'n_per_group': n_per_group,
211         'n_total': 2 * n_per_group,
212         'p1': p1,
213         'p2': p2,
214         'effect_absolute': p2 - p1,
215         'effect_relative': (p2 - p1) / p1,
216         'alpha': self.alpha,
217         'power': power
218     }
219
220     def format_results(self, result: ABTestResult) -> str:
221         """Format results for reporting."""
222         report = f"""
223 A/B Test Results
224 =====
225 Control Mean: {result.control_mean:.4f}
226 Treatment Mean: {result.treatment_mean:.4f}
227
228 Difference: {result.difference:.4f}
229 Relative Lift: {result.relative_lift:.2%}
230 Standard Error: {result.standard_error:.4f}
231
232 95% Confidence Interval: [{result.confidence_interval[0]:.4f}, {result.
233     confidence_interval[1]:.4f}]
234
235 Test Statistic: {result.test_statistic:.4f}
236 P-value: {result.p_value:.4f}
237 Statistically Significant: {'Yes' if result.statistically_significant
238     else 'No'}
239
240 Sample Sizes:
241   Control: {result.sample_size_control:,}
242   Treatment: {result.sample_size_treatment:,}
243 """
244         return report
245
246 # Example usage
247 if __name__ == "__main__":
248     analyzer = ABTestAnalyzer(alpha=0.05)
249
250     # Test proportions
251     result_prop = analyzer.test_proportions(
252         control_conversions=80,
253         control_total=1000,
254         treatment_conversions=95,
255         treatment_total=1000
256     )
257
258     print(analyzer.format_results(result_prop))
259
260     # Calculate sample size
261     sample_size = analyzer.calculate_sample_size_proportions(
262         p1=0.08,

```

```

261     p2=0.09,
262     power=0.80
263 )
264
265 print("\nSample_Size_Calculation:")
266 print(f"Required_per_group:{sample_size['n_per_group'],}")
267 print(f"Total_required:{sample_size['n_total'],}")

```

.2.2 Experiment Analysis Pipeline (R)

Listing 3: Automated Ethical Guardrails

```

1 library(tidyverse)
2 library(broom)
3
4 # Comprehensive experiment analysis pipeline
5 analyze_experiment <- function(data, metric, variant_col = "variant",
6                               alpha = 0.05) {
7   #' Analyze A/B test experiment
8   #'
9   #' @param data DataFrame with experiment data
10  #' @param metric Character, name of metric column
11  #' @param variant_col Character, name of variant column
12  #' @param alpha Numeric, significance level
13  #'
14  #' @return List with analysis results
15
16  # 1. Pre-analysis checks
17  cat("===Pre-AnalysisChecks===\n")
18
19  # Sample ratio mismatch
20  variant_counts <- table(data[[variant_col]])
21  srm_test <- chisq.test(variant_counts)
22  cat(sprintf("SRM_p-value: %.4f%s\n", srm_test$p.value,
23            ifelse(srm_test$p.value < 0.01, "(WARNING)", "(OK)")))
24
25  # 2. Summary statistics
26  cat("\n===SummaryStatistics===\n")
27  summary_stats <- data %>%
28    group_by(!!sym(variant_col)) %>%
29    summarise(
30      n = n(),
31      mean = mean(!!sym(metric), na.rm = TRUE),
32      median = median(!!sym(metric), na.rm = TRUE),
33      sd = sd(!!sym(metric), na.rm = TRUE),
34      se = sd / sqrt(n),
35      .groups = 'drop'
36    )
37
38  print(summary_stats)
39
40  # 3. Statistical test
41  cat("\n===StatisticalTest===\n")
42
43  variants <- unique(data[[variant_col]])
44  control <- data %>% filter(!!sym(variant_col) == variants[1]) %>%
45    pull(!!sym(metric))

```

[illegible]

```

103         n_sim = 100000) {
104   #' Bayesian A/B test for proportions
105   #'
106   #' @param x_control Conversions in control
107   #' @param n_control Sample size control
108   #' @param x_treatment Conversions in treatment
109   #' @param n_treatment Sample size treatment
110   #' @param prior_alpha Beta prior alpha
111   #' @param prior_beta Beta prior beta
112   #' @param n_sim Number of Monte Carlo simulations
113   #'
114   #' @return List with Bayesian results
115
116   # Posterior parameters
117   post_alpha_control <- prior_alpha + x_control
118   post_beta_control <- prior_beta + n_control - x_control
119   post_alpha_treatment <- prior_alpha + x_treatment
120   post_beta_treatment <- prior_beta + n_treatment - x_treatment
121
122   # Simulate from posteriors
123   p_control <- rbeta(n_sim, post_alpha_control, post_beta_control)
124   p_treatment <- rbeta(n_sim, post_alpha_treatment, post_beta_treatment
125   )
126
127   # Probability treatment > control
128   prob_treatment_better <- mean(p_treatment > p_control)
129
130   # Expected lift
131   lift <- p_treatment / p_control - 1
132   expected_lift <- mean(lift)
133   lift_ci <- quantile(lift, c(0.025, 0.975))
134
135   # Expected loss
136   loss_control <- mean(pmax(0, p_treatment - p_control))
137   loss_treatment <- mean(pmax(0, p_control - p_treatment))
138
139   list(
140     prob_treatment_better = prob_treatment_better,
141     expected_lift = expected_lift,
142     lift_ci = lift_ci,
143     loss = list(
144       control = loss_control,
145       treatment = loss_treatment,
146       recommendation = ifelse(loss_control < loss_treatment,
147                               "Control", "Treatment")
148     )
149   )
150 }
151
152 # Example usage:
153 # data <- read_csv("experiment_data.csv")
154 # results <- analyze_experiment(data, metric = "conversion")
155 # print(results$plots$boxplot)

```

.2.3 Experiment Simulation Framework

Listing 4: Automated Ethical Guardrails

```

1 import numpy as np
2 import pandas as pd
3 from typing import Callable
4
5 class ExperimentSimulator:
6     """Simulate A/B tests for power analysis and methodology validation
7     """
8
9     def __init__(self, seed=None):
10         if seed is not None:
11             np.random.seed(seed)
12
13     def simulate_proportions(self, n_control, n_treatment,
14                             p_control, p_treatment,
15                             n_simulations=1000, alpha=0.05):
16         """
17         Simulate A/B tests with binary outcomes.
18
19         Returns:
20         dict: Simulation results including empirical power and Type I
21             error
22         """
23         results = []
24
25         for _ in range(n_simulations):
26             # Generate data
27             control = np.random.binomial(1, p_control, n_control)
28             treatment = np.random.binomial(1, p_treatment, n_treatment)
29
30             # Test
31             p1 = np.mean(control)
32             p2 = np.mean(treatment)
33
34             p_pooled = (np.sum(control) + np.sum(treatment)) / (n_
35                 control + n_treatment)
36             se = np.sqrt(p_pooled * (1 - p_pooled) * (1/n_control + 1/n_
37                 treatment))
38
39             z = (p2 - p1) / se
40             p_value = 2 * (1 - stats.norm.cdf(abs(z)))
41
42             results.append({
43                 'p_value': p_value,
44                 'reject': p_value < alpha,
45                 'p1_hat': p1,
46                 'p2_hat': p2,
47                 'effect_estimate': p2 - p1
48             })
49
50         results_df = pd.DataFrame(results)
51
52         # Calculate metrics
53         empirical_power = results_df['reject'].mean()
54         type_i_error = empirical_power if p_control == p_treatment else
55             None

```

```
53     return {
54         'results': results_df,
55         'empirical_power': empirical_power,
56         'type_i_error': type_i_error,
57         'mean_p_value': results_df['p_value'].mean(),
58         'significant_count': results_df['reject'].sum()
59     }
60
61 # Example: Verify power calculation
62 simulator = ExperimentSimulator(seed=42)
63 sim_result = simulator.simulate_proportions(
64     n_control=1000,
65     n_treatment=1000,
66     p_control=0.10,
67     p_treatment=0.12,
68     n_simulations=10000,
69     alpha=0.05
70 )
71
72 print(f"Empirical_power: {sim_result['empirical_power']:.3f}")
73 print(f"Significant_tests: {sim_result['significant_count']}/{10000}")
```

Statistical Tables

.3 Statistical Tables

This appendix provides reference tables for common statistical distributions and critical values used in A/B testing.

.3.1 Standard Normal Distribution (Z) Critical Values

Confidence Level	Significance (α)	$z_{\alpha/2}$
90%	0.10	1.645
95%	0.05	1.960
98%	0.02	2.326
99%	0.01	2.576
99.5%	0.005	2.807
99.9%	0.001	3.291

.3.2 Student's t-Distribution Critical Values ($t_{\alpha/2,df}$)

For two-sided tests at $\alpha = 0.05$:

Degrees of Freedom	$t_{0.025,df}$
5	2.571
10	2.228
15	2.131
20	2.086
25	2.060
30	2.042
40	2.021
50	2.009
60	2.000
80	1.990
100	1.984
∞	1.960

.3.3 Chi-Squared Distribution Critical Values ($\chi^2_{\alpha,df}$)

For significance level $\alpha = 0.05$:

Degrees of Freedom	$\chi^2_{0.05,df}$
1	3.841
2	5.991
3	7.815
4	9.488
5	11.070
6	12.592
7	14.067
8	15.507
9	16.919
10	18.307
15	24.996
20	31.410
25	37.652
30	43.773

.3.4 Sample Size Requirements (Two Proportions)

For $\alpha = 0.05$ (two-sided), power = 0.80, and baseline proportion p_1 :

Relative Lift	Baseline Proportion (p_1)				
	0.05	0.10	0.15	0.20	0.25
5%	62,448	30,081	19,718	14,616	11,540
10%	15,396	7,337	4,784	3,540	2,790
15%	6,772	3,213	2,092	1,547	1,218
20%	3,788	1,793	1,166	862	678
25%	2,405	1,137	739	546	429
30%	1,650	780	507	374	294

Note: Values show sample size per variant.

.3.5 Minimum Detectable Effect (MDE)

For $\alpha = 0.05$, power = 0.80, and sample size n per variant:

Baseline (p_1)	Sample Size per Variant			
	1,000	5,000	10,000	25,000
0.05	3.95%	1.76%	1.24%	0.79%
0.10	5.32%	2.38%	1.68%	1.06%
0.15	6.31%	2.82%	2.00%	1.26%
0.20	7.08%	3.17%	2.24%	1.42%
0.25	7.66%	3.42%	2.42%	1.53%

Note: Values show absolute minimum detectable effect (e.g., 0.05 = 5 percentage points).

.3.6 Bonferroni Corrected Significance Levels

For overall $\alpha = 0.05$:

Number of Tests (m)	Corrected $\alpha^* = \alpha/m$
2	0.0250
3	0.0167
4	0.0125
5	0.0100
6	0.0083
8	0.0063
10	0.0050
15	0.0033
20	0.0025
25	0.0020
50	0.0010
100	0.0005

.3.7 Power vs. Sample Size

For two-proportion test with $p_1 = 0.10$, $p_2 = 0.12$ (20% relative lift), $\alpha = 0.05$:

Sample Size per Variant	Statistical Power
1,000	0.17
2,000	0.27
3,000	0.35
4,000	0.42
5,000	0.49
7,500	0.62
10,000	0.72
15,000	0.85
20,000	0.92
25,000	0.96
30,000	0.98

.3.8 Beta Distribution Quantiles

For Beta(a, b) distributions commonly used as priors:

Distribution	Mean	Mode	2.5% Quantile	97.5% Quantile
Beta(1, 1)	0.500	-	0.025	0.975
Beta(2, 2)	0.500	0.500	0.082	0.918
Beta(5, 5)	0.500	0.500	0.186	0.814
Beta(1, 9)	0.100	0.000	0.003	0.352
Beta(10, 90)	0.100	0.099	0.050	0.165
Beta(50, 450)	0.100	0.099	0.071	0.133

.3.9 Conversion Rate Effect Size Examples

Real-world context for effect sizes at different baseline rates:

Baseline	5% Lift	10% Lift	20% Lift	50% Lift
1%	1.05%	1.10%	1.20%	1.50%
5%	5.25%	5.50%	6.00%	7.50%
10%	10.50%	11.00%	12.00%	15.00%
20%	21.00%	22.00%	24.00%	30.00%
50%	52.50%	55.00%	60.00%	75.00%

.3.10 Design Effect for Cluster Randomization

Design effect = $1 + (m - 1)\rho$ where m is cluster size and ρ is ICC:

Cluster Size (m)	Intraclass Correlation (ρ)			
	0.01	0.05	0.10	0.20
5	1.04	1.20	1.40	1.80
10	1.09	1.45	1.90	2.80
20	1.19	1.95	2.90	4.80
50	1.49	3.45	5.90	10.80
100	1.99	5.95	10.90	20.80

.3.11 Quick Reference Formulas

Standard Error for Proportion:

$$SE(\hat{p}) = \sqrt{\frac{p(1-p)}{n}}$$

Confidence Interval for Proportion:

$$\hat{p} \pm z_{\alpha/2} \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$$

Z-Test Statistic for Proportions:

$$Z = \frac{\hat{p}_2 - \hat{p}_1}{\sqrt{\bar{p}(1-\bar{p})(1/n_1 + 1/n_2)}}$$

Sample Size for Proportions:

$$n = \frac{(z_{\alpha/2} + z_{\beta})^2 [p_1(1-p_1) + p_2(1-p_2)]}{(p_2 - p_1)^2}$$

Relative Lift:

$$\text{Relative Lift} = \frac{p_2 - p_1}{p_1} = \frac{p_2}{p_1} - 1$$

Minimum Detectable Effect:

$$\text{MDE} \approx (z_{\alpha/2} + z_{\beta}) \sqrt{\frac{2p(1-p)}{n}}$$

3.12 Decision Matrix

Scenario	Recommended Approach
Small sample ($n < 100$)	Fisher’s exact test or permutation test
Large sample, proportions	Z-test or chi-squared test
Continuous metric, normal	T-test (Welch’s preferred)
Continuous metric, skewed	Transform data, use robust statistics, or permutation test
Multiple metrics ($m < 10$)	Bonferroni or Holm-Bonferroni correction
Multiple metrics ($m > 10$)	Benjamini-Hochberg FDR control
Sequential testing	Group sequential with spending function
Rare events ($p < 0.01$)	Longer duration, proxy metrics, or exact tests
Network effects	Cluster randomization with appropriate analysis